COBHAM

GR712RC

Dual-Core LEON3FT SPARC V8 Processor

User's Manual

GR712RC

Table of contents

14.4   Signal definitions 99

15   UART Serial Interface 100
15.1   Overview 100
15.2   Operation 100
15.3   Baud-rate generation 101
15.4   Loop back mode 101
15.5   FIFO debug mode 102
15.6   Interrupt generation 102
15.7   Registers 102
15.8   Signal definitions 104

16   SpaceWire Interface with RMAP support 105
16.1   Overview 105
16.2   Operation 105
16.3   Link interface 106
16.4   Receiver DMA channels 108
16.5   Transmitter DMA channels 114
16.6   RMAP 117
16.7   AMBA interface 122
16.8   SpaceWire clock generation 123
16.9   Registers 124
16.10   Signal definitions 129

17   Ethernet Media Access Controller (MAC) 130
17.1   Overview 130
17.2   Operation 130
17.3   Tx DMA interface 131
17.4   Rx DMA interface 132
17.5   MDIO Interface 134
17.6   Reduced Media Independent Interfaces (RMII) 134
17.7   Registers 134
17.8   Signal definitions 137

18   CAN Interface 138
18.1   Overview 138
18.2   CAN controller overview 138
18.3   AHB interface 138
18.4   BasicCAN mode 139
18.5   PeliCAN mode 143
18.6   Common registers 153
18.7   Design considerations 154
18.8   Signal definitions 154

19   Obsolete 155
19.1   Signal definitions 155

20   CAN Bus multiplexer 156
20.1   Overview 156
20.2   Operation 156
20.3   Registers 156

1 Introduction

1.1 Scope

This document is a user's manual for the radiation-hard GR712RC Dual-Core LEON3FT SPARC V8 processor. GR712RC is based on LEON3FT and IP cores from the GRLIB IP library, implemented with the Ramon RadSafe^{TM} 180 nm cell library on Tower Semiconductors 180nm CMOS process.

1.2 GR712RC Architecture

The GR712RC has the following on-chip features:

- 2 x LEON3FT processor cores with 32KiB cache and 32-entry TLB SRMMU
- 2 x Double-precision high-performance floating-point units (GRFPU V2)
- Debug Support Unit (DSU) with 256-line instruction and AMBA AHB trace buffers
- PROM/SRAM/SDRAM memory controller with BCH and Reed-Solomon error-correction
  - up to 32 MiB PROM over two 16 MiB banks
  - up to 32 MiB SRAM over two 16 MiB banks
  - up to 1 GiB SDRAM over two 512 MiB banks
- 192 KiB on-chip RAM with BCH error-correction (uncacheable)
- 6 x SpaceWire links (two with RMAP support, one DMA channel per link, one port per link)
- 6 x UARTs
- 6 x General Purpose Timers (2 with time latch capability)
- Multi-processor Interrupt Controller with support for 31 interrupts
- 2 x 32 bits General Purpose I/O
- JTAG debug link
- 10/100 Ethernet MAC with RMII interface
- MIL-STD-1553B BC/RT/BM controller
- 2 x CAN 2.0 controller
- High-speed CCSDS Telecommand decoder and Telemetry encoder
- SPI master controller
- I2C master controller
- SLINK and ASCS16 controllers
- Clock gating unit
- I/O switch matrix
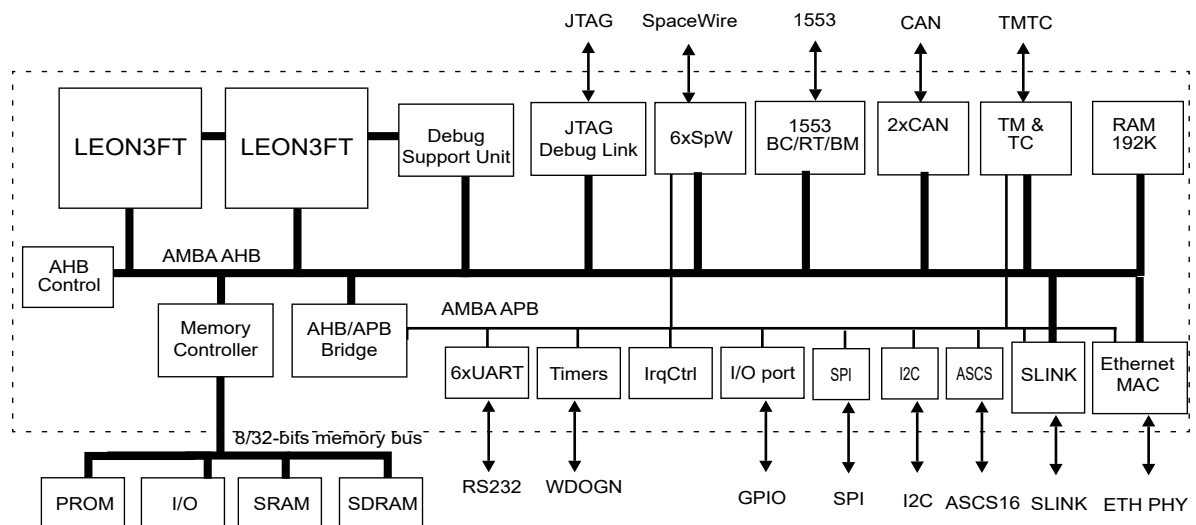
GR712RC Block diagram

GR712RC

1.3 Memory map

The memory map of the internal AHB/APB buses is specified as shown below.

TABLE 1. Internal memory map

Core	Address range		Bus
FTMCTRL	0x00000000 - 0x20000000 0x20000000 - 0x40000000 0x40000000 - 0x80000000	PROM area I/O area SRAM/SDRAM area	AHB
APBCTRL - 1	0x80000000 - 0x80100000	APB bridge	AHB
FTMCTRL	0x80000000 - 0x80000100	Registers	APB
APBUART - 0	0x80000100 - 0x80000200	Registers	APB
IRQMP	0x80000200 - 0x80000300	Registers	APB
GPTIMER	0x80000300 - 0x80000400	Registers	APB
SPICTRL	0x80000400 - 0x80000500	Registers	APB
CANMUX	0x80000500 - 0x80000600	Registers	APB
GRGPREG	0x80000600 - 0x80000700	Registers	APB
GRASCS	0x80000700 - 0x80000800	Registers	APB
GRSLINK	0x80000800 - 0x80000900	Registers	APB
GRGPIO - 1	0x80000900 - 0x80000A00	Registers	APB
GRGPIO - 2	0x80000A00 - 0x80000B00	Registers	APB
GRTM	0x80000B00 - 0x80000C00	Registers	APB
I2CMST	0x80000C00 - 0x80000D00	Registers	APB
CLKGATE	0x80000D00 - 0x80000E00	Registers	APB
GRETH	0x80000E00 - 0x80000F00	Registers	APB
AHBSTAT	0x80000F00 - 0x80001000	Registers	APB
APB1 plug&play	0x800FF000 - 0x80100000	Plug&play configuration	APB
APBCTRL - 2	0x80100000 - 0x80200000	APB bridge	AHB
APBUART - 1	0x80100100 - 0x80100200	Registers	APB
APBUART - 2	0x80100200 - 0x80100300	Registers	APB
APBUART - 3	0x80100300 - 0x80100400	Registers	APB
APBUART - 4	0x80100400 - 0x80100500	Registers	APB
APBUART - 5	0x80100500 - 0x80100600	Registers	APB
GRTIMER	0x80100600 - 0x80100700	Registers	APB
GRSPW2 - 0	0x80100800 - 0x80100900	Registers	APB
GRSPW2 - 1	0x80100900 - 0x80100A00	Registers	APB
GRSPW2 - 2	0x80100A00 - 0x80100B00	Registers	APB
GRSPW2 - 3	0x80100B00 - 0x80100C00	Registers	APB
GRSPW2 - 4	0x80100C00 - 0x80100D00	Registers	APB
GRSPW2 - 5	0x80100D00 - 0x80100E00	Registers	APB
APB2 plug&play	0x801FF000 - 0x80200000	Plug&play configuration	APB
DSU3	0x90000000 - 0xA0000000	Registers	AHB
FTAHBRAM	0xA0000000 - 0xA0100000 0x80100000 - 0x80100100	RAM area Registers	AHB APB
B1553BRM	0xFFF00000 - 0xFFF01000	Registers	AHB
GRTC	0xFFF10000 - 0xFFF10100	Registers	AHB
Proprietary, do not access	0xFFF20000 - 0xFFF20100	Registers	AHB
CANOC	0xFFF30000 - 0xFFF31000	Registers	AHB
AHB plug&play	0xFFFFF000 - 0xFFFFFFFF	Plug&play configuration	AHB

Access to addresses outside the ranges described above returns an AMBA AHB error response.

1.4   Interrupts

The following table specifies the interrupt assignment for the GR712RC.

TABLE 2. Interrupt assignment

Core	Interrupt #	Function
AHBSTAT	1	AHB bus error
APBUART - 0	2	UART1 RX/TX interrupt
GRGPIO - 1 to 2	1-15	External I/O interrupt (3 & 4 are GPIO only)
CANOC	5-6	OCCAN interrupt (core 1 - 2)
GRTIMER	7	GRTIMER timer underflow interrupt
GPTIMER	8-11	GPTIMER timer underflow interrupts
IRQMP	12	IRQMP extended interrupt
SPICTRL, SLINK	13	SPI and SLINK interrupt
B1553BRM, GRETH, GRTC	14	1553, Ethernet, and Telecommand interrupt
ASCS	16	ASCS interrupt
APBUART - 1to 5	17-21	UART2-6 RX/TX interrupt
GRSPW2 - 0 to 5	22-27	SpaceWire 0-5 RX/TX data interrupt
I2CMST	28	I2C master interrupt
GRTM	29	Telemetry encoder interrupt
	30	Telemetry encoder time strobe interrupt

The interrupts are routed to the IRQMP interrupt controller and forwarded to the LEON3FT processors. Interrupt 16-31 are generated as interrupt 12 and for those interrupts a chained interrupt handler is used.

1.5   GRLIB IP cores

The GR712RC uses the following cores from the GRLIB IP library:

TABLE 3. Used IP cores

Core	Function	Vendor ID	Device ID
LEON3FT	Fault tolerant SPARC V8 32-bit processor	0x01	0x053
DSU3	Debug support unit	0x01	0x004
IRQMP	Interrupt controller	0x01	0x00D
APBCTRL	AHB/APB Bridge	0x01	0x006
FTMCTRL	8/32-bit Memory controller with EDAC	0x01	0x054
AHBSTAT	AHB failing address register	0x01	0x052
FTAHBRAM	Fault tolerant on chip memory	0x01	0x050
AHBJTAG	JTAG/AHB debug interface	0x01	0x01C
GRSPW2	SpaceWire link with DMA	0x01	0x029
B1553BRM	MIL-STD1553B BC/RT/BM controller	0x01	0x072
CANOC	OC CAN controller	0x01	0x019
GRETH	10/100 Ethernet MAC	0x01	0x01D
APBUART	8-bit UART with FIFO	0x01	0x00C
GPTIMER	Modular timer unit	0x01	0x011
GRTIMER	Modular timer unit with time latch capability	0x01	0x038
GRGPIO	General purpose I/O port	0x01	0x01A
GRTC	CCSDS Telecommand Decoder	0x01	0x031
GRTM	CCSDS Telemetry Encoder	0x01	0x030
-	Proprietary	0x01	0x080
CANMUX	CAN bus multiplexer	0x01	0x081

TABLE 3. Used IP cores

Core	Function	Vendor ID	Device ID
GRASCS	ASCS16 controller	0x01	0x043
GRSLINK	SLINK controller	0x01	0x02F
SPICTRL	SPI controller (master only)	0x01	0x02D
I2CMST	I2C master	0x01	0x028
CLKGATE	Clock gating unit	0x01	0x02C
GRGPREG	General purpose register	0x01	0x087

1.6   References

[DS]	GR712RC - Dual-Core LEON3FT SPARC V8 Processor, Data Sheet, GR712RC-DS, Issue 2.1, June 2014, Cobham Gaisler, www.cobham.com/gaisler
[SPARC]	The SPARC Architecture Manual, Version 8, Revision SAV080SI9308, SPARC International Inc.
[SPW]	ECSS - Space Engineering, SpaceWire - Links, Nodes, Routers and Networks, ECSS-E-ST-50-12C, 31 July 2008
[RMAPID]	ECSS - Space Engineering, SpaceWire Protocol Identification, ECSS-E-ST-50-51C, February 2010
[RMAP]	ECSS - Space Engineering, Remote Memory Access Protocol, ECSS-E-ST-50-52C, February 2010
[1553BRM]	Core1553BRM Product Handbook, 50200040-0/11-04, November 2004, Actel Corporation
	Core1553BRM MIL-STD-1553 BC, RT, and MT, 51700052-4/12.05, v 5.0, December 2005, Actel Corporation
	Core1553BRM User's Guide, 50200023-0/06.04, June 2004, Actel Corporation
	Core1553BRM v2.16 Release Notes, 51300019-8/6.06, June 2006, Actel Corporation
[MMU]	Technical Note on LEON SRMMU Behaviour, GRLIB-TN-0002, Issue 1, Revision 1, February 2015, Cobham Gaisler, www.cobham.com/gaisler
[GRETH]	Technical Note on GRETH Controller Behaviour, GRLIB-TN-0008, Issue 1, October 2015, Cobham Gaisler, www.cobham.com/gaisler
[B2B]	Technical Note on Stale Cache Entry After Store with Data Tag Parity Error, GRLIB-TN-0009, Issue 1, September 2016, Cobham Gaisler, www.cobham.com/gaisler
[TN0011]	Technical Note on LEON3/FT AHB Lock Release during Atomic Operation, GRLIB-TN-0011, Issue 1, November 2017, Cobham Gaisler, www.cobham.com/gaisler
[TN0012]	Technical Note on GR712RC Incorrect Annulation of Floating-point Operation on Instruction Cache Parity Error, GRLIB-TN-0012, Issue 1, November 2017 Cobham Gaisler, www.cobham.com/gaisler
[TN0013]	Technical Note on GRFPU Floating-point controller: Missing FDIV/FSQRT Result, GRLIB-TN-0013, Issue 1, November 2017 Cobham Gaisler, www.cobham.com/gaisler
[AN003]	FTMCTRL: BCH EDAC with multiple 8-bit wide PROM and SRAM banks, GRLIB-AN-0003, Issue 1, April 2017 Cobham Gaisler, www.cobham.com/gaisler

1.7 Errata

1.7.1 FTAHBRAM: On-chip Memory not cacheable

The 192 KiB on-chip memory at address 0xA0000000 is not cacheable in the L1 caches of the processors. This means that executing from the on-chip memory will be slow, approximately 0.25 MIPS/MHz. The on-chip memory should preferably be used for data storage for on-chip IP cores, such as the Mil-Std-1553B controller, Ethernet controller and SpaceWire links.

See sections 4.2.17 and 6.1 for details.

1.7.2 CAN OC: interrupt can be cleared before read

The Transmit Interrupt can be cleared before read, if a read access is made to the Interrupt Register in the same clock cycle as the interrupt is set by the transmitter logic.

Instead of looking at the Transmit Interrupt flag, the Transmit Buffer Status flag in the Status Register can bee looked at when something has been scheduled for transmission.

See section 18.4.5.

1.7.3 GRSPW2: interrupt can be lost

If an AMBA AHB error occurs in the same clock cycle as a SpaceWire link error or Time-Code Tick-Out interrupt is generated and the AMBA AHB error interrupts are disabled then the corresponding interrupt will be lost (i.e. SpaceWire link error or Time-Code Tick-Out interrupt)

The workaround is to also enable the AMBA AHB error interrupt when any of the link error or Time-Code Tick-Out interrupts are enabled. An AMBA AHB error interrupt will cause the corresponding DMA channel to stop, and therefore interrupt handling will be required anyway.

See section 16.9.

1.7.4 GRSPW2: CRC calculation partially incorrect

RMAP CRC calculation for the DMA receiver does not work correctly under all conditions. The detected header length is not used correctly in cases when two characters are received on consecutive cycles and one of them is the header CRC. This results in an incorrect header CRC indication in the receiver descriptor (i.e. Header CRC (HC) field). The calculation is continued however (as described in the manual) and the data CRC indication (i.e. Data CRC (DC) field) is still correct and covers the whole packet and can thus be used to determine that there are no errors in the complete packet.

See section 16.4.5.

1.7.5 SPICTRL: transfers in progress bit not cleared

The Transfer in Progress bit (TIP the core's Event Register is not cleared automatically by the core if an overrun occurs. The core needs to be disabled in order to clear the bit. The problem does not affect most software drivers.

Workaround 1: Read receive queue so that no overrun occurs.

Workaround 2: Disable and re-enable core if overrun occurs.

Workaround 3: Do not use Transfer in Progress (TIP) interrupts.

See section 23.3.

1.7.6 SPICTRL: back-to-back transfers

Back-to-back transfers where a word is written to the core's transmit queue at the very end of the transfer of a previous word may cause the core to essentially change the clock phase (CPHA) of the SPI communication. This leads to transmission and reception of wrong data.

The SPI controller regards a transfer as completed when the last bit has been sampled. With clock phase 1 there is a window between sampling of the last bit and when the SCK clock line returns to its idle state. If a new transfer is started (a word is written to the core's transmit queue) in this window, then the bug will be triggered.

This scenario can occur when writing a series of words to the SPI controller and the delay between two writes matches the time it takes for the core to transfer the current transmit queue contents. A possible scenario is that a processor wants to send two words over SPI, after the first word has been written the processor gets an interrupt, when the processor returns from serving the interrupt and writes the second word the second write could hit the window.

Software drivers that use SPI clock phase 0 (data sampling starts on the first transition of SCK) are potentially affected.

Workaround 0: Use a SPI clock phase of 1 and adjust the CPOL value instead.

Workaround 1: Ensure that all words of a transfer are written to the transmit queue before the last bit of a preceding word has been sent.

See section 23.3.

1.7.7 FTMCTLR: EDAC usage with 8-bit wide memory

When multiple SRAM or PROM memory banks are used in 8-bit mode with EDAC, the configured bank size must be set to at least four times the actual device size.

Refer to [AN003] for further details.

1.7.8 LEON3FT Cache Controller: Incorrect Bus Access After Power-Down

A LEON with support for clock gating has one clock that is kept running and one clock that is gated off in power-down mode. Due to a design error, the gated clock was used for logic that keeps track of the state of the AMBA bus. Due to this error, the first instruction or data (whichever is first) access to the bus after leaving power-down might be performed incorrectly.

The errata triggers a failure when all of the following conditions are met:

• The processor enters power-down mode while it has received grant to the bus
• For data cache: The instructions to be executed after the power-down instruction contain a load or store instruction. The memory position accessed is not present in the data cache, or is located in a noncachable memory area.

  For instruction cache: The instructions following the power-down instruction are not present in the instruction cache or the instruction cache is disabled.
• The processor leaves power-down mode and has lost grant to the bus. For this condition to be met, another master in the system has made a bus access simultaneously with the processor leaving power-down mode.

The default behaviour in GRLIB systems is to grant master 0 by default, this master is typically the LEON processor.

On systems where the errata is applicable and not using any of the workarounds described in this document, the impact is that:

• The first data access that goes to the AHB bus may not be performed correctly. If it was a load it will return invalid data.
• If the instruction cache is disabled, or if the processor is executing from a memory area marked as noncachable, then the instruction fetch when the processor leaves power-down mode may return erroneous data that will be interpreted as an instruction and executed by the processor.

Workaround 1: Do not use power-down

The issues described in this document are avoided by not entering power-down mode. Refrain from using the wr %asr19 sequence.

Workaround 2: Keep instruction cache enabled, execute power-down sequence from cacheable memory.

When executing from cacheable memory with the instruction cache enabled, the sequence below is immune to the errata and is suitable for systems that have been implemented with a MMU:

```
unsigned int address = (unsigned int)<any physical memory address>;
```

```
__asm__ __volatile__ (
 "mov %%g0, %%asr19\n"
 "lda [%0] 0x1C, %%g0\n"

 :
 : "r"(address));
```

The power-down sequence that makes use of ASI 0x1C will only work on systems that have been implemented with a MMU. For systems without MMU, a noncacheable APB register can be used as source for the load:

```
unsigned int address = (unsigned int)<address of noncacheable memory. like APB, PnP>;
__asm__ __volatile__ (
 "mov %%g0, %%asr19\n"
 "ld [%0], %%g0\n"

 :
 : "r"(address));
```

Note that the stack pointer or other cachable memory should never be used as the source operand, even if using ASI 1, since this may lead to faulty data being loaded into the data cache. Also note that FIFOs should not be accessed via the LD. The recommended address to use for the load in both cases above is the AMBA plug&play area, typically at top of RAM (address 0xFFFFFFF0) 1 .

The load instruction after the write to %asr19 provides immunity against a data cache failure due to the error described by this errata. After coming back from power-down with a pending interrupt, the instruction following the %asr19 access will be executed prior to taking the interrupt trap. The load will be executed on the bus before trying to fetch the first instruction of the trap handler. So even if the trap handler is not currently in cache, the load will occur before fetching it. The load will also be executed to memory before trying to fetch any instruction from the calling function.

Both the write to asr19 and the load, plus the three following instructions, are guaranteed to be in the I-cache when leaving power-down since they are the last to be fetched into the pipeline before entering power-down.

Software packages with workarounds

The Cobham Gaisler provided run time environments, from versions shown below, have integrated a workaround for the data cache for this bug. The software is assumed to be run with cache enabled (default behaviour in bootladers) and from cacheable memory and does not include any workaround for the instruction cache, since this is not required with these assumptions.

BCC (bare-C)	Not applicable, does not use power-down instruction
Linux	2.6.21 (SnapGear p42)
	All 2.6.36 releases and later that have been distributed via Cobham Gaisler
	In mainline since 2.6.39
RCC (RTEMS)	RCC-1.2.x based on RTEMS-4.10 and later
VxWorks	VxWorks 6.3 and 6.5 are not affected since power-down is not used
	Fixed in 6.7 since release 1.0.3

1.7.9 Failing SDRAM Access After Uncorrectable EDAC Error

When the memory controller of the GR712RC LEON system detects an uncorrectable EDAC error, it should respond with an AMBA ERROR response and then return to normal operation. Due to an incomplete condition check for starting new SDRAM accesses, the memory controller may perform a read access following an uncorrectable error even if there is no incoming access on the AMBA bus. The result will be discarded unless a AMBA read access to the SDRAM memory area is performed before the SDRAM read operation has finished. The extra read access will not occur if there is a SDRAM refresh operation pending.

The memory controller will return to normal operation after the extra read access has been performed. If a AMBA read is performed to the SDRAM area before this unintended read access has completed then the result of the incoming AMBA read access may be erroneous. The result can be an AMBA

ERROR response or the memory controller may deliver data from the wrong memory location without a AMBA ERROR response (note that the first access that read a location which had an uncorrectable error will always receive an AMBA ERROR response).

The erratum can be triggered when:

- The FTMCTRL has been configured with minimum t_{RP} SDRAM timing, and
- A read access to SDRAM results in an uncorrectable EDAC error, and
- A second AMBA read access is performed to the SDRAM memory area in the window zero to five system clock cycles after the AMBA ERROR response given due to the uncorrectable EDAC error.

If the incoming read access occurs during the last cycle of the vulnerable window in time then the controller will return data from the memory location of the first access, which will may trigger an AMBA ERROR response if the uncorrectable error remains at that memory location. For incoming accesses during the other cycles in the vulnerable window, the access will malfunction but the data read by the memory controller may still have valid check bits. If the check bits are valid then the erroneous data will be delivered without any AMBA ERROR for the second access.

Uncorrectable SDRAM errors are expected to be a rare occurrence, provided that scrubbing of correctable errors has been implemented in the system.

The uncorrected error response may be the result of a processor access which leads to the processor taking a trap. The read of the trap table may then occur during the window when the FTMCTRL performs the extra read access which leads to erroneous data or a AMBA ERROR response when the processor fetches the trap table. In multi-master systems, one master may trigger the uncorrectable error and a second master may perform a read access in the window where the FTMCTRL is malfunctioning.

There are cases where the unintended read access is prevented. A SDRAM command issue or a refresh operation are cases that prevent triggering the erratum.

It is recommended that all GR712RC systems implement the following workaround.

Workaround

The erratum cannot be triggered when the MCFG2.TRP field (bit 30 of MCFG2 register) is set to 1 (SDRAM t_{RP} parameter is three clock cycles).

VxWorks board support package for GR712RC will be updated to default MCFG2.TRP=1. For other boot loaders, please refer to the software package's documentation for information on how to set MCFG2.TRP.

GRMON2 will be updated to automatically set MCFG2.TRP for GR712RC.

Please contact Cobham Gaisler support for information on specific software versions.

Other information

The failure will not occur if the second incoming access is to the PROM, SRAM or IO area. The memory controller may perform the unintended read access to SDRAM but software will not be affected.

The bug does not affect SRAM and PROM areas.

The controller will always perform the unintended SDRAM read access when MCFG2.TRP=0, unless there is a refresh operation pending.

It takes exactly five system clock cycles after the AMBA ERROR response until the memory controller is back in normal state.

1.7.10 MIL-STD-1553B core duplicate interrupt assertion

Under certain conditions interrupts from the MIL-STD-1553B core can appear to be duplicate. The Interrupt Pending status bit in the Multiprocessor Interrupt Controller might be asserted an extra time after the interrupt event is successfully processed. The problem appears more frequently with a higher ratio of the AHB system bus frequency to the MIL-STD-1553B core frequency, but can be present for any combination thereof. A possible consequence of this problem is that the Interrupt Service Routine associated with the MIL-STD-1553B core, IRQ 14, gets called twice.

The reason for this spurious behavior lies in how the Actel Core1553BRM controller propagates the IRQ to the interrupt controller. The IRQ line is propagated from the Interrupt Request pulse from the

Core1553BRM and is re-synchronized into the system clock domain and fed as an IRQ source to the interrupt controller. Assuming, for instance, an AHB system bus frequency four times higher than the MIL-STD-1553B core frequency, since the Interrupt Request line output of the Core1553BRM is high for 3 cycles when INTLEVEL=0, this will result in a 12 cycle pulse coming in to the interrupt controller. When this pulse is longer than it takes for the CPU to handle the interrupt, this will lead to a second interrupt being stored into the interrupt controller, and this will then cause the interrupt handler to be called twice.

Workaround

A possible workaround is to manually clear the interrupt pending status in the IRQ handler via the Interrupt Clear Register of the IRQ controller. This assumes that there are no other other sources (hardware or software) that can assert interrupt 14. This needs to be done early, before reading the BRM IRQ status/log to eliminate the risk of losing any IRQ events from the BRM.

The other workaround is simply to accept the extra IRQs and make the IRQ handler capable of handling this case.

1.7.11 Technical Note on LEON SRMMU Behaviour

The MMU behavior when updating the FSR and FAR registers in this device does not fully comply to the SPARC V8 Manual. Special care should be taken by users that utilize the MMU. This is addressed in Linux release 3.10.58-1.0.4 and onwards. Other users who implement MMU handling should consult [MMU] for further information.

1.7.12 Technical Note on GRETH Ethernet Controller Behaviour

The state in which the GRETH receiver control finite state machine discards packets does not take overrun into account in all cases. The discard state is entered when a received packet is determined to be dropped. The size of the packet is checked and that amount of data is read from the FIFO. When an overrun has occurred, the amount written to the FIFO is not the same as the length of the packet and in those cases the FIFO should be emptied completely. This was not done if:

1 The packet in question had a MAC address which the core should not receive, or
2 if the receiver was enabled when the packet was received but the descriptor was not enabled and the descriptor read was so slow that an overrun occurred during that time.

These conditions cause the Ethernet receiver to hang. No further traffic can be received and the GRETH must receive a hardware reset to recover. This hardware reset can be attained by either resetting the whole device or, in devices that have individual clock gating of the Ethernet controller, forcing a reset of the GRETH via the system's clock gating unit.

Workaround

As described in the previous section, the problem occurs under two conditions. The first condition can be avoided by enabling promiscuous mode for the Ethernet controller. This mode is enabled via the GRETH control register. Promiscuous mode may lead to additional software load of the systems since more packets may need to be processed. In applications where the GRETH is used in a limited network with few members, the addition of promiscuous mode should have less negative effects.

Use of switches instead of hubs in the network since the GRETH controller will in this case receive fewer packets with MAC addresses that will be discarded.

The second trigger condition is system and application specific. The condition is not expected to happen in systems like the affected components listed earlier in this document. Other systems with memory controllers that have high latency and systems that introduce several bridges between the GRETH controller and main memory may encounter condition 2.

There is no status bit in the controller that will indicate that the receiver has hanged. The condition can be detected by the Ethernet controller no longer receiving any packets. This could require implementation at a higher level of a heart beat function so that system software can detect that packets are missing and correct the situation. The only way to correct the condition is to perform a hardware reset of the GRETH controller. The soft reset functionality that can be triggered via the GRETH control register is not enough to resolve the receiver hang.

Users of the GRETH Ethernet Controller should consult [GRETH] for further information and updates.

1.7.13 Technical Note on Stale Cache Entry After Store with Data Tag Parity Error

Under certain conditions, a data coherency issue can occur when the processor has executed a back-to-back store operation and the first store's address hits a cache set that has a data cache tag parity error. The second store operation successfully stores its data content into the main memory but the data cache controller fails to update the old data in the data cache with the new data from the store operation. This leads to an incoherent state, data in the data cache is not the same as in the main memory. A subsequent read from the second location will get the old data instead of the new data if the data cache line is still valid in the cache.

Refer to [B2B] for further details and workaround.

1.7.14 Never disable the TLB when the MMU is enabled

When the MMU is enabled but TLB disabled, the MMU can in certain cases perform incorrect address translation. Therefore it is advised to never set the TLB disable bit in the MMU control register. The TLB disable bit is not set by any operating system provided by Cobham Gaisler.

1.7.15 Technical Note on LEON3/FT AHB Lock Release During Atomic Operation

In certain cases, where the MMU is enabled, a LEON3/FT core can lose the bus lock before completing the store operation required by an atomic instruction (swap, ldstub, casa). The atomic instruction will complete atomically in terms of instruction order but due to losing bus lock before store, the atomicity on the bus level is not guaranteed. As a result, if before being able to make the store required by the atomic instruction, if another master gains access to the bus and makes a store to the same address (used by the atomic instruction) it may cause erroneous behavior since atomicity is lost.

Refer to [TN0011] for further details and workaround.

1.7.16 Technical Note on GR712RC Incorrect Annulation of Floating-point Operation on Instruction Cache Parity Error

In certain cases, an incorrect update of a floating point register or floating-point condition register due to execution and commit of an extra floating point instruction outside of the correct instruction flow can occur. This can in turn cause erroneous behavior of software. The error can occur when there is a tag or data parity error in the instruction cache.

Refer to [TN0012] for further details and workaround.

1.7.17 Technical Note on GRFPU Floating-point controller: Missing FDIV/FSQRT Result

This document describes a corner case present in the GRFPC floating point controller, which is the hardware used to interface the GRFPU floating point unit with the LEON3/FT and LEON4/FT processors. The corner case described in this document can cause a FDIV/FSQRT operation result not to be committed to the floating point register file. The issue manifests if;

*   two instructions exist in the instruction flow between two FDIV/FSQRT instructions, and
*   at least one of them is a floating point operation, and
*   the processor encounters certain number of hold cycles after the first FDIV/FSQRT operation reaches execute stage of the pipeline.

Refer to [TN0013] for further details and workaround.

1.7.18 Don't break into debug mode on RETT instructions

Avoid breaking into debug mode on RETT instructions because this causes incorrect behavior when resuming.

1.7.19 Stores to ASI 0x1C may update data cache

When a store is made using the store-alternate instruction with the address space identifier set to MMU/cache bypass (STA with ASI set to 0x1C), the processor's data cache controller handles this as if it was a regular store instruction with the same address and data. This may lead to unexpected data cache updates.

This is particularly an issue when the MMU is enabled and used for virtual-to-physical address translation. In this case, the given store address is used as a physical address when written to memory, as intended, but gets used as a virtual address in the cache update.

It is recommended that applications avoid using stores with ASI 0x1C and reimplement such code using ordinary store instructions instead. Alternatively, a data cache flush can be used to get a consistent cache state after STA ASI=0x1C instructions have been performed (however note that traps occurring between the STA and the flush may see the invalid cache state).

1.8 Document revision history

TABLE 4. Revision history

Issue	Date	Sections	Description
2.12	2018-07-19	4.5.4	Corrected naming of DDE in Cache Control Register field descriptions
		4.5.7	Clarified which fields of the Cache Control Register that are 0 after reset
		5.10.2	Corrected code generation polynomial for Reed-Solomon EDAC
		12.3	Corrected GRTIMER Configuration register reset value
		15.5	Added note about debug mode in APBUART not being affected by system reset
		18.3	Corrected address range for CAN core 2
		18.6	Corrected spelling of BasicCan
		20	Updated spelling of multiplexer to be consistent
		20.1	Clarification of connection between registers and BUS A mux
		16.3.3	Remove incorrect statement about clock domains in receiver.
2.11	2018-03-19	1.7.19	Added errata about ASI 0x1C
		4.5.3	Clarified data cache snooping with MMU
2.10	2017-11-21	1.7.7, 5.1, 5.10	Clarified FTMCTLR EDAC usage with 8-bit wide memory and added a reference to GRLIB-AN-0003
		1.7.15	Added errata about LEON3/FT AHB lock release during atomic operation
		1.7.16	Added errata about incorrect annulation of floating-point operation on instruction cache parity error
		1.7.17	Added errata about floating-point controller: missing FDIV/FSQRT result
		1.7.18	Added errata about breaking on RETT instructions
		4.1, 4.2.5, 4.2.10	Clarification that GR712RC doesn't have a co-processor
		4.2.10, 4.9	Clarified that ERRORN is only connected to CPU0
2.9	2017-06-22	1.7.14	Added errata about TLB disable when MMU is enabled

TABLE 4. Revision history

Issue	Date	Sections	Description
2.8	2017-02-28	1.7.3	Updated error handling for GRSPW2
		1.7.13	Added errata for back-to-back storage
		4.2.6	Added description of CASA instruction
		4.2.9	FPU option value corrected
		4.2.12, 4.5.1, 4.6.1	ASI clarification for ASI 0x10, 0x11 and 0x13.
		4.2.15	Updated multiprocessor wording
		4.5.2, 4.5.5, 4.5.6	Clarified cache access descriptions
		4.5.4	FD bit interpretation corrected in Cache Control Register
		5.14.2	TRP bit description corrected in Memory configuration register 2
		15.6	Clarified APBUART receive interrupts
		16.1	Further clarified number of DMA channels per GRSPW2
		16.8	Clarified relationship between SPW receive data rate and transmit clock
		16.9	Clarified reset value of AS bit in GRSPW control register
		18.1	Added reference to CAN multiplexer to clarify the need to activate OC-CAN
		23.3	LSD bit description corrected in SPI controller Command register
2.7	2016-01-26	7.2	AHBSTAT: Remove incorrect note about AMBA ERROR overriding correctable errors.
		15.2.2	APBUART documentation mentioned receive ready instead of data ready field.
			Update footer
2.6	2015-11-19	21.1, 21.4	MIL-STD-1553B, specify AHB slave interface and update graphics
		7.2	AHB Status Register needs NE bit to be cleared for error monitoring
		12.3	GRTIMER register interface overhauled and reset values corrected
		13.2	GPREG register typo for 1553DIVH
		16.9	Spacewire Register address space correction
		8.3.7, 8.3.9	IRQMP register figure captions
		11.3	GPTIMER register fix for CTRL.IP bit
		5.4	Anticipated *ramsn* signal in Figure 25, which was one clock late
		26.1, 26.6	Clarify that there is no hardware link between GRTC and GRTM
		15.7.3	UART FA field read-only and read-value specified
		16.9	Added condition to bits TI and RI of DMA control register
		1.7.12	Technical Note on GRETH Ethernet Controller
		23.1	SPI controller does not support periodic transfers

TABLE 4. Revision history

Issue	Date	Sections	Description
2.5	2015-04-27	1.7.11	Technical Note on LEON SRMMU Behaviour
		1.6	Updated DS and MMU References
		1.2, 5.1, 5.2, 5.4, 5.8.1,	Corrected the max SRAM and PROM bank sizes and the maximum supported SRAM and PROM size
		5.14.1, 5.14.2	
		3, 3.4, 13.2	Telemetry clock input name inconsistency
		3.8	SDCLK always driven by Delay line output
		11.3	Refactored GPTIMER register description with error fixes
		12.2, 12.3	GRTIMER number of timers and scaler bits corrected
		8.3.5	Document BA bit in Multiprocessor Status Register
		4.8	Single Error Correction/Double Error Detection
		16.1	SpaceWire documentation references update
		9.4	Removed incorrect statements about multiply, divide and FPU operations
		13.2	Specify SDCLK delay value meaning and fix TMCLKI name
		4.7	FPU trap queue has 8 entries
		6.3	Corrected Table 36 layout with MEMSIZE MSb
2.4	2014-09-04	1.7.9	Added "Failing SDRAM Access After Uncorrectable EDAC Error" errata
		1.7.10	Added "MIL-STD-1553B core duplicate interrupt assertion" errata

TABLE 4. Revision history

Issue	Date	Sections	Description
2.3	2013-05-13	1.2	Memory controller maximum SRAM/SDRAM size supported
		1.7.8	Added "LEON3FT Cache Controller: Incorrect Bus Access After Power-Down" errata
		2	Behavior of switch matrix output when CAN core is enabled
		3.2	SpaceWire clock max frequency and startup bandwidth
		5.1	Added maximum capacities for SRAM and SDRAM
		5.8.1	Units and SDRAM + RS EDAC bus width
		5.10.3	EDAC Error reporting behavior clarifications
		5.14.2	SDRAM Column Size clarification
		6.2	On-chip memory performance statistics update
		11.3	Register layout fixes for General Purpose Timer
		12.3	Register layout fixes for Latching General Purpose Timer
		14.3	Documented interrupt mask register
		15.3	Clarified scaler behavior with respect to clocking
		16.3.4	Removed unsupported tick-in signal reference
		16.6	RMAP support only on GRSPW2-0 and GRSPW2-1
		16.8	SpaceWire clock max frequency and startup bandwidth
		17.7	Clarified reset PHY address value
		25.1	Specified number of slaves and data word length for GRASCS
		25.3	Documented Output Enable bit
			Documented TMD bit
			Fixed ETR scale register width
2.2	2013-02-20	1.2, 1.3, 1.4, 1.5, 2, 3.6, 3.9, 8.3, 13.2, 19, 20, 28.2	Obsolete proprietary function removed
		1.3	APB addresses corrected
		1.3, 18.3	CANOC addresses corrected
		1.7.2	CANOC errata updated
		1.7.7, 5.1, 5.10	Note on EDAC usage with 8-bit wide memory introduced
		3.3, 21.3	Mil-Std-1553B clocking clarified
		3.8	DLLBPN polarity corrected
		3.8	Delay line explained, figure added
		4.2.13, 4.6.1, 4.6.3	MMU and cache handling clarified
		5.8	SDRAM example programming clarified
		15.3	UART baud rate generation clarified

TABLE 4. Revision history

Issue	Date	Sections	Description
2.1	2012-12-21	1.3	Updated GRTC address mapping
		1.7	Added errata section
		2	Updated I/O switch matrix conflict table 8
		4.2.17, 6.1	On-chip memory is not cacheable
		5	Added SRAM read-modify-write timing diagram and various memory controller clarifications.
		5.10.2	Clarified 8-bit EDAC limitation to a single bank
		5.15	Clarified that busy ready signal also applies to PROM area
		8.3	Clarified interrupt controller register definitions
		10.4	Clarified handling of unused JTAG interface
		11.2	Watchdog dependent on interrupt setting for timer
		12.2	Clarified timer latching limitations
		12.3	Corrected memory map
		13.2	Typo in register definition
		1.2, 16.1	SpaceWire core configuration information added
		16.4.8	SpaceWire read descriptor handling clarified
		16.9	SpaceWire register reset values clarified
		18.3	CAN register mapping corrected
		21.4	Reset value of ssysfn bit corrected for Mil-Std-1553B core
		23	Updated SPI controller capability register descriptions with implemented values, removed functions not implemented.
2.0	2011-10-18		Clarifications, see change bars
1.9	2011-02-25		Updated all sections with additional details
	2010-11-09		Corrected typos in address map
1.8	2010-07-30		Re-ordered chapters, improved description of clocking
1.7	2010-03-31		Refined text to describe only the GR712RC config
1.6	2009-06-23		Added GRTIMER, updated GRTM and IRQs
1.5	2009-05-19		Corrected bus timing register
1.4	2009-04-07		Added separate chapters for clocking and I/O switch matrix. Updated GPREG, BRM.
1.3	2009-04-06		Added BTR registers and GPREG section
1.2	2009-04-01		Changes in clocking options
1.1	2009-03-09		Incorporated feedback from RC
1.0	2009-03-02		Initial specification

2   Signals and I/O Switch Matrix

The GR712RC is housed in a 240-pin ceramic quad-flat pack package (CQFP-240). To fit in this package, some of the on-chip peripheral functions have to share I/O pins. A programmable I/O switch matrix provides access to several I/O units. When an interface is not activated, its pins automatically become general purpose I/O. After reset, all I/O switch matrix pins are defined as I/O until programmed otherwise. The programmable I/O switch matrix consists of 67 pins.

Figure 1 shows how the various I/O units are connected to the I/O switch matrix. Table 5 shows examples of possible configurations using the I/O switch matrix. Note that two SpaceWire interfaces are always available outside the I/O switch matrix.

Table 6 shows a listing of all pins in the I/O switch matrix, indicating the priority amongst them. Note that some pins are input only, some are output only, and the rest are both input and output, as described in table 6.

Table 7 shows a listing of pins in the I/O switch matrix grouped per function (GPIO is not listed). Table 8 shows a complete listing of conflicts between I/O units (GPIO is not listed).
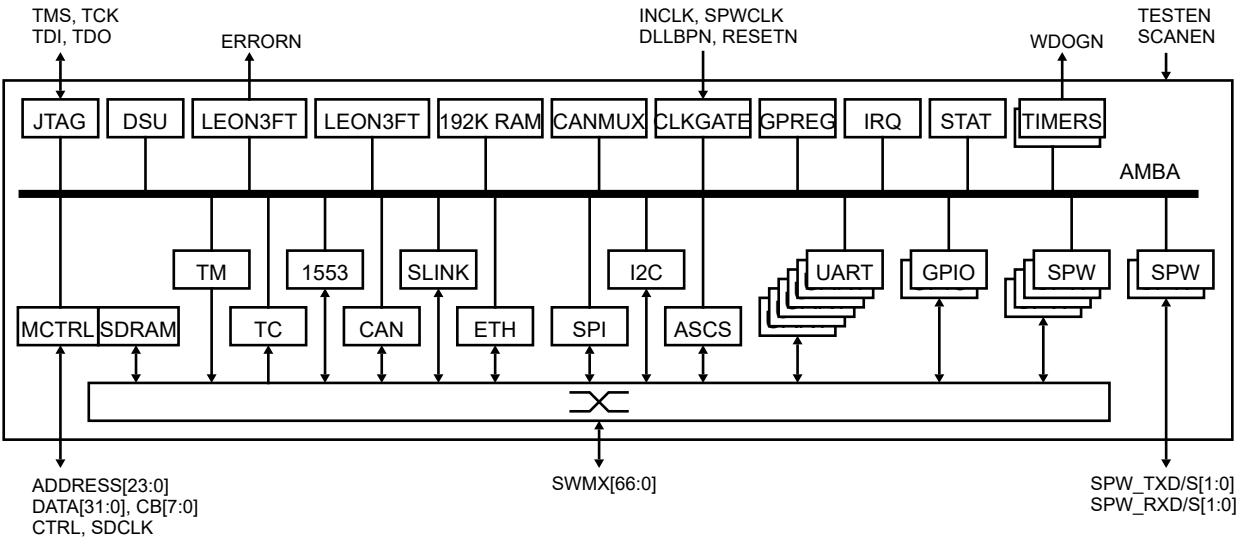
Figure 1. Architectural block diagram showing connections to the I/O switch matrix

Table 5. Example of possible configurations using the I/O switch matrix

Interface type	Example configuration					
	CF0	CF1	CF2	CF3	CF4	CF5
SDRAM with or without Reed-Solomon		1	1	1	1	1
UART	6	4	6	6	6	6
SpaceWire	6	4	2	2	4	3
Ethernet					1	1
CAN						
MIL-STD-1553B BC/RT/BM	1					
I2C	1		1	1	1	1
SPI	1				1	1
SLINK			1	1		
ASCS16			1			
CCSDS/ECSS TC & TM		1				

Table 6. I/O switch matrix pin description, defining the order of priority for outputs and input/outputs, with the highest priority for each pin listed first.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
4	SWMX[0]	UART_TX[0]	-	High	Out	UART Transmit 0
3	SWMX[1]	UART_RX[0]	-		In	UART Receive 0
2	SWMX[2]	UART_TX[1]	-	High	Out	UART Transmit 1
1	SWMX[3]	UART_RX[1]	-		In	UART Receive 1
		GPIO[0]	-		In	GPIO 1 Register, bit 0 (input only)
240	SWMX[4]	UART_TX[2]	-	High-Z	Out	UART Transmit 2
		GPIO[1]	-	High-Z	In/Out	GPIO 1 Register, bit 1
		MCFG3[8]	-		In	At reset, bit 8 in MCFG3 register in the memory controller is set from this input.
239	SWMX[5]	UART_RX[2]	-		In	UART Receive 2
		GPIO[2]	-		In	GPIO 1 Register, bit 2 (input only)
238	SWMX[6]	UART_TX[3]	-	High-Z	Out	UART Transmit 3
		GPIO[3]	-	High-Z	In/Out	GPIO 1 Register, bit 3
		MCFG1[9]	-		In	At reset, bit 9 in MCFG1 register in the memory controller is set from this input
233	SWMX[7]	UART_RX[3]	-		In	UART Receive 3
		GPIO[4]	-		In	GPIO 1 Register, bit 4 (input only)
232	SWMX[8]	UART_TX[4]	-	High-Z	Out	UART Transmit 4
		TMDO	-	High-Z	Out	Telemetry Data Out
		GPIO[5]	-	High-Z	In/Out	GPIO 1 Register, bit 5
231	SWMX[9]	UART_RX[4]	-		In	UART Receive 4
		TMCLKI	Rising		In	Telemetry Clock Input
		GPIO[6]	-		In	GPIO 1 Register, bit 6 (input only)
230	SWMX[10]	UART_TX[5]	-	High-Z	Out	UART Transmit 5
		TMCLKO	-	High-Z	Out	Telemetry Clock Output
		GPIO[7]	-	High-Z	In/Out	GPIO 1 Register, bit 7
229	SWMX[11]	UART_RX[5]	-		In	UART Receive 5
		TCACT[0]	High		In	Telecommand Active 0
		GPIO[8]	-		In	GPIO 1 Register, bit 8 (input only)
228	SWMX[12]	SPW_TXS[4]	High	High-Z	Out	SpaceWire Transmit Strobe 4
		SDCSN[0]	Low	High-Z	Out	SDRAM Select 0
		GPIO[9]	-	High-Z	In/Out	GPIO 1 Register, bit 9
227	SWMX[13]	SPW_TXD[4]	High	High-Z	Out	SpaceWire Transmit Data 4
		SDCSN[1]	Low	High-Z	Out	SDRAM Select 1
		GPIO[10]	-	High-Z	In/Out	GPIO 1 Register, bit 10
226	SWMX[14]	SPW_RXS[4]	High		In	SpaceWire Receive Strobe 4
		TCCLK[0]	Rising		In	Telecommand Clock 0
		A16DASA	-		In	ASCS DAS A - Slave data in
		GPIO[11]	-		In	GPIO 1 Register, bit 11 (input only)
225	SWMX[15]	SPW_RXD[4]	High		In	SpaceWire Receive Data 4
		TCD[0]	-		In	Telecommand Data 0
		A16DASB	-		In	ASCS DAS B - Slave data in
		GPIO[12]	-		In	GPIO 1 Register, bit 12 (input only)
220	SWMX[16]	SPW_TXS[2]	High	High-Z	Out	SpaceWire Transmit Strobe 2
		CANTXA	-	High-Z	Out	CAN Transmit A
		GPIO[13]	-	High-Z	In/Out	GPIO 1 Register, bit 13
219	SWMX[17]	SPW_TXD[2]	High	High-Z	Out	SpaceWire Transmit Data 2
		CANTXB	-	High-Z	Out	CAN Transmit B
		GPIO[14]	-	High-Z	In/Out	GPIO 1 Register, bit 14

Table 6. I/O switch matrix pin description, defining the order of priority for outputs and input/outputs, with the highest priority for each pin listed first.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
218	SWMX[18]	SPW_RXS[2]	High		In	SpaceWire Receive Strobe 2
		CANRXA	-		In	CAN Receive A
		GPIO[15]	-		In	GPIO 1 Register, bit 15 (input only)
217	SWMX[19]	SPW_RXD[2]	High		In	SpaceWire Receive Data 2
		CANRXB	-		In	CAN Receive B
		GPIO[16]	-		In	GPIO 1 Register, bit 16 (input only)
203	SWMX[20]	SPW_TXS[3]	High	High-Z	Out	SpaceWire Transmit Strobe 3
		SLSYNC	High	High-Z	Out	SLINK SYNC
		GPIO[17]	-	High-Z	In/Out	GPIO 1 Register, bit 17
202	SWMX[21]	SPW_TXD[3]	High	High-Z	Out	SpaceWire Transmit Data 3
		A16ETR	High	High-Z	Out	ASCS ETR - Synchronization signal
		GPIO[18]	-	High-Z	In/Out	GPIO 1 Register, bit 18
201	SWMX[22]	SPW_RXS[3]]	High		In	SpaceWire Receive Strobe 3
		GPIO[19]	-		In	GPIO 1 Register, bit 19 (input only)
200	SWMX[23]	SPW_RXD[3]	High		In	SpaceWire Receive Data 3
		GPIO[20]	-		In	GPIO 1 Register, bit 20 (input only)
197	SWMX[24]	SPW_TXD[5]	High	High-Z	Out	SpaceWire Transmit Data 5
		SDDQM[0]	High	High-Z	Out	SDRAM Data Mask 0, corresponds to DATA[7:0]
		GPIO[21]	-	High-Z	In/Out	GPIO 1 Register, bit 21
196	SWMX[25]	SPW_TXS[5]	High	High-Z	Out	SpaceWire Transmit Strobe 5
		SDDQM[1]	High	High-Z	Out	SDRAM Data Mask 1, corresponds to DATA[15:8]
		GPIO[22]	-	High-Z	In/Out	GPIO 1 Register, bit 22
193	SWMX[26]	SPW_RXS[5]	High		In	SpaceWire Receive Strobe 5
		TCRFAVL[0]	High		In	Telecommand RF Available 0
		GPIO[23]	-		In	GPIO 1 Register, bit 23 (input only)
192	SWMX[27]	SPW_RXD[5]	High		In	SpaceWire Receive Data 5
		TCCLK[1]	Rising		In	Telecommand Clock 1
		GPIO[24]	-		In	GPIO 1 Register, bit 24 (input only)
191	SWMX[28]	1553RXENA	High	High-Z	Out	MIL-STD-1553B Receive Enable A
		-		High-Z	Out	Proprietary, enabled by CAN
		RMTXD[0]	-	High-Z	Out	Ethernet Transmit Data 0
		GPIO[25]	-	High-Z	In/Out	GPIO 1 Register, bit 25
190	SWMX[29]	1553TXA	High	High-Z	Out	MIL-STD-1553B Transmit Positive A
		-		High-Z	Out	Proprietary, enabled by CAN
		RMTXD[1]	-	High-Z	Out	Ethernet Transmit Data 1
		GPIO[26]	-	High-Z	In/Out	GPIO 1 Register, bit 26
189	SWMX[30]	1553RXA	High		In	MIL-STD-1553B Receive Positive A
		TCD[1]	-		In	Telecommand Data 1
		RMRXD[0]	-		In	Ethernet Receive Data 0
		GPIO[27]	-		In	GPIO 1 Register, bit 27 (input only)
188	SWMX[31]	1553RXNA	Low		In	MIL-STD-1553B Receive Negative A
		TCACT[1]	High		In	Telecommand Active 1
		RMRXD[1]	-		In	Ethernet Receive Data 1
		GPIO[28]	-		In	GPIO 1 Register, bit 28 (input only)
185	SWMX[32]	1553TXNA	Low	High-Z	Out	MIL-STD-1553B Transmit Negative A
		-		High-Z	Out	Proprietary, enabled by CAN
		RMTXEN	High	High-Z	Out	Ethernet Transmit Enable
		GPIO[29]	-	High-Z	In/Out	GPIO 1 Register, bit 29
184	SWMX[33]	1553TXINHA	High	High-Z	Out	MIL-STD-1553B Transmit Inhibit A

Table 6. I/O switch matrix pin description, defining the order of priority for outputs and input/outputs, with the highest priority for each pin listed first.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
		-		High-Z	Out	Proprietary, enabled by CAN
		GPIO[30]	-	High-Z	In/Out	GPIO 1 Register, bit 30
183	SWMX[34]	1553RXB	High		In	MIL-STD-1553B Receive Positive B
		TCRFAVL[1]	High		In	Telecommand RF Available 1
		RMCRSDV	High		In	Ethernet Carrier Sense / Data Valid
		GPIO[31]	-		In	GPIO 1 Register, bit 31 (input only)
182	SWMX[35]	1553RXNB	Low		In	MIL-STD-1553B Receive Negative B
		TCCLK[2]	Rising		In	Telecommand Clock 2
		RMINTN	Low		In	Ethernet Management Interrupt
		GPIO[32]	-		In	GPIO 2 Register, bit 0 (input only)
179	SWMX[36]	1553RXENB	High	High-Z	Out	MIL-STD-1553B Receive Enable B
		A16MCS	High	High-Z	Out	ASCS MCS - TM start/stop signal
		RMMDIO	-	High-Z	In/Out	Ethernet Media Interface Data
		GPIO[33]	-	High-Z	In/Out	GPIO 2 Register, bit 1
178	SWMX[37]	1553TXB	High	High-Z	Out	MIL-STD-1553B Transmit Positive B
		A16HS	High	High-Z	Out	ASCS HS - TM/TC serial clock
		RMMDC	-	High-Z	Out	Ethernet Media Interface Clock
		GPIO[34]	-	High-Z	In/Out	GPIO 2 Register, bit 2
		SpaceWire clock divisor registers	-		In	At reset, bits 8 and 0 in the clock divisor register of the SpaceWire interfaces are set from this input
177	SWMX[38]	1553CK	-		In	MIL-STD-1553B Clock
		TCD[2]	-		In	Telecommand Data 2
		RMRFCLK	-		In	Ethernet Reference Clock
		GPIO[35]	-		In	GPIO 2 Register, bit 3 (input only)
176	SWMX[39]	TCACT[2]	High		In	Telecommand Active 2
		GPIO[36]	-		In	GPIO 2 Register, bit 4 (input only)
175	SWMX[40]	1553TXNB	Low	High-Z	Out	MIL-STD-1553B Transmit Negative B
		A16DCS	-	High-Z	Out	ASCS DCS - Slave data out
		GPIO[37]	-	High-Z	In/Out	GPIO 2 Register, bit 5
		SpaceWire clock divisor registers	-		In	At reset, bits 9 and 1 in the clock divisor register of the SpaceWire interfaces are set from this input
174	SWMX[41]	1553TXINHB	High	High-Z	Out	MIL-STD-1553B Transmit Inhibit B
		A16MAS	High	High-Z	Out	ASCS MAS - TM start/stop signal
		GPIO[38]	-	High-Z	In/Out	GPIO 2 Register, bit 6
173	SWMX[42]	TCRFAVL[2]	High		In	Telecommand RF Available 2
		GPIO[39]	-		In	GPIO 2 Register, bit 7 (input only)
172	SWMX[43]	-		High-Z	Out	Proprietary, enabled by CAN
		GPIO[40]	-	High-Z	In/Out	GPIO 2 Register, bit 8
		SpaceWire clock divisor registers	-		In	At reset, bits 10 and 2 in the clock divisor register of the SpaceWire interfaces are set from this input
169	SWMX[44]	SPICLK		High-Z	Out	SPI Clock
		SLO	-	High-Z	Out	SLINK Data Out
		GPIO[41]	-	High-Z	In/Out	GPIO 2 Register, bit 9
166	SWMX[45]	SPIMOSI	-	High-Z	Out	SPI Master Out Slave In
		SLCLK	High	High-Z	Out	SLINK Clock
		GPIO[42]	-	High-Z	In/Out	GPIO 2 Register, bit 10
		SpaceWire clock divisor registers	-		In	At reset, bits 11 and 3 in the clock divisor register of the SpaceWire interfaces are set from this input
165	SWMX[46]	TCCLK[3]	Rising		In	Telecommand Clock 3
		GPIO[43]	-		In	GPIO 2 Register, bit 11 (input only)

Table 6. I/O switch matrix pin description, defining the order of priority for outputs and input/outputs, with the highest priority for each pin listed first.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
164	SWMX[47]	TCD[3]	-		In	Telecommand Data 3
		GPIO[44]	-		In	GPIO 2 Register, bit 12 (input only)
163	SWMX[48]	SDCASN	Low	High-Z	Out	SDRAM Column Address Strobe
		GPIO[45]	-	High-Z	In/Out	GPIO 2 Register, bit 13
162	SWMX[49]	SDRASN	Low	High-Z	Out	SDRAM Row Address Strobe
		GPIO[46]	-	High-Z	In/Out	GPIO 2 Register, bit 14
161	SWMX[50]	TCACT[3]	High		In	Telecommand Active 3
		GPIO[47]	-		In	GPIO 2 Register, bit 15 (input only)
160	SWMX[51]	SPIMISO			In	SPI Master In Slave Out
		TCRFAVL[3]	High		In	Telecommand RF Available 3
		SLI	-		In	SLINK Data In
		GPIO[48]	-		In	GPIO 2 Register, bit 16 (input only)
157	SWMX[52]	SDWEN	Low	High-Z	Out	SDRAM Write Enable
		GPIO[49]	-	High-Z	In/Out	GPIO 2 Register, bit 17
155	SWMX[53]	SDDQM[2]	High	High-Z	Out	SDRAM Data Mask 2, corresponds to DATA[23:16]
		GPIO[50]	-	High-Z	In/Out	GPIO 2 Register, bit 18
154	SWMX[54]	SDDQM[3]	High	High-Z	Out	SDRAM Data Mask 3, corresponds to DATA[31:24]
		GPIO[51]	-	High-Z	In/Out	GPIO 2 Register, bit 19
153	SWMX[55]	TCACT[4]	High		In	Telecommand Active 4
		GPIO[52]	-		In	GPIO 2 Register, bit 20 (input only)
144	SWMX[56]	TCRFAVL[4]	High		In	Telecommand RF Available 4
		GPIO[53]	-		In	GPIO 2 Register, bit 21 (input only)
143	SWMX[57]	I2CSDA		High-Z	In/Out	I2C Serial Data
		GPIO[54]	-	High-Z	In/Out	GPIO 2 Register, bit 22
		TCCLK[4]	Rising		In	Telecommand Clock 4
142	SWMX[58]	I2CSCL		High-Z	In/Out	I2C Serial Clock
		GPIO[55]	-	High-Z	In/Out	GPIO 2 Register, bit 23
		TCD[4]	-		In	Telecommand Data 4
140	SWMX[59]	CB[8]	-	High-Z	In/Out	Reed-Solomon Check Bit 8
		GPIO[56]	-	High-Z	In/Out	GPIO 2 Register, bit 24
137	SWMX[60]	CB[9]	-	High-Z	In/Out	Reed-Solomon Check Bit 9
		GPIO[57]	-	High-Z	In/Out	GPIO 2 Register, bit 25
136	SWMX[61]	CB[10]	-	High-Z	In/Out	Reed-Solomon Check Bit 10
		GPIO[58]	-	High-Z	In/Out	GPIO 2 Register, bit 26
135	SWMX[62]	CB[11]	-	High-Z	In/Out	Reed-Solomon Check Bit 11
		GPIO[59]	-	High-Z	In/Out	GPIO 2 Register, bit 27
132	SWMX[63]	CB[12]	-	High-Z	In/Out	Reed-Solomon Check Bit 12
		GPIO[60]	-	High-Z	In/Out	GPIO 2 Register, bit 28
129	SWMX[64]	CB[13]	-	High-Z	In/Out	Reed-Solomon Check Bit 13
		GPIO[61]	-	High-Z	In/Out	GPIO 2 Register, bit 29
128	SWMX[65]	CB[14]	-	High-Z	In/Out	Reed-Solomon Check Bit 14
		GPIO[62]	-	High-Z	In/Out	GPIO 2 Register, bit 30
127	SWMX[66]	CB[15]	-	High-Z	In/Out	Reed-Solomon Check Bit 15
		GPIO[63]	-	High-Z	In/Out	GPIO 2 Register, bit 31

Table 7. I/O switch matrix pins listed per function, including supporting signals outside the I/O switch matrix.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
157	SWMX[52]	SDWEN	Low	High-Z	Out	SDRAM Write Enable
163	SWMX[48]	SDCASN	Low	High-Z	Out	SDRAM Column Address Strobe
162	SWMX[49]	SDRASN	Low	High-Z	Out	SDRAM Row Address Strobe
228	SWMX[12]	SDCSN[0]	Low	High-Z	Out	SDRAM Select 0
227	SWMX[13]	SDCSN[1]	Low	High-Z	Out	SDRAM Select 1
197	SWMX[24]	SDDQM[0]	High	High-Z	Out	SDRAM Data Mask 0, corresponds to DATA[7:0]
196	SWMX[25]	SDDQM[1]	High	High-Z	Out	SDRAM Data Mask 1, corresponds to DATA[15:8]
155	SWMX[53]	SDDQM[2]	High	High-Z	Out	SDRAM Data Mask 2, corresponds to DATA[23:16]
154	SWMX[54]	SDDQM[3]	High	High-Z	Out	SDRAM Data Mask 3, corresponds to DATA[31:24]
140	SWMX[59]	CB[8]	-	High-Z	In/Out	Check Bit 8, Reed-Solomon
137	SWMX[60]	CB[9]	-	High-Z	In/Out	Check Bit 9, Reed-Solomon
136	SWMX[61]	CB[10]	-	High-Z	In/Out	Check Bit 10, Reed-Solomon
135	SWMX[62]	CB[11]	-	High-Z	In/Out	Check Bit 11, Reed-Solomon
132	SWMX[63]	CB[12]	-	High-Z	In/Out	Check Bit 12, Reed-Solomon
129	SWMX[64]	CB[13]	-	High-Z	In/Out	Check Bit 13, Reed-Solomon
128	SWMX[65]	CB[14]	-	High-Z	In/Out	Check Bit 14, Reed-Solomon
127	SWMX[66]	CB[15]	-	High-Z	In/Out	Check Bit 15, Reed-Solomon
	ADDRESS[16:2]	ADDRESS[16:2]	-	Low	Out	Memory address
	DATA[31:0]	DATA[31:0]	-	High-Z	In/Out	Memory data bus
	CB[7:0]	CB[7:0]	-	High-Z	In/Out	Memory checkbits
240	SWMX[4]	MCFG3[8]	-		In	At reset, bit 8 in MCFG3 register in the memory controller is set from this input.
238	SWMX[6]	MCFG1[9]	-		In	At reset, bit 9 in MCFG1 register in the memory controller is set from this input
4	SWMX[0]	UART_TX[0]	-	High	Out	UART Transmit 0
3	SWMX[1]	UART_RX[0]	-		In	UART Receive 0
2	SWMX[2]	UART_TX[1]	-	High	Out	UART Transmit 1
1	SWMX[3]	UART_RX[1]	-		In	UART Receive 1
240	SWMX[4]	UART_TX[2]	-	High-Z	Out	UART Transmit 2
239	SWMX[5]	UART_RX[2]	-		In	UART Receive 2
238	SWMX[6]	UART_TX[3]	-	High-Z	Out	UART Transmit 3
233	SWMX[7]	UART_RX[3]	-		In	UART Receive 3
232	SWMX[8]	UART_TX[4]	-	High-Z	Out	UART Transmit 4
231	SWMX[9]	UART_RX[4]	-		In	UART Receive 4
230	SWMX[10]	UART_TX[5]	-	High-Z	Out	UART Transmit 5
229	SWMX[11]	UART_RX[5]	-		In	UART Receive 5
213	SPW_RXD[0]	SPW_RXD[0]	High		In	SpaceWire Receive Data 0
214	SPW_RXS[0]	SPW_RXS[0]	High		In	SpaceWire Receive Strobe 0
215	SPW_TXD[0]	SPW_TXD[0]	High	Low	Out	SpaceWire Transmit Data 0
216	SPW_TXS[0]	SPW_TXS[0]	High	Low	Out	SpaceWire Transmit Strobe 0
204	SPW_RXD[1]	SPW_RXD[1]	High		In	SpaceWire Receive Data 1
205	SPW_RXS[1]	SPW_RXS[1]	High		In	SpaceWire Receive Strobe 1
206	SPW_TXD[1]	SPW_TXD[1]	High	Low	Out	SpaceWire Transmit Data 1
209	SPW_TXS[1]	SPW_TXS[1]	High	Low	Out	SpaceWire Transmit Strobe 1
217	SWMX[19]	SPW_RXD[2]	High		In	SpaceWire Receive Data 2
218	SWMX[18]	SPW_RXS[2]	High		In	SpaceWire Receive Strobe 2
219	SWMX[17]	SPW_TXD[2]	High	High-Z	Out	SpaceWire Transmit Data 2
220	SWMX[16]	SPW_TXS[2]	High	High-Z	Out	SpaceWire Transmit Strobe 2
200	SWMX[23]	SPW_RXD[3]	High		In	SpaceWire Receive Data 3
201	SWMX[22]	SPW_RXS[3]	High		In	SpaceWire Receive Strobe 3

Table 7. I/O switch matrix pins listed per function, including supporting signals outside the I/O switch matrix.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
202	SWMX[21]	SPW_TXD[3]	High	High-Z	Out	SpaceWire Transmit Data 3
203	SWMX[20]	SPW_TXS[3]	High	High-Z	Out	SpaceWire Transmit Strobe 3
225	SWMX[15]	SPW_RXD[4]	High		In	SpaceWire Receive Data 4
226	SWMX[14]	SPW_RXS[4]	High		In	SpaceWire Receive Strobe 4
227	SWMX[13]	SPW_TXD[4]	High	High-Z	Out	SpaceWire Transmit Data 4
228	SWMX[12]	SPW_TXS[4]	High	High-Z	Out	SpaceWire Transmit Strobe 4
192	SWMX[27]	SPW_RXD[5]	High		In	SpaceWire Receive Data 5
193	SWMX[26]	SPW_RXS[5]	High		In	SpaceWire Receive Strobe 5
197	SWMX[24]	SPW_TXD[5]	High	High-Z	Out	SpaceWire Transmit Data 5
196	SWMX[25]	SPW_TXS[5]	High	High-Z	Out	SpaceWire Transmit Strobe 5
178	SWMX[37]	SpaceWire clock divisor registers values at reset, all other bits are zero.	-		In	At reset, bits 8 & 0 are set from this input
175	SWMX[40]		-		In	At reset, bits 9 & 1 are set from this input
172	SWMX[43]		-		In	At reset, bits 10 & 2 are set from this input
166	SWMX[45]		-		In	At reset, bits 11 & 3 are set from this input
185	SWMX[32]	RMTXEN	High	High-Z	Out	Ethernet Transmit Enable
191	SWMX[28]	RMTXD[0]	-	High-Z	Out	Ethernet Transmit Data 0
190	SWMX[29]	RMTXD[1]	-	High-Z	Out	Ethernet Transmit Data 1
189	SWMX[30]	RMRXD[0]	-		In	Ethernet Receive Data 0
188	SWMX[31]	RMRXD[1]	-		In	Ethernet Receive Data 1
183	SWMX[34]	RMCRSDV	High		In	Ethernet Carrier Sense / Data Valid
182	SWMX[35]	RMINTN	Low		In	Ethernet Management Interrupt
179	SWMX[36]	RMMDIO	-	High-Z	In/Out	Ethernet Media Interface Data
178	SWMX[37]	RMMDC	-	High-Z	Out	Ethernet Media Interface Clock
177	SWMX[38]	RMRFCLK	-		In	Ethernet Reference Clock
220	SWMX[16]	CANTXA	-	High-Z	Out	CAN Transmit A
218	SWMX[18]	CANRXA	-		In	CAN Receive A
219	SWMX[17]	CANTXB	-	High-Z	Out	CAN Transmit B
217	SWMX[19]	CANRXB	-		In	CAN Receive B
191	SWMX[28]	-		High-Z	Out	Proprietary, enabled by CAN
190	SWMX[29]	-		High-Z	Out	Proprietary, enabled by CAN
185	SWMX[32]	-		High-Z	Out	Proprietary, enabled by CAN
184	SWMX[33]	-		High-Z	Out	Proprietary, enabled by CAN
172	SWMX[43]	-		High-Z	Out	Proprietary, enabled by CAN
177	SWMX[38]	1553CK	-		In	MIL-STD-1553B Clock
184	SWMX[33]	1553TXINHA	High	High-Z	Out	MIL-STD-1553B Transmit Inhibit A
190	SWMX[29]	1553TXA	High	High-Z	Out	MIL-STD-1553B Transmit Positive A
185	SWMX[32]	1553TXNA	Low	High-Z	Out	MIL-STD-1553B Transmit Negative A
191	SWMX[28]	1553RXENA	High	High-Z	Out	MIL-STD-1553B Receive Enable A
189	SWMX[30]	1553RXA	High		In	MIL-STD-1553B Receive Positive A
188	SWMX[31]	1553RXNA	Low		In	MIL-STD-1553B Receive Negative A
174	SWMX[41]	1553TXINHB	High	High-Z	Out	MIL-STD-1553B Transmit Inhibit B
178	SWMX[37]	1553TXB	High	High-Z	Out	MIL-STD-1553B Transmit Positive B
175	SWMX[40]	1553TXNB	Low	High-Z	Out	MIL-STD-1553B Transmit Negative B
179	SWMX[36]	1553RXENB	High	High-Z	Out	MIL-STD-1553B Receive Enable B
183	SWMX[34]	1553RXB	High		In	MIL-STD-1553B Receive Positive B
182	SWMX[35]	1553RXNB	Low		In	MIL-STD-1553B Receive Negative B
143	SWMX[57]	I2CSDA		High-Z	In/Out	I2C Serial Data
142	SWMX[58]	I2CSCL		High-Z	In/Out	I2C Serial Clock
169	SWMX[44]	SPICLK		High-Z	Out	SPI Clock
166	SWMX[45]	SPIMOSI	-	High-Z	Out	SPI Master Out Slave In
160	SWMX[51]	SPIMISO			In	SPI Master In Slave Out

Table 7. I/O switch matrix pins listed per function, including supporting signals outside the I/O switch matrix.

Pin no.	Pin name	Pin function	Polarity	Reset value	Dir.	Description
203	SWMX[20]	SLSYNC	High	High-Z	Out	SLINK SYNC
166	SWMX[45]	SLCLK	High	High-Z	Out	SLINK Clock
160	SWMX[51]	SLI	-		In	SLINK Data In
169	SWMX[44]	SLO	-	High-Z	Out	SLINK Data Out
226	SWMX[14]	A16DASA	-		In	ASCS DAS A - Slave data in
225	SWMX[15]	A16DASB	-		In	ASCS DAS B - Slave data in
202	SWMX[21]	A16ETR	High	High-Z	Out	ASCS ETR - Synchronization signal
175	SWMX[40]	A16DCS	-	High-Z	Out	ASCS DCS - Slave data out
174	SWMX[41]	A16MAS	High	High-Z	Out	ASCS MAS - TM start/stop signal
179	SWMX[36]	A16MCS	High	High-Z	Out	ASCS MCS - TC start/stop signal
178	SWMX[37]	A16HS	High	High-Z	Out	ASCS HS - TM/TC serial clock
229	SWMX[11]	TCACT[0]	High		In	Telecommand Active 0
226	SWMX[14]	TCCLK[0]	Rising		In	Telecommand Clock 0
225	SWMX[15]	TCD[0]	-		In	Telecommand Data 0
193	SWMX[26]	TCRFAVL[0]	High		In	Telecommand RF Available 0
188	SWMX[31]	TCACT[1]	High		In	Telecommand Active 1
192	SWMX[27]	TCCLK[1]	Rising		In	Telecommand Clock 1
189	SWMX[30]	TCD[1]	-		In	Telecommand Data 1
183	SWMX[34]	TCRFAVL[1]	High		In	Telecommand RF Available 1
176	SWMX[39]	TCACT[2]	High		In	Telecommand Active 2
182	SWMX[35]	TCCLK[2]	Rising		In	Telecommand Clock 2
177	SWMX[38]	TCD[2]	-		In	Telecommand Data 2
173	SWMX[42]	TCRFAVL[2]	High		In	Telecommand RF Available 2
161	SWMX[50]	TCACT[3]	High		In	Telecommand Active 3
165	SWMX[46]	TCCLK[3]	Rising		In	Telecommand Clock 3
164	SWMX[47]	TCD[3]	-		In	Telecommand Data 3
160	SWMX[51]	TCRFAVL[3]	High		In	Telecommand RF Available 3
153	SWMX[55]	TCACT[4]	High		In	Telecommand Active 4
143	SWMX[57]	TCCLK[4]	Rising		In	Telecommand Clock 4
142	SWMX[58]	TCD[4]	-		In	Telecommand Data 4
144	SWMX[56]	TCRFAVL[4]	High		In	Telecommand RF Available 4
231	SWMX[9]	TMCLKI	Rising		In	Telemetry Clock Input
230	SWMX[10]	TMCLKO	-	High-Z	Out	Telemetry Clock Output
232	SWMX[8]	TMDO	-	High-Z	Out	Telemetry Data Out

Table 8. Conflicting interfaces in the I/O switch matrix are marked with an X, with duplicates shown in bold typeface.

	SDRAM (with RS)	UART 0	UART 1	UART 2	UART 3	UART 4	UART 5	SpaceWire 0	SpaceWire 1	SpaceWire 2	SpaceWire 3	SpaceWire 4	SpaceWire 5	Ethernet	CAN	MIL-STD-1553B	I2C	SPI	SLINK	ASCS	CCSDS TC 0	CCSDS TC 1	CCSDS TC 2	CCSDS TC 3	CCSDS TC 4	CCSDS TM
SDRAM (with RS)	-											**X**	**X**													
UART 0		-																								
UART 1			-																							
UART 2				-																						
UART 3					-																					
UART 4						-																				X
UART 5							-												X							X
SpaceWire 0								-																		
SpaceWire 1									-																	
SpaceWire 2										-						X			X							
SpaceWire 3											-							X								
SpaceWire 4	X											-							X	X						
SpaceWire 5	X												-						X	X						
Ethernet														-	X	X						*X*	*X*			
CAN										X				X	-											
MIL-STD-1553B														X		-			X	X	X					
I2C																	-								X	
SPI																		-	X					X		
SLINK											X							X	-					X		
ASCS										X		X				X				-	*X*					
CCSDS TC 0						X									X	X				*X*	-					
CCSDS TC 1														*X*		X						-				
CCSDS TC 2														*X*		X							-			
CCSDS TC 3																			X	X				-		
CCSDS TC 4																X									-	
CCSDS TM					X	X																				-

Which peripheral function that drives a shared I/O pin is determined through a set of enabling conditions. Peripheral functions that are clock gated use the clock enable bit as I/O matrix select signal. Enabling conditions for peripherals without clock gating are defined in table 9.

TABLE 9. I/O switch matrix enables for non-gated devices

Device	Enable
UART	TX enable bit in APBUART control register
SPI	Enable bit in SPICTRL control register
I2C	Enable bit in I2CMST control register
ASCS	Output enable bit in GRASCS control register
SLINK	Enable bit in SLINKMST control register
SDRAM	SDRAM enable in MCFG2
RS checkbits	Reed-Solomon enable in MCFG3

The peripherals are connected to the switch matrix I/O pin through a priority structure that makes sure only one peripheral is granted the pin at a time. If no device is enabled, the pin defaults to the respective GPIO[63:0]. Note that there are two GRGPIO cores each with 32 bits, thus only 64 of the 67 pins have a GPIO pin alternative. The first GRGPIO core is mapped on GPIO[31:0], and the second GPIO core is mapped on GPIO[63:32]. Also note that when the CAN core is enabled, the output of pins 172,

184, 185, 190 and 191 will be driven with random values. These pins are shared with the Mil-Std-1553 interface, but since this interface has a higher priority, the two interfaces are not considered conflicting in the case they are both enabled simultaneously.

The pins listed in table 10 have fixed functions that are not connected to the switch matrix.

TABLE 10. GR712RC fixed pins

Interface	GR712RC pin	In	Out	In/Out	Total
Miscellaneous	ERRORN		1		1
	INCLK	1			1
	RESETN	1			1
	DLLBPN	1			1
	TESTEN	1			1
	SCANEN	1			1
	WDOGN		1		1
JTAG	TDI	1			1
	TDO		1		1
	TMS	1			1
	TCK	1			1
Memory interface	ADDRESS[23:0]		24		24
	DATA[31:0]			32	32
	Control (READ, BRDYN, BEXCN)	2	1		3
	Check bits (CB[7:0])			8	8
	SRAM control (RAMSN[1:0], RAMOEN, RAMWEN)		4		4
	PROM control (ROMSN[1:0], OEN, WRITEN)		4		4
	I/O control (IOSN)		1		1
	SDRAM clock (SDCLK)		1		1
SpaceWire 0-1	SpaceWire receiver and transmitter clock (SPWCLK)	1			1
(with RMAP)	SpaceWire 0-1 receive/transmit data/strobe	4	4		8
Number of pins		*15*	*42*	*40*	**97**

3   Clocking

The table below specifies the clock inputs to the GR712RC.

TABLE 11. Clock inputs

Clock input	Description
INCLK	System clock input
SPWCLK	SpaceWire clock
TMCLKI	Telemetry transponder clock (input shared with UART4)
1553CK	1553 clock (input shared with Ethernet and TC)
RMRFCLK	Ethernet RMII clock (input shared with 1553 and TC)

The design makes use of two MFDLL clock multipliers to create the system clock and the SpaceWire transmitter clock. The system clock (AHB clock), the SpaceWire transmitter clock and the 1553 clock is connected through clock multiplexers in order to select between a number of clock options described below.

3.1   System clock

The system clock is used to clock the processors, the AMBA buses, and all on-chip cores. The system clock is either taken from the INCLK input, or from a MFDLL clock multiplier providing 2x INCLK. The DLLBPN input is used to select system clock source.

3.2   SpaceWire clock

The clock source for SpaceWire is selectable through a clock multiplexer, whose inputs are SPW-CLK, and INCLK. The selected clock is then used either as 1X, 2X or 4X as shown in the figure below. The system clock (CLK) is also selectable as SpaceWire transmit clock.

After reset the selected SpaceWire transmit clock is SPWCLK without any multiplication. The Space-Wire clock multiplexers and the DLL reset are controlled through the GRGPREG register.

The SpaceWire transmit clock must be a multiple of 10 MHz in order to achieve 10 Mbps start up bit rate. The division to 10 MHz is done internally in the GRSPW2 core. During reset the clock divider register in GRSPW2 gets its value from 4 I/O pins that must be pulled up/down to set the divider correctly. Pins (MSB-LSB) SWMX45, 43, 40 and 37 are used for this. Thus it is possible to use a SPW-

CLK which is any multiple of 10 between 10-100 MHz (note that the required precision is 10 MHz +/ - 1 MHz).

The SpaceWire transmitter uses SDR output registers meaning that the bitrate will be equal to the transmit clock. The SpaceWire input signals are sampled synchronously with DDR registers using the transmit clock.

3.3 MIL-STD-1553

The B1553BRM core needs a 24, 20 or 16 MHz clock. The frequency used can be configured through a register in the core. This clock can either be supplied directly to the B1553BRM core through the 1553CK pin, the system clock, or it can be generated through a clock divider that divides the system clock and is programmable through the GPREG.

If the system clock is used to generate the BRM clock it must be one of the frequencies in the table below:

Table 12. MIL-STD-1553 BRM frequencies

System clock (MHz)	Division	BRM frequency (MHz)
32	2	16 [2]
40	2	20
48	2	24
64	4	16
80	4	20
96	4	24
120	6	20
144	6	24

Note 1: The system clock frequency must always be higher than BRM clock frequency, regardless of how the BRM clock is generated.

Note 2: This specific case has not been validated.

3.4 Telemetry

The telemetry input clock frequency must be less than or equal to the system clock frequency. It is possible to select either TMCLKI or system clock as input clock to the GRTM core.

The TMDO data signal is timed relative to TMCLKO.

3.5 Telecommand

The telecommand input signals are sampled using the system clock and the telecommand input clock frequency must be less than or equal to the system clock frequency divided by eight.

3.6 Obsolete

Proprietary function not supported.

3.7 SLINK

The SLINK core is clocked by the system clock and has a built in clock divider. When using this interface the system clock must be a multiple of 12 MHz if a 50% duty cycle is necessary. If the SLINK clock duty cycle does not need to be 50% any multiple of 6 MHz is possible.

3.8 SDRAM clock

The SDRAM clock (SDCLK) output is driven from the internal system clock. A programmable delay line allows tuning the SDCLK phase relative to the internal clock. The output is tuned to be in phase with the internal clock tree when the delay line is set to zero. If the system clock is generated by the 2x MFDLL, SDCLK is not active when the reset input (RESETN) is asserted. When the system clock is generated directly from INCLK (i.e. when DLLBPN = 0), then SDCLK is driven directly from INCLK passing through the Delay line.
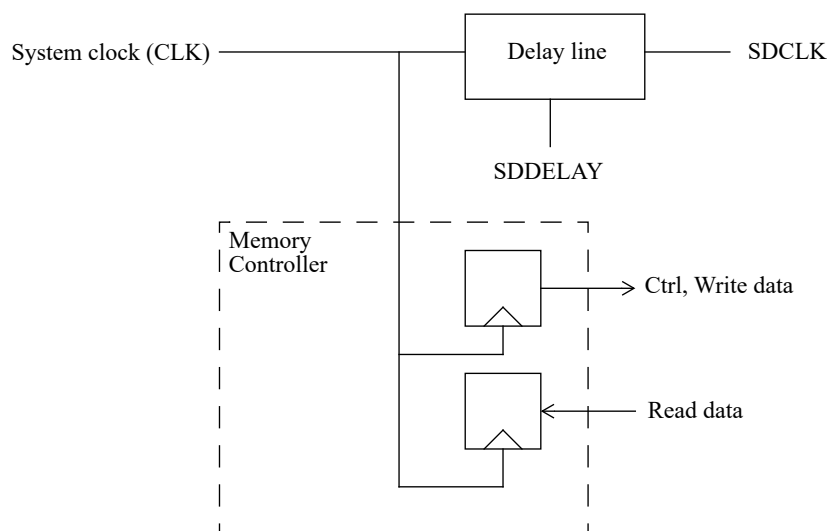
3.9 Clock gating unit

The GR712RC device has a clock gating unit through which individual cores can have their AHB clocks enabled/disabled and resets driven. The cores connected to the clock gating unit are listed in the table below:

Table 13. Devices with gatable clock

Device
Ethernet
SpaceWire 0
SpaceWire 1
SpaceWire 2
SpaceWire 3
SpaceWire 4
SpaceWire 5
CAN 0 - 1
Proprietary
CCSDS Telemetry
CCSDS Telecommand
MIL-STD-1553
LEON3FT (0)
LEON3FT (1)
GRFPU(0)
GRFPU(1)

The LEON3FT processor cores will automatically be clock gated when the processor enter power-down or halt state. The floating-point units (GRFPU) will be clock gated when the corresponding processor have disabled FPU operations by setting the %psr.ef bit to zero, or when the processor have entered power-down/halt mode. For more information see the chapter about the clock gating unit.

3.10   Test mode clocking

When in test mode (TESTEN signal = 1) all clocks in the design are connected to the INCLK test clock.

4 LEON3FT - High-performance SPARC V8 32-bit Processor

4.1 Overview

The GR712RC implements two LEON3FT processor cores in SMP configuration. LEON3FT is a 32-bit processor core conforming to the IEEE-1754 (SPARC V8) architecture. It is designed for embedded applications, combining high performance with low complexity and low power consumption.

The LEON3FT core has the following main features: 7-stage pipeline with Harvard architecture, separate instruction and data caches, hardware multiplier and divider, on-chip debug support and multiprocessor extensions.
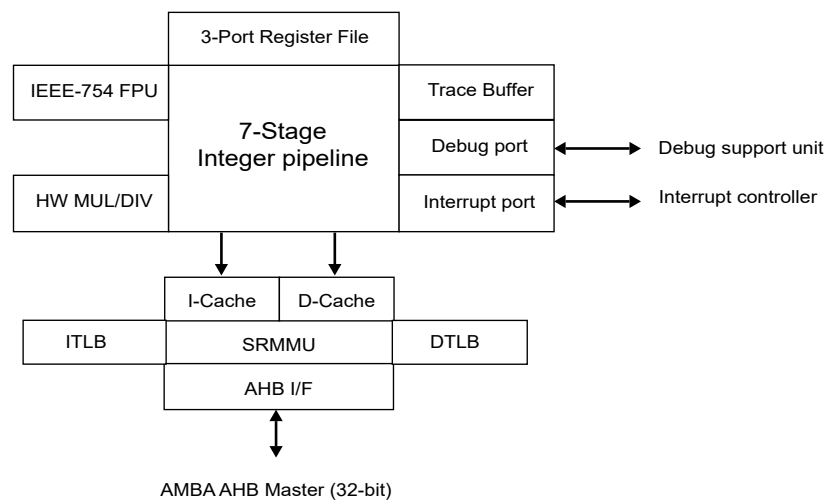
Figure 2. LEON3 processor core block diagram

4.1.1 Integer unit

The LEON3 integer unit implements the full SPARC V8 standard, including hardware multiply and divide instructions. The number of register windows is 8. The pipeline consists of 7 stages with a separate instruction and data cache interface (Harvard architecture).

4.1.2 Cache sub-system

Each processor is configured with a 16 KiB instruction cache and a 16 KiB data cache. The instruction cache uses streaming during line-refill to minimize refill latency. The data cache uses write-through policy and implements a double-word write-buffer. The data cache can also perform bus-snooping on the AHB bus.

4.1.3 Floating-point unit

Each LEON3 processor is connected to a unique GRFPU floating-point unit. The GRFPU implements the IEEE-754 floating-point format and supports both single and double-precision operands.

4.1.4 Memory management unit

A SPARC V8 Reference Memory Management Unit (SRMMU) is provided per processor. The SRMMU implements the full SPARC V8 MMU specification, and provides mapping between multiple 32-bit virtual address spaces and 32-bit physical memory. A three-level hardware table-walk is implemented, and the MMU is configured with 32 fully associative TLB entries.

4.1.5 On-chip debug support

The LEON3 pipeline includes functionality to allow non-intrusive debugging on target hardware. Through the JTAG debug support interface, full access to all processor registers and caches is provided. The debug interfaces also allows single stepping, instruction tracing and hardware breakpoint/watchpoint control. An internal trace buffer can monitor and store executed instructions, which can later be read out over the debug interface.

4.1.6 Interrupt interface

LEON3 supports the SPARC V8 interrupt model with a total of 15 asynchronous interrupts. The interrupt interface provides functionality to both generate and acknowledge interrupts. The GR712RC contains an interrupt controller that support 31 system interrupts, mapped on the 15 processor interrupts.

4.1.7 AMBA interface

The cache system implements an AMBA AHB master to load and store data to/from the caches. The interface is compliant with the AMBA-2.0 standard. During line refill, incremental burst are generated to optimise the data transfer.

4.1.8 Power-down mode

The LEON3 processor core implements a power-down mode, which halts the pipeline and caches until the next interrupt using clock gating. This is an efficient way to minimize power-consumption when the application is idle or when one of the processor cores is not used.

4.1.9 Multi-processor support

LEON3 is designed to be use in multi-processor systems and the GR712RC contains two cores. Each processor has a unique index to allow processor enumeration. The write-through caches and snooping mechanism guarantees memory coherency in shared-memory systems.

4.1.10 Fault-tolerance

LEON3FT contains logic to correct up to 4 bit errors per 32-bit cache word and associated tag. The processor registers are protected by a SEC/DED BCH EDAC.

4.1.11 Performance

Using gcc-4.4.2, a dhrystone figure of 1.34 DMIPS/MHz can be achieved.

4.2 LEON3 integer unit

4.2.1 Overview

The LEON3 integer unit implements the integer part of the SPARC V8 instruction set. The implementation is focused on high performance and low complexity. The LEON3 integer unit has the following main features:

- 7-stage instruction pipeline
- Separate instruction and data cache interface
- 8 register windows
- 16-bit Hardware multiplier and Radix-2 divider (non-restoring)
- Single-vector trapping for reduced code size

Figure 3 shows a block diagram of the integer unit.
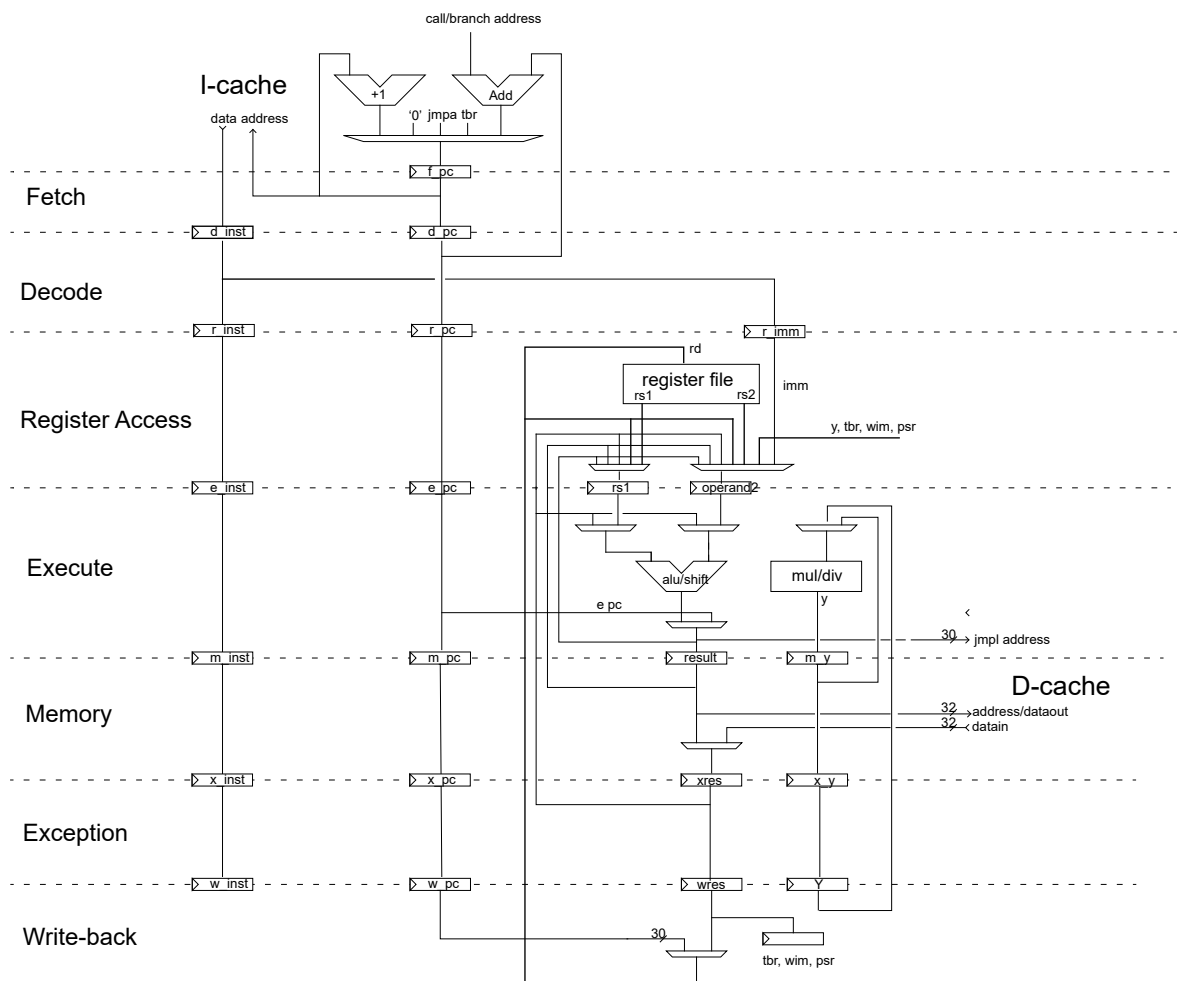
Figure 3. LEON3 integer unit datapath diagram

4.2.2 Instruction pipeline

The LEON integer unit uses a single instruction issue pipeline with 7 stages:

1. FE (Instruction Fetch): If the instruction cache is enabled, the instruction is fetched from the instruction cache. Otherwise, the fetch is forwarded to the memory controller. The instruction is valid at the end of this stage and is latched inside the IU.

2. DE (Decode): The instruction is decoded and the CALL/Branch target addresses are generated.

3. RA (Register access): Operands are read from the register file or from internal data bypasses.

4. EX (Execute): ALU operations are performed. For memory operations (e.g., LD) and for JMPL/RETT, the address is generated.

5. ME (Memory): Data cache is accessed. Store data is written to the data cache at this time.

6. XC (Exception) Traps and interrupts are resolved. For cache reads, the data is aligned.

7. WR (Write): The result of any ALU or cache operations are written back to the register file.

Table 14 lists the cycles per instruction (assuming cache hit and no icc or load interlock):

Table 14. Instruction timing

Instruction	Clock Cycles
JMPL, RETT	3
Double load	2
Single store	2
Double store	3
SMUL/UMUL	4
SDIV/UDIV	35
Taken Trap	5
Atomic load/store	3
All other instructions	**1**

4.2.3   SPARC Implementor's ID

Gaisler Research is assigned number 15 (0xF) as SPARC implementor's identification. This value is hard-coded into bits 31:28 in the %psr register. The version number for LEON3 is 3, which is hard-coded in to bits 27:24 of the %psr.

4.2.4   Divide instructions

Full support for SPARC V8 divide instructions is provided (SDIV, UDIV, SDIVCC & UDIVCC). The divide instructions perform a 64-by-32 bit divide and produce a 32-bit result. Rounding and overflow detection is performed as defined in the SPARC V8 standard.

4.2.5   Co-processor

This implementation does not have a co-processor connected to the integer unit. The %psr.EC bit is writable by software but does not have any effect on processor operation.

4.2.6 Multiply instructions

The LEON processor supports the SPARC integer multiply instructions UMUL, SMUL UMULCC and SMULCC. These instructions perform a 32x32-bit integer multiply, producing a 64-bit result. SMUL and SMULCC performs signed multiply while UMUL and UMULCC performs unsigned multiply. UMULCC and SMULCC also set the condition codes to reflect the result. The multiply instructions are performed using a 16x16 hardware multiplier which is iterated four times. To improve the timing, the 16x16 multiplier is provided with a pipeline stage.

4.2.7 Compare and Swap instruction (CASA)

GR712RC implements the SPARC V9 Compare and Swap Alternative (CASA) instruction. The CASA operates as described in the SPARC V9 manual. The instruction is privileged but setting ASI = 0xA (user data) will allow it to be used in user mode.

4.2.8 Hardware breakpoints

The integer unit includes two hardware breakpoints. Each breakpoint consists of a pair of application-specific registers (%asr24/25 and %asr26/27), one with the break address and one with a mask:

	31		2	1	0
%asr24, %asr26		WADDR[31:2]			IF

	31		2		0
%asr25, %asr27		WMASK[31:2]		DL	DS

Figure 4. Watch-point registers

Any binary aligned address range can be watched - the range is defined by the WADDR field, masked by the WMASK field (WMASK[x] = 1 enables comparison). On a breakpoint hit, trap 0x0B is generated. By setting the IF, DL and DS bits, a hit can be generated on instruction fetch, data load or data store. Clearing these three bits will effectively disable the breakpoint function.

4.2.9 Instruction trace buffer

The instruction trace buffer consists of a circular buffer that stores executed instructions. The trace buffer operation is controlled through the debug support interface, and does not affect processor operation. The size of the trace buffer is 4 KiB, equivalent to 256 lines. The trace buffer is 128 bits wide, and stores the following information:

*   Instruction address and opcode
*   Instruction result
*   Load/store data and address
*   Trap information
*   30-bit time tag

The operation and control of the trace buffer is further described in section 9.4. Note that in GR712RC, each processor has its own trace buffer allowing simultaneous tracing of both instruction streams.

4.2.10 Processor configuration register

The application specific register 17 (%asr17) provides information on how various configuration options were set during synthesis. This can be used to enhance the performance of software, or to support enumeration in multi-processor systems. The register can be accessed through the RDASR instruction, and has the following layout:

	31	28		14	13	12	11 10	9	8	7	5	4	0
%asr17	INDEX		RESERVED	DW	SV	0	10	0	1	010		00111	

Figure 5. LEON3 configuration register (%asr17)

Field Definitions:

[31:28]: Processor index. In multi-processor systems, each LEON core gets a unique index to support enumeration. The value in this field is 0 for core0 and 1 for core1.
[14]: Disable write error trap (DWT). When set, a write error trap (tt = 0x2b) will be ignored. Set to zero after reset.
[13]: Single-vector trapping (SVT) enable. If set, will enable single-vector trapping. Set to zero after reset.
[11:10]: FPU option. Hardcoded to 1 (GRFPU)
[8]: If set, the SPARC V8 multiply and divide instructions are available. Hardcoded to 1.
[7:5]: Number of implemented watchpoints. Hardcoded to 2
[4:0]: Number of implemented registers windows corresponds to NWIN+1. Hardcoded to 7.

4.2.11 Exceptions

LEON3 adheres to the general SPARC V8 trap model. The table below shows the implemented traps and their individual priority. When PSR (processor status register) bit ET=0, an exception trap causes the processor to halt execution and enter error mode. If CPU0 enters error mode, then the external signal ERRORN is asserted.

Table 15. Trap allocation and priority

Trap	TT	Pri	Description
reset	0x00	1	Power-on reset
write error	0x2b	2	write buffer error
instruction_access_error	0x01	3	Error during instruction fetch
illegal_instruction	0x02	5	UNIMP or other un-implemented instruction
privileged_instruction	0x03	4	Execution of privileged instruction in user mode
fp_disabled	0x04	6	FP instruction while FPU disabled
cp_disabled	0x24	6	CP instruction while Co-processor disabled. The GR712RC does not implement a co-processor and CP instructions will trigger this trap
watchpoint_detected	0x0B	7	Hardware breakpoint match
window_overflow	0x05	8	SAVE into invalid window
window_underflow	0x06	8	RESTORE into invalid window
register_hadrware_error	0x20	9	Uncorrectable register file EDAC error
mem_address_not_aligned	0x07	10	Memory access to un-aligned address
fp_exception	0x08	11	FPU exception
data_access_exception	0x09	13	Access error during load or store instruction
tag_overflow	0x0A	14	Tagged arithmetic overflow
divide_exception	0x2A	15	Divide by zero
interrupt_level_1	0x11	31	Asynchronous interrupt 1
interrupt_level_2	0x12	30	Asynchronous interrupt 2
interrupt_level_3	0x13	29	Asynchronous interrupt 3
interrupt_level_4	0x14	28	Asynchronous interrupt 4
interrupt_level_5	0x15	27	Asynchronous interrupt 5
interrupt_level_6	0x16	26	Asynchronous interrupt 6
interrupt_level_7	0x17	25	Asynchronous interrupt 7
interrupt_level_8	0x18	24	Asynchronous interrupt 8
interrupt_level_9	0x19	23	Asynchronous interrupt 9
interrupt_level_10	0x1A	22	Asynchronous interrupt 10
interrupt_level_11	0x1B	21	Asynchronous interrupt 11
interrupt_level_12	0x1C	20	Asynchronous interrupt 12
interrupt_level_13	0x1D	19	Asynchronous interrupt 13
interrupt_level_14	0x1E	18	Asynchronous interrupt 14
interrupt_level_15	0x1F	17	Asynchronous interrupt 15
trap_instruction	0x80 - 0xFF	16	Software trap instruction (TA)

4.2.12 Single vector trapping (SVT)

Single-vector trapping (SVT) is an SPARC V8e option to reduce code size for embedded applications. When enabled, any taken trap will always jump to the reset trap handler (%tbr.tba + 0). The trap type

will be indicated in %tbr.tt, and must be decoded by the shared trap handler. SVT is enabled by setting bit 13 in %asr17.

4.2.13 Address space identifiers (ASI)

In addition to the address, a SPARC processor also generates an 8-bit address space identifier (ASI), providing up to 256 separate, 32-bit address spaces. During normal operation, the LEON3 processor accesses instructions and data using ASI 0x8 - 0xB as defined in the SPARC standard. Using the LDA/STA instructions, alternative address spaces can be accessed. The table shows the ASI usage for LEON3. Only ASI[5:0] are used for the mapping, ASI[7:6] have no influence on operation.

Table 16. ASI usage

ASI	Usage
0x01	Forced cache miss
0x02	System control registers (cache control register)
0x08, 0x09, 0x0A, 0x0B	Normal cached access (replace if cacheable)
0x0C	Instruction cache tags
0x0D	Instruction cache data
0x0E	Data cache tags
0x0F	Data cache data
0x10	Flush instruction and data caches
0x11	Flush data cache
0x13	Flush instruction and data caches
0x14	MMU diagnostic data cache context access
0x15	MMU diagnostic instruction cache context access
0x18	Flush MMU TLB and instruction and data caches
0x19	MMU registers
0x1C	MMU bypass
0x1D	MMU diagnostic access
0x1E	MMU snoop tags diagnostic access

For historical reasons there are multiple ASIs that flush the cache in different ways.

Writing to ASI 0x10 will flush the data cache and the instruction cache.

Writing to ASI 0x11 will flush the data cache only.

Writing to ASI 0x13 will flush the data cache and the instruction cache.

Writing to ASI 0x18, will flush both the MMU TLB, the I-cache, and the D-cache. This will block execution for a few cycles while the TLB is flushed and then continue asynchronously with the cache flushes continuing in the background.

4.2.14 Power-down

The processor can be configured to include a power-down feature to minimize power consumption during idle periods. The power-down mode is entered by performing a WRASR instruction to %asr19:

```
wr %g0, %asr19
```

During power-down, the pipeline is halted until the next interrupt occurs. The processor clock is gated, reducing power consumption from dynamic switching of logic and clock net.

4.2.15 Processor reset operation

The processor is reset by asserting the RESET input for at least 4 clock cycles. The following table indicates the reset values of the registers which are affected by the reset. All other registers maintain their value (or are undefined). Since PC is set to 0, execution start at address 0 in the PROM area.

Table 17. Processor reset values

Register	Reset value
PC (program counter)	0x0
nPC (next program counter)	0x4
PSR (processor status register)	ET=0, S=1

4.2.16 Multi-processor support

The GR712RC contains two LEON3 processor support symmetric multiprocessing (SMP) configuration. After reset, only the first processor will start while the second processor will remain halted in power-down mode. After the system has been initialized, the second processor can be started by writing to the 'MP status register', located in the multi-processor interrupt controller. The halted processors will start executing from address 0. Cache snooping should always be enabled in SMP systems to maintain data cache coherency between the processors.

4.2.17 Cache sub-system

The LEON3 processor implements a Harvard architecture with separate instruction and data buses, connected to two independent cache controllers. The cacheable memory areas are shown in the table below.

Table 18. Cacheable areas

Address range	Usage
0x00000000 – 0x20000000	External PROM
0x40000000 – 0x80000000	External SRAM/SDRAM

4.3 Instruction cache

4.3.1 Operation

The instruction cache is as a 4x4 KiB multi-way cache with 32 byte/line, using true LRU replacement policy. Each line has a cache tag associated with it consisting of a tag field and a valid field with one valid bit for each 4-byte sub-block. On an instruction cache miss to a cachable location, the instruction is fetched and the corresponding tag and data line updated.

If instruction burst fetch is enabled in the cache control register (CCR), the cache line is filled starting at the missed address and until the end of the line. At the same time, the instructions are forwarded to the IU (streaming). If the IU cannot accept the streamed instructions due to internal dependencies or multi-cycle instruction, the IU is halted until the line fill is completed. If the IU executes a control transfer instruction (branch/CALL/JMPL/RETT/TRAP) during the line fill, the line fill will be terminated on the next fetch. If instruction burst fetch is enabled, instruction streaming is enabled even when the cache is disabled. In this case, the fetched instructions are only forwarded to the IU and the cache is not updated. During cache line refill, incremental burst are generated on the AHB bus.

If a memory access error occurs during a line fill with the IU halted, the corresponding valid bit in the cache tag will not be set. If the IU later fetches an instruction from the failed address, a cache miss will occur, triggering a new access to the failed address. If the error remains, an instruction access error trap (tt=0x1) will be generated.

4.3.2 Instruction cache tag

A instruction cache tag entry consists of several fields as shown in figure 6:

Tag for 4 KiB way, 32 bytes/line

31	12	7	0
ATAG	0000	VALID	

Figure 6. Instruction cache tag layout examples

Field Definitions:

[31:12]: Address Tag (ATAG) - Contains the tag address of the cache line.

[7:0]: Valid (V) - When set, the corresponding sub-block of the cache line contains valid data. These bits are set when a sub-block is filled due to a successful cache miss; a cache fill which results in a memory error will leave the valid bit unset. A FLUSH instruction will clear all valid bits. V[0] corresponds to address 0 in the cache line, V[1] to address 1, V[2] to address 2 and so on.

4.4 Data cache

4.4.1 Operation

The data cache is implemented as a 4x4 KiB multi-way cache with 16 bytes/line, using true LRU replacement policy. Each line has a cache tag associated with it consisting of a tag field and valid field with one valid bit for each 4-byte sub-block. On a data cache read-miss to a cachable location 4 bytes of data are loaded into the cache from main memory. The write policy for stores is write-through with no-allocate on write-miss. Locked AHB transfers are generated for LDD, STD, LDST and SWAP instructions. If a memory access error occurs during a data load, the corresponding valid bit in the cache tag will not be set. and a data access error trap (tt=0x9) will be generated.

4.4.2 Write buffer

The write buffer (WRB) consists of three 32-bit registers used to temporarily hold store data until it is sent to the destination device. For half-word or byte stores, the stored data replicated into proper byte alignment for writing to a word-addressed device, before being loaded into one of the WRB registers. The WRB is emptied prior to a load-miss cache-fill sequence to avoid any stale data from being read in to the data cache.

Since the processor executes in parallel with the write buffer, a write error will not cause an exception to the store instruction. Depending on memory and cache activity, the write cycle may not occur until several clock cycles after the store instructions has completed. If a write error occurs, the currently executing instruction will take trap 0x2b.

Note: the 0x2b trap handler should flush the data cache, since a write hit would update the cache while the memory would keep the old value due the write error.

4.4.3 Data cache tag

A data cache tag entry consists of several fields as shown in figure 7:

Tag for 4 KiB way, 16 bytes/line

31	12	3	0
ATAG	000000	VALID	

Figure 7. Data cache tag layout

Field Definitions:

[31:12]:   Address Tag (ATAG) - Contains the address of the data held in the cache line.

[3:0]:   Valid (V) - When set, the corresponding sub-block of the cache line contains valid data. These bits are set when a sub-block is filled due to a successful cache miss; a cache fill which results in a memory error will leave the valid bit unset. V[0] corresponds to address 0 in the cache line, V[1] to address 1, V[2] to address 2 and V[3] to address 3.

4.5 Additional cache functionality

4.5.1 Cache flushing

Both instruction and data cache are flushed by executing the FLUSH instruction, or by writing to certain ASI address spaces. The instruction cache is also flushed by setting the FI bit in the cache control register and the data cache is also flushed by setting the FD bit in the cache control register.

Cache flushing takes one cycle per cache line, during which the IU will not be halted, but during which the caches are disabled. When the flush operation is completed, the cache will resume the state (disabled, enabled or frozen) indicated in the cache control register. Diagnostic access to the cache is not possible during a FLUSH operation and will cause a data exception (trap=0x09) if attempted.

4.5.2 Diagnostic cache access

Tags and data in the instruction and data cache can be accessed through ASI address space 0xC, 0xD, 0xE and 0xF by executing LDA and STA instructions. Address bits making up the cache offset will be used to index the tag to be accessed while the least significant bits of the bits making up the address tag will be used to index the cache set.

Diagnostic read of tags is possible by executing an LDA instruction with ASI=0xC for instruction cache tags and ASI=0xE for data cache tags. A cache line and set are indexed by the address bits making up the cache offset and the least significant bits of the address bits making up the address tag. Similarly, the data sub-blocks may be read by executing an LDA instruction with ASI=0xD for instruction cache data and ASI=0xF for data cache data. The sub-block to be read in the indexed cache line and set is selected by A[4:2].

The tags can be directly written by executing a STA instruction with ASI=0xC for the instruction cache tags and ASI=0xE for the data cache tags. The cache line and set are indexed by the address bits making up the cache offset and the least significant bits of the address bits making up the address tag. D[31:12] is written into the ATAG field (see figure 6 and figure 7) and the valid bits are written with the D[7:0] of the write data. The data sub-blocks can be directly written by executing a STA instruction with ASI=0xD for the instruction cache data and ASI=0xF for the data cache data. The sub-block to be read in the indexed cache line and set is selected by A[4:2].

In multi-way caches, the address of the tags and data of the ways are concatenated. The address of a tag or data is thus:

ADDRESS = WAY & LINE & DATA & "00"

Examples: the tag for line 2 in way 1 of the 4x4 KiB data cache with 16 byte line would be:

A[13:12]   = 1   (WAY)
A[11:5]   = 2   (TAG)
=> TAG ADDRESS = 0x1040

The data of this line would be at addresses 0x1040 - 0x104C

4.5.3 Data Cache snooping

To keep the data cache synchronized with external memory, cache snooping can be enabled by setting the DS bit in the cache control register. When enabled, the data cache monitors write accesses on the AHB bus to cacheable locations. If an other AHB master writes to a cacheable location which is currently cached in the data cache, the corresponding cache line is marked as invalid.

Note that the bus snooping reacts on accesses from other masters on the bus. The processor implementation in GR712RC does not snoop on its own accesses. If the MMU is enabled and there are mul-

tiple cached mappings to the same physical address that is being written to by the processor itself, then only the cache location corresponding to the virtual address that is being written will be updated. Cached locations belonging to aliased addresses will contain the old data.

4.5.4 Cache Control Register

The operation of the instruction and data caches is controlled through a common Cache Control Register (CCR) (figure 8). Each cache can be in one of three modes: disabled, enabled and frozen. If disabled, no cache operation is performed and load and store requests are passed directly to the memory controller. If enabled, the cache operates as described above. In the frozen state, the cache is accessed and kept in sync with the main memory as if it was enabled, but no new lines are allocated on read misses.

31	29 28	27 24	23	22	21	20 19	16	15	14	13 12	11 10	9 8	7 6	5	4	3 2	1 0
	PS	TB	DS	FD	FI	FT	IB	IP	DP	ITE	IDE	DTE	DDE	DF	IF	DCS	ICS

Figure 8. Cache control register

[28]: Parity Select [PS] - if set diagnostic read will return 4 check bits in the lsb bits, otherwise tag or data word is returned.

[27:24]: Test Bits [TB] - if set, check bits will be xored with test bits TB during diagnostic write

[23]: Data cache snoop enable [DS] - if set, will enable data cache snooping.

[22]: Flush data cache (FD). If set, will flush the data cache. Always reads as zero.

[21]: Flush Instruction cache (FI). If set, will flush the instruction cache. Always reads as zero.

[20:19]: FT scheme: "00" = no FT, "01" = 4-bit checking implemented

[16]: Instruction burst fetch (IB). This bit enables burst fill during instruction fetch.

[15]: Instruction cache flush pending (IP). This bit is set when an instruction cache flush operation is in progress.

[14]: Data cache flush pending (DP). This bit is set when an data cache flush operation is in progress.

[13:12]: Instruction Tag Errors (ITE) - Number of detected parity errors in the instruction tag cache.

[11:10]: Instruction Data Errors (IDE) - Number of detected parity errors in the instruction data cache.

[9:8]: Data Tag Errors (DTE) - Number of detected parity errors in the data tag cache.

[7:6]: Data Data Errors (DDE) - Number of detected parity errors in the data data cache.

[5]: Data Cache Freeze on Interrupt (DF) - If set, the data cache will automatically be frozen when an asynchronous interrupt is taken.

[4]: Instruction Cache Freeze on Interrupt (IF) - If set, the instruction cache will automatically be frozen when an asynchronous interrupt is taken.

[3:2]: Data Cache state (DCS) - Indicates the current data cache state according to the following: X0= disabled, 01 = frozen, 11 = enabled.

[1:0]: Instruction Cache state (ICS) - Indicates the current data cache state according to the following: X0= disabled, 01 = frozen, 11 = enabled.

If the DF or IF bit is set, the corresponding cache will be frozen when an asynchronous interrupt is taken. This can be beneficial in real-time system to allow a more accurate calculation of worst-case execution time for a code segment. The execution of the interrupt handler will not evict any cache lines and when control is returned to the interrupted task, the cache state is identical to what it was before the interrupt. If a cache has been frozen by an interrupt, it can only be enabled again by enabling it in the CCR. This is typically done at the end of the interrupt handler before control is returned to the interrupted task.

4.5.5 Cache configuration registers

The configuration of the two caches if defined in two registers: the instruction and data configuration registers. These registers are read-only and indicate the size and configuration of the caches.

31	30 29	28 27	26 25 24	23 20	19	18 16	15 12	11 4	3	0
CL		REPL	SN SETS	SSIZE	LR	LSIZE	LRSIZE	LRSTART	M	

Figure 9. Cache configuration register

[31]: Cache locking (CL). Set if cache locking is implemented.

[29:28]: Cache replacement policy (REPL). 00 - no replacement policy (direct-mapped cache), 01 - least recently used (LRU), 10 - least recently replaced (LRR), 11 - random

[27]: Cache snooping (SN). Set if snooping is implemented.

[26:24]: Cache associativity (SETS). Number of sets in the cache: 000 - direct mapped, 001 - 2-way associative, 010 - 3-way associative, 011 - 4-way associative

[23:20]: Set size (SSIZE). Indicates the size (KiB) of each cache set. Size = 2^{SIZE}

[19]: Local ram (LR). Set if local scratch pad ram is implemented.

[18:16]: Line size (LSIZE). Indicated the size (words) of each cache line. Line size = 2^{LSZ}

[15:12]: Local ram size (LRSZ). Indicates the size (KiB) of the implemented local scratch pad ram. Local ram size = 2^{LRSZ}

[11:4]: Local ram start address. Indicates the 8 most significant bits of the local ram start address.

[3]: MMU present. This bit is set to '1' if an MMU is present.

The value of the instruction cache configuration register is 0x132308e8 and the value of the data cache configuration register is 0x1b2208f8.

4.5.6 Cache registers accesses

All cache registers are accessed through load/store operations to the alternate address space (LDA/STA), using ASI = 2. The table below shows the register addresses:

Table 19. ASI 2 (system registers) address map

Address	Register
0x00	Cache control register
0x04	Reserved
0x08	Instruction cache configuration register
0x0C	Data cache configuration register

4.5.7 Software consideration

After reset, the caches are disabled and the PS, TB, DS, IB, DCS, ICS fields of the cache control register (CCR) are 0. Other writable fields in the cache control register are not affected by reset. Before the caches may be enabled, a flush operation must be performed to initialized (clear) the tags and valid bits. A suitable assembly sequence could be:

```
flush
set 0x81000f, %g1
sta%g1, [%g0] 2
```

4.6 Memory management unit

A memory management unit (MMU) compatible with the SPARC V8 reference MMU is provided in the cache control module. For details on operation, see the SPARC V8 manual.

4.6.1 ASI mappings

See table 16 for MMU ASI mappings.

4.6.2 MMU/Cache operation

When the MMU is disabled, the caches operate as normal with physical address mapping. When the MMU is enabled, the caches tags store the virtual address and also include an 8-bit context field. On cache miss or access to an uncacheable location, the virtual address is translated to a physical address before the access appears on the AHB bus.

4.6.3 MMU registers

The following MMU registers are implemented:

Table 20. MMU registers (ASI = 0x19)

Address	Register
0x000	MMU control register
0x100	Context pointer register
0x200	Context register
0x300	Fault status register
0x400	Fault address register

The definition of the registers can be found in the SPARC V8 manual.

31	28 27	24 23	21 20	18 17 16 15 14		2 1 0
IMPL	VER	ITLB	DTLB	ST - TD	RESERVED	NF E

Figure 10. MMU control register

[31:28]: MMU Implementation ID. Hardcoded to "0000".

[27:24]: MMU Version ID. Hardcoded to "0000".

[23:21]: Number of ITLB entries. The number of ITLB entries is calculated as 2^{ITLB}.

[20:18]: Number of DTLB entries. The number of DTLB entries is calculated as 2^{DTLB}.

[17]: Separate TLB. Set to 1 to indicate that separate instructions and data TLBs are implemented.

[15]: TLB disable. When set to 1, the TLB will be disabled and each data access will generate an MMU page table walk.

[1]: No Fault. When NF= 0, any fault detected by the MMU causes FSR and FAR to be updated and causes a fault to be generated to the processor. When NF= 1, a fault on an access to ASI 9 is handled as when NF= 0; a fault on an access to any other ASI causes FSR and FAR to be updated but no fault is generated to the processor.

[0]: Enable MMU. 0 = MMU disabled, 1 = MMU enabled.

4.6.4 Translation look-aside buffer (TLB)

The MMU uses a separate 16-entry TLB for instructions and data accesses. The organisation of the TLB and number of entries is not visible to the software and does thus not require any modification to the operating system.

4.6.5 Snoop tag diagnostic access

To implement cache snooping, the MMU contains a separate snoop tag for each virtual tag. The tags can be accessed via ASI 0x1E. This is primarily useful for RAM testing, and should not be performed during normal operation. The figure below shows the layout of the snoop tag for the 4 KiB way in data cache:

31	12 11	2 1 0
ATAG	"0000000000"	PAR HIT

Figure 11. Snoop cache tag layout

[31:12] Address tag. The physical address tag of the cache line.

[1]: Parity. The odd parity over the snoop tag (ATAG field.

[0]: Snoop hit. When set, the cache line is not valid and will cause a cache miss if accessed by the processor.

4.7 Floating-point unit

Each of the LEON3 processor in GR712RC contains a high-performance GRFPU floating-point unit. The GRFPU operates on single- and double-precision operands, and implements all SPARC V8 FPU instructions. The FPU is interfaced to the LEON3 pipeline using a LEON3-specific FPU controller (GRFPC) that allows FPU instructions to be executed simultaneously with integer instructions. Only in case of a data or resource dependency is the integer pipeline held. The GRFPU is fully pipelined and allows the start of one instruction each clock cycle, with the exception is FDIV and FSQRT which can only be executed one at a time. The FDIV and FSQRT are however executed in a separate divide unit and do not block the FPU from performing all other operations in parallel.

All instructions except FDIV and FSQRT has a latency of three cycles, but to improve timing, the LEON3 FPU controller inserts an extra pipeline stage in the result forwarding path. This results in a latency of four clock cycles at instruction level. The table below shows the GRFPU/GRFPC instruction timing:

Table 21. GRFPU instruction timing with GRFPC

Instruction	Throughput	Latency
FADDS, FADDD, FSUBS, FSUBD,FMULS, FMULD, FSMULD, FITOS, FITOD, FSTOI, FDTOI, FSTOD, FDTOS, FCMPS, FCMPD, FCMPES. FCMPED	1	4
FDIVS	14	16
FDIVD	15	17
FSQRTS	22	24
FSQRTD	23	25

The GRFPC controller implements the SPARC deferred trap model, and the FPU trap queue (FQ) can contain up to 8 queued instructions when an FPU exception is taken. The version field in %fsr has the value of 2 to indicate the presence of the GRFPU

4.8 Error detection and correction

4.8.1 Register file error detection and correction

The data in the integer register file is protected using BCH coding, capable of correcting one error and detecting two errors (SEC/DED). The error detection has no impact on normal operation, but a correction cycle will delay the current instruction with 6 clock cycles. An uncorrectable error in the IU register file will cause trap 0x20 (*register_access_error*).

For test purposes, the IU register file checkbits can be modified by software. This is done by setting the ITE bit to '1'. In this mode, the checkbits are first XORed with the contents of %asr16.TB before written to the register file.

4.8.2 ASR16 register

ASR register 16 (%asr16) is used to control the IU register file error protection. It is possible to disable the error protection by setting the IDI bit, and to inject errors using the ITE bits. Corrected errors in the register file are counted, and available in ICNT field. The counter saturate at their maximum value (7), and should be reset by software after read-out.

31	15	14 13	11	10	3	2	1	0
RESERVED	IUFT	ICNT		TB[7:0]		DP	ITE	IDI

Figure 12. %asr16 - Register protection control register

[15:14]: IU FT ID - Defines which error protection is implemented. Hardcoded to "11".

[13:11]: IU RF error counter - Number of detected errors in the IU register file.

[10:3]: RF Test bits (FTB) - In test mode, these bits are xored with correct checkbits before written to the register file.

[2]: DP ram select (DP) - See table below for details.

[1]: IU RF Test Enable - Enables register file test mode. Checkbits are xored with TB before written to the register file.

[0]: IU RF protection disable (IDI) - Disables IU register file protection when set.

Table 22. DP ram select usage

ITE	DP	Function
1	0	Write to IU register (%i, %l, %o, %g) will only write location of %rs1
1	1	Write to IU register (%i, %l, %o, %g) will only write location of %rs2
0	X	IU registers written nominally

4.8.3 Register file EDAC/parity bits diagnostic read-out

The register file checkbits can be read out through the DSU address space at 0x90300800. The check-bits are read out for both read ports simultaneously as defined in the figure below:

31	16	8	7	0
RESERVED		RF ECC Port 2	RF ECC port 1	

Figure 13. Register file checkbits read-out layout

4.8.4 Cache error protection

Each word in the cache tag and data memories is protected by four check bits. An error during cache access will cause a cache line flush, and a re-execution of the failing instruction. This will insure that the complete cache line (tags and data) is refilled from external memory. For every detected error, a counter in the cache control register is incremented. The counters saturate at their maximum value (3), and should be reset by software after read-out.

The context and parity bits for data and instruction caches can be read out via ASI 0xC - 0xF when the PS bit in the cache control register is set. The data will be organized as shown below:

	31	23	16 15	4 3	0
ASI = 0xC ASI = 0xE		MMU CTX [7:0]		TAG PAR[3:0]	
ASI = 0xD ASI = 0xF				DATA PAR[3:0]	

Figure 14. Data cache tag diagnostic access when CCR.PS = '1'

4.8.5 Data scrubbing

There is generally no need to perform data scrubbing on either IU register file or the cache memory. During normal operation, the active part of the IU register files will be flushed to memory on each task switch. This will cause all registers to be checked and corrected if necessary. Since most real-time operating systems performs several task switches per second, the active data in the register file will be frequently refreshed.

The similar situation arises for the cache memory. In most applications, the cache memory is significantly smaller than the full application image, and the cache contents is gradually replaced as part of normal operation. For very small programs, the only risk of error build-up is if a part of the applica-

tion is resident in the cache but not executed for a long period of time. In such cases, executing a cache flush instruction periodically (e.g. once per minute) is sufficient to refresh the cache contents.

4.8.6 Initialisation

After power-on, the check bits in the IU register file are not initialized. This means that access to an uninitialized (un-written) register could cause a register access trap (tt = 0x20). Such behaviour is considered as a software error, as the software should not read a register before it has been written. It is recommended that the boot code for the processor writes all registers in the IU register files before launching the main application.

The check bits in the cache memories do not need to be initialized as this is done automatically during cache line filling. However, a cache flush must be executed before the caches are enabled.

4.9 Signal definitions

When CPU0 enters error mode, the *ERRORN* output is driven active low, else it is in tri-state and therefore requires an additional external pull-up.

The signals are described in table 23.

Table 23. Signal definitions

Signal name	Type	Function	Active
ERRORN	Tri-state output	Processor error mode indicator	Low

5   Fault Tolerant Memory Controller

5.1   Overview

The combined 8/32-bit memory controller provides a bridge between external memory and the AHB bus. The memory controller can handle four types of devices: PROM, asynchronous static RAM (SRAM), synchronous dynamic RAM (SDRAM) and memory mapped I/O devices (IO). The PROM, SRAM and SDRAM areas can be EDAC-protected using a (39,7) BCH code. The BCH code provides single-error correction and double-error detection for each 32-bit memory word.

The SDRAM area can optionally also be protected using Reed-Solomon coding. In this case a 16-bit checksum is used for each 32-bit word, and any two adjacent 4-bit (nibble) errors can be corrected.

The memory controller is configured through three configuration registers accessible via an APB bus interface. The external data bus can be configured in 8-, or 32-bit mode, depending on application requirements. The controller decodes three address spaces on the AHB bus (PROM, IO, and SRAM/ SDRAM). The IO area is marked as non-cacheable in the core's AMBA plug'n'play information record.

External chip-selects are provided for two PROM banks, one IO bank, two SRAM banks and two SDRAM banks. Refer to section 1.7.7 for information about configuration when enabling EDAC in 8-bit bus mode.

The maximum PROM capacity supported is 32 MiB + 25% for EDAC checksum over 2 banks.

The maximum SRAM capacity supported is 32 MiB + 25% for EDAC checksum over 2 banks.

The maximum SDRAM capacity supported is 1 GiB + 25% or 50% for EDAC checksum over 2 banks when not using SRAM. When also using SRAM the maximum amount of supported SDRAM is 512 MiB + 25% or 50% for EDAC checksum over 1 or 2 banks.

Figure 15 below shows how the connection to the different device types is made.
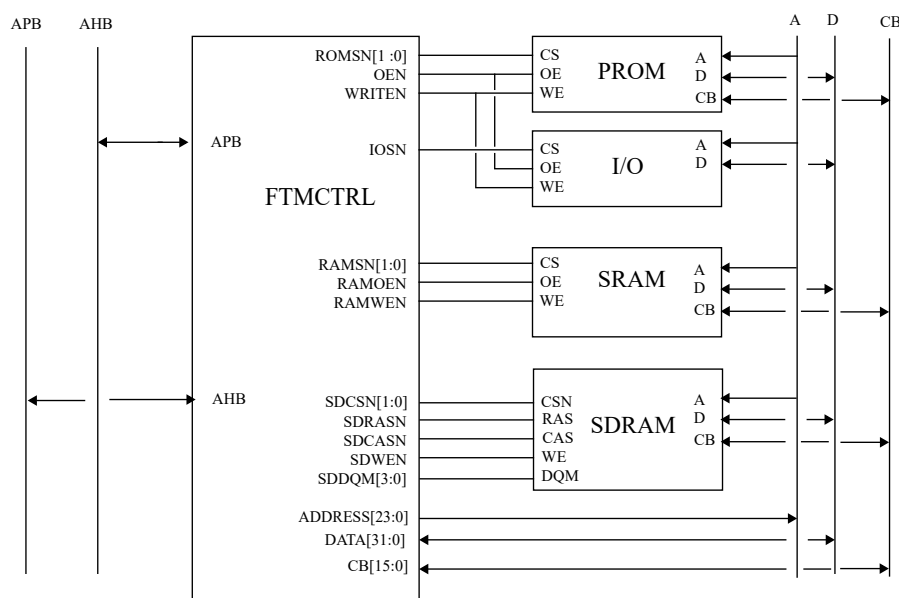
Figure 15. FTMCTRL connected to different types of memory devices

5.2   PROM access

Up to two PROM chip-select signals are provided for the PROM area, ROMSN[1:0]. The size of the banks can be set in binary steps from 16KiB to 16MiB. The total maximum supported PROM capacity is 32 MiB.

A read access to PROM consists of two data cycles and between 0 and 30 waitstates. The read data (and optional EDAC check-bits) are latched on the rising edge of the clock on the last data cycle. On non-consecutive accesses, a lead-out cycle is added after a read cycle to prevent bus contention due to slow turn-off time of PROM devices. Figure 16 shows the basic read cycle waveform (zero waitstate) for non-consecutive PROM reads. Note that the address is undefined in the lead-out cycle. Figure 17 shows the timing for consecutive cycles (zero waitstate). Waitstates are added by extending the data2 phase. This is shown in figure 18 and applies to both consecutive and non-consecutive cycles. Only an even number of waitstates can be assigned to the PROM area.

The PROM area will wrap back to the first bank after the end of the last decoded bank. As an example, if the ROMBANKSZ field is set to 13 the following banks will be decoded:

```
bank 0: 0x00000000 - 0x03FFFFFF
bank 1: 0x04000000 - 0x07FFFFFF
bank 2: 0x08000000 - 0x0BFFFFFF (no external select signal provided)
bank 3: 0x0C000000 - 0x0FFFFFFF (no external select signal provided)
...
```

bank 0 starting again at 0x10000000 (the same pattern applies for other values less than 13, addresses will wrap after the last decoded bank).
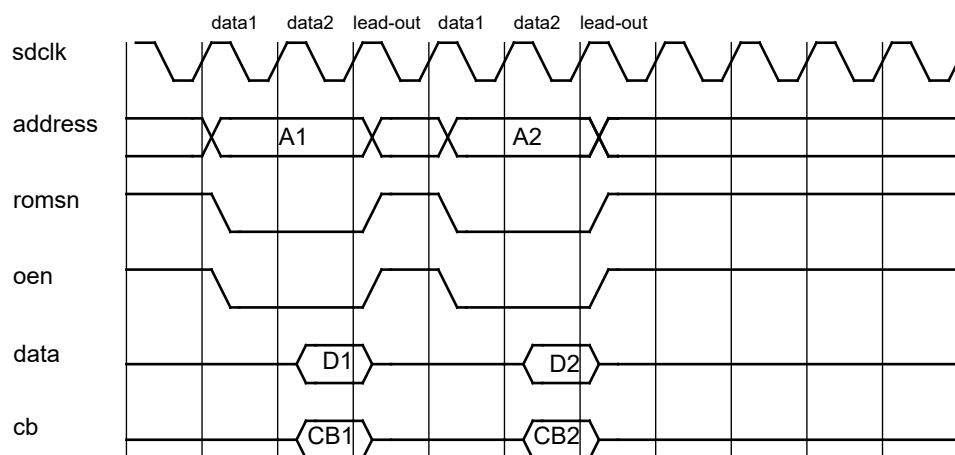
Figure 16. Prom non-consecutive read cyclecs.
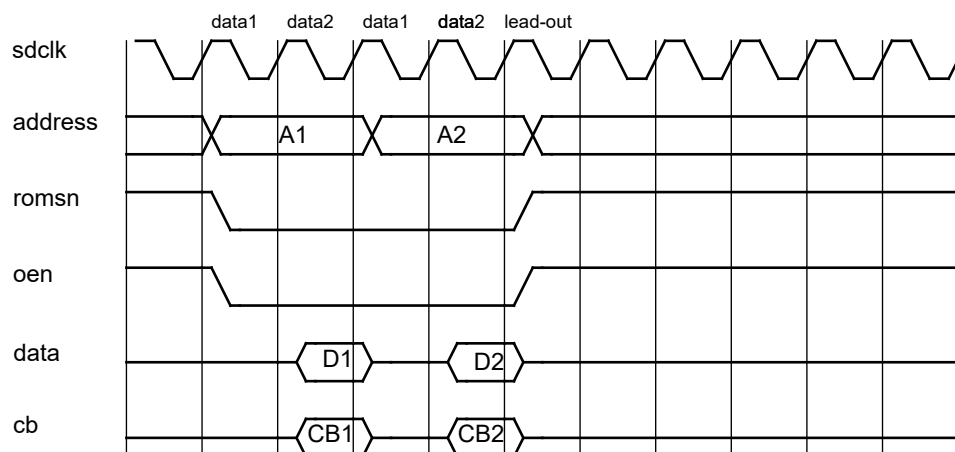
Figure 17. Prom consecutive read cyclecs.
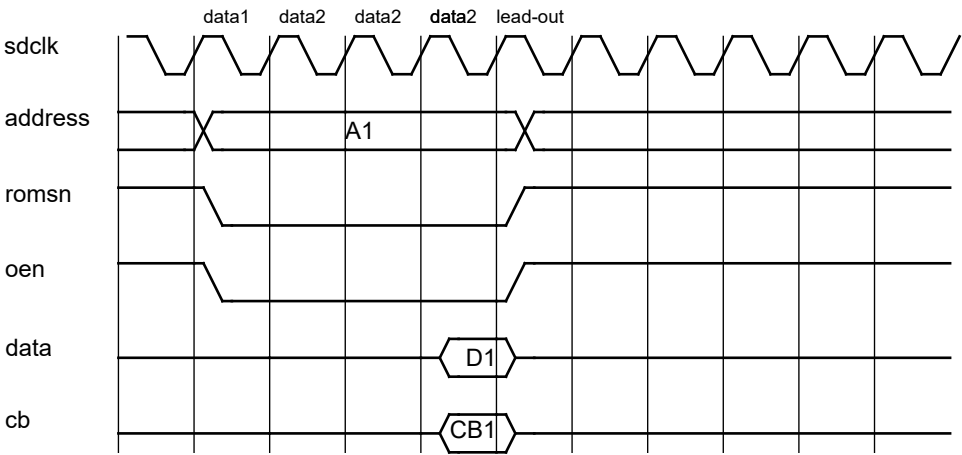
Figure 18. Prom read access with two waitstates.
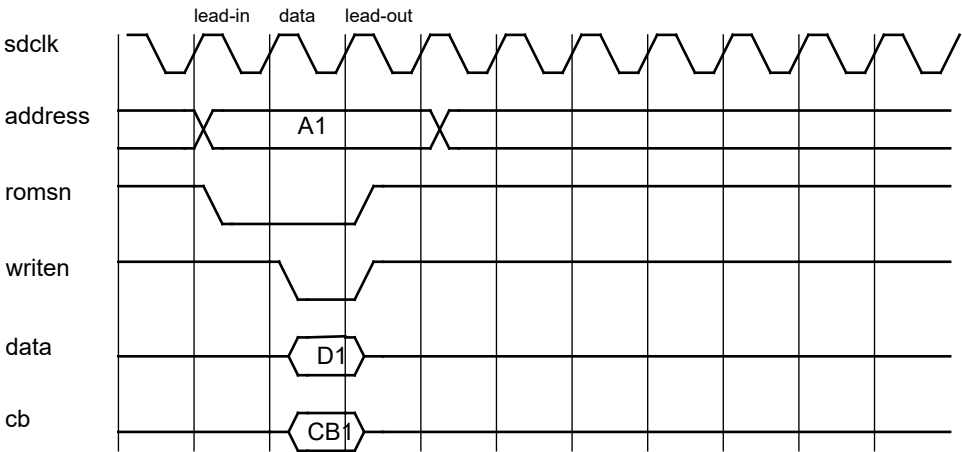
Figure 19. Prom write cycle (0-waitstates)
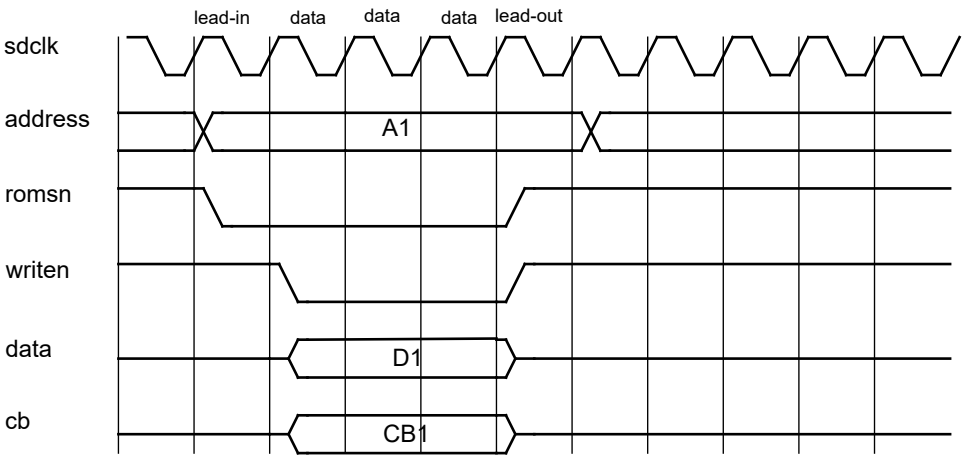
Figure 20. Prom write cycle (2-waitstates)

5.3 Memory mapped IO

Accesses to IO have similar timing as PROM accesses. The IO select (IOSN) and output enable (OEN) signals are delayed one clock to provide stable address before IOSN is asserted. All accesses are performed as non-consecutive accesses as shown in figure 21. The data2 phase is extended when waitstates are added.
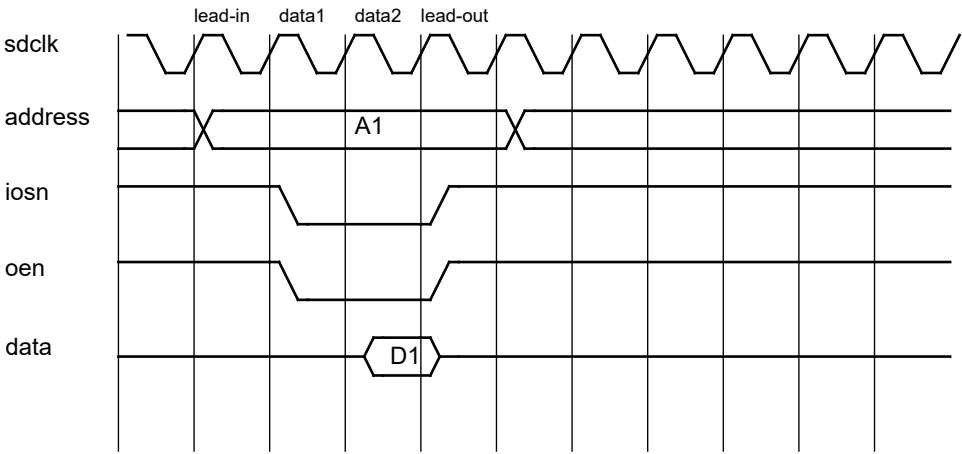
Figure 21. I/O read cycle (0-waitstates)
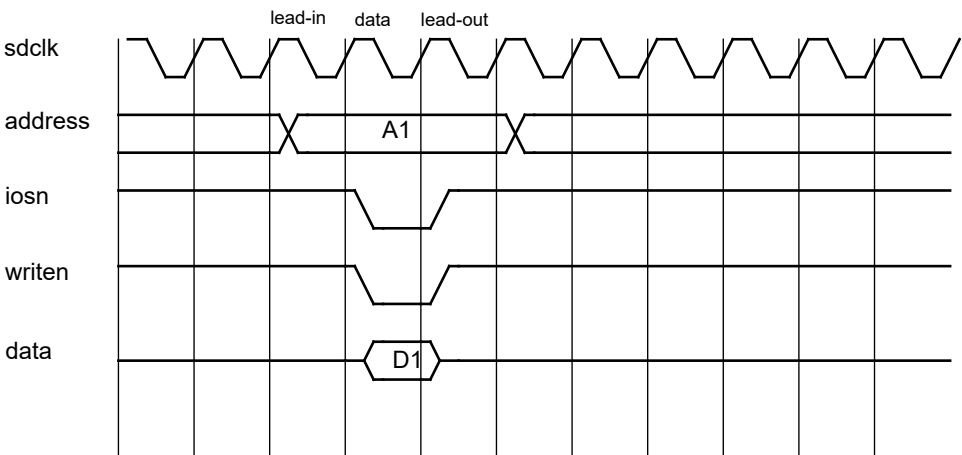
Figure 22. I/O write cycle (0-waitstates)

5.4   SRAM access

The SRAM area is divided on two RAM banks. The size of the banks (RAMSN[1:0]) is programmed in the RAM bank-size field (MCFG2[12:9]) and can be set in binary steps from 8KiB to 16MiB. The total maximum supported SRAM capacity is 32 MiB. A read access to SRAM consists of two data cycles and between zero and three waitstates. The read data (and optional EDAC check-bits) are latched on the rising edge of the clock on the last data cycle. On non-consecutive accesses, a lead-out cycle is added after a read cycle to prevent bus contention due to slow turn-off time of memories. Figure 23 shows the basic read cycle waveform (zero waitstate). Waitstates are added in the same way as for PROM in figure 18.
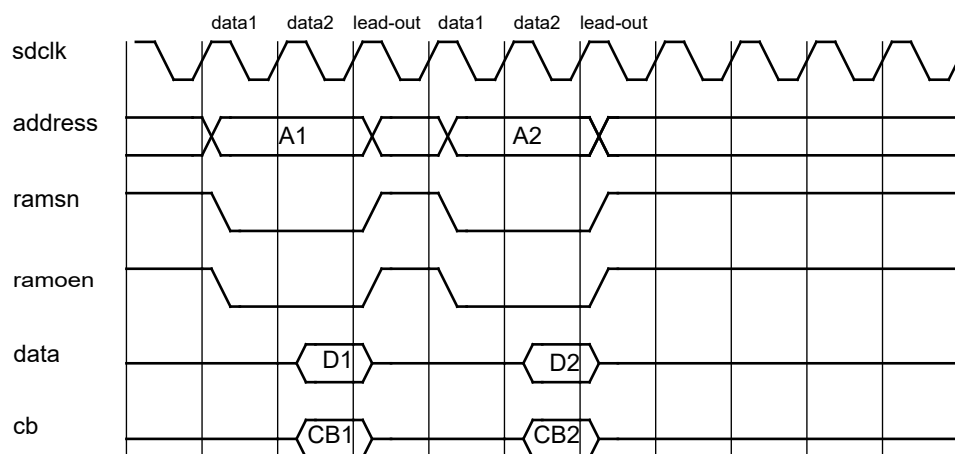
Figure 23.   Sram non-consecutive read cyclecs.

For read accesses to RAMSN[1:0], a shared output enable signal (RAMOEN) is provided. A write access is similar to the read access but takes a minimum of three cycles. Waitstates are added in the same way as for PROM.

The GR712RC uses a common SRAM write strobe (RAMWEN), and the read-modify-write bit MCFG2 must be set to enable read-modify-write cycles for sub-word writes.
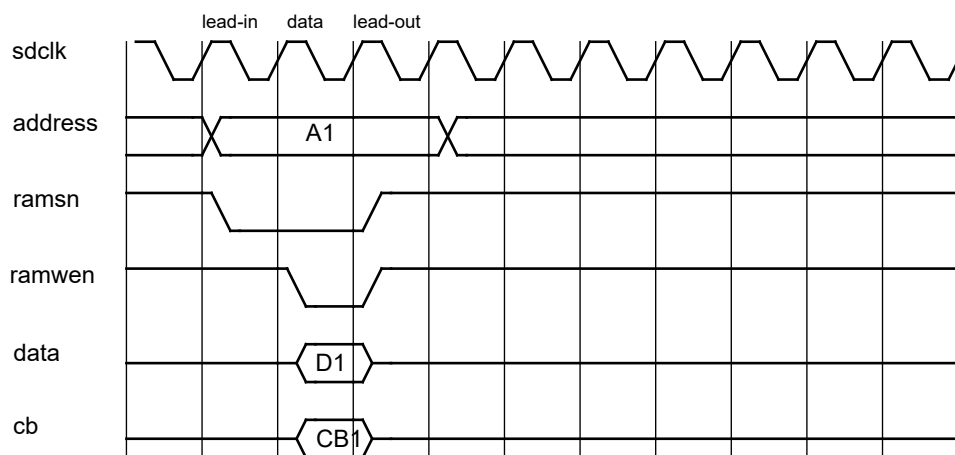
Figure 24.   Sram write cycle (0-waitstates)
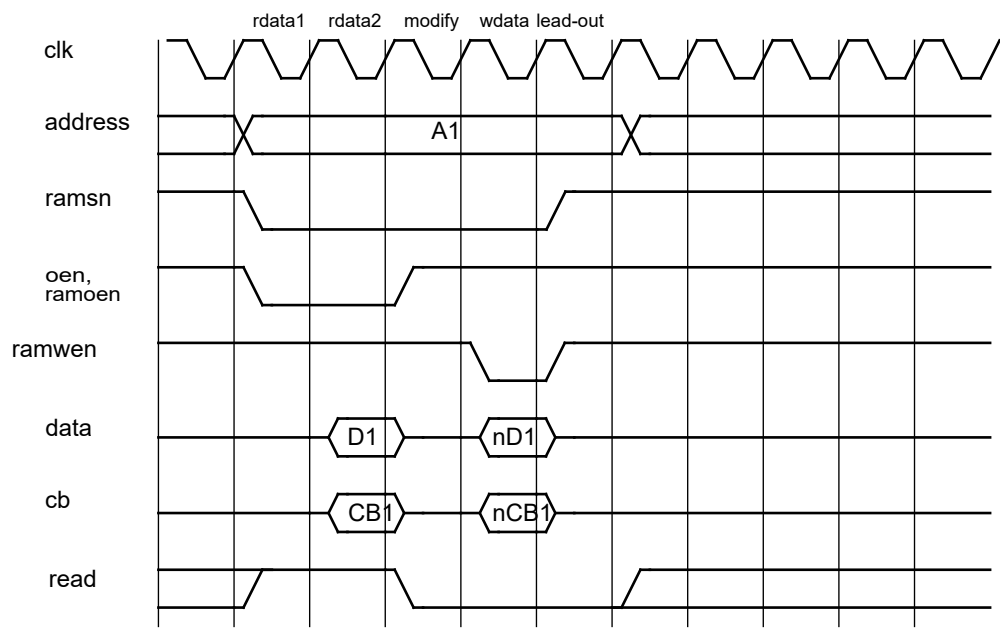
Figure 25. Sram read-modify-write cycle (0-waitstates)

5.5   8-bit PROM and SRAM access

To support applications with low memory and performance requirements efficiently, the SRAM and PROM areas can be individually configured for 8-bit operation by programming the ROM and RAM width fields in the memory configuration registers. Since reads to memory are always done on 32-bit word basis, read access to 8-bit memory will be transformed in a burst of four read cycles. During writes, only the necessary bytes will be written. Figure 26 shows an interface example with 8-bit PROM and 8-bit SRAM.

The supported combinations of width, EDAC and RMW for the PROM and SRAM areas are given in table 24.

Table 24. Supported SRAM and PROM width, EDAC and RMW combinations

Bus width	EDAC	RMW bit (SRAM only)
8	None	0
8	BCH	1
32	None	1
32+7	BCH	1

The RMW bit in MCFG2 must **not** be set if RAM EDAC is dissembled when RAM width is set to 8-bit.
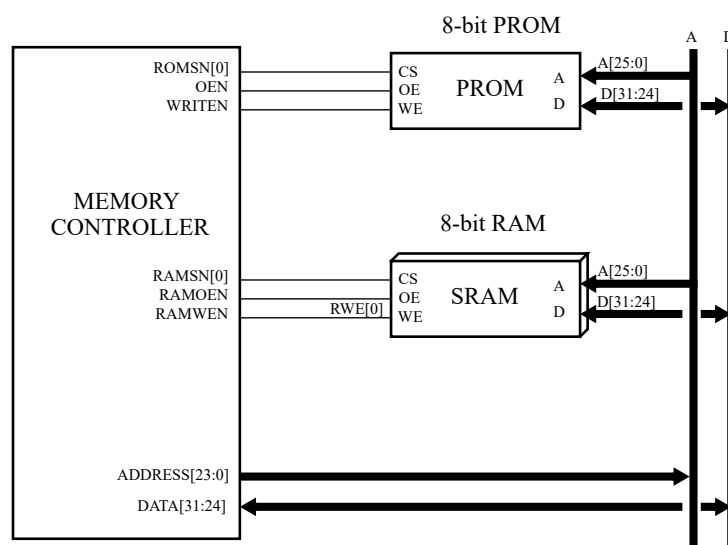
Figure 26. 8-bit memory interface example

In 8-bit mode, the PROM/SRAM devices should be connected to the most significant byte of the data bus (DATA[31:24]). The least significant part of the address bus should be used for addressing (ADDRESS[23:0]).

5.6   8- bit I/O access

Similar to the PROM/SRAM areas, the IO area can also be configured to 8-bits mode. However, the I/O device will NOT be accessed by multiple 8/16 bits accesses as the memory areas, but only with one single access just as in 32-bit mode. To access an IO device on an 8-bit bus, only byte accesses should be used (LDUB/STB instructions for the CPU).

5.7   Burst cycles

To improve the bandwidth of the memory bus, accesses to consecutive addresses can be performed in burst mode. Burst transfers will be generated when the memory controller is accessed using an AHB burst request. These includes instruction cache-line fills, double loads and double stores. The timing of a burst cycle is identical to the programmed basic cycle with the exception that during read cycles, the lead-out cycle will only occurs after the last transfer. Burst cycles will not be generated to the IO area.

5.8   SDRAM access

5.8.1   General

Synchronous dynamic RAM (SDRAM) access is supported to two banks of PC100 compatible devices. The SDRAM controller supports 64Mibit, 256Mibit and 512Mibit devices with 8 - 12 column-address bits, and up to 13 row-address bits. The size of the two banks can be programmed in binary steps between 4MiB and 512MiB. The total maximum supported SDRAM capacity is 1 GiB. The operation of the SDRAM controller is controlled through MCFG2 and MCFG3 (see below).

The supported combinations of width and EDAC for the SDRAM is given in table 25.

Table 25. Supported SDRAM width and EDAC combinations

SDRAM bus width	EDAC
32	None
32+7	BCH
32+16	RS

5.8.2 Address mapping

The two SDRAM chip-select signals are decoded. The SDRAM is mapped in address range 0x60000000 - 0x80000000 when the SRAM is enabled. When the SRAM disable bit is set, all access to SRAM is disabled and the SDRAM banks are mapped in address range 0x40000000 - 0x80000000.

5.8.3 Initialisation

When the SDRAM controller is enabled, it automatically performs the SDRAM initialisation sequence of PRECHARGE, 8x AUTO-REFRESH and LOAD-MODE-REG on both banks simultaneously. The controller programs the SDRAM mode register to use single location access on write and length 8 line burst on read.

5.8.4 Configurable SDRAM timing parameters

To provide optimum access cycles for different SDRAM devices (and at different frequencies), three SDRAM parameters can be programmed through memory configuration register 2 (MCFG2): TCAS, TRP and TRFCD. The value of these field affects the SDRAM timing as described in table 26.

Table 26. SDRAM programmable minimum timing parameters

SDRAM timing parameter	Minimum timing (clocks)
CAS latency, RAS/CAS delay (t_{CAS}, t_{RCD})	TCAS + 2
Precharge to activate (t_{RP})	TRP + 2
Auto-refresh command period (t_{RFC})	TRFC + 3
Activate to precharge (t_{RAS})	TRFC + 1
Activate to Activate (t_{RC})	TRP + TRFC + 4

If the TCAS, TRP and TRFC are programmed such that the PC100/133 specifications are fulfilled, the remaining SDRAM timing parameters will also be met. The table below shows typical settings for 100 and 133 MHz operation and the resulting SDRAM timing (in ns).

It should be noted that the maximum SDRAM frequency of 100 and 133 MHz is not possible with GR712RC.

Table 27. SDRAM example programming (clock periods)

SDRAM settings	t_{CAS}	t_{RC}	t_{RP}	t_{RFC}	t_{RAS}
CL=2; TRP=0, TCAS=0, TRFC=4	2	8	2	7	5
CL=3; TRP=0, TCAS=1, TRFC=4	3	8	2	7	5

5.9 Refresh

The SDRAM controller contains a refresh function that periodically issues an AUTO-REFRESH command to both SDRAM banks. The period between the commands (in clock periods) is pro-

grammed in the refresh counter reload field in the MCFG3 register. Depending on SDRAM type, the required period is typically 7.8 or 15.6 μs (corresponding to 780 or 1560 clocks at 100 MHz). The generated refresh period is calculated as (reload value+1)/sysclk. The refresh function is enabled by setting bit 31 in MCFG2.

5.9.1 SDRAM commands

The controller can issue three SDRAM commands by writing to the SDRAM command field in MCFG2: PRE-CHARGE, AUTO-REFRESH and LOAD-MODE-REG (LMR). If the LMR command is issued, the CAS delay as programmed in MCFG2 will be used. Remaining fields are fixed: 8-word sequential read burst, single location write. The command field will be cleared after a command has been executed. When changing the value of the CAS delay, a LOAD-MODE-REGISTER command should be generated at the same time. NOTE: when issuing SDRAM commands, the SDRAM refresh must be disabled.

5.9.2 Read and write cycles

A read cycle is started by performing an ACTIVATE command to the desired bank and row, followed by a READ command after the programmed CAS delay. A read burst is performed if a burst access has been requested on the AHB bus. The read cycle is terminated with a PRE-CHARGE command, no banks are left open between two accesses.

Write cycles are performed similarly to read cycles, with the difference that WRITE commands are issued after activation. A write burst on the AHB bus will generate a burst of write commands without idle cycles in-between.

5.9.3 Read-modify-write cycles

If EDAC is enabled and a byte or half-word write is performed, the controller will perform a read-modify-write cycle to update the checkbits correctly. This is done by performing an ACTIVATE command, followed by READ, WRITE and PRE-CHARGE. The write command interrupts the read burst and the data mask signals will be raised two cycles before this happens as required by the SDRAM standard.

5.9.4 Address bus

The address bus of the SDRAM devices should be connected to ADDRESS[14:2], the bank address to ADDRESS[16:15]. The MSB part of ADDRESS[14:2] can be left unconnected if not used.

5.9.5 Initialisation

Each time the SDRAM is enabled (bit 14 in MCFG2), an SDRAM initialisation sequence will be sent to both SDRAM banks. The sequence consists of one PRECHARGE, eight AUTO-REFRESH and one LOAD-COMMAND-REGISTER command.

5.10 Memory EDAC

The memory controller contains two separate EDACs, a BCH-based for all memory types and a Reed-Solomon based for SDRAM only.

SDRAM is possible to use on both banks in 32-bit mode with EDAC.

Both SRAM and PROM are possible to use on both banks in 32-bit mode with or without EDAC.

Both SRAM and PROM are possible to use on both banks in 8-bit mode without EDAC.

When multiple SRAM or PROM memory banks are used in 8-bit mode with EDAC, the configured bank size must be set to at least four times the actual device size. See section 1.7.7 for details.

5.10.1 BCH EDAC

The FTMCTRL is provided with an BCH EDAC that can correct one error and detect two errors in a 32-bit word. For each word, a 7-bit checksum is generated according to the equations below. A correctable error will be handled transparently by the memory controller, but adding one waitstate to the access. If an un-correctable error (double-error) is detected, the current AHB cycle will end with an error response. The EDAC can be used during access to PROM, SRAM and SDRAM areas by setting the corresponding EDAC enable bits in the MCFG3 register. The equations below show how the EDAC checkbits are generated (where Dn corresponds to DATA[n]:

```
CB[0]=D0 ^ D4 ^ D6 ^ D7 ^ D8 ^ D9 ^ D11 ^ D14 ^ D17 ^ D18 ^ D19 ^ D21 ^ D26 ^ D28 ^ D29 ^ D31
CB[1]=D0 ^ D1 ^ D2 ^ D4 ^ D6 ^ D8 ^ D10 ^ D12 ^ D16 ^ D17 ^ D18 ^ D20 ^ D22 ^ D24 ^ D26 ^ D28
CB[2]=D0 ^ D3 ^ D4 ^ D7 ^ D9 ^ D10 ^ D13 ^ D15 ^ D16 ^ D19 ^ D20 ^ D23 ^ D25 ^ D26 ^ D29 ^ D31
CB[3]=D0 ^ D1 ^ D5 ^ D6 ^ D7 ^ D11 ^ D12 ^ D13 ^ D16 ^ D17 ^ D21 ^ D22 ^ D23 ^ D27 ^ D28 ^ D29
CB[4]=D2 ^ D3 ^ D4 ^ D5 ^ D6 ^ D7 ^ D14 ^ D15 ^ D18 ^ D19 ^ D20 ^ D21 ^ D22 ^ D23 ^ D30 ^ D31
CB[5]=D8 ^ D9 ^ D10 ^ D11 ^ D12 ^ D13 ^ D14 ^ D15 ^ D24 ^ D25 ^ D26 ^ D27 ^ D28 ^ D29 ^ D30^ D31
CB[6]=D0 ^ D1 ^ D2 ^ D3 ^ D4 ^ D5 ^ D6 ^ D7 ^ D24 ^ D25 ^ D26 ^ D27 ^ D28 ^ D29 ^ D30 ^ D31
```

If the SRAM is configured in 8-bit mode, the EDAC checkbit bus (CB[7:0]) is not used but it is still possible to use EDAC protection. Data is always accessed as words (4 bytes at a time) and the corresponding checkbits are located at the address acquired by inverting the word address (bits 2 to 27) and using it as a byte address. The same chip-select is kept active. A word written as four bytes to addresses 0, 1, 2, 3 will have its checkbits at address 0xFFFFFFFF, addresses 4, 5, 6, 7 at 0xFFFFFFFE and so on. All the bits up to the maximum bank size will be inverted while the same chip-select is always asserted. This way all the bank sizes can be supported and no memory will be unused (except for a maximum of 4 byte in the gap between the data and checkbit area). A read access will automatically read the four data bytes individually from the nominal addresses and the EDAC checkbit byte from the top part of the bank. A write cycle is performed the same way. Byte or half-word write accesses will result in an automatic read-modify-write access where 4 data bytes and the checkbit byte are firstly read, and then 4 data bytes and the newly calculated checkbit byte are writen back to the memory. This 8-bit mode applies to SRAM while SDRAM always uses 32-bit accesses. The size of the memory bank is determined from the settings in MCFG2.

If the ROM is configured in 8-bit mode, EDAC protection is provided in a similar way as for the SRAM memory described above. The difference is that write accesses are not being handled automatically. Instead, write accesses must only be performed as individual byte accesses by the software, writing one byte at a time, and the corresponding checkbit byte must be calculated and be written to the correct location by the software.

The operation of the EDAC can be tested trough the MCFG3 register. If the WB (write bypass) bit is set, the value in the TCB field will replace the normal checkbits during memory write cycles. If the RB (read bypass) is set, the memory checkbits of the loaded data will be stored in the TCB field during memory read cycles. NOTE: when the EDAC is enabled, the RMW bit in memory configuration register 2 must be set.

The 8-bit EDAC protection scheme works independently of the actual size of the physical memory device that is mounted in the first bank, since the physical memory would simply be replicated over

the complete 256MiB bank space, always locating the upper part of the physical memory device at the upper part of the bank space. The EDAC checkbytes will be located automatically at power up (especially for the first word access from address 0x00000000). Thus, it is sufficient to enable the EDAC protection setting the SWMX[4] input at reset, and then later in the boot sequence one may constrain the PROM bank size.

5.10.2 Reed-Solomon EDAC

The Reed-Solomon EDAC provides block error correction, and is capable of correcting up to two 4-bit nibble errors in a 32-bit data word or 16-bit checksum. The Reed-Solomon EDAC can be enabled for the SDRAM area only, and uses a 16-bit checksum. The Reed-Solomon EDAC is enabled by setting the RSE and RE bits in MCFG3, and the RMW bit in MCFG2.

The Reed-Solomon data symbols are 4-bit wide, represented as $GF(2^4)$. The basic Reed-Solomon code is a shortened RS(15, 13, 2) code, represented as RS(6, 4, 2). It has the capability to detect and correct a single symbol error anywhere in the codeword. The EDAC implements an interleaved RS(6, 4, 2) code where the overall data is represented as 32 bits and the overall checksum is represented as 16 bits. The codewords are interleaved nibble-wise. The interleaved code can correct two 4-bit errors when each error is located in a nibble and not in the same original RS(6, 4, 2) codeword.

The Reed-Solomon RS(15, 13, 2) code has the following definition:

* there are 4 bits per symbol;
* there are 15 symbols per codeword;
* the code is systematic;
* the code can correct one symbol error per codeword;
* the field polynomial is

$$f(x) = x^4 + x + 1$$

* the code generator polynomial is

$$g(x) = \prod_{i=7}^{8} (x + \alpha^i) = \sum_{j=0}^{2} g_j \cdot x^j$$

  for which the highest power of x is stored first;

* a codeword is defined as 15 symbols:

  $c_0, c_1, c_2, c_3, c_4, c_5, c_6, c_7, c_8, c_9, c_{10}, c_{11}, c_{12}, c_{13}, c_{14}$

  where c_0 to c_{12} represent information symbols and c_{13} to c_{14} represent check symbols.

The shortened and interleaved RS(6, 4, 2) code has the following definition:

* the codeword length is shortened to 4 information symbols and 2 check symbols and as follows:

  $c_0 = c_1 = c_2 = c_3 = c_4 = c_5 = c_6 = c_7 = c_8 = 0$

  where the above information symbols are suppressed or virtually filled with zeros;

* two codewords are interleaved (i.e. interleaved depth I=2) with the following mapping to the 32-bit data and 16-bit checksum, were $c_{i,j}$ is a symbol with codeword index i and symbol index j:

  $c_{0,9}$ = DATA[31:28]

  $c_{1,9}$ = DATA[27:24]

  $c_{0,10}$ = DATA[23:20]

$c_{1,10}$ = DATA[19:16]

$c_{0,11}$ = DATA[15:12]

$c_{1,11}$ = DATA[11:8]

$c_{0,12}$ = DATA[7:4]

$c_{1,12}$ = DATA[3:0]

$c_{0,13}$ = CB[15:12]

$c_{1,13}$ = CB[11:8]

$c_{0,14}$ = CB[7:4]

$c_{1,14}$ = CB[3:0]

5.10.3 EDAC Error reporting

An un-correctable EDAC error will results in an AHB error response. The initiating AHB master will be notified of the error and take corresponding action. If the master was the LEON3FT processor, an instruction or data error trap will be taken (see LEON3 section). The AHB error response will also be registered in the AHB status register (see section 7). Correctable EDAC errors are amended on the fly so that correct data is provided on the AHB bus, but data in memory is left unchanged. The presence of a correctable error is howewer registered in the AHB Status register by raising the NE and CE flags (see section 7). This information can be used for providing a scrubbing mechanism and/or error statistics in software.

5.11 Bus Ready signalling

The BRDYN signal can be used to stretch accesses to the PROM and I/O area. The accesses will always have at least the pre-programmed number of waitstates as defined in memory configuration registers 1, but will be further stretched until BRDYN is asserted. BRDYN should be asserted in the cycle preceding the last one. If bit 29 in MCFG1 is set, BRDYN can be asserted asynchronously with the system clock. In this case, the read data must be kept stable until the de-assertion of OEN and BRDYN must be asserted for at least 1.5 clock cycle. The use of BRDYN can be enabled separately for the PROM and I/O areas. It is recommended that BRDYN is asserted until the corresponding chip select signal is de-asserted, to ensure that the access has been properly completed and avoiding the system to stall.
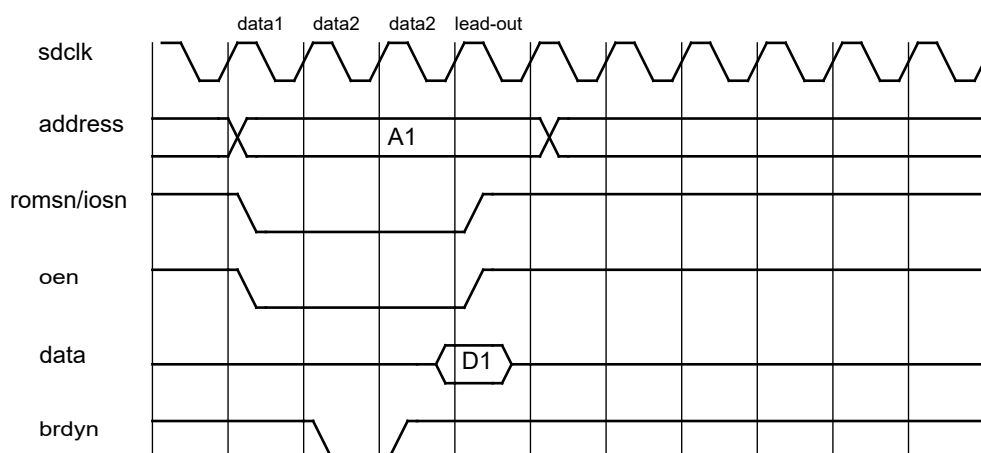
Figure 27. READ cycle with one extra data2 cycle added with BRDYN (synchronous sampling). Lead-out cycle is only applicable for I/O accesses.

Figure 28 shows the use of BRDYN with asynchronous sampling. BRDYN is kept asserted for more than 1.5 clock-cycle. Two synchronization registers are used so it will take at least one additional

cycle from when BRDYN is first asserted until it is visible internally. In figure 28, one cycle is added to the data2 phase.
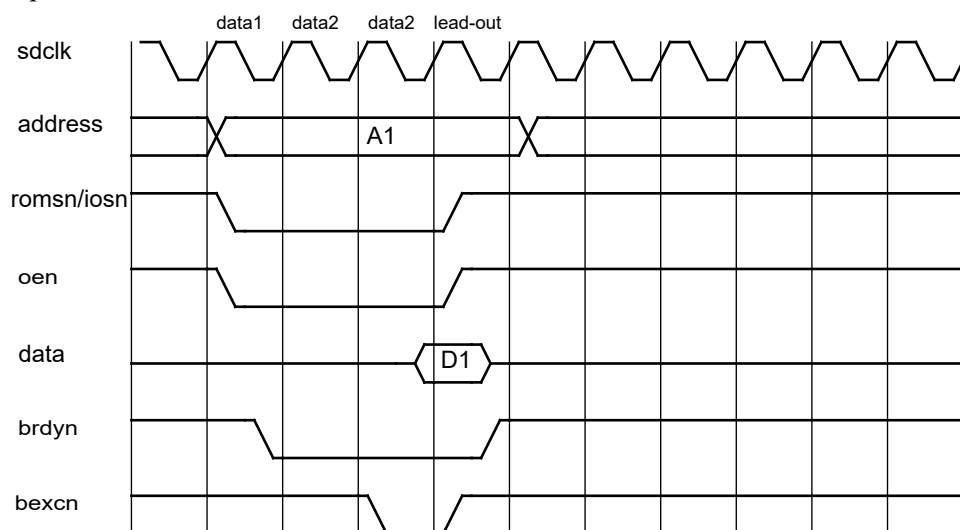
Figure 28. BRDYN (asynchronous) sampling and BEXCN timing. Lead-out cycle is only applicable for I/O-accesses.
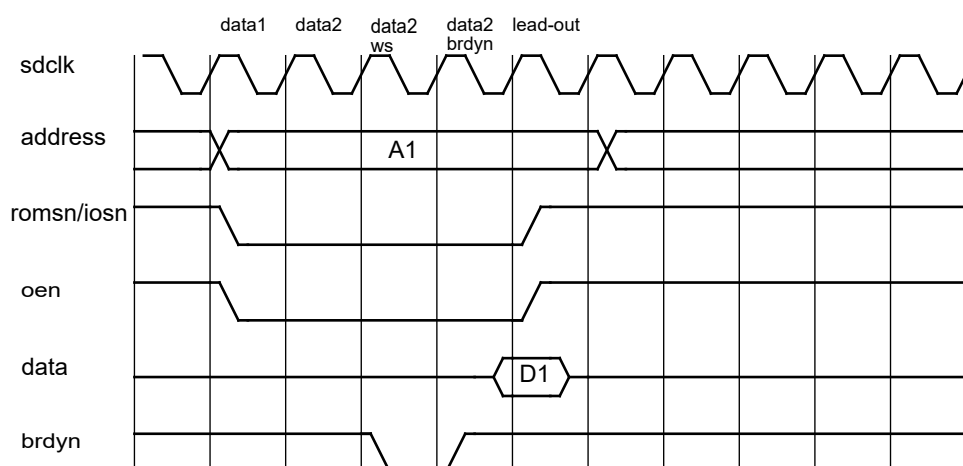
Figure 29. Read cycle with one waitstate (configured) and one BRDYN generated waitstate (synchronous sampling).

If burst accesses and BRDYN signaling are to be used together, special care needs to be taken to make sure BRDYN is raised between the separate accesses of the burst. The controller does not raise the select and OEN signals (in the read case) between accesses during the burst so if BRDYN is kept asserted until the select signal is raised, all remaining accesses in the burst will finish with just the configured fixed number of wait states.

5.12   External bus errors

An access error on the external memory bus can be signalled by asserting the BEXCN signal for read and write accesses. For reads it is sampled together with the read data. For writes it is sampled on the last rising edge before chip select is de-asserted, which is controlled by means of waitstates or bus ready signalling. If the usage of BEXCN is enabled in memory configuration register 1, an error response will be generated on the internal AHB bus. BEXCN can be enabled or disabled through memory configuration register 1, and is active for all areas (PROM, IO and RAM). BEXCN is only sampled in the last access for 8-bit mode for RAM and PROM. That is, when four bytes are written for a word access to 8-bit wide memory BEXCN is only sampled in the last access with the same timing as a single access in 32-bit mode.
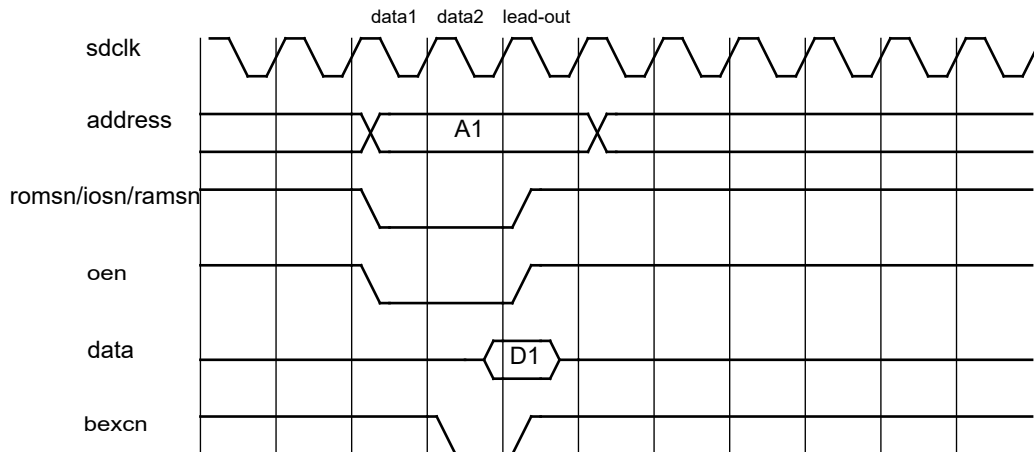
Figure 30. Read cycle with BEXCN.
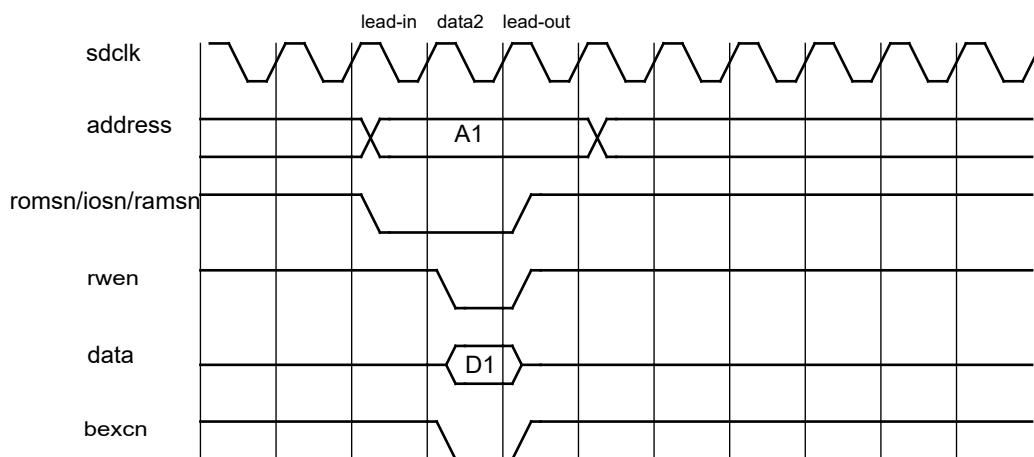
Figure 31. Write cycle with BEXCN. Chip-select (iosn) is not asserted in lead-in cycle for io-accesses.

5.13 Read strobe

The READ output signal indicates the direction of the current PROM,SRAM,IO or SDRAM transfer, and it can be used to drive external bi-directional buffers on the data bus. It always is valid at least one cycle before and after the bus is driven, at other times it is held either constant high or low.

5.14 Registers

The core is programmed through registers mapped into APB address space.

Table 28. FTMCTRL memory controller registers

APB Address offset	Register
0x80000000	Memory configuration register 1 (MCFG1)
0x80000004	Memory configuration register 2 (MCFG2)
0x80000008	Memory configuration register 3 (MCFG3)
0x8000000C	Memory configuration register 4 (MCFG4)

5.14.1 Memory configuration register 1 (MCFG1)

Memory configuration register 1 is used to program the timing of rom and IO accesses.

Table 29. Memory configuration register 1

31	30	29	28	27	26	25	24	23	20	19	18	17	
	PBRDY	ABRDY	IOBUSW		IBRDY	BEXCN		IO WAITSTATES		IOEN		ROMBANKSZ	

14	13	12	11	10	9	8	7	4	3	0
	RESERVED		PWEN		PROM WIDTH		PROM WRITE WS		PROM READ WS	

31	RESERVED
30	PROM area bus ready enable (PBRDY) - Enables bus ready (BRDYN) signalling for the PROM area. Reset to '0'.
29	Asynchronous bus ready (ABRDY) - Enables asynchronous bus ready.
28: 27	I/O bus width (IOBUSW) - Sets the data width of the I/O area ("00"=8, "10" =32).
26	I/O bus ready enable (IBRDY) - Enables bus ready (BRDYN) signalling for the I/O area. Reset to '0'.
25	Bus error enable (BEXCN) - Enables bus error signalling for all areas. Reset to '0'.
24	RESERVED
23: 20	I/O waitstates (IO WAITSTATES) - Sets the number of waitstates during I/O accesses ("0000"=0, "0001"=1, "0010"=2,..., "1111"=15).
19	I/O enable (IOEN) - Enables accesses to the memory bus I/O area.
18	RESERVED
17: 14	PROM bank size (ROMBANKSZ) - Returns current PROM bank size when read. "0000" is a special case and corresponds to a bank size of 256 MiB. All other values give the bank size in binary steps: "0001"=16KiB, "0010"=32KiB, "0011"=64KiB,... , "1011"=16 MiB (i.e. 8 KiB * 2**ROMBANKSZ).
	For any value, two ROM chip select signals are available.
	Programmable bank sizes can be changed by writing to this register field. The written values correspond to the bank sizes and number of chip-selects as above. Reset to "0000".
13:12	RESERVED
11	PROM write enable (PWEN) - Enables write cycles to the PROM area.
10	RESERVED
9: 8	PROM width (PROM WIDTH) - Sets the data width of the PROM area ("00"=8, "10"=32).
7: 4	PROM write waitstates (PROM WRITE WS) - Sets the number of wait states for PROM write cycles ("0000"=0, "0001"=2, "0010"=4,..., "1111"=30).
3: 0	PROM read waitstates (PROM READ WS) - Sets the number of wait states for PROM read cycles ("0000"=0, "0001"=2, "0010"=4,...,"1111"=30). Reset to "1111".

During reset, the MSB PROM width bit (MCFG1[9]) is set to the value of SWMX[6], while the LSB bit (MCFG1[8]) is set to 0. The prom waitstates fields are set to 15 (maximum). External bus error and bus ready are disabled. All other fields are undefined.

5.14.2 Memory configuration register 2 (MCFG2)

Memory configuration register 2 is used to control the timing of the SRAM and SDRAM.

Table 30. Memory configuration register 2

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16
SDRF	TRP		SDRAM TRFC		TCAS		SDRAM BANKSZ		SDRAM COLSZ		SDRAM CMD		D64	SDPB	

15	14	13	12	9	8	7	6	5	4	3	2	1	0
	SE	SI	RAM BANK SIZE			RBRDY	RMW	RAM WIDTH		RAM WRITE WS		RAM READ WS	

31	SDRAM refresh (SDRF) - Enables SDRAM refresh.
30	SDRAM TRP parameter (TRP) - t_{RP} will be equal to 2 or 3 system clocks (0/1).

Table 30. Memory configuration register 2

29: 27	SDRAM TRFC parameter (SDRAM TRFC) - t_{RFC} will be equal to 3+field-value system clocks.
26	SDRAM TCAS parameter (TCAS) - Selects 2 or 3 cycle CAS delay (0/1). When changed, a LOAD-COMMAND-REGISTER command must be issued at the same time. Also sets RAS/CAS delay (t_{RCD}).
25: 23	SDRAM bank size (SDRAM BANKSZ) - Sets the bank size for SDRAM chip selects ("000"=4MiB, "001"=8MiB, "010"=16MiB.... "111"=512MiB).
22: 21	SDRAM column size (SDRAM COLSZ) - "00"=256, "01"=512, "10"=1024, "11"=2048 except when bit[25:23]="111" then "11"=4096
20: 19	SDRAM command (SDRAM CMD) - Writing a non-zero value will generate a SDRAM command. "01"=PRECHARGE, "10"=AUTO-REFRESH, "11"=LOAD-COMMAND-REGISTER. The field is reset after the command has been executed.
18	0
17	0
16: 15	RESERVED
14	SDRAM enable (SE) - Enables the SDRAM controller
13	SRAM disable (SI) - Disables accesses to SRAM bank if bit 14 (SE) is set to '1'.
12: 9	RAM bank size (RAM BANK SIZE) - Sets the size of each RAM bank ("0000"=8KiB, "0001"=16KiB, "0010"=32KiB, "0011"=64KiB,.., "1011"=16MiB) (i.e. 8KiB * 2**RAM BANK SIZE).
8	RESERVED
7	0
6	Read-modify-write enable (RMW) - Enables read-modify-write cycles for sub-word writes 32-bit RAM areas. NOTE: **must** be set if RAM area is 32-bit.
5: 4	RAM width (RAM WIDTH) - Sets the data width of the RAM area ("00"=8, "1X"=32).
3: 2	RAM write waitstates (RAM WRITE WS) - Sets the number of wait states for RAM write cycles ("00"=0, "01"=1, "10"=2, "11"=3).
1: 0	RAM read waitstates (RAM READ WS) - Sets the number of wait states for RAM read cycles ("00"=0, "01"=1, "10"=2, "11"=3).

5.14.3 Memory configuration register 3 (MCFG3)

MCFG3 controls the SDRAM refresh counter and memory EDAC.

Table 31. Memory configuration register 3

31		28	27	26					
RESERVED		RSE	ME	SDRAM REFRESH COUNTER					
		12	11	10	9	8	7		0
			WB	RB	RE	PE	TCB		

31: 29	RESERVED
28	Reed-Solomon EDAC enable (RSE) - if set, will enable Reed-Solomon protection of SDRAM area when implemented
27	1
26: 12	SDRAM refresh counter reload value (SDRAM REFRESH COUNTER)
11	EDAC diagnostic write bypass (WB) - Enables EDAC write bypass.
10	EDAC diagnostic read bypass (RB) - Enables EDAC read bypass.
9	RAM EDAC enable (RE) - Enable EDAC checking of the RAM area (including SDRAM).
8	PROM EDAC enable (PE) - Enable EDAC checking of the PROM area. Ar reset, this bit is initialized with the value of SWMX[4].
7: 0	Test checkbits (TCB) - This field replaces the normal checkbits during write cycles when WB is set. It is also loaded with the memory checkbits during read cycles when RB is set.

The period between each AUTO-REFRESH command is calculated as follows:

$$t_{REFRESH} = ((\text{reload value}) + 1) / \text{SYSCLK}$$

5.14.4 Memory configuration register 4 (MCFG4)

MCFG4 provides means to insert Reed-Solomon EDAC errors into memory for diagnostic purposes.

Table 32. Memory configuration register 4

31	16
RESERVED	WB

15	0
TCB[15:0]	

31: 17	RESERVED
16	EDAC diagnostic write bypass (WB) - Enables EDAC write bypass. Identical to WB in MCFG3.
15: 0	Test checkbits (TCB) - This field replaces the normal checkbits during write cycles when WB is set. It is also loaded with the memory checkbits during read cycles when RB is set. Note that TCB[7:0] are identical to TCB[7:0] in MCFG3

5.15 Signal definitions

The signals are described in table 33.

Table 33. Signal definitions

Signal name	Type	Function	Active
ADDRESS[23:0]	Output	Memory address	High
DATA[31:0]	Input/Output	Memory data	High
CB[15:0]	Input/Output	Check bits	High
RAMSN[1:0]	Output	SRAM chip select	Low
RAMOEN	Output	SRAM output enable	Low
RAMWEN	Output	SRAM write enable	Low
OEN	Output	Output enable	Low
WRITEN	Output	Write strobe	Low
READ	Output	Read strobe	High
IOSN	Output	I/O area chip select	Low
ROMSN[1:0]	Output	PROM chip select	Low
BRDYN	Input	Bus ready. Extends accesses to the PROM and IO areas.	Low
BEXCN	Input	Bus exception.	Low
SDWEN	Output	SDRAM write enable	Low
SDRASN	Output	SDRAM row address strobe	Low
SDCASN	Output	SDRAM column address strobe	Low
SDDQM[3:0]	Output	SDRAM data mask: SDDQM[3] corresponds to DATA[31:24], SDDQM[2] corresponds to DATA[23:16], SDDQM[1] corresponds to DATA[15:8], SDDQM[0] corresponds to DATA[7:0]. Any SDDQM[ ] signal can be used for CB[ ].	High
GPIO[3] / SWMX[6]	Input	Configuring PROM data width at reset. 8-bit data width when "0", and 32-bit when "1".	-
GPIO[1] / SWMX[4]	Input	Enabling EDAC usage for PROM at reset when "1".	-

6 On-chip Memory with EDAC Protection

6.1 Overview

The GR712RC is provided with 192 KiB on-chip RAM, based on the FTAHBRAM core from GRLIB. The RAM is protected with an error detection and correction unit (EDAC), capable of correcting one error per word, and detecting two errors per word. The on-chip memory is not cacheable.

6.2 Operation

The 192 KiB of on-chip memory is accessed through the AHB bus at address 0xA0000000 - 0xA0030000. The configuration register is accessible via the APB bus at address 0x80100000.

The on-chip RAM implements a general-purpose storage and can be accessed by any AHB master. The on-chip RAM can be protected against soft errors by enabling the EDAC protection in the configuration register. When enabled, the EDAC will transparently correct any single-bit error, and return an AHB error response in case of an uncorrectable multi-bit error.

If EDAC is enabled by setting the EN bit in configuration register. When enabled for the first time after reset, the full RAM contents should be written with a pre-defined value to initialize the EDAC checkbits using word (32-bit) writes.

Reads are performed with 3 AHB waitstates while (32-bit) writes are fully buffered and do not generate any waitstates. Sub-word writes require 4 waitstates.

Table 34. Summary of the number of waitstates for the different operations for the memory.

Operation	Waitstates
Read	3
Word write	0
Subword write	4

Note: the FTAHBRAM is not cached in the L1 caches of the LEON3FT processor(s). It is therefore better suited to be used for data buffers than containing program instructions. If the processors executes out of the FTAHBRAM, the performance is typically 0.25 MIPS/MHz.

6.2.1 EDAC checkbits diagnostic access

The EDAC checkbits can be accessed for diagnostic purposes using read and write bypass mode. In read bypass mode, the checkbits are stored in the TCB field of the configuration register after each read to the RAM. In write bypass mode, the checkbits in the RAM array are written with the value of the TCB field instead of the automatically calculated checkbits. The bypass modes can be enabled by setting the RB and/or WB bits in the configuration register.

6.2.2 Error reporting

An 8-bit single error counter (SEC) field is available in the configuration register. It is incremented each time a single data bit error is encountered (reads or subword writes), saturation at 255. The counter can be reset by writing the value 255 to it. A single-error will also be reported in the AHB status register, see section 7 for details.

6.3   Registers

The on-chip RAM is configured through a configuration register mapped at address 0x80100000:

Table 35. FTAHBRAM registers

APB Address offset	Register
0x80100000	Configuration Register

Table 36. Configuration Register

31 30 29		24 23	21 20	13 12	10 9	8	7	6	0
DIAG		MEMSIZE	SEC	MEMSIZE	WB	RB	EN	TCB	

31:	30	RAM timing adjust. Must be written with "00".
23	21	3 MSB bits of log2 of the memory size (read-only). Hardcoded to "001".
20:	13	Single error counter (SEC): Incremented each time a single error is corrected (includes errors on checkbits). Each bit can be set to zero by writing a one to it.
12:	10	3 LSB bits of log2 of the memory size (read-only). Hardcoded to "000".
9		Write Bypass (WB): When set, the TCB field is stored as check bits when a write is performed to the memory.
8		Read Bypass (RB): When set during a read or subword write, the check bits loaded from memory are stored in the TCB field.
7		EDAC Enable (EN): When set, the EDAC is used otherwise it is bypassed during read and write operations.
6:	0	Test Check Bits (TCB): Used as checkbits when the WB bit is set during writes and loaded with the check bits during a read operation when the RB bit is set.

Any unused most significant bits are reserved. Always read as '000...0'.

All fields except TCB are initialized to 0 at reset.

7   AHB Status Registers

7.1   Overview

The AHB status registers store information about AMBA AHB accesses returning an error response. The AHB status registers consists of a failing address register capturing the address of the failed access, and a status register providing information about the access type.

7.2   Operation

The registers monitor AMBA AHB bus transactions and store the current HADDR, HWRITE, HMASTER and HSIZE internally. The monitoring is activated after reset and until an error response (HRESP = "01") is detected. When an error is detected, the status and address register contents are frozen and the New Error (NE) bit is set to one. At the same time, interrupt 1 is generated. To restart error monitoring, the NE bit must be cleared by software.

To be able to monitor error rates in on-chip and off-chip memories, the AHB status registers also latch the address and control signals when a correctable error occurs in the off-chip memory controller, or in the on-chip RAM module. Such errors do not cause an error response on the AHB bus, but can still be detected through dedicated sideband signals. When such a correctable error is detected, the effect will be the same as for an AHB error response, with the only difference that the Correctable Error (CE) bit in the status register is set to one. When the CE bit is set, the interrupt routine can either scrub the failing address or just register the error for statistical use. The NE bit should then be cleared to restart error monitoring.

7.3   Registers

The core is programmed through registers mapped into APB address space.

Table 37. AHB Status registers

APB address	Registers
0x80000F00	AHB Status register
0x80000F04	AHB Failing address register

Table 38. AHB Status register

31					10	9	8	7	6	3	2	0
		RESERVED				CE	NE	HWRITE	HMASTER		HSIZE	

31: 10	RESERVED
9	CE: Correctable Error. Set if the detected error was caused by a single error and zero otherwise.
8	NE: New Error. Deasserted at start-up and after reset. Asserted when an error is detected. Reset by writing a zero to it.
7	The HWRITE signal of the AHB transaction that caused the error.
6: 3	The HMASTER signal of the AHB transaction that caused the error.
2: 0	The HSIZE signal of the AHB transaction that caused the error

Table 39. AHB Failing address register

31	0
AHB FAILING ADDRESS	

31: 0	The HADDR signal of the AHB transaction that caused the error.

8 Multiprocessor Interrupt Controller

8.1 Overview

The interrupts generated by the GR712RC peripherals are forwarded to the IRQMP multi-processor interrupt controller. The interrupt controller prioritizes, masks and propagates the interrupt with the highest priority to the two processors.

8.2 Operation

8.2.1 Interrupt routing

The GR712RC device has an internal interrupt bus consisting of 31 interrupts, which are mapped on the 15 SPARC interrupts by the IRQMP controller.

The lowest 15 interrupts (1 - 15) are mapped directly on the 15 SPARC interrupts. When any of these lines are asserted high, the corresponding bit in the interrupt pending register is set. The pending bits will stay set even if the PIRQ line is de-asserted, until cleared by software or by an interrupt acknowledge from the processor. Interrupt 16-31 are generated as interrupt 12, and for those interrupts a chained interrupt handler is typically used by the software.

8.2.2 Interrupt prioritization

Each interrupt can be assigned to one of two levels (0 or 1) as programmed in the interrupt level register. Level 1 has higher priority than level 0. The interrupts are prioritised within each level, with interrupt 15 having the highest priority and interrupt 1 the lowest. The highest interrupt from level 1 will be forwarded to the processor. If no unmasked pending interrupt exists on level 1, then the highest unmasked interrupt from level 0 will be forwarded.

Interrupts are prioritised at system level, while masking and forwarding of interrupts in done for each processor separately. Each processor in an multiprocessor system has separate interrupt mask and force registers. When an interrupt is signalled on the interrupt bus, the interrupt controller will prioritize interrupts, perform interrupt masking for each processor according to the mask in the corresponding mask register and forward the interrupts to the processors.
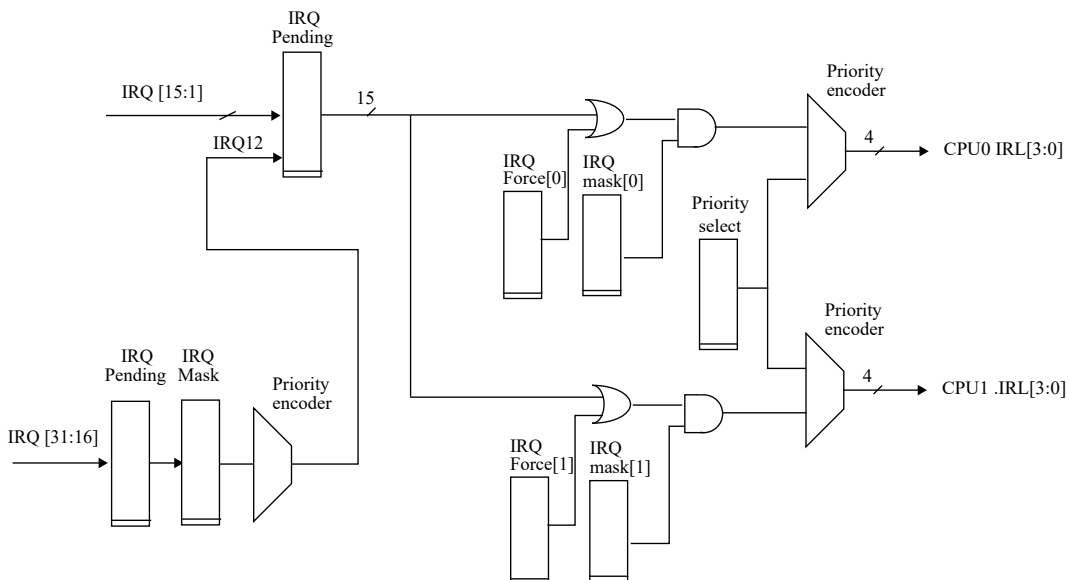
Figure 32. Interrupt controller block diagram

When a processor acknowledges the interrupt, the corresponding pending bit will automatically be cleared. Interrupt can also be forced by setting a bit in the interrupt force register. In this case, the pro-

cessor acknowledgement will clear the force bit rather than the pending bit. After reset, the interrupt mask register is set to all zeros while the remaining control registers are undefined. Note that interrupt 15 cannot be maskable by the LEON3 processor and should be used with care - most operating systems do not safely handle this interrupt.

8.2.3   Interrupt assignment

The following table specifies the interrupt assignment for the GR712RC:

TABLE 40. Interrupt assignment

Core	Interrupt #	Function
AHBSTAT	1	AHB bus error
APBUART - 0	2	UART0 RX/TX interrupt
GRGPIO1-2	1-15	External I/O interrupt (3 & 4 are GPIO only)
CANOC	5-6	OCCAN interrupt (core 1 - 2)
GRTIMER	7	GRTIMER timer underflow interrupt
GPTIMER	8-11	GPTIMER timer underflow interrupts
IRQMP	12	IRQMP Extended interrupt
SPICTRL, SLINK	13	SPI and SLINK interrupt
B1553BRM, GRETH, GRTC	14	1553, Ethernet, and Telecommand interrupt
ASCS	16	ASCS interrupt
APBUART - 1 to 5	17-21	UART1-5 RX/TX interrupt
GRSPW2 - 0 to 5	22-27	SpaceWire 0-5 RX/TX data interrupt
I2CMST	28	I2C master interrupt
GRTM	29	Telemetry encoder interrupt
	30	Telemetry encoder time strobe interrupt

8.2.4   Processor status monitoring

The processor status can be monitored through the Multiprocessor Status Register. The STATUS field in this register indicates if a processor is power-down ('1') or running ('0'). A processor in power-down can be made running by writing a '1' to its status field, which makes the processor continue execution where it was powered down. After reset, all processors except processor 0 are powered down. When the system is properly initialized, processor 0 can start the remaining processors by writing to their STATUS bits.

8.2.5   Interrupt broadcasting

An incoming interrupt that has its bit set in the Broadcast Register is propagated to the force register of *all* processors rather than only to the Pending Register. This can be used to implement a timer that fires to all processors with that same interrupt.

8.3 Registers

The core is controlled through registers mapped into APB address space. The number of implemented registers depend on number of processor in the multiprocessor system.

Table 41. Interrupt Controller registers

APB address offset	Register
0x80000200	Interrupt level register
0x80000204	Interrupt pending register
0x80000208	Interrupt force register, processor 0
0x8000020C	Interrupt clear register
0x80000210	Multiprocessor status register
0x80000214	Broadcast register
0x80000240	Processor 0 interrupt mask register
0x80000244	Processor 1 interrupt mask register
0x80000280	Processor 0 interrupt force register
0x80000284	Processor 1 interrupt force register
0x800002C0	Processor 0 extended interrupt identification register
0x800002C4	Processor 1 extended interrupt identification register

8.3.1 Interrupt level register

31	17	16	1	0
"000..0"		IL[15:1]		0

Figure 33. Interrupt level register

[31:16]   Reserved.
[15:1]   Interrupt Level *n* (IL[*n*]): Interrupt level for interrupt *n*.
[0]   Reserved.

8.3.2 Interrupt pending register

31	16	15	1	0
EIP[31:16]		IP[15:1]		0

Figure 34. Interrupt pending register

[31:17]   Extended Interrupt Pending n (EIP[n]).
[15:1]   Interrupt Pending n (IP[n]): Interrupt pending for interrupt n.
[0]   Reserved

8.3.3 Interrupt force register, processor 0

31	16	15	1	0
"000...0"		IF[15:1]		0

Figure 35. Processor interrupt force register

[31:16]   Reserved.
[15:1]   Interrupt Force *n* (IF[*n*]): Force interrupt no. *n*. The resulting IF[*n*] is the value of the written value to IF[*n*].
[0]   Reserved.

8.3.4 Interrupt clear register

31	16	15	1	0
"000...0"		IC[15:1]		0

Figure 36. Interrupt clear register

[31:16]   Reserved.
[15:1]   Interrupt Clear *n* (IC[*n*]): Writing '1' to IC*n* will clear interrupt *n*. Write only, reads zeros.
[0]   Reserved.

8.3.5 Multiprocessor status register

31	28	27	26	20	19	16	15	0
NCPU		BA	"000...0"		EIRQ		STATUS[15:0]	

Figure 37. Multiprocessor status register

[31:28]   NCPU. Number of CPU's in the system -1. Read only. Fixed to 1, i.e. 2 processors.
[27]   Broadcast Available (BA). Set to '1'.
[19:16]   EIRQ. Interrupt number (1 - 15) used for extended interrupts. Read only. Fixed to 12.
[15:1]   Power-down status of CPU [*n*]: reads '1' = power-down, '0' = running. Write to start processor *n*: '1'=to start, '0=has no effect.

8.3.6 Processor interrupt mask register

31	16	15	1	0
EIM[31:16]		IM[15:1]		0

Figure 38. Processor interrupt mask register

[31:16]   Interrupt mask for extended interrupts
[15:1]   Interrupt Mask *n* (IM[*n*]): If IMn = 0 the interrupt *n* is masked, otherwise it is enabled.
[0]   Reserved.

8.3.7 Broadcast register

```
 31                                              16 15                              1  0
┌──────────────────────────────────────────────┬──────────────────────────────┬────┐
│                  "000...0"                     │           IM[15:1]           │ 0  │
└──────────────────────────────────────────────┴──────────────────────────────┴────┘
```

Figure 39. Broadcast register

[31:16]   Reserved.
[15:1]   Broadcast Mask n (BM[n]): If BMn = 1 the interrupt n is broadcasted (written to the Force Register of all CPUs), otherwise standard semantic applies (Pending Register).
[0]   Reserved.

8.3.8 Processor interrupt force register

```
 31                                      17 16 15                              1  0
┌──────────────────────────────────────┬────┬──────────────────────────────┬────┐
│              IFC[15:1]                 │ 0  │           IF[15:1]           │ 0  │
└──────────────────────────────────────┴────┴──────────────────────────────┴────┘
```

Figure 40. Processor interrupt force register

[31:17]   Interrupt Force Clear n (IFC[n]). Write only, reads zeros.
[15:1]   Interrupt Force n (IF[n]): Force interrupt no. n. The resulting IF[n] is a combination of the written value to IF[n] and the written value of IFC[n] and the previous state of IF[n].
[0]   Reserved.

8.3.9 Extended interrupt identification register

```
 31                                                            5  4           0
┌─────────────────────────────────────────────────────────────┬─────────────┐
│                                                               │   EID[4:0]  │
└─────────────────────────────────────────────────────────────┴─────────────┘
```

Figure 41. Extended interrupt identification register

[4:0]   ID (16 - 31) of the acknowledged extended interrupt. Read only.

9 Hardware Debug Support Unit

9.1 Overview

To simplify debugging on target hardware, the LEON3 processor implements a debug mode during which the pipeline is idle and the processor is controlled through a Debug Support Unit (DSU). The DSU acts as an AHB slave and can be accessed by any AHB master. In the GR712RC, an external debug host can access the DSU through JTAG or SpaceWire links 0 and 1 (using RMAP).
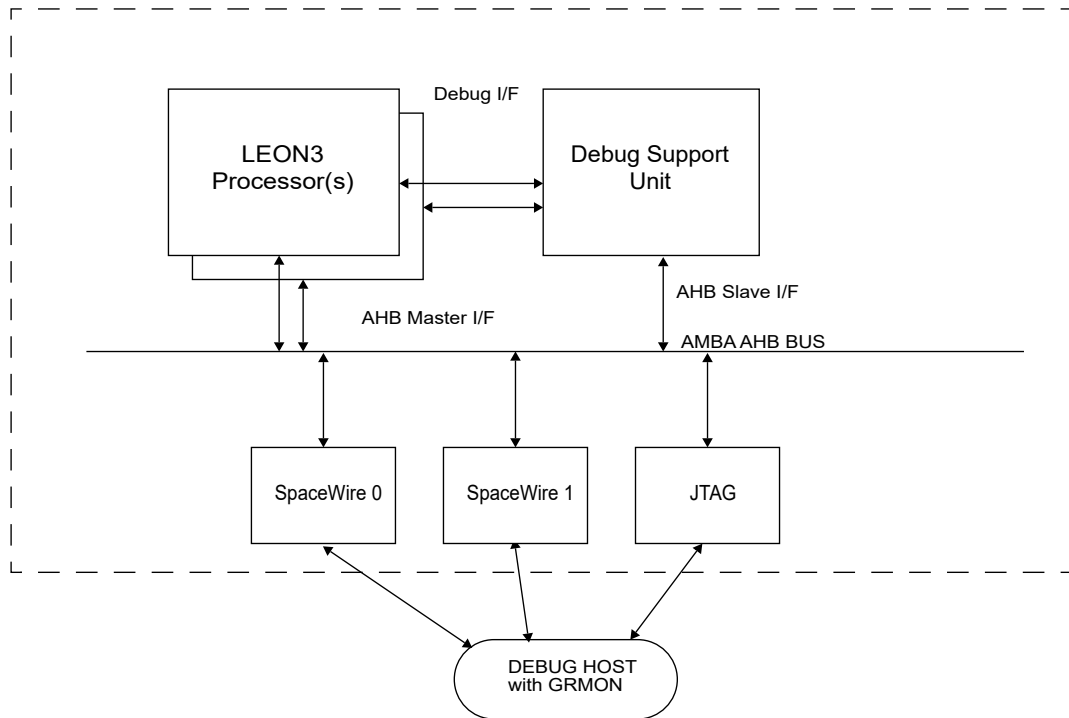
Figure 42. LEON3/DSU Connection

9.2 Operation

Through the DSU AHB slave interface, the debug link AHB master can access the processor registers and the contents of the instruction trace buffer. The DSU control registers can be accessed at any time, while the processor registers, caches and trace buffer can only be accessed when the processor has entered debug mode. In debug mode, the processor pipeline is held and the processor state can be accessed by the DSU. Entering the debug mode can occur on the following events:

- executing a breakpoint instruction (ta 1)
- integer unit hardware breakpoint/watchpoint hit (trap 0xb)
- setting the break-now (BN) bits in the DSU break and single step register
- a trap that would cause the processor to enter error mode
- occurrence of any, or a selection of traps as defined in the DSU control register
- after a single-step operation
- one of the processors in a multiprocessor system has entered the debug mode
- DSU AHB breakpoint hit

When the debug mode is entered, the following actions are taken:

- PC and nPC are saved in temporary registers (accessible by the debug unit)
- the timer unit is (optionally) stopped to freeze the LEON timers and watchdog

The instruction that caused the processor to enter debug mode is not executed, and the processor state is kept unmodified. Execution is resumed by clearing the BN bits in the DSU break and single step register. The timer unit will be re-enabled and execution will continue from the saved PC and nPC. Debug mode can also be entered after the processor has entered error mode, for instance when an application has terminated and halted the processor. The error mode can be reset and the processor restarted at any address.

When a processor is in the debug mode, an access to ASI diagnostic area is forwarded to the IU which performs access with ASI equal to value in the DSU ASI register and address consisting of 20 LSB bits of the original address.

9.3   AHB Trace Buffer

The AHB trace buffer consists of a circular buffer that stores AHB data transfers. The address, data and various control signals of the AHB bus are stored and can be read out for later analysis. The trace buffer is 128 bits wide and 256 lines deep. The information stored is indicated in the table below:

Table 42. AHB Trace buffer data allocation

Bits	Name	Definition
127	AHB breakpoint hit	Set to '1' if a DSU AHB breakpoint hit occurred.
126	-	Not used
125:96	Time tag	DSU time tag counter
95	-	Not used
94:80	Hirq	AHB HIRQ[15:1]
79	Hwrite	AHB HWRITE
78:77	Htrans	AHB HTRANS
76:74	Hsize	AHB HSIZE
73:71	Hburst	AHB HBURST
70:67	Hmaster	AHB HMASTER
66	Hmastlock	AHB HMASTLOCK
65:64	Hresp	AHB HRESP
63:32	Load/Store data	AHB HRDATA or HWDATA
31:0	Load/Store address	AHB HADDR

In addition to the AHB signals, the DSU time tag counter is also stored in the trace.

The trace buffer is enabled by setting the enable bit (EN) in the trace control register. Each AHB transfer is then stored in the buffer in a circular manner. The address to which the next transfer is written is held in the trace buffer index register, and is automatically incremented after each transfer. Tracing is stopped when the EN bit is reset, or when a AHB breakpoint is hit. Tracing is temporarily suspended when the processor enters debug mode. Note that neither the trace buffer memory nor the breakpoint registers (see below) can be read/written by software when the trace buffer is enabled.

9.4 Instruction trace buffer

The instruction trace buffer consists of a circular buffer that stores executed instructions. The instruction trace buffer is located in the processor, and read out via the DSU. The trace buffer is 128 bits wide and 256 lines deep. The information stored is indicated in the table below:

Table 43. Instruction trace buffer data allocation

Bits	Name	Definition
127	-	Unused
126	Multi-cycle instruction	Set to '1' on the second and third instance of a multi-cycle instruction (LDD, ST or FPOP)
125:96	Time tag	The value of the DSU time tag counter
95:64	Load/Store parameters	Instruction result, Store address or Store data
63:34	Program counter	Program counter (2 lsb bits removed since they are always zero)
33	Instruction trap	Set to '1' if traced instruction trapped
32	Processor error mode	Set to '1' if the traced instruction caused processor error mode
31:0	Opcode	Instruction opcode

During tracing, one instruction is stored per line in the trace buffer with the exception of multi-cycle instructions. Multi-cycle instructions are entered two or three times in the trace buffer. For store instructions, bits [63:32] correspond to the store address on the first entry and to the stored data on the second entry (and third in case of STD). Bit 126 is set on the second and third entry to indicate this. A double load (LDD) is entered twice in the trace buffer, with bits [63:32] containing the loaded data. Bit 126 is set for the second entry.

When the processor enters debug mode, tracing is suspended. The trace buffer and the trace buffer control register can be read and written while the processor is in the debug mode. During the instruction tracing (processor in normal mode) the trace buffer and the trace buffer control register can not be accessed.

9.5   DSU memory map

The DSU memory map can be seen in table 44 below. In a multiprocessor systems, the register map is duplicated and address bits 27 - 24 are used to index the processor.

Table 44. DSU memory map

Address offset	Register
0x90000000	DSU control register
0x90000008	Time tag counter
0x90000020	Break and Single Step register
0x90000024	Debug Mode Mask register
0x90000040	AHB trace buffer control register
0x90000044	AHB trace buffer index register
0x90000050	AHB breakpoint address 1
0x90000054	AHB mask register 1
0x90000058	AHB breakpoint address 2
0x9000005c	AHB mask register 2
0x90100000 - 0x9010FFFF	Instruction trace buffer (..0: Trace bits 127 - 96, ..4: Trace bits 95 - 64, ..8: Trace bits 63 - 32, ..C : Trace bits 31 - 0)
0x90110000	Instruction Trace buffer control register
0x90200000 - 0x90210000	AHB trace buffer (..0: Trace bits 127 - 96, ..4: Trace bits 95 - 64, ..8: Trace bits 63 - 32, ..C : Trace bits 31 - 0)
0x90300000 - 0x903007FC	IU register file
0x90300800 - 0x90300FFC	IU register file check bits (LEON3FT only)
0x90301000 - 0x9030107C	FPU register file
0x90400000 - 0x904FFFFC	IU special purpose registers
0x90400000	Y register
0x90400004	PSR register
0x90400008	WIM register
0x9040000C	TBR register
0x90400010	PC register
0x90400014	NPC register
0x90400018	FSR register
0x9040001C	CPSR register
0x90400020	DSU trap register
0x90400024	DSU ASI register
0x90400040 - 0x9040007C	ASR16 - ASR31 (when implemented)
0x90700000 - 0x907FFFFC	ASI diagnostic access (ASI = value in DSU ASI register, address = address[19:0]) ASI = 0x9 : Local instruction RAM ASI = 0xB : Local data RAM ASI = 0xC : Instruction cache tags ASI = 0xD : Instruction cache data ASI = 0xE : Data cache tags ASI = 0xF : Data cache data ASI = 0x1E : Separate snoop tags

The addresses of the IU registers are formed as follows:

* %on : 0x300000 + (((psr.cwp * 64) + 32 + n*4) mod (NWINDOWS*64))
* %ln : 0x300000 + (((psr.cwp * 64) + 64 + n*4) mod (NWINDOWS*64))
* %in : 0x300000 + (((psr.cwp * 64) + 96 + n*4) mod (NWINDOWS*64))
* %gn : 0x300000 + (NWINDOWS*64)
* %fn : 0x301000 + n*4

9.6   DSU registers

9.6.1   DSU control register

The DSU is controlled by the DSU control register:

.

31		11	10	9	8	7	6	5	4	3	2	1	0
		PW	HL	PE	0	1	DM	BZ	BX	BS	BW	BE	TE

Figure 43. DSU control register

[0]:   Trace enable (TE). Enables instruction tracing. If set the instructions will be stored in the trace buffer. Remains set when then processor enters debug or error mode.

[1]:   Break on error (BE) - if set, will force the processor to debug mode when the processor would have entered error condition (trap in trap).

[2]:   Break on IU watchpoint (BW)- if set, debug mode will be forced on a IU watchpoint (trap 0xb).

[3]:   Break on S/W breakpoint (BS) - if set, debug mode will be forced when an breakpoint instruction (ta 1) is executed.

[4]:   Break on trap (BX) - if set, will force the processor into debug mode when any trap occurs.

[5]:   Break on error traps (BZ) - if set, will force the processor into debug mode on all *except* the following traps: priviledged_instruction, fpu_disabled, window_overflow, window_underflow, asynchronous_interrupt, ticc_trap.

[6]:   Debug mode (DM). Indicates when the processor has entered debug mode (read-only).

[7]:   Unused, hardcoded to 1.

[8]:   Unused, hardcoded to 0.

[9]:   Processor error mode (PE) - returns '1' on read when processor is in error mode, else '0'. If written with '1', it will clear the error and halt mode.

[10]:   Processor halt (HL). Returns '1' on read when processor is halted. If the processor is in debug mode, setting this bit will put the processor in halt mode.

[11]:   Power down (PW). Returns '1' when processor in power-down mode.

9.6.2   DSU Break and Single Step register

This register is used to break or single step the processor(s). This register controls all processors in a multi-processor system, and is only accessible in the DSU memory map of processor 0.

31		17	16			1	0
. . .		SS1	SS0		. . .	BN1	BN0

Figure 44. DSU Break and Single Step register

[15:0]:   Break now (BNx) -Force processor x into debug mode if the Break on watchpoint (BW) bit in the processors DSU control register is set. If cleared, the processor x will resume execution.

[31:16]: Single step (SSx) - if set, the processor x will execute one instruction and return to debug mode. The bit remains set after the processor goes into the debug mode.

9.6.3   DSU Debug Mode Mask Register

When one of the processors in a multiprocessor LEON3 system enters the debug mode the value of the DSU Debug Mode Mask register determines if the other processors are forced in the debug mode. This register controls all processors in a multi-processor system, and is only accessible in the DSU memory map of processor 0.

31		18	17	16	15		2	1	0
. . .			DM1	DM0		. . .		ED1	ED0

Figure 45. DSU Debug Mode Mask register

GR712RC

[15:0]: Enter debug mode (EDx) - Force processor x into debug mode if any of processors in a multiprocessor system enters the debug mode. If 0, the processor x will not enter the debug mode.

[31:16]: Debug mode mask. If set, the corresponding processor will not be able to force running processors into debug mode even if it enters debug mode.

9.6.4   DSU trap register

The DSU trap register is a read-only register that indicates which SPARC trap type that caused the processor to enter debug mode. When debug mode is force by setting the BN bit in the DSU break and single step register, the trap type will be 0xb (hardware watchpoint trap).

31	13	12	11	4	3	0
RESERVED		EM	TRAP TYPE		0000	

Figure 46. DSU trap register

[11:4]: 8-bit SPARC trap type
[12]: Error mode (EM). Set if the trap would have cause the processor to enter error mode.

9.6.5   Trace buffer time tag counter

The trace buffer time tag counter is incremented each clock as long as the processor is running. The counter is stopped when the processor enters debug mode, and restarted when execution is resumed.

31	29	0
00	DSU TIME TAG VALUE	

Figure 47. Trace buffer time tag counter

The value is used as time tag in the instruction and AHB trace buffer.

9.6.6   DSU ASI register

The DSU can perform diagnostic accesses to different ASI areas. The value in the ASI diagnostic access register is used as ASI while the address is supplied from the DSU.

31	7	0
	ASI	

Figure 48. ASI diagnostic access register

[7:0]: ASI to be used on diagnostic ASI access

9.6.7   AHB Trace buffer control register

The AHB trace buffer is controlled by the AHB trace buffer control register:

31	16	2	1	0
DCNT	RESERVED	BR	DM	EN

Figure 49. AHB trace buffer control register

[0]: Trace enable (EN). Enables the trace buffer.
[1]: Delay counter mode (DM). Indicates that the trace buffer is in delay counter mode.
[2]: Break (BR). If set, the processor will be put in debug mode when AHB trace buffer stops due to AHB breakpoint hit.
[31:16] Trace buffer delay counter (DCNT). Note that the number of bits actually implemented depends on the size of the trace buffer.

GR712RC-UM, Jul 2018, Version 2.12   85   www.cobham.com/gaisler

9.6.8 AHB trace buffer index register

The AHB trace buffer index register contains the address of the next trace line to be written.

31	4	3	0
INDEX		0000	

Figure 50. AHB trace buffer index register

31:4   Trace buffer index counter (INDEX). Note that the number of bits actually implemented depends on the size of the trace buffer.

9.6.9 AHB trace buffer breakpoint registers

The DSU contains two breakpoint registers for matching AHB addresses. A breakpoint hit is used to freeze the trace buffer by automatically clearing the enable bit. Freezing can be delayed by programming the DCNT field in the trace buffer control register to a non-zero value. In this case, the DCNT value will be decremented for each additional trace until it reaches zero, after which the trace buffer is frozen. A mask register is associated with each breakpoint, allowing breaking on a block of addresses. Only address bits with the corresponding mask bit set to '1' are compared during breakpoint detection. To break on AHB load or store accesses, the LD and/or ST bits should be set.

Break address reg.

31	2	1	0
BADDR[31:2]		0	0

Break mask reg.

31	2	1	0
BMASK[31:2]		LD	ST

Figure 51. Trace buffer breakpoint registers

[31:2]:   Breakpoint address (bits 31:2)
[31:2]:   Breakpoint mask (see text)
[1]:   LD - break on data load address
[0]:   ST - beak on data store address

9.6.10 Instruction trace control register

The instruction trace control register contains a pointer that indicates the next line of the instruction trace buffer to be written.

31	16	0
RESERVED	IT POINTER	

Figure 52. Instruction trace control register

[15:0]   Instruction trace pointer. Note that the number of bits actually implemented depends on the size of the trace buffer.

10 JTAG Debug Interface

10.1 Overview

The JTAG debug interface provides access to on-chip AMBA AHB bus through JTAG. The JTAG debug interface implements a simple protocol which translates JTAG instructions to AHB transfers. Through this link, a read or write transfer can be generated to any address on the AHB bus. This is typically used by an external debug tool such as *grmon*, to load and execute programs on the system.
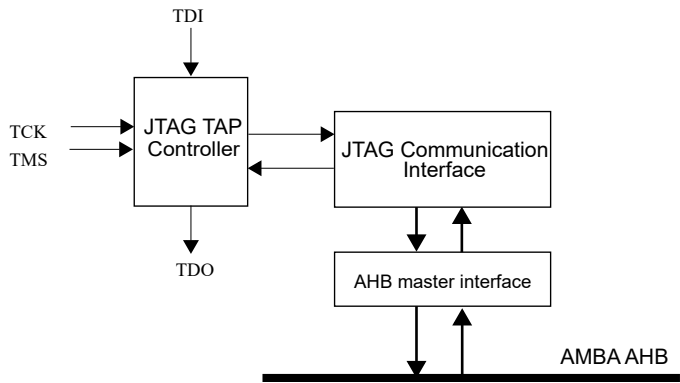
Figure 53. JTAG Debug link block diagram

10.2 Operation

10.2.1 Transmission protocol

The JTAG Debug link decodes two JTAG instructions and implements two JTAG data registers: the command/address register and data register. A read access is initiated by shifting in a command consisting of read/write bit, AHB access size and AHB address into the command/address register. The AHB read access is performed and data is ready to be shifted out of the data register. Write access is performed by shifting in command, AHB size and AHB address into the command/data register followed by shifting in write data into the data register. Sequential transfers can be performed by shifting in command and address for the transfer start address and shifting in SEQ bit in data register for following accesses. The SEQ bit will increment the AHB address for the subsequent access. Sequential transfers should not cross a 1 kB boundary. Sequential transfers are always word based.

Table 45. JTAG debug link Command/Address register

34	33 32 31		0
W	SIZE	AHB ADDRESS	

34	Write (W) - '0' - read transfer, '1' - write transfer
33 32	AHB transfer size - "00" - byte, "01" - half-word, "10" - word, "11"- reserved
31 30	AHB address

Table 46. JTAG debug link Data register

32	31		0
SEQ	AHB DATA		

Table 46. JTAG debug link Data register

32	Sequential transfer (SEQ) - If '1' is shifted in this bit position when read data is shifted out or write data shifted in, the subsequent transfer will be to next word address.
31 30	AHB Data - AHB write/read data. For byte and half-word transfers data is aligned according to big-endian order where data with address offset 0 data is placed in MSB bits.

10.3 Registers

The core does not implement any registers mapped in the AMBA AHB or APB address space.

10.4 Signal definitions

The signals are described in table 47.

Table 47. Signal definitions

Signal name	Type	Function	Active
TCK	Input	JTAG clock	-
TMS	Input	JTAG TMS	High
TDI	Input	JTAG TDI	High
TDO	Output	JTAG TDO	High

NOTE: for correct operation, all JTAG signals should be pulled-up externally with 10kOhm. The GR712RC does not include any internal pull-ups. This is in line with the TAP specification where TMS and TDI implementation should be such that if an external signal fails (e.g. open circuit) then the behavior of TMS and TDI should be equivalent to a logical 1 input. Since TCK is not switching, the JTAG interface will effectively be disabled.

11 General Purpose Timer Unit

11.1 Overview

The General Purpose Timer Unit provides a common 16-bit prescaler and four 32-bit decrementing timers. Each timer can generate an unique interrupt on underflow.
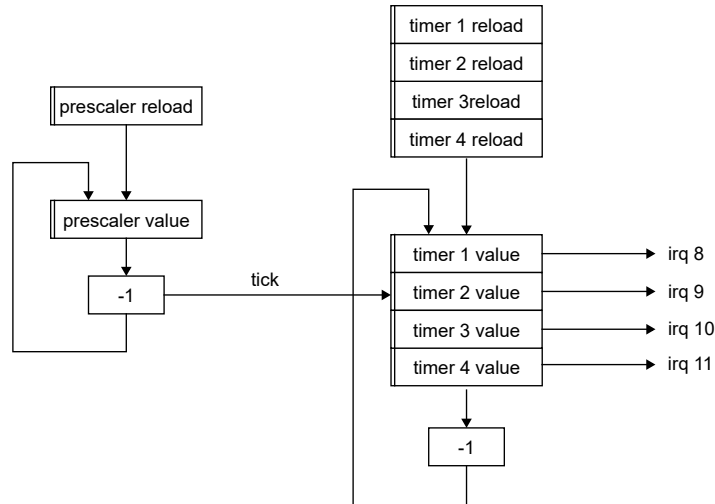
Figure 54. General Purpose Timer Unit block diagram

11.2 Operation

The prescaler is decremented on each system clock cycle. When the prescaler underflows, it is reloaded from the prescaler reload register and a timer tick is generated. The four timers are decremented on each tick. To minimize complexity, timers share the same decrementer. This means that the minimum allowed prescaler division factor is *5*.

The operation of each timer is controlled through its control register. A timer is enabled by setting the enable bit in the control register. The timer value is then decremented on each prescaler tick. When a timer underflows, it will automatically be reloaded with the value of the corresponding timer reload register if the restart bit in the control register is set, otherwise it will stop at -1 and clear the enable bit. Each timer can generate a unique interrupt when it underflows, if the interrupt enable bit is set.

By setting the chain bit in the control register, timer *n* can be chained with preceding timer *n*-1, creating a 64-bit (or larger) timer. Timer *n* will then be decremented once each time timer *n*-1 underflows.

Each timer can be reloaded with the value in its reload register at any time by writing a 'one' to the load bit in the control register.

Timer 4 also operates as a watchdog. After reset, the timer is enabled and pre-set to value 0xFFFF. When timer 4 underflows, the external WDOGN signals is asserted. This can be used to trigger a reset or other system level action. Watchdog output is equivalent to the interrupt pending bit of the last timer. The interrupt pending bit is only set when the interrupt is enabled for that timer.

Note: The WDOGN signal **cannot** be used if the system is clocked using the 2x DLL, but only when the system is clocked directly from INCLK.

11.3   Registers

The core is programmed through registers mapped at address 0x80000300.

Table 48. GPTIMER unit registers

APB address	Register	Reset value
0x80000300	Scaler value	0xFFFF
0x80000304	Scaler reload value	0xFFFF
0x80000308	Configuration register	0x0044
0x80000310	Timer 1 counter value register	-
0x80000314	Timer 1 reload value register	-
0x80000318	Timer 1 control register	IE=0, EN=0
0x80000320	Timer 2 counter value register	-
0x80000324	Timer 2 reload value register	-
0x80000328	Timer 2 control register	IE=0, EN=0
0x80000330	Timer 3 counter value register	-
0x80000334	Timer 3 reload value register	-
0x80000338	Timer 3control register	IE=0, EN=0
0x80000340	Timer 4 counter value register	0xFFFF
0x80000344	Timer 4 reload value register	0xFFFF
0x80000348	Timer 4 control register	0x0009

Table 49. Scaler value

31	16	15	0
"000..0"		SCALER VALUE	

15:    0   Scaler value

Any unused most significant bits are reserved. Always reads as '000...0'.

Table 50. Scaler reload value

31	16	15	0
"000..0"		SCALER RELOAD VALUE	

15:    0   Scaler reload value

Any unused most significant bits are reserved. Always read as '000...0'.

Table 51. GPTIMER Configuration register

31	10	9	8	7	3	2	0
"000..0"		DF	SI	IRQ		TIMERS	

31: 10   Reserved. Write only with zeroes.

9   Disable timer freeze (DF). If set the timer unit can not be freezed, otherwise timers will halt when the processor enters debug mode.

8   Separate interrupts (SI). Reads '1' to indicate the timer unit generates separate interrupts for each timer.

7: 3   Interrupt ID of first timer. Set to 8. Read-only.

2: 0   Number of implemented timers. Set to 4. Read-only.

Table 52. Timer counter value register

31	0
TIMER COUNTER VALUE	

Table 52. Timer counter value register

31:	0	Timer Counter value. Decremented by 1 for each prescaler tick.

Table 53. Timer reload value register

31	0
TIMER RELOAD VALUE	

31:	0	Timer Reload value. This value is loaded into the timer counter value register when '1' is written to load bit in the timers control register.

Table 54. Timer control register

31		7	6	5	4	3	2	1	0
"000..0"			DH	CH	IP	IE	LD	RS	EN

31: 7	Reserved. Always reads as '000...0'.
6	Debug Halt (DH): Value of GPTI.DHALT signal which is used to freeze counters (e.g. when a system is in debug mode). Read-only.
5	Chain (CH): Chain with preceding timer. If set for timer *n*, timer *n* will be decremented each time when timer (*n*-1) underflows.
4	Interrupt Pending (IP): The core sets this bit to '1' when an interrupt is signalled. This bit remains '1' until cleared by writing '0' to this bit.
3	Interrupt Enable (IE): If set the timer signals interrupt when it underflows.
2	Load (LD): Load value from the timer reload register to the timer counter value register.
1	Restart (RS): If set, the timer counter value register is reloaded with the value of the reload register when the timer underflows
0	Enable (EN): Enable the timer.

11.4 Signal definitions

The timer unit signals are described in table 55.

Table 55. Signal definitions

Signal name	Type	Function	Active
wdogn	Tri-state output	Watchdog output. Equivalent to interrupt pending bit of last timer. Can NOT be used when system clock is generated from 2x DLL.	Low

12 General Purpose Timer Unit with Time Latch Capability

12.1 Overview

The General Purpose Timer Unit with time latch capability provides a common 8-bit prescaler and two 32-bit decrementing timers. The timer value can be latched when a system interrupt occurs. Each timer can also generate an interrupt on underflow.
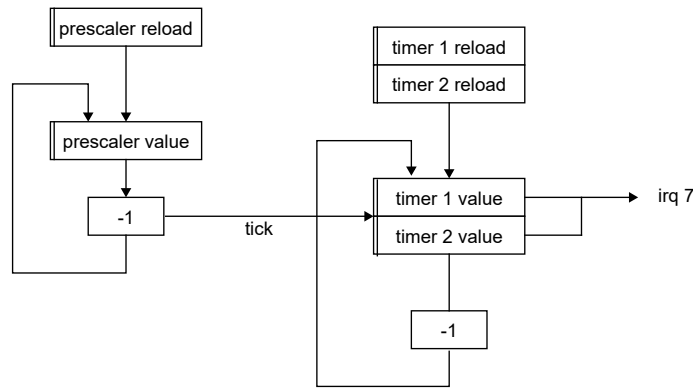
Figure 55. General Purpose Timer Unit block diagram

12.2 Operation

The prescaler is decremented on each system clock cycle. When the prescaler underflows, it is reloaded from the prescaler reload register and a timer tick is generated. The two timers are decremented each on each tick. To minimize complexity, timers share the same decrementer. This means that the minimum allowed prescaler division factor is *3*.

The operation of each timer is controlled through its control register. A timer is enabled by setting the enable bit in the control register. The timer value is then decremented on each prescaler tick. When a timer underflows, it will automatically be reloaded with the value of the corresponding timer reload register if the restart bit in the control register is set, otherwise it will stop at -1 and clear the enable bit. Each timer can generate a unique interrupt when it underflows, if the interrupt enable bit is set.

By setting the chain bit in the control register, timer 2 can be chained with timer 1, creating a 64-bit timer. Timer 2 will be decremented once each time timer 1 underflows.

Each timer can be reloaded with the value in its reload register at any time by writing a 'one' to the load bit in the control register.

Each timer can latch its value to dedicated registers when any of the 31 system interrupts occur. A dedicated mask register is provided to select the interrupts. Note that incorrect values can be latched if an incoming interrupt is occurring in the middle of the decrementing process. The likelihood is approximately 1/(prescaler reload value). Note that the latch function cannot be used with the core's own interrupt number 7.

12.3 Registers

The core is programmed through registers mapped at address 0x80100600.

Table 56. GRTIMER unit registers

APB address	Register	Reset value
0x80100600	Scaler value	0x000000FF
0x80100604	Scaler reload value	0x000000FF
0x80100608	Configuration register	0x0000003A
0x8010060C	Timer latch configuration register	0x00000000
0x80100610	Timer 1 counter value register	-
0x80100614	Timer 1 reload value register	-
0x80100618	Timer 1 control register	IE=0, EN=0
0x8010061C	Timer 1 latch register	-
0x80100620	Timer 2 counter value register	-
0x80100624	Timer 2 reload value register	-
0x80100628	Timer 2 control register	IE=0, EN=0
0x8010062C	Timer 2 latch register	-

Table 57. Scaler value

31	8	7	0
"000..0"		SCALER VALUE	

31:	8	Reserved. Always reads as '000...0'
7:	0	Scaler value

Table 58. Scaler reload value

31	8	7	0
"000..0"		SCALER RELOAD VALUE	

31:	8	Reserved. Always reads as '000...0'
7:	0	Scaler reload value

Table 59. GRTIMER Configuration register

31	12	11	10	9	8	7	3	2	0
"000..0"		EL	"0"	DF	SI	IRQ		TIMERS	

31:	12	Reserved. Always reads as '000...0'
	11	Enable latching (EL). If set, on the next matching interrupt, the latches will be loaded with the corresponding timer values. The bit is then automatically cleared, not to load a timer value until set again.
	10	Reserved. Always reads as '0'
	9	Disable timer freeze (DF). If set the timer unit can not be frozen, otherwise timers will halt when the processor enters debug mode.
	8	Separate interrupts (SI). Set to 0 to indicate common interrupt (7) for each timer. Read-only.
7:	3	Interrupt ID of first timer. Set to 7. Read-only.
2:	0	Number of implemented timers. Set to 2. Read-only.

Table 60. Timer latch configuration register

31	0
SELECT	

Table 60. Timer latch configuration register

31:	0	Specifies what bits of the AMBA interrupt shall cause the Timer Latch Register to latch the timer values. Bit 31 - 1 indicates interrupts 31 - 1

Table 61. Timer counter value registers

31	0
TIMER COUNTER VALUE	

31:	0	Timer Counter value. Decremented by 1 for each prescaler tick.

Table 62. Timer reload value registers

31	0
TIMER RELOAD VALUE	

31:	0	Timer Reload value. This value is loaded into the timer counter value register when '1' is written to load bit in the timers control register.

Table 63. Timer control registers

31	7	6	5	4	3	2	1	0
"000..0"		DH	CH	IP	IE	LD	RS	EN

31:	7	Reserved. Always reads as '000...0'
	6	Debug Halt (DH): Value of GPTI.DHALT signal which is used to freeze counters (e.g. when a system is in debug mode). Read-only.
	5	Chain (CH): Chain with preceding timer. If set for timer n, decrementing timer n begins when timer (n-1) underflows.
	4	Interrupt Pending (IP): Sets when an interrupt is signalled. Remains '1' until cleared by writing '0' to this bit.
	3	Interrupt Enable (IE): If set the timer signals interrupt when it underflows.
	2	Load (LD): Load value from the timer reload register to the timer counter value register.
	1	Restart (RS): If set the value from the timer reload register is loaded to the timer counter value register and decrementing the timer is restarted.
	0	Enable (EN): Enable the timer.

Table 64. Timer latch registers

31	0
LTCV	

31:	0	Latched Timer Counter Value (LTCV). Value latch from corresponding timer.

13 General Purpose Register

13.1 Operation

The GRGPREG provides a programmable register that controls clock generation and multiplexing.

13.2 Registers

The core is programmed through registers mapped into APB address space.

Table 65. General Purpose Register

APB address offset	Register
0x80000600	General purpose register

Table 66. I/O port data register

26 19	18	17 16	15 14	13 6	5	4	3	2	1	0
SDCLK Del.	DIVO	1553DIVL	1553DIVH	RESERVED	1553 clock	SpW rst	TM clk	SpW clk	SpW src

31 - 27	Reserved
26 - 19	Programable SDCLK delay, number of stages in the programmable delay line (0 - 255)
18	MIL-STD-1553B clock divider output: 1- system clock, 0 = divided clock
17 - 16	MIL-STD-1553B clock divider low count, clock is low DIVL+1 system clock cycles
15 - 14	MIL-STD-1553B clock divider high count, clock is high DIVH+1 system clock cycles
13 - 6	Reserved, always write all zeros to this field
5	MIL-STD-1553B clock select: 0 - generated clock, 1 - 1553CK input
4	SpaceWire DLL reset (active low)
3	TM clock select: 0 - TMCLKI, 1 - System clock
2 - 1:	SpaceWire clock select: 0 - 1X, 1 - System clock, 2 - 2X, 3 - 4X
0	SpaceWire clock source: 0 - SPWCLK, 1 - INCLK

This register resets to 0.

14 General Purpose I/O Port

14.1 Overview

The GR712RC contains two 32-bit GPIO ports, providing up to 64 controllable I/O signals.

The I/O signals also have alternative functions when connected to the various on-chip peripherals, and can also serve as external interrupts inputs.

Some bits in the I/O ports can only work as inputs, while other bits are bi-directional (see table 67). Bits 1 - 15 in each port can generate an interrupt.

Figure 56 shows a diagram for one I/O line.
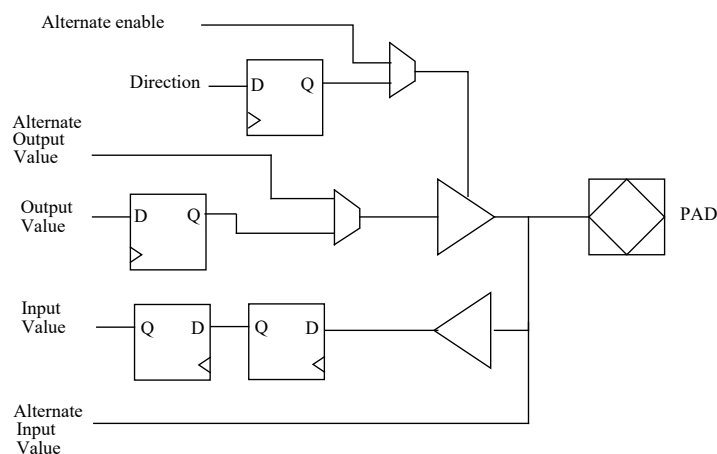
Figure 56. General Purpose I/O Port diagram

14.2 Operation

The I/O ports are implemented as bi-directional buffers with programmable output enable. The input from each buffer is synchronized by two flip-flops in series to remove potential meta-stability. The synchronized values can be read-out from the I/O port data register. The output enable is controlled by the I/O port direction register. A '1' in a bit position will enable the output buffer for the corresponding I/O line. The output value driven is taken from the I/O port output register.

Bits 1 - 15 on each I/O port can drive a separate interrupt line on the APB interrupt bus. The interrupt number is equal to the I/O line index (GPIO[1] and GPIO[33] correspond to interrupt 1, etc.). The interrupt generation is controlled by three registers: interrupt mask, polarity and edge registers. To enable an interrupt, the corresponding bit in the interrupt mask register must be set. If the edge register is '0', the interrupt is treated as level sensitive. If the polarity register is '0', the interrupt is active low. If the polarity register is '1', the interrupt is active high. If the edge register is '1', the interrupt is edge-triggered. The polarity register then selects between rising edge ('1') or falling edge ('0').

GR712RC

TABLE 67. GPIO ports and signals

GPIO port	Register bit no.	Pin function	External pin name	Interrupt	Direction
GRGPIO - 1	0	GPIO[0]	SWMX[3]		In
GRGPIO - 1	1	GPIO[1]	SWMX[4]	1	In/Out
GRGPIO - 1	2	GPIO[2]	SWMX[5]	2	In
GRGPIO - 1	3	GPIO[3]	SWMX[6]	3	In/Out
GRGPIO - 1	4	GPIO[4]	SWMX[7]	4	In
GRGPIO - 1	5	GPIO[5]	SWMX[8]	5	In/Out
GRGPIO - 1	6	GPIO[6]	SWMX[9]	6	In
GRGPIO - 1	7	GPIO[7]	SWMX[10]	7	In/Out
GRGPIO - 1	8	GPIO[8]	SWMX[11]	8	In
GRGPIO - 1	9	GPIO[9]	SWMX[12]	9	In/Out
GRGPIO - 1	10	GPIO[10]	SWMX[13]	10	In/Out
GRGPIO - 1	11	GPIO[11]	SWMX[14]	11	In
GRGPIO - 1	12	GPIO[12]	SWMX[15]	12	In
GRGPIO - 1	13	GPIO[13]	SWMX[16]	13	In/Out
GRGPIO - 1	14	GPIO[14]	SWMX[17]	14	In/Out
GRGPIO - 1	15	GPIO[15]	SWMX[18]	15	In
GRGPIO - 1	16	GPIO[16]	SWMX[19]		In
GRGPIO - 1	17	GPIO[17]	SWMX[20]		In/Out
GRGPIO - 1	18	GPIO[18]	SWMX[21]		In/Out
GRGPIO - 1	19	GPIO[19]	SWMX[22]		In
GRGPIO - 1	20	GPIO[20]	SWMX[23]		In
GRGPIO - 1	21	GPIO[21]	SWMX[24]		In/Out
GRGPIO - 1	22	GPIO[22]	SWMX[25]		In/Out
GRGPIO - 1	23	GPIO[23]	SWMX[26]		In
GRGPIO - 1	24	GPIO[24]	SWMX[27]		In
GRGPIO - 1	25	GPIO[25]	SWMX[28]		In/Out
GRGPIO - 1	26	GPIO[26]	SWMX[29]		In/Out
GRGPIO - 1	27	GPIO[27]	SWMX[30]		In
GRGPIO - 1	28	GPIO[28]	SWMX[31]		In
GRGPIO - 1	29	GPIO[29]	SWMX[32]		In/Out
GRGPIO - 1	30	GPIO[30]	SWMX[33]		In/Out
GRGPIO - 1	31	GPIO[31]	SWMX[34]		In
GRGPIO - 2	0	GPIO[32]	SWMX[35]		In
GRGPIO - 2	1	GPIO[33]	SWMX[36]	1	In/Out
GRGPIO - 2	2	GPIO[34]	SWMX[37]	2	In/Out
GRGPIO - 2	3	GPIO[35]	SWMX[38]	3	In
GRGPIO - 2	4	GPIO[36]	SWMX[39]	4	In
GRGPIO - 2	5	GPIO[37]	SWMX[40]	5	In/Out
GRGPIO - 2	6	GPIO[38]	SWMX[41]	6	In/Out
GRGPIO - 2	7	GPIO[39]	SWMX[42]	7	In
GRGPIO - 2	8	GPIO[40]	SWMX[43]	8	In/Out
GRGPIO - 2	9	GPIO[41]	SWMX[44]	9	In/Out
GRGPIO - 2	10	GPIO[42]	SWMX[45]	10	In/Out
GRGPIO - 2	11	GPIO[43]	SWMX[46]	11	In
GRGPIO - 2	12	GPIO[44]	SWMX[47]	12	In
GRGPIO - 2	13	GPIO[45]	SWMX[48]	13	In/Out
GRGPIO - 2	14	GPIO[46]	SWMX[49]	14	In/Out
GRGPIO - 2	15	GPIO[47]	SWMX[50]	15	In
GRGPIO - 2	16	GPIO[48]	SWMX[51]		In
GRGPIO - 2	17	GPIO[49]	SWMX[52]		In/Out
GRGPIO - 2	18	GPIO[50]	SWMX[53]		In/Out

TABLE 67. GPIO ports and signals

GPIO port	Register bit no.	Pin function	External pin name	Interrupt	Direction
GRGPIO - 2	19	GPIO[51]	SWMX[54]		In/Out
GRGPIO - 2	20	GPIO[52]	SWMX[55]		In
GRGPIO - 2	21	GPIO[53]	SWMX[56]		In
GRGPIO - 2	22	GPIO[54]	SWMX[57]		In/Out
GRGPIO - 2	23	GPIO[55]	SWMX[58]		In/Out
GRGPIO - 2	24	GPIO[56]	SWMX[59]		In/Out
GRGPIO - 2	25	GPIO[57]	SWMX[60]		In/Out
GRGPIO - 2	26	GPIO[58]	SWMX[61]		In/Out
GRGPIO - 2	27	GPIO[59]	SWMX[62]		In/Out
GRGPIO - 2	28	GPIO[60]	SWMX[63]		In/Out
GRGPIO - 2	29	GPIO[61]	SWMX[64]		In/Out
GRGPIO - 2	30	GPIO[62]	SWMX[65]		In/Out
GRGPIO - 2	31	GPIO[63]	SWMX[66]		In/Out

14.3 Registers

The core is programmed through registers mapped into APB address space. GRGPIO port 1 registers are mapped at address 0x80000900, while GRGPIO port 2 registers are at 0x80000A00.

Table 68. General Purpose I/O Port registers

APB address port 1	APB address port 2	Register
0x80000900	0x80000A00	I/O port data register
0x80000904	0x80000A04	I/O port output register
0x80000908	0x80000A08	I/O port direction register
0x8000090C	0x80000A0C	Interrupt mask register
0x80000910	0x80000A10	Interrupt polarity register
0x80000914	0x80000A14	Interrupt edge register

Table 69. I/O port data register

31   0

I/O port input value

31: 0   I/O port input value. Bits that are configured as outputs will always read zero from this register.

Table 70. I/O port output register

31   0

I/O port output value

31: 0   I/O port output value

Table 71. I/O port direction register

31   0

I/O port direction value

31: 0   I/O port direction value (0=output disabled, 1=output enabled). Input-only bits are don't care.

Table 72. Interrupt mask register

31   0

Interrupt mask

GR712RC

Table 72. Interrupt mask register

31:    0   Interrupt mask (0=interrupt masked, 1=interrupt enabled)

Table 73. Interrupt polarity register

31	0
Interrupt polarity	

31:    0   Interrupt polarity (0=low/falling, 1=high/rising)

Table 74. Interrupt edge register

31	0
Interrupt edge	

31:    0   Interrupt edge (0=level, 1=edge)

14.4   Signal definitions

The signals are described in table 75.

Table 75. Signal definitions

Signal name	Type	Function	Active
GPIO[63:0]	Input/Output	General purpose input output	-

15 UART Serial Interface

15.1 Overview

The GR712RC contains six UART interfaces for asynchronous serial communications. The UARTs supports data frames with 8 data bits, one optional parity bit and one stop bit. To generate the bit-rate, each UART has a programmable 12-bit clock divider. Two 8-byte FIFOs are used for data transfer between the APB bus and UART. Both odd and even parity is supported.
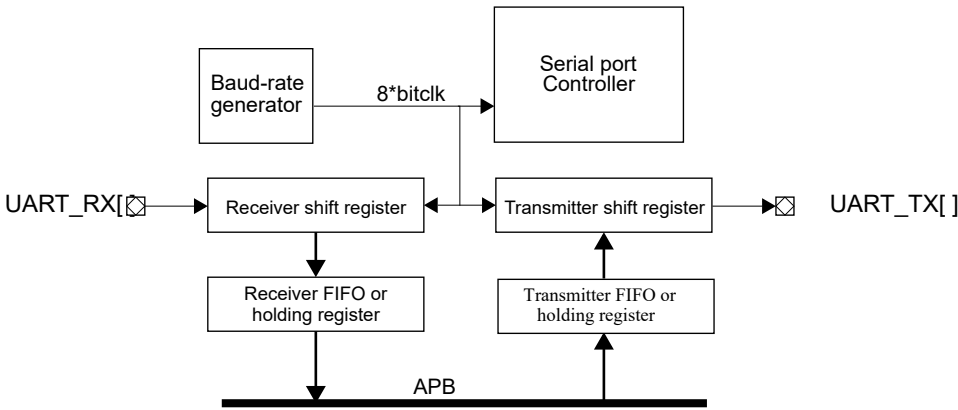
Figure 57. Block diagram

15.2 Operation

15.2.1 Transmitter operation

The transmitter is enabled through the TE bit in the UART control register. Data that is to be transferred is stored in the FIFO/holding register by writing to the data register. When ready to transmit, data is transferred from the transmitter FIFO/holding register to the transmitter shift register and converted to a serial stream on the transmitter serial output pin (TX). It automatically sends a start bit followed by eight data bits, an optional parity bit, and one stop bit (figure 58). The least significant bit of the data is sent first.

Data frame, no parity:

Start	D0	D1	D2	D3	D4	D5	D6	D7	Stop

Data frame with parity:

Start	D0	D1	D2	D3	D4	D5	D6	D7	Parity	Stop

Figure 58. UART data frames

Following the transmission of the stop bit, if a new character is not available in the transmitter FIFO, the transmitter serial data output remains high and the transmitter shift register empty bit (TS) will be

set in the UART status register. Transmission resumes and the TS is cleared when a new character is loaded into the transmitter FIFO. When the FIFO is empty the TE bit is set in the status register. If the transmitter is disabled, it will immediately stop any active transmissions including the character currently being shifted out from the transmitter shift register. The transmitter holding register may not be loaded when the transmitter is disabled or when the FIFO (or holding register) is full. If this is done, data might be overwritten and one or more frames are lost.

The TF status bit (not to be confused with the TF control bit) is set if the transmitter FIFO is currently full and the TH bit is set as long as the FIFO is *less* than half-full (less than half of entries in the FIFO contain data). The TF control bit enables FIFO interrupts when set. The status register also contains a counter (TCNT) showing the current number of data entries in the FIFO.

15.2.2 Receiver operation

The receiver is enabled for data reception through the receiver enable (RE) bit in the UART control register. The receiver looks for a high to low transition of a start bit on the receiver serial data input pin. If a transition is detected, the state of the serial input is sampled a half bit clocks later. If the serial input is sampled high the start bit is invalid and the search for a valid start bit continues. If the serial input is still low, a valid start bit is assumed and the receiver continues to sample the serial input at one bit time intervals (at the theoretical centre of the bit) until the proper number of data bits and the parity bit have been assembled and one stop bit has been detected. The serial input is shifted through an 8-bit shift register where all bits have to have the same value before the new value is taken into account, effectively forming a low-pass filter with a cut-off frequency of 1/8 system clock.

The receiver also has a configurable FIFO which is identical to the one in the transmitter. As mentioned in the transmitter part, both the holding register and FIFO will be referred to as FIFO.

During reception, the least significant bit is received first. The data is then transferred to the receiver FIFO and the data ready (DR) bit is set in the UART status register as soon as the FIFO contains at least one data frame. The parity, framing and overrun error bits are set at the received byte boundary, at the same time as the data ready would have been set. The data frame is not stored in the FIFO if an error is detected. Also, the new error status bits are or:ed with the old values before they are stored into the status register. Thus, they are not cleared until written to with zeros from the AMBA APB bus. If both the receiver FIFO and shift registers are full when a new start bit is detected, then the character held in the receiver shift register will be lost and the overrun bit will be set in the UART status register.

The RF status bit (not to be confused with the RF control bit) is set when the receiver FIFO is full. The RH status bit is set when the receiver FIFO is half-full (at least half of the entries in the FIFO contain data frames). The RF control bit enables receiver FIFO interrupts when set. A RCNT field is also available showing the current number of data frames in the FIFO.

15.3 Baud-rate generation

Each UART contains a 12-bit down-counting scaler to generate the desired baud-rate. The scaler is clocked by the system clock and generates a UART tick each time it underflows. It is reloaded with the value of the UART scaler reload register after each underflow. The resulting UART tick frequency should be 8 times the desired baud-rate.

The scalear reload value is calculated as follows (where F_{CLK} is the sytem clock frequency):

$$SCALER\_RELOAD\_VALUE = F_{CLK} / (BAUD\_RATE * 8) - 1$$

15.4 Loop back mode

If the LB bit in the UART control register is set, the UART will be in loop back mode. In this mode, the transmitter output is internally connected to the receiver input. It is then possible to perform loop back tests to verify operation of receiver, transmitter and associated software routines. In this mode, the outputs remain in the inactive state, in order to avoid sending out data.

15.5 FIFO debug mode

FIFO debug mode is entered by setting the debug mode bit in the control register. In this mode it is possible to read the transmitter FIFO and write the receiver FIFO through the FIFO debug register. The transmitter output is held inactive when in debug mode. A write to the receiver FIFO generates an interrupt if receiver interrupts are enabled.

Note that the debug mode bit is not affected by system reset and should be considered undefined by the bootloader.

15.6 Interrupt generation

For FIFOs, two different kinds of interrupts are available: normal interrupts and FIFO interrupts. For the transmitter, normal interrupts are generated when transmitter interrupts are enabled (TI), the transmitter is enabled and the transmitter FIFO goes from containing data to being empty. FIFO interrupts are generated when the FIFO interrupts are enabled (TF), transmissions are enabled (TE) and the UART is less than half-full (that is, whenever the TH status bit is set). This is a level interrupt and the interrupt signal is continuously driven high as long as the condition prevails. The receiver interrupts work in the same way. FIFO interrupts are generated when receiver FIFO interrupts are enabled, the receiver is enabled and the FIFO is half-full. The interrupt signal is continuously driven high as long as the receiver FIFO is half-full (at least half of the entries contain data frames).

15.7 Registers

The core is controlled through registers mapped into APB address space.

Table 76. UART register start address

APB address	UART device	Interrupt
0x80000100	UART 0	2
0x80100100	UART 1	17
0x80100200	UART 2	18
0x80100300	UART 3	19
0x80100400	UART 4	20
0x80100500	UART 5	21

Table 77. UART registers

APB address offset	Register
0x0	UART Data register
0x4	UART Status register
0x8	UART Control register
0xC	UART Scaler register
0x10	UART FIFO debug register

15.7.1 UART Data Register

Table 78. UART data register

31	8	7	0
RESERVED		DATA	

7: 0   Receiver holding register or FIFO (read access)

7: 0   Transmitter holding register or FIFO (write access)

15.7.2 UART Status Register

Table 79. UART status register

31	26	25	20	19	11	10	9	8	7	6	5	4	3	2	1	0
RCNT		TCNT		RESERVED		RF	TF	RH	TH	FE	PE	OV	BR	TE	TS	DR

31: 26   Receiver FIFO count (RCNT) - shows the number of data frames in the receiver FIFO.

25: 20   Transmitter FIFO count (TCNT) - shows the number of data frames in the transmitter FIFO.

10   Receiver FIFO full (RF) - indicates that the Receiver FIFO is full.

9   Transmitter FIFO full (TF) - indicates that the Transmitter FIFO is full.

8   Receiver FIFO half-full (RH) -indicates that at least half of the FIFO is holding data.

7   Transmitter FIFO half-full (TH) - indicates that the FIFO is less than half-full.

6   Framing error (FE) - indicates that a framing error was detected.

5   Parity error (PE) - indicates that a parity error was detected.

4   Overrun (OV) - indicates that one or more character have been lost due to overrun.

3   Break received (BR) - indicates that a BREAK has been received.

2   Transmitter FIFO empty (TE) - indicates that the transmitter FIFO is empty.

1   Transmitter shift register empty (TS) - indicates that the transmitter shift register is empty.

0   Data ready (DR) - indicates that new data is available in the receiver holding register

15.7.3 UART Control Register

Table 80. UART control register

31	30								13	12	11	10	9	8	7	6	5	4	3	2	1	0
FA	RESERVED									0	DB	RF	TF	0	LB	0	PE	PS	TI	RI	TE	RE

31	FIFOs available (FA) - Set to 1, read-only. Receiver and transmitter FIFOs are available.
30: 13	RESERVED
11	FIFO debug mode enable (DB) - when set, it is possible to read and write the FIFO debug register
10	Receiver FIFO interrupt enable (RF) - when set, Receiver FIFO level interrupts are enabled
9	Transmitter FIFO interrupt enable (TF) - when set, Transmitter FIFO level interrupts are enabled.
7	Loop back (LB) - if set, loop back mode will be enabled
5	Parity enable (PE) - if set, enables parity generation and checking (when implemented)
4	Parity select (PS) - selects parity polarity (0 = even parity, 1 = odd parity) (when implemented)
3	Transmitter interrupt enable (TI) - if set, interrupts are generated when a frame is transmitted
2	Receiver interrupt enable (RI) - if set, interrupts are generated when a frame is received
1	Transmitter enable (TE) - if set, enables the transmitter.
0	Receiver enable (RE) - if set, enables the receiver.

15.7.4 UART Scaler Register

Table 81. UART scaler reload register

31	12	11	0
RESERVED		SCALER RELOAD VALUE	

11: 0	Scaler reload value

15.7.5 UART FIFO Debug Register

Table 82. UART FIFO debug register

31	8	7	0
RESERVED		DATA	

7: 0	Transmitter holding register or FIFO (read access)
7: 0	Receiver holding register or FIFO (write access)

15.8 Signal definitions

The signals are described in table 83.

Table 83. Signal definitions

Signal name	Type	Function	Active
UART_TX[5:0]	Output	UART transmit data line	-
UART_RX[5:0]	Input	UART receive data line	-

16 SpaceWire Interface with RMAP support

16.1 Overview

The SpaceWire core provides an interface between the AHB bus and a SpaceWire network. It implements the SpaceWire standard [SPW] with the protocol identification extension [RMAPID]. The Remote Memory Access Protocol (RMAP) command handler implements the ECSS standard [RMAP].

The SpaceWire interface is configured through a set of registers accessed through an APB interface. Data is transferred through DMA channels using an AHB master interface.

GR712RC includes six GRSPW2 cores. Only cores GRSPW2-0 and GRSPW2-1 include RMAP support. Each core includes a single DMA channel. Each core supports a single port.

Note that the subsequent sections are written in a generic manner as if there were several DMA channels, even though there is only one channel per core.
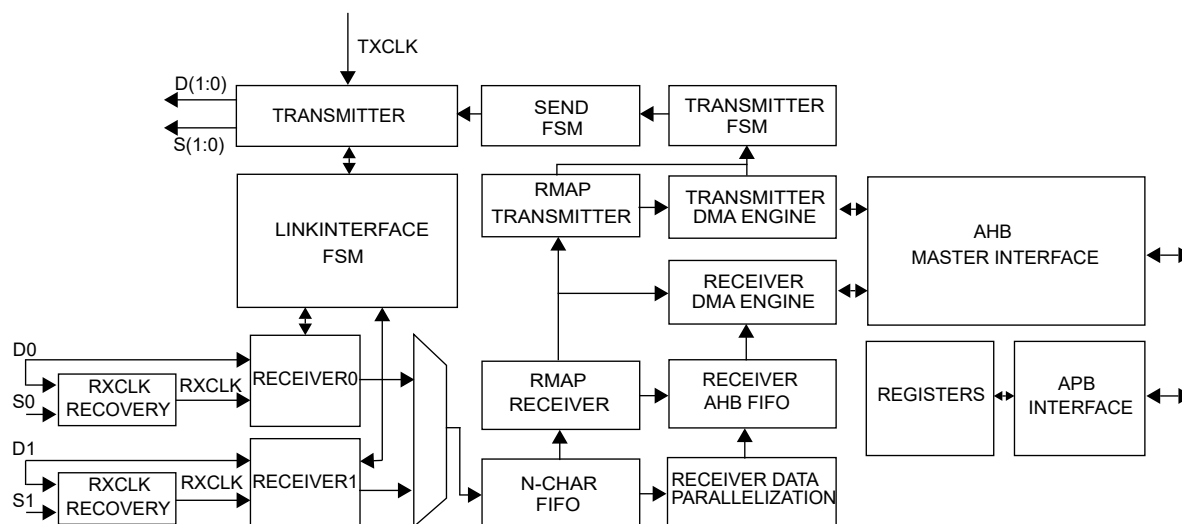
Figure 59. Block diagram

16.2 Operation

16.2.1 Overview

The GRSPW can be split into three main parts: the link interface, the AMBA interface and the RMAP handler. A block diagram of the internal structure can be found in figure 59.

The link interface consists of the receiver, transmitter and the link interface FSM. They handle communication on the SpaceWire network. The AMBA interface consists of the DMA engines, the AHB master interface and the APB interface. The link interface provides FIFO interfaces to the DMA engines. These FIFOs are used to transfer N-Chars between the AMBA and SpaceWire domains during reception and transmission.

The RMAP handler handles incoming packets which are determined to be RMAP commands instead of the receiver DMA engine. The RMAP command is decoded and if it is valid, the operation is performed on the AHB bus. If a reply was requested it is automatically transmitted back to the source by the RMAP transmitter.

16.2.2 Protocol support

The GRSPW only accepts packets with a valid destination address in the first received byte. Packets with address mismatch will be silently discarded (except in promiscuous mode which is covered in section 16.4.10).

The second byte is always interpreted as a protocol ID. The only protocol handled separately in hardware is the RMAP protocol (ID=0x1) while other packets are stored to a DMA channel. If the RMAP command handler is present and enabled all RMAP commands will be processed, executed and replied automatically in hardware. Otherwise RMAP commands are stored to a DMA channel in the same way as other packets. RMAP replies are always stored to a DMA channel. there is no need to include a protocol ID in the packets and the data can start immediately after the address.

All packets arriving with the extended protocol ID (0x00) are stored to a DMA channel. This means that the hardware RMAP command handler will not work if the incoming RMAP packets use the extended protocol ID. Note also that packets with the reserved extended protocol identifier (ID = 0x000000) are not ignored by the GRSPW. It is up to the client receiving the packets to ignore them.

When transmitting packets, the address and protocol-ID fields must be included in the buffers from where data is fetched. They are *not* automatically added by the GRSPW.

Figure 60 shows a packet with a normal protocol identifier. The GRSPW also allows reception and transmission with extended protocol identifiers but then the hardware RMAP and RMAP CRC calculations will not work.

| Addr | ProtID | D0 | D1 | D2 | D3 | .. | Dn-2 | Dn-1 | EOP |

Figure 60. The SpaceWire packet with protocol ID that is expected by the GRSPW.

16.3 Link interface

The link interface handles the communication on the SpaceWire network and consists of a transmitter, receiver, a FSM and FIFO interfaces. An overview of the architecture is found in figure 59.

16.3.1 Link interface FSM

The FSM controls the link interface (a more detailed description is found in the SpaceWire standard). The low-level protocol handling (the signal and character level of the SpaceWire standard) is handled by the transmitter and receiver while the FSM handles the exchange level.

The link interface FSM is controlled through the control register. The link can be disabled through the link disable bit, which depending on the current state, either prevents the link interface from reaching the started state or forces it to the error-reset state. When the link is not disabled, the link interface FSM is allowed to enter the started state when either the link start bit is set or when a NULL character has been received and the autostart bit is set.

The current state of the link interface determines which type of characters are allowed to be transmitted which together with the requests made from the host interfaces determine what character will be sent.

Time-codes are sent when the FSM is in the run-state and a request is made through the time-interface (described in section 16.3.4).

When the link interface is in the connecting- or run-state it is allowed to send FCTs. FCTs are sent automatically by the link interface when possible. This is done based on the maximum value of 56 for the outstanding credit counter and the currently free space in the receiver N-Char FIFO. FCTs are sent as long as the outstanding counter is less than or equal to 48 and there are at least 8 more empty FIFO entries than the counter value.

N-Chars are sent in the run-state when they are available from the transmitter FIFO and there are credits available. NULLs are sent when no other character transmission is requested or the FSM is in a state where no other transmissions are allowed.

The credit counter (incoming credits) is automatically increased when FCTs are received and decreased when N-Chars are transmitted. Received N-Chars are stored to the receiver N-Char FIFO for further handling by the DMA interface. Received Time-codes are handled by the time-interface.

16.3.2 Transmitter

The state of the FSM, credit counters, requests from the time-interface and requests from the DMA-interface are used to decide the next character to be transmitted. The type of character and the character itself (for N-Chars and Time-codes) to be transmitted are presented to the low-level transmitter which is located in a separate clock-domain.

This is done because one usually wants to run the SpaceWire link on a different frequency than the host system clock. The GRSPW has a separate clock input which is used to generate the transmitter clock. Since the transmitter often runs on high frequency clocks (> 100 MHz) as much logic as possible has been placed in the system clock domain to minimize power consumption and timing issues.

The transmitter logic in the host clock domain decides what character to send next and sets the proper control signal and presents any needed character to the low-level transmitter as shown in figure 61. The transmitter sends the requested characters and generates parity and control bits as needed. If no requests are made from the host domain, NULLs are sent as long as the transmitter is enabled. Most of the signal and character levels of the SpaceWire standard is handled in the transmitter. External LVDS drivers are needed for the data and strobe signals.

Figure 61. Schematic of the link interface transmitter.

A transmission FSM reads N-Chars for transmission from the transmitter FIFO. It is given packet lengths from the DMA interface and appends EOPs/EEPs and RMAP CRC values if requested. When it is finished with a packet the DMA interface is notified and a new packet length value is given.

16.3.3 Receiver

The receiver detects connections from other nodes and receives characters as a bit stream on the data and strobe signals.

The receiver is activated as soon as the link interface leaves the error reset state. Then after a NULL is received it can start receiving any characters. It detects parity, escape and credit errors which causes the link interface to enter the error reset state. Disconnections are handled in the link interface part in the tx clock domain because no receiver clock is available when disconnected.

Received Characters are flagged to the host domain and the data is presented in parallel form. The interface to the host domain is shown in figure 62. L-Chars are the handled automatically by the host domain link interface part while all N-Chars are stored in the receiver FIFO for further handling. If two or more consecutive EOPs/EEPs are received all but the first are discarded.

Figure 62. Schematic of the link interface receiver.

16.3.4 Time interface

The time interface is used for sending Time-codes over the SpaceWire network and consists of a time-counter register, time-ctrl register, tick-in signal, tick-out signal, tick-in register field and a tick-out register field. There are also two control register bits which enable the time receiver and transmitter respectively.

Each Time-code sent from the GRSPW2 is a concatenation of the time-ctrl and the time-counter register. There is a timetxen bit which is used to enable Time-code transmissions. It is not possible to send time-codes if this bit is zero.

Received Time-codes are stored to the same time-ctrl and time-counter registers which are used for transmission. The timerxen bit in the control register is used for enabling time-code reception. No time-codes will be received if this bit is zero.

The two enable bits are used for ensuring that a node will not (accidentally) both transmit and receive time-codes which violates the SpaceWire standard. It also ensures that a the master sending time-codes on a network will not have its time-counter overwritten if another (faulty) node starts sending time-codes.

The time-counter register is set to 0 after reset and is incremented each time the tick-in signal is asserted for one clock-period and the timetxen bit is set. This also causes the link interface to send the new value on the network. Tick-in can be generated by writing a one to the register field. A Tick-in should not be generated too often since if the time-code after the previous Tick-in has not been sent the register will not be incremented and no new value will be sent. The tick-in field is automatically cleared when the value has been sent and thus no new ticks should be generated until this field is zero. If the tick-in signal is used there should be at least 4 system-clock and 25 transmit-clock cycles between each assertion.

A tick-out is generated each time a valid time-code is received and the timerxen bit is set. When the tick-out is generated the tick-out signal will be asserted one clock-cycle and the tick-out register field is asserted until it is cleared by writing a one to it.

The current time counter value can be read from the time register. It is updated each time a Time-code is received and the timerxen bit is set. The same register is used for transmissions and can also be written directly from the APB interface.

The control bits of the Time-code are stored to the time-ctrl register when a Time-code is received whose time-count is one more than the nodes current time-counter register. The time-ctrl register can be read through the APB interface. The same register is used during time-code transmissions.

It is possible to have both the time-transmission and reception functions enabled at the same time.

16.4 Receiver DMA channels

The receiver DMA engine handles reception of data from the SpaceWire network to different DMA channels.

16.4.1 Address comparison and channel selection

Packets are received to different channels based on the address and whether a channel is enabled or not. When the receiver N-Char FIFO contains one or more characters, N-Chars are read by the receiver DMA engine. The first character is interpreted as the logical address and is compared with the addresses of each channel starting from 0. The packet will be stored to the first channel with an matching address. The complete packet including address and protocol ID but excluding EOP/EEP is stored to the memory address pointed to by the descriptors (explained later in this section) of the channel.

Each SpaceWire address register has a corresponding mask register. Only bits at an index containing a zero in the corresponding mask register are compared. This way a DMA channel can accept a range of addresses. There is a default address register which is used for address checking of RMAP commands in the RMAP command handler and all implemented DMA channels that do not have separate addressing enabled. With separate addressing enabled the DMA channels' own address/mask register pair is used instead.

If an RMAP command is received it is only handled by the command handler if the default address register (including mask) matches the received address. Otherwise the packet will be stored to a DMA channel if one or more of them has a matching address. If the address does not match neither the default address nor one of the DMA channels' separate register, the packet is still handled by the RMAP command handler if enabled since it has to return the invalid address error code. The packet is only discarded (up to and including the next EOP/EEP) if an address match cannot be found and the RMAP command handler is disabled or not present. Figure 63 shows a flowchart of packet reception.

At least 2 non EOP/EEP N-Chars needs to be received for a packet to be stored to the DMA channel. If it is an RMAP packet with hardware RMAP enabled 3 N-Chars are needed since the command byte determines where the packet is processed. Packets smaller than these sizes are discarded.
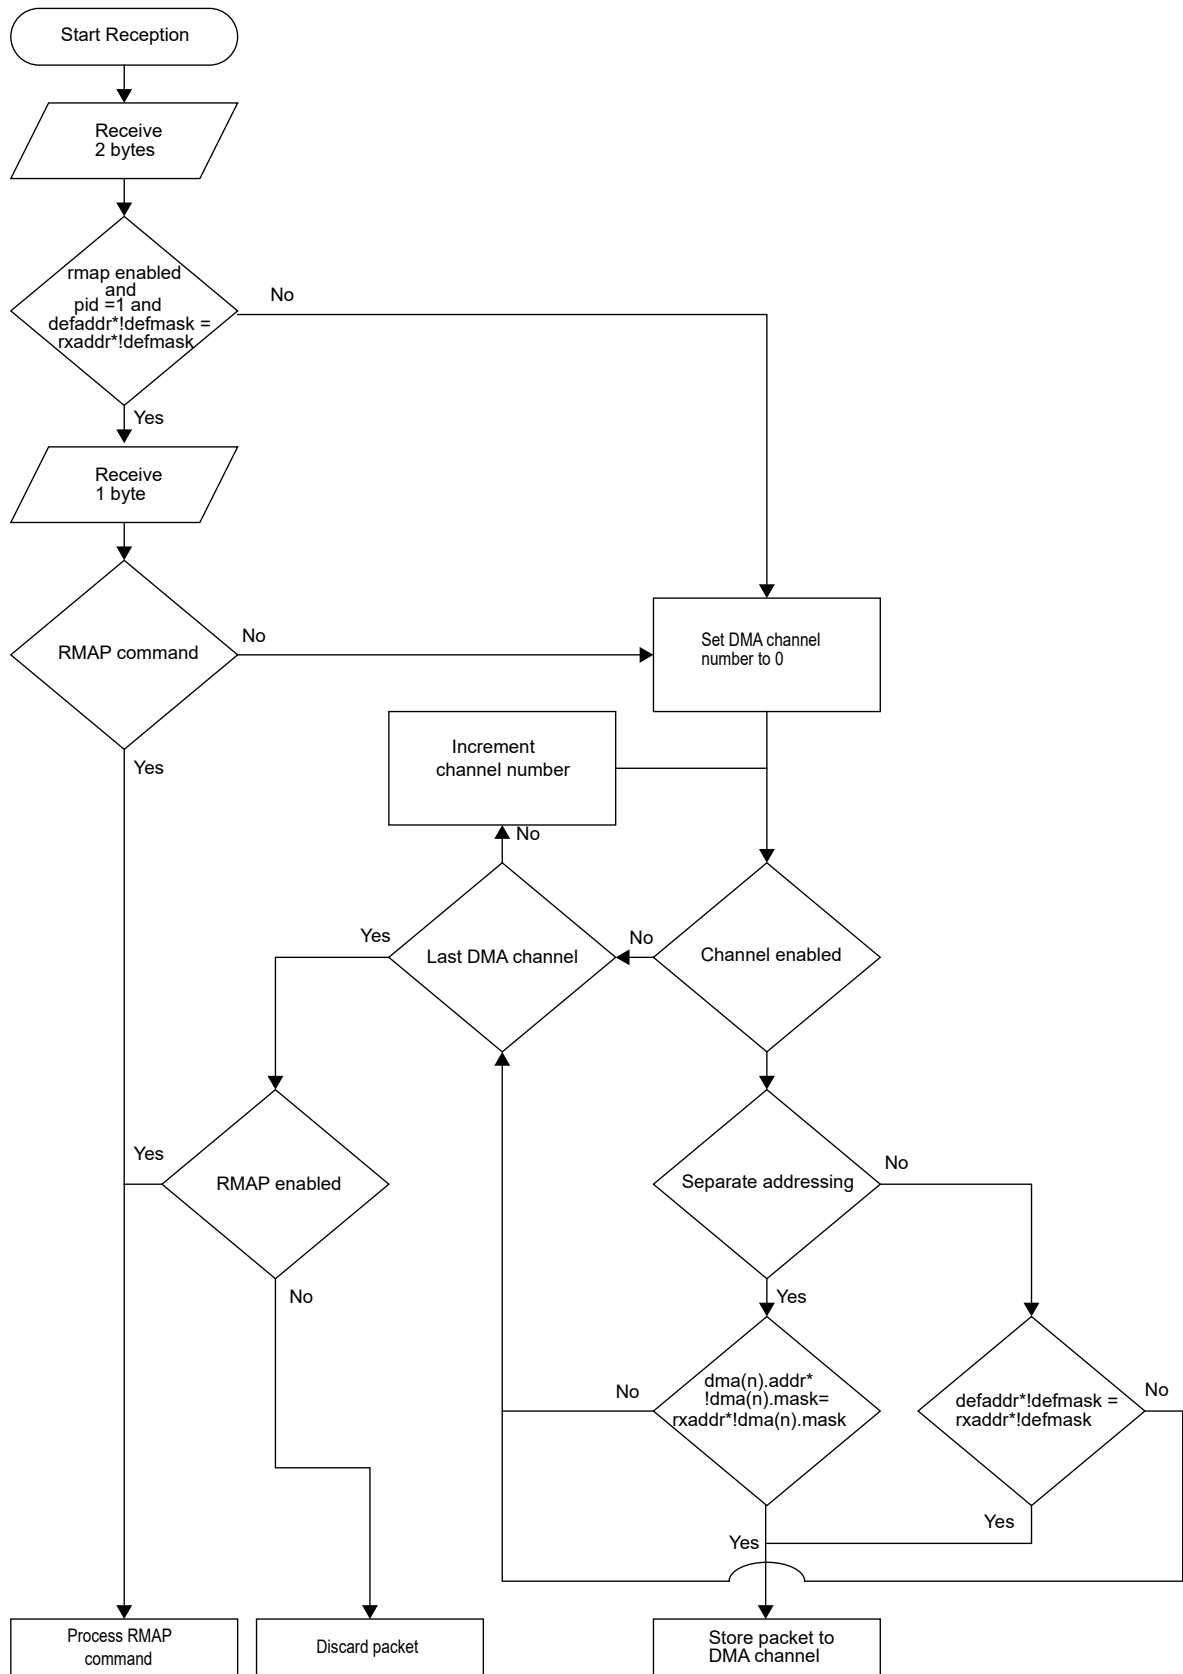
Figure 63. Flow chart of packet reception.

16.4.2 Basic functionality of a channel

Reception is based on descriptors located in a consecutive area in memory that hold pointers to buffers where packets should be stored. When a packet arrives at the GRSPW the channel which should receive it is first determined as described in the previous section. A descriptor is then read from the channels' descriptor area and the packet is stored to the memory area pointed to by the descriptor. Lastly, status is stored to the same descriptor and increments the descriptor pointer to the next one. The following sections will describe DMA channel reception in more detail.

16.4.3 Setting up the GRSPW for reception

A few registers need to be initialized before reception to a channel can take place. First the link interface need to be put in the run state before any data can be sent. The DMA channel has a maximum length register which sets the maximum packet size in bytes that can be received to this channel. Larger packets are truncated and the excessive part is spilled. If this happens an indication will be given in the status field of the descriptor. The minimum value for the receiver maximum length field is 4 and the value can only be incremented in steps of four bytes up to the maximum value 33554428. If the maximum length is set to zero the receiver will *not* function correctly.

Either the default address register or the channel specific address register (the accompanying mask register must also be set) needs to be set to hold the address used by the channel. A control bit in the DMA channel control register determines whether the channel should use default address and mask registers for address comparison or the channel's own registers. Using the default register the same address range is accepted as for RMAP commands while the separate provides the channel it own range. If all channels use the default registers they will accept the same address range and the enabled channel with the lowest number will receive the packet.

Finally, the descriptor table and control register must be initialized. This will be described in the two following sections.

16.4.4 Setting up the descriptor table address

The GRSPW reads descriptors from an area in memory pointed to by the receiver descriptor table address register. The register consists of a base address and a descriptor selector. The base address points to the beginning of the area and must start on a 1024 bytes aligned address. It is also limited to be 1024 bytes in size which means the maximum number of descriptors is 128 since the descriptor size is 8 bytes.

The descriptor selector points to individual descriptors and is increased by 1 when a descriptor has been used. When the selector reaches the upper limit of the area it wraps to the beginning automatically. It can also be set to wrap at a specific descriptor before the upper limit by setting the wrap bit in the descriptor. The idea is that the selector should be initialized to 0 (start of the descriptor area) but it can also be written with another 8 bytes aligned value to start somewhere in the middle of the area. It will still wrap to the beginning of the area.

If one wants to use a new descriptor table the receiver enable bit has to be cleared first. When the rxactive bit for the channel is cleared it is safe to update the descriptor table register. When this is finished and descriptors are enabled the receiver enable bit can be set again.

16.4.5 Enabling descriptors

As mentioned earlier one or more descriptors must be enabled before reception can take place. Each descriptor is 8 byte in size and the layout can be found in the tables below. The descriptors should be written to the memory area pointed to by the receiver descriptor table address register. When new descriptors are added they must always be placed after the previous one written to the area. Otherwise they will not be noticed.

A descriptor is enabled by setting the address pointer to point at a location where data can be stored and then setting the enable bit. The WR bit can be set to cause the selector to be set to zero when

reception has finished to this descriptor. IE should be set if an interrupt is wanted when the reception has finished. The DMA control register interrupt enable bit must also be set for an interrupt to be generated.

The descriptor packet address should be word aligned. All accesses on the bus are word accesses so complete words will always be overwritten regardless of whether all 32-bit contain received data. Also if the packet does not end on a word boundary the complete word containing the last data byte will be overwritten.

Table 84. GRSPW receive descriptor word 0 (address offset 0x0)

31	30	29	28	27	26	25	24	0
TR	DC	HC	EP	IE	WR	EN	PACKETLENGTH	

31	Truncated (TR) - Packet was truncated due to maximum length violation.
30	Data CRC (DC) - 1 if a CRC error was detected for the data and 0 otherwise.
29	Header CRC (HC) - 1 if a CRC error was detected for the header and 0 otherwise.
28	EEP termination (EP) - This packet ended with an Error End of Packet character.
27	Interrupt enable (IE) - If set, an interrupt will be generated when a packet has been received if the receive interrupt enable bit in the DMA channel control register is set.
26	Wrap (WR) - If set, the next descriptor used by the GRSPW will be the first one in the descriptor table (at the base address). Otherwise the descriptor pointer will be increased with 0x8 to use the descriptor at the next higher memory location. The descriptor table is limited to 1 KiB in size and the pointer will be automatically wrap back to the base address when it reaches the 1 KiB boundary.
25	Enable (EN) - Set to one to activate this descriptor. This means that the descriptor contains valid control values and the memory area pointed to by the packet address field can be used to store a packet.
24: 0	Packet length (PACKETLENGTH) - The number of bytes received to this buffer. Only valid after EN has been set to 0 by the GRSPW.

Table 85. GRSPW receive descriptor word 1 (address offset 0x4)

31	0
PACKETADDRESS	

31: 0	Packet address (PACKETADDRESS) - The address pointing at the buffer which will be used to store the received packet.

16.4.6 Setting up the DMA control register

The final step to receive packets is to set the control register in the following steps: The receiver must be enabled by setting the rxen bit in the DMA control register. This can be done anytime and before this bit is set nothing will happen. The rxdescav bit in the DMA control register is then set to indicate that there are new active descriptors. This must always be done after the descriptors have been enabled or the GRSPW might not notice the new descriptors. More descriptors can be activated when reception has already started by enabling the descriptors and writing the rxdescav bit. When these bits are set reception will start immediately when data is arriving.

16.4.7 The effect to the control bits during reception

When the receiver is disabled all packets going to the DMA-channel are discarded if the packet's address does not fall into the range of another DMA channel. If the receiver is enabled and the address falls into the accepted address range, the next state is entered where the rxdescav bit is checked. This bit indicates whether there are active descriptors or not and should be set by the external application using the DMA channel each time descriptors are enabled as mentioned above. If the rxdescav bit is '0' and the nospill bit is '0' the packets will be discarded. If nospill is one the GRSPW waits until rxdescav is set and the characters are kept in the N-Char fifo during this time. If the fifo becomes full further N-char transmissions are inhibited by stopping the transmission of FCTs.

When rxdescav is set the next descriptor is read and if enabled the packet is received to the buffer. If the read descriptor is not enabled, rxdescav is set to '0' and the packet is spilled depending on the value of nospill.

The receiver can be disabled at any time and will stop packets from being received to this channel. If a packet is currently received when the receiver is disabled the reception will still be finished. The rxdescav bit can also be cleared at any time. It will not affect any ongoing receptions but no more descriptors will be read until it is set again. Rxdescav is also cleared by the GRSPW when it reads a disabled descriptor.

16.4.8 Status bits

When the reception of a packet is finished the enable (EN) bit in the current descriptor is cleared to zero. Notet that alos the interrupt enable (IE) and Wrap (WR) bits are cleared to zero. When enable is zero, the status bits are also valid and the number of received bytes is indicated in the length field. The DMA control register contains a status bit which is set each time a packet has been received. The GRSPW can also be made to generate an interrupt for this event as mentioned in section.

RMAP CRC is always checked for all packets when CRC logic is included in the implementation . All bytes in the packet except the EOP/EEP is included in the CRC calculation. The packet is always assumed to be a RMAP packet and the length of the header is determined by checking byte 3 which should be the command field. The calculated CRC value is then checked when the header has been received (according to the calculated number of bytes) and if it is non-zero the HC bit is set indicating a header CRC error.

The CRC value is not set to zero after the header has been received, instead the calculation continues in the same way until the complete packet has been received. Then if the CRC value is non-zero the DC bit is set indicating a data CRC error. This means that the GRSPWcan indicate a data CRC error even if the data field was correct when the header CRC was incorrect. However, the data should not be used when the header is corrupt and therefore the DC bit is unimportant in this case. When the header is not corrupted the CRC value will always be zero when the calculation continues with the data field and the behaviour will be as if the CRC calculation was restarted

If the received packet is not of RMAP type the header CRC error indication bit cannot be used. It is still possible to use the DC bit if the complete packet is covered by a CRC calculated using the RMAP CRC definition. This is because the GRSPW does not restart the calculation after the header has been received but instead calculates a complete CRC over the packet. Thus any packet format with one CRC at the end of the packet calculated according to RMAP standard can be checked using the DC bit.

If the packet is neither of RMAP type nor of the type above with RMAP CRC at the end, then both the HC and DC bits should be ignored.

16.4.9 Error handling

If a packet reception needs to be aborted because of congestion on the network, the suggested solution is to set link disable to '1'. Unfortunately, this will also cause the packet currently being transmitted to be truncated but this is the only safe solution since packet reception is a passive operation depending on the transmitter at the other end. A channel reset bit could be provided but is not a satisfactory solution since the untransmitted characters would still be in the transmitter node. The next character (somewhere in the middle of the packet) would be interpreted as the node address which would probably cause the packet to be discarded but not with 100% certainty. Usually this action is performed when a reception has stuck because of the transmitter not providing more data. The channel reset would not resolve this congestion.

If an AHB error occurs during reception the current packet is spilled up to and including the next EEP/EOP and then the currently active channel is disabled and the receiver enters the idle state. A bit in the channels control/status register is set to indicate this condition.

16.4.10 Promiscuous mode

The GRSPW supports a promiscuous mode where all the data received is stored to the first DMA channel enabled regardless of the node address and possible early EOPs/EEPs. This means that all non-eop/eep N-Chars received will be stored to the DMA channel. The rxmaxlength register is still checked and packets exceeding this size will be truncated.

If the RMAP handler is present, RMAP commands will still be handled by it when promiscuous mode is enabled if the rmapen bit is set. If it is cleared, RMAP commands will also be stored to a DMA channel.

16.5   Transmitter DMA channels

The transmitter DMA engine handles transmission of data from the DMA channels to the SpaceWire network. Each receive channel has a corresponding transmit channel which means there can be up to 4 transmit channels. It is however only necessary to use a separate transmit channel for each receive channel if there are also separate entities controlling the transmissions. The use of a single channel with multiple controlling entities would cause them to corrupt each other's transmissions. A single channel is more efficient and should be used when possible.

Multiple transmit channels with pending transmissions are arbitrated in a round-robin fashion.

16.5.1   Basic functionality of a channel

A transmit DMA channel reads data from the AHB bus and stores them in the transmitter FIFO for transmission on the SpaceWire network. Transmission is based on the same type of descriptors as for the receiver and the descriptor table has the same alignment and size restrictions. When there are new descriptors enabled the GRSPW reads them and transfer the amount data indicated.

16.5.2   Setting up the GRSPW for transmission

Four steps need to be performed before transmissions can be done with the GRSPW. First the link interface must be enabled and started by writing the appropriate value to the ctrl register. Then the address to the descriptor table needs to be written to the transmitter descriptor table address register and one or more descriptors must also be enabled in the table. Finally, the txen bit in the DMA control register is written with a one which triggers the transmission. These steps will be covered in more detail in the next sections.

16.5.3   Enabling descriptors

The descriptor table address register works in the same way as the receiver's corresponding register which was covered in section 16.4. The maximum size is 1024 bytes as for the receiver but since the descriptor size is 16 bytes the number of descriptors is 64.

To transmit packets one or more descriptors have to be initialized in memory which is done in the following way: The number of bytes to be transmitted and a pointer to the data has to be set. There are two different length and address fields in the transmit descriptors because there are separate pointers for header and data. If a length field is zero the corresponding part of a packet is skipped and if both are zero no packet is sent. The maximum header length is 255 bytes and the maximum data length is 16 MiB - 1. When the pointer and length fields have been set the enable bit should be set to enable the descriptor. This must always be done last. The other control bits must also be set before enabling the descriptor.

The transmit descriptors are 16 bytes in size so the maximum number in a single table is 64. The different fields of the descriptor together with the memory offsets are shown in the tables below.

The HC bit should be set if RMAP CRC should be calculated and inserted for the header field and correspondingly the DC bit should be set for the data field. This field is only used by the GRSPW when the CRC logic is available. The header CRC will be calculated from the data fetched from the

header pointer and the data CRC is generated from data fetched from the data pointer. The CRCs are appended after the corresponding fields. The NON-CRC bytes field is set to the number of bytes in the beginning of the header field that should not be included in the CRC calculation.

The CRCs are sent even if the corresponding length is zero, but when both lengths are zero no packet is sent not even an EOP.

16.5.4 Starting transmissions

When the descriptors have been initialized, the transmit enable bit in the DMA control register has to be set to tell the GRSPW to start transmitting. New descriptors can be activated in the table on the fly (while transmission is active). Each time a set of descriptors is added the transmit enable register bit should be set. This has to be done because each time the GRSPW encounters a disabled descriptor this register bit is set to 0.

Table 86. GRSPW transmit descriptor word 0 (address offset 0x0)

31 ... 18	17	16	15	14	13	12	11 ... 8	7 ... 0
RESERVED	DC	HC	LE	IE	WR	EN	NONCRCLEN	HEADERLEN

31: 18	RESERVED
17	Append data CRC (DC) - Append CRC calculated according to the RMAP specification after the data sent from the data pointer. The CRC covers all the bytes from this pointer. A null CRC will be sent if the length of the data field is zero.
16	Append header CRC (HC) - Append CRC calculated according to the RMAP specification after the data sent from the header pointer. The CRC covers all bytes from this pointer except a number of bytes in the beginning specified by the non-crc bytes field. The CRC will not be sent if the header length field is zero.
15	Link error (LE) - A Link error occurred during the transmission of this packet.
14	Interrupt enable (IE) - If set, an interrupt will be generated when the packet has been transmitted and the transmitter interrupt enable bit in the DMA control register is set.
13	Wrap (WR) - If set, the descriptor pointer will wrap and the next descriptor read will be the first one in the table (at the base address). Otherwise the pointer is increased with 0x10 to use the descriptor at the next higher memory location.
12	Enable (EN) - Enable transmitter descriptor. When all control fields (address, length, wrap and crc) are set, this bit should be set. While the bit is set the descriptor should not be touched since this might corrupt the transmission. The GRSPW clears this bit when the transmission has finished.
11: 8	Non-CRC bytes (NONCRCLEN)- Sets the number of bytes in the beginning of the header which should not be included in the CRC calculation. This is necessary when using path addressing since one or more bytes in the beginning of the packet might be discarded before the packet reaches its destination.
7: 0	Header length (HEADERLEN) - Header Length in bytes. If set to zero, the header is skipped.

Table 87. GRSPW transmit descriptor word 1 (address offset 0x4)

31 ... 0
HEADERADDRESS

31: 0	Header address (HEADERADDRESS) - Address from where the packet header is fetched. Does not need to be word aligned.

Table 88. GRSPW transmit descriptor word 2 (address offset 0x8)

31 ... 24	23 ... 0
RESERVED	DATALEN

Table 88. GRSPW transmit descriptor word 2 (address offset 0x8)

31: 24	RESERVED
23: 0	Data length (DATALEN) - Length of data part of packet. If set to zero, no data will be sent. If both data- and header-lengths are set to zero no packet will be sent.

Table 89. GRSPW transmit descriptor word 3(address offset 0xC)

31	0
DATAADDRESS	

31: 0	Data address (DATAADDRESS) - Address from where data is read. Does not need to be word aligned.

16.5.5 The transmission process

When the txen bit is set the GRSPW starts reading descriptors immediately. The number of bytes indicated are read and transmitted. When a transmission has finished, status will be written to the first field of the descriptor and a packet sent bit is set in the DMA control register. If an interrupt was requested it will also be generated. Then a new descriptor is read and if enabled a new transmission starts, otherwise the transmit enable bit is cleared and nothing will happen until it is enabled again.

16.5.6 The descriptor table address register

The internal pointer which is used to keep the current position in the descriptor table can be read and written through the APB interface. This pointer is set to zero during reset and is incremented each time a descriptor is used. It wraps automatically when the 1024 bytes limit for the descriptor table is reached or it can be set to wrap earlier by setting a bit in the current descriptor.

The descriptor table register can be updated with a new table anytime when no transmission is active. No transmission is active if the transmit enable bit is zero and the complete table has been sent or if the table is aborted (explained below). If the table is aborted one has to wait until the transmit enable bit is zero before updating the table pointer.

16.5.7 Error handling

Abort Tx

The DMA control register contains a bit called Abort TX which if set causes the current transmission to be aborted, the packet is truncated and an EEP is inserted. This is only useful if the packet needs to be aborted because of congestion on the SpaceWire network. If the congestion is on the AHB bus this will not help (This should not be a problem since AHB slaves should have a maximum of 16 wait-states). The aborted packet will have its LE bit set in the descriptor. The transmit enable register bit is also cleared and no new transmissions will be done until the transmitter is enabled again.

AHB error

When an AHB error is encountered during transmission the currently active DMA channel is disabled and the transmitter goes to the idle mode. A bit in the DMA channel's control/status register is set to indicate this error condition and, if enabled, an interrupt will also be generated. Further error handling depends on what state the transmitter DMA engine was in when the AHB error occurred. If the descriptor was being read the packet transmission had not been started yet and no more actions need to be taken.

If the AHB error occurs during packet transmission the packet is truncated and an EEP is inserted. Lastly, if it occurs when status is written to the descriptor the packet has been successfully transmitted but the descriptor is not written and will continue to be enabled (this also means that no error bits are set in the descriptor for AHB errors).

The client using the channel has to correct the AHB error condition and enable the channel again. No more AHB transfers are done again from the same unit (receiver or transmitter) which was active during the AHB error until the error state is cleared and the unit is enabled again.

Link error

When a link error occurs during the transmission the remaining part of the packet is discarded up to and including the next EOP/EEP. When this is done status is immediately written (with the LE bit set) and the descriptor pointer is incremented. The link will be disconnected when the link error occurs but the GRSPW will automatically try to connect again provided that the link-start bit is asserted and the link-disabled bit is deasserted. If the LE bit in the DMA channel's control register is not set the transmitter DMA engine will wait for the link to enter run-state and start a new transmission immediately when possible if packets are pending. Otherwise the transmitter will be disabled when a link error occurs during the transmission of the current packet and no more packets will be transmitted until it is enabled again.

16.6 RMAP

The Remote Memory Access Protocol (RMAP) is used to implement access to resources in the node via the SpaceWire Link. Some common operations are reading and writing to memory, registers and FIFOs. The GRSPW has an optional hardware RMAP command handler which on the GR712RC is

enabled only on core GRSPW2-0 and on core GRSPW-1. This section describes the basics of the RMAP protocol and the command handler implementation.

16.6.1 Fundamentals of the protocol

RMAP is a protocol which is designed to provide remote access via a SpaceWire network to memory mapped resources on a SpaceWire node. It has been assigned protocol ID 0x01. It provides three operations write, read and read-modify-write. These operations are posted operations which means that a source does not wait for an acknowledge or reply. It also implies that any number of operations can be outstanding at any time and that no timeout mechanism is implemented in the protocol. Time-outs must be implemented in the user application which sends the commands. Data payloads of up to 16 MiB - 1 is supported in the protocol. A destination can be requested to send replies and to verify data before executing an operation. A complete description of the protocol is found in the RMAP standard.

16.6.2 Implementation

The GRSPW includes an handler for RMAP commands which processes all incoming packets with protocol ID = 0x01, type field (bit 7 and 6 of the 3rd byte in the packet) equal to 01b and an address falling in the range set by the default address and mask register. When such a packet is detected it is not stored to the DMA channel, instead it is passed to the RMAP receiver.

The GRSPW implements all three commands defined in the standard with some restrictions. The implementation is based on draft F of the RMAP standard (the only exception being that error code 12 is not implemented). Support is only provided for 32-bit big-endian systems. This means that the first byte received is the msb in a word. The command handler will not receive RMAP packets using the extended protocol ID which are always dumped to the DMA channel.

The RMAP receiver processes commands. If they are correct and accepted the operation is performed on the AHB bus and a reply is formatted. If an acknowledge is requested the RMAP transmitter automatically send the reply. RMAP transmissions have priority over DMA channel transmissions.

Packets with a mismatching destination logical address are never passed to the RMAP handler. There is a user accessible destination key register which is compared to destination key field in incoming packets. If there is a mismatch and a reply has been requested the error code in the reply is set to 3. Replies are sent if and only if the ack field is set to '1'.

When a failure occurs during a bus access the error code is set to 1 (General Error). There is predetermined order in which error-codes are set in the case of multiple errors in the GRSPW. It is shown in table 90.

Table 90. The order of error detection in case of multiple errors in the GRSPW. The error detected first has number 1.

Detection Order	Error Code	Error
1	12	Invalid destination logical address
2	2	Unused RMAP packet type or command code
3	3	Invalid destination key
4	9	Verify buffer overrun
5	11	RMW data length error
6	10	Authorization failure
7*	1	General Error (AHB errors during non-verified writes)
8	5/7	Early EOP / EEP (if early)
9	4	Invalid Data CRC
10	1	General Error (AHB errors during verified writes or RMW)
11	7	EEP
12	6	Cargo Too Large
*The AHB error is not guaranteed to be detected before Early EOP/EEP or Invalid Data CRC. For very long accesses the AHB error detection might be delayed causing the other two errors to appear first.		

Read accesses are performed on the fly, that is they are not stored in a temporary buffer before transmitting. This means that the error code 1 will never be seen in a read reply since the header has already been sent when the data is read. If the AHB error occurs the packet will be truncated and ended with an EEP.

Errors up to and including Invalid Data CRC (number 8) are checked before verified commands. The other errors do not prevent verified operations from being performed.

The details of the support for the different commands are now presented. All defined commands which are received but have an option set which is not supported in this specific implementation will not be executed and a possible reply is sent with error code 10.

16.6.3 Write commands

The write commands are divided into two subcategories when examining their capabilities: verified writes and non-verified writes. Verified writes have a length restriction of 4 bytes and the address must be aligned to the size. That is 1 byte writes can be done to any address, 2 bytes must be halfword aligned, 3 bytes are not allowed and 4 bytes writes must be word aligned. Since there will always be only on AHB operation performed for each RMAP verified write command the incrementing address bit can be set to any value.

Non-verified writes have no restrictions when the incrementing bit is set to 1. If it is set to 0 the number of bytes must be a multiple of 4 and the address word aligned. There is no guarantee how many words will be written when early EOP/EEP is detected for non-verified writes.

16.6.4 Read commands

Read commands are performed on the fly when the reply is sent. Thus if an AHB error occurs the packet will be truncated and ended with an EEP. There are no restrictions for incrementing reads but non-incrementing reads have the same alignment restrictions as non-verified writes. Note that the "Authorization failure" error code will be sent in the reply if a violation was detected even if the length field was zero. Also note that no data is sent in the reply if an error was detected i.e. if the status field is non-zero.

16.6.5 RMW commands

All read-modify-write sizes are supported except 6 which would have caused 3 B being read and written on the bus. The RMW bus accesses have the same restrictions as the verified writes. As in the verified write case, the incrementing bit can be set to any value since only one AHB bus operation will be performed for each RMW command. Cargo too large is detected after the bus accesses so this error will not prevent the operation from being performed. No data is sent in a reply if an error is detected i.e. the status field is non-zero.

16.6.6 Control

The RMAP command handler mostly runs in the background without any external intervention, but there are a few control possibilities.

There is an enable bit in the control register of the GRSPW which can be used to completely disable the RMAP command handler. When it is set to '0' no RMAP packets will be handled in hardware, instead they are all stored to the DMA channel.

There is a possibility that RMAP commands will not be performed in the order they arrive. This can happen if a read arrives before one or more writes. Since the command handler stores replies in a buffer with more than one entry several commands can be processed even if no replies are sent. Data for read replies is read when the reply is sent and thus writes coming after the read might have been performed already if there was congestion in the transmitter. To avoid this the RMAP buffer disable bit can be set to force the command handler to only use one buffer which prevents this situation.

The last control option for the command handler is the possibility to set the destination key which is found in a separate register.

Table 91. GRSPW hardware RMAP handling of different packet type and command fields.

Bit 7 Reserved	Bit 6 Command / Response	Bit 5 Write / Read	Bit 4 Verify data before write	Bit 3 Acknow- ledge	Bit 2 Increment Address	Command	Action
0	0	-	-	-	-	Response	Stored to DMA-channel.
0	1	0	0	0	0	Not used	Does nothing. No reply is sent.
0	1	0	0	0	1	Not used	Does nothing. No reply is sent.
0	1	0	0	1	0	Read single address	Executed normally. Address has to be word aligned and data size a multiple of four. Reply is sent. If alignment restrictions are violated error code is set to 10.
0	1	0	0	1	1	Read incre- menting address.	Executed normally. No restrictions. Reply is sent.
0	1	0	1	0	0	Not used	Does nothing. No reply is sent.
0	1	0	1	0	1	Not used	Does nothing. No reply is sent.
0	1	0	1	1	0	Not used	Does nothing. Reply is sent with error code 2.
0	1	0	1	1	1	Read-Mod- ify-Write increment- ing address	Executed normally. If length is not one of the allowed rmw values nothing is done and error code is set to 11. If the length was correct, alignment restrictions are checked next. 1 byte can be rmw to any address. 2 bytes must be halfword aligned. 3 bytes are not allowed. 4 bytes must be word aligned. If these restrictions are violated nothing is done and error code is set to 10. If an AHB error occurs error code is set to 1. Reply is sent.
0	1	1	0	0	0	Write, sin- gle-address, do not verify before writ- ing, no acknowledge	Executed normally. Address has to be word aligned and data size a multiple of four. If alignment is violated nothing is done. No reply is sent.
0	1	1	0	0	1	Write, incre- menting address, do not verify before writ- ing, no acknowledge	Executed normally. No restric- tions. No reply is sent.
0	1	1	0	1	0	Write, sin- gle-address, do not verify before writ- ing, send acknowledge	Executed normally. Address has to be word aligned and data size a multiple of four. If alignment is violated nothing is done and error code is set to 10. If an AHB error occurs error code is set to 1. Reply is sent.

Table 91. GRSPW hardware RMAP handling of different packet type and command fields.

Bit 7 Reserved	Bit 6 Command / Response	Bit 5 Write / Read	Bit 4 Verify data before write	Bit 3 Acknow-ledge	Bit 2 Increment Address	Command	Action
0	1	1	0	1	1	Write, incrementing address, do not verify before writing, send acknowledge	Executed normally. No restrictions. If AHB error occurs error code is set to 1. Reply is sent.
0	1	1	1	0	0	Write, single address, verify before writing, no acknowledge	Executed normally. Length must be 4 or less. Otherwise nothing is done. Same alignment restrictions apply as for rmw. No reply is sent.
0	1	1	1	0	1	Write, incrementing address, verify before writing, no acknowledge	Executed normally. Length must be 4 or less. Otherwise nothing is done. Same alignment restrictions apply as for rmw. If they are violated nothing is done. No reply is sent.
0	1	1	1	1	0	Write, single address, verify before writing, send acknowledge	Executed normally. Length must be 4 or less. Otherwise nothing is done and error code is set to 9. Same alignment restrictions apply as for rmw. If they are violated nothing is done and error code is set to 10. If an AHB error occurs error code is set to 1. Reply is sent.
0	1	1	1	1	1	Write, incrementing address, verify before writing, send acknowledge	Executed normally. Length must be 4 or less. Otherwise nothing is done and error code is set to 9. Same alignment restrictions apply as for rmw. If they are violated nothing is done and error code is set to 10. If an AHB error occurs error code is set to 1. Reply is sent.
1	0	-	-	-	-	Unused	Stored to DMA-channel.
1	1	-	-	-	-	Unused	Stored to DMA-channel.

16.7 AMBA interface

The AMBA interface consists of an APB interface, an AHB master interface and DMA FIFOs. The APB interface provides access to the user registers. The DMA engines have 32-bit wide FIFOs to the AHB master interface which are used when reading and writing to the bus.

The transmitter DMA engine reads data from the bus in bursts which are half the FIFO size in length. A burst is always started when the FIFO is half-empty or if it can hold the last data for the packet. The burst containing the last data might have shorter length if the packet is not an even number of bursts in size.

The receiver DMA works in the same way except that it checks if the FIFO is half-full and then performs a burst write to the bus which is half the fifo size in length. The last burst might be shorter. There might be 1 to 3 single byte writes when writing the beginning and end of the received packets.

16.7.1 APB slave interface

As mentioned above, the APB interface provides access to the user registers which are 32-bits in width. The accesses to this interface are required to be aligned word accesses. The result is undefined if this restriction is violated.

16.7.2 AHB master interface

The GRSPW contains a single master interface which is used by both the transmitter and receiver DMA engines. The arbitration algorithm between the channels is done so that if the current owner requests the interface again it will always acquire it. This will not lead to starvation problems since the DMA engines always deassert their requests between accesses.

The AHB accesses are always word accesses (HSIZE = 0x010) of type incremental burst with unspecified length (HBURST = 0x001) if rmap and rxunaligned are disabled. Otherwise the accesses can be of size byte, halfword and word (HSIZE = 0x000, 0x001, 0x010). Byte and halfword accesses are always NONSEQ.

The burst length will be half the AHB FIFO size except for the last transfer for a packet which might be smaller. Shorter accesses are also done during descriptor reads and status writes.

The AHB master also supports non-incrementing accesses where the address will be constant for several consecutive accesses. HTRANS will always be NONSEQ in this case while for incrementing accesses it is set to SEQ after the first access. This feature is included to support non-incrementing reads and writes for RMAP.

If the GRSPW does not need the bus after a burst has finished there will be one idle cycle (HTRANS = IDLE).

BUSY transfer types are never requested and the core provides full support for ERROR, RETRY and SPLIT responses.

16.8 SpaceWire clock generation

The clock source for SpaceWire is selectable through a clock multiplexer, which inputs are the SPW-CLK, and INCLK inputs. The selected clock is then used either as 1X, 2X or 4X as shown in the figure below. The system clock (CLK) is also selectable as SpaceWire transmit clock.

After reset the selected SpaceWire transmit clock is SPWCLK without any multiplication. The Space-Wire clock multiplexers and the DLL reset are controlled through the GRGPREG register.

The SpaceWire transmit clock must be a multiple of 10 MHz in order to achieve 10 Mbps start up bit rate. The division to 10 MHz is done internally in the GRSPW2 core. During reset the clock divider register in GRSPW2 gets its value from 4 I/O pins that must be pulled up/down to set the divider correctly. Pins (MSB-LSB) SWMX45, 43, 40 and 37 are used for this. Thus it is possible to use a SPW-CLK which is any multiple of 10 between 10-100 MHz (note that the required precision is 10 MHz +/- 1 MHz).

The SpaceWire transmitter uses SDR output registers meaning that the bitrate will be equal to the transmit clock. The SpaceWire input signals are sampled synchronously with DDR registers using the transmit clock. This allows the GR712RC to receive SpaceWire data at a bitrate equal to the transmit clock.

16.9   Registers

The core is programmed through registers mapped into APB address space.The table below shows the base address and interrupt number for each GRSWP core.

Table 92. GRSPW core base addresses and interrupts

APB base address	GRSWP core	Interrupt
0x80100800	GRSPW2 - 0	22
0x80100900	GRSPW2 - 1	23
0x80100A00	GRSPW2 - 2	24
0x80100B00	GRSPW2 - 3	25
0x80100C00	GRSPW2 - 4	26
0x80100D00	GRSPW2 - 5	27

Table 93. GRSPW registers

APB address offset	Register
0x0	Control
0x4	Status/Interrupt-source
0x8	Node address
0xC	Clock divisor
0x10	Destination key
0x14	Time
0x20	DMA channel 1 control/status
0x24	DMA channel 1 rx maximum length
0x28	DMA channel 1 transmit descriptor table address.
0x2C	DMA channel 1 receive descriptor table address.
0x30	DMA channel 1 address register
0x34	Unused
0x38	Unused
0x3C	Unused
0x40 - 0x9C	Unused

Table 94. GRSPW control register

31	30	29	28 27 26	26	25 24 23 22	21	20	19 18	17	16	15 14 13 12	11	10	9	8	7	6	5	4	3	2	1	0
RA	RX	RC	NCH	PO	RESERVED	PS	NP		RD	RE	RESERVED	TR	TT	LI	TQ		RS	PM	TI	IE	AS	LS	LD

31	RMAP available (RA) - Set to one if the RMAP command handler is available. Only readable. Reset value '1' for links 0 and 1. Reset value '0' for links 2, 3, 4 and 5.
30	RX unaligned access (RX) - Set to one if unaligned writes are available for the receiver. Only readable.
29	RMAP CRC available (RC) - Set to one if RMAP CRC is enabled in the core. Only readable.
28: 27	Number of DMA channels (NCH) - The number of available DMA channels minus one (Number of channels = NCH+1). Only readable. Reset value 0.
26	Number of ports (PO) - The number of available SpaceWire ports minus one. Only readable. Reset value 0.
25: 22	RESERVED
21	Port select (PS) - Selects the active port when the no port force bit is zero. '0' selects the port connected to data and strobe on index 0 while '1' selects index 1. Fixed to '0'.
20	No port force (NP) - Disable port force. When disabled the port select bit cannot be used to select the active port. Instead, it is automatically selected by checking the activity on the respective receive links. Reset value: '0'. Fixed to '0'.
19: 18	RESERVED
17	RMAP buffer disable (RD) - If set only one RMAP buffer is used. This ensures that all RMAP commands will be executed consecutively. Reset value: '0'.
16	RMAP Enable (RE) - Enable RMAP command handler. Reset value: '1' on links 0 - 1, '0' on links 2 - 5.
15: 12	RESERVED
11	Time Rx Enable (TR) - Enable time-code receptions. Reset value: '0'.
10	Time Tx Enable (TT) - Enable time-code transmissions. Reset value: '0'.
9	Link error IRQ (LI) - Generate interrupt when a link error occurs. Not reset.
8	Tick-out IRQ (TQ) - Generate interrupt when a valid time-code is received. Not reset.
7	RESERVED
6	Reset (RS) - Make complete reset of the SpaceWire node. Self clearing. Reset value: '0'.
5	Promiscuous Mode (PM) - Enable Promiscuous mode. Reset value: '0'.
4	Tick In (TI) - The host can generate a tick by writing a one to this field. This will increment the timer counter and the new value is transmitted after the current character is transferred. A tick can also be generated by asserting the tick_in signal. Reset value: '0'.
3	Interrupt Enable (IE) - If set, an interrupt is generated when one of bit 8 to 9 is set and its corresponding event occurs. Reset value: '0'.
2	Autostart (AS) - Automatically start the link when a NULL has been received. Not reset on links 0-1. Reset value '0' on links 2-5.
1	Link Start (LS) - Start the link, i.e. allow a transition from ready to started state. Reset value: '1' on links 0 - 1, '0' on links 2 - 5.
0	Link Disable (LD) - Disable the SpaceWire codec. Reset value: '0'.

Table 95. GRSPW status register

31 30 29 28 27 26 25 24	23 22 21	20 19 18 17 16 15 14 13 12 11 10	9	8	7	6 5	4	3	2	1	0
RESERVED	LS	RESERVED	AP	EE	IA		PE	DE	ER	CE	TO

31: 24	RESERVED
23: 21	Link State (LS) - The current state of the start-up sequence. 0 = Error-reset, 1 = Error-wait, 2 = Ready, 3 = Started, 4 = Connecting, 5 = Run. Reset value: 0.
20: 10	RESERVED
9	Active port (AP) - Shows the currently active port. '0' = Port 0 and '1' = Port 1 where the port numbers refer to the index number of the data and strobe signals.
8	Early EOP/EEP (EE) - Set to one when a packet is received with an EOP after the first byte for a non-rmap packet and after the second byte for a RMAP packet. Cleared when written with a one. Reset value: '0'.
7	Invalid Address (IA) - Set to one when a packet is received with an invalid destination address field, i.e it does not match the nodeaddr register. Cleared when written with a one. Reset value: '0'.
6: 5	RESERVED
4	Parity Error (PE) - A parity error has occurred. Cleared when written with a one. Reset value: '0'.
3	Disconnect Error (DE) - A disconnection error has occurred. Cleared when written with a one. Reset value: '0'.
2	Escape Error (ER) - An escape error has occurred. Cleared when written with a one. Reset value: '0'.
1	Credit Error (CE) - A credit has occurred. Cleared when written with a one. Reset value: '0'.
0	Tick Out (TO) - A new time count value was received and is stored in the time counter field. Cleared when written with a one. Reset value: '0'.

Table 96. GRSPW default address register

31 16	15 8	7 0
RESERVED	DEFMASK	DEFADDR

31: 8	RESERVED
15: 8	Default mask (DEFMASK) - Default mask used for node identification on the SpaceWire network. This field is used for masking the address before comparison. Both the received address and the DEFADDR field are anded with the inverse of DEFMASK before the address check.
7: 0	Default address (DEFADDR) - Default address used for node identification on the SpaceWire network. Reset value: 254.

Table 97. GRSPW clock divisor register

31 16	15 8	7 0
RESERVED	CLKDIVSTART	CLKDIVRUN

31: 16	RESERVED
15: 8	Clock divisor startup (CLKDIVSTART) - Clock divisor value used for the clock-divider during startup (link-interface is in other states than run). The actual divisor value is Clock Divisor register + 1. Reset value: "0000" & SWMX[45] & SWMX[43] & SWMX[40] & SWMX[37].
7: 0	Clock divisor run (CLKDIVRUN) - Clock divisor value used for the clock-divider when the link-interface is in the run-state. The actual divisor value is Clock Divisor register + 1. Reset value: "0000" & SWMX[45] & SWMX[43] & SWMX[40] & SWMX[37]

Table 98. GRSPW destination key

31		8	7		0
RESERVED			DESTKEY		

31: 8	RESERVED
7: 0	Destination key (DESTKEY) - RMAP destination key. Reset value: 0.

Table 99. GRSPW time register

31		8	7	6	5		0
RESERVED			TCTRL		TIMECNT		

31: 8	RESERVED
7: 6	Time control flags (TCTRL) - The current value of the time control flags. Sent with time-code resulting from a tick-in. Received control flags are also stored in this register. Reset value: '0'.
5: 0	Time counter (TIMECNT) - The current value of the system time counter. It is incremented for each tick-in and the incremented value is transmitted. The register can also be written directly but the written value will not be transmitted. Received time-counter values are also stored in this register. Reset value: '0'.

Table 100. GRSPW dma control register

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
RESERVED															LE	SP	SA	EN	NS	RD	RX	AT	RA	TA	PR	PS	AI	RI	TI	RE	TE

31: 17	RESERVED
16	Link error disable (LE) - Disable transmitter when a link error occurs. No more packets will be transmitted until the transmitter is enabled again. Reset value: '0'.
15	Strip pid (SP) - Remove the pid byte (second byte) of each packet. The address byte (first byte) will also be removed when this bit is set independent of the SA bit. Reset value: '0'.
14	Strip addr (SA) - Remove the addr byte (first byte) of each packet. Reset value: '0'.
13	Enable addr (EN) - Enable separate node address for this channel. Reset value: '0'.
12	No spill (NS) - If cleared, packets will be discarded when a packet is arriving and there are no active descriptors. If set, the GRSPW will wait for a descriptor to be activated.
11	Rx descriptors available (RD) - Set to one, to indicate to the GRSPW that there are enabled descriptors in the descriptor table. Cleared by the GRSPW when it encounters a disabled descriptor: Reset value: '0'.
10	RX active (RX) - Is set to '1' if a reception to the DMA channel is currently active otherwise it is '0'. Only readable.
9	Abort TX (AT) - Set to one to abort the currently transmitting packet and disable transmissions. If no transmission is active the only effect is to disable transmissions. Self clearing. Reset value: '0'.
8	RX AHB error (RA) - An error response was detected on the AHB bus while this receive DMA channel was accessing the bus. Cleared when written with a one. Reset value: '0'.
7	TX AHB error (TA) - An error response was detected on the AHB bus while this transmit DMA channel was accessing the bus. Cleared when written with a one. Reset value: '0'.
6	Packet received (PR) - This bit is set each time a packet has been received. never cleared by the SW-node. Cleared when written with a one. Reset value: '0'.
5	Packet sent (PS) - This bit is set each time a packet has been sent. Never cleared by the SW-node. Cleared when written with a one. Reset value: '0'.
4	AHB error interrupt (AI) - If set, an interrupt will be generated each time an AHB error occurs when this DMA channel is accessing the bus. Not reset.
3	Receive interrupt (RI) - If set, an interrupt will be generated each time a packet has been received, if the interrupt enable (IE) bit in the corresponding receive descriptor is set as well. This happens both if the packet is terminated by an EEP or EOP. Not reset.
2	Transmit interrupt (TI) - If set, an interrupt will be generated each time a packet is transmitted, if the interrupt enable (IE) bit in the corresponding transmit descriptor is set as well. The interrupt is generated regardless of whether the transmission was successful or not. Not reset.

Table 100. GRSPW dma control register

1	Receiver enable (RE) - Set to one when packets are allowed to be received to this channel. Reset value: '0'.
0	Transmitter enable (TE) - Write a one to this bit each time new descriptors are activated in the table. Writing a one will cause the SW-node to read a new descriptor and try to transmit the packet it points to. This bit is automatically cleared when the SW-node encounters a descriptor which is disabled. Reset value: '0'.

Table 101. GRSPW RX maximum length register.

31	25	24	0
RESERVED		RXMAXLEN	

31: 25	RESERVED
24: 0	RX maximum length (RXMAXLEN) - Receiver packet maximum length in bytes. Only bits 24 - 2 are writable. Bits 1 - 0 are always 0. Not reset.

Table 102. GRSPW transmitter descriptor table address register.

31	10	9	4	3	0
DESCBASEADDR		DESCSEL		RESERVED	

31: 10	Descriptor table base address (DESCBASEADDR) - Sets the base address of the descriptor table. Not reset.
9: 4	Descriptor selector (DESCSEL) - Offset into the descriptor table. Shows which descriptor is currently used by the GRSPW. For each new descriptor read, the selector will increase with 16 and eventually wrap to zero again. Reset value: 0.
3: 0	RESERVED

Table 103. GRSPW receiver descriptor table address register.

31	10	9	3	2	0
DESCBASEADDR		DESCSEL		RESERVED	

31: 10	Descriptor table base address (DESCBASEADDR) - Sets the base address of the descriptor table. Not reset.
9: 3	Descriptor selector (DESCSEL) - Offset into the descriptor table. Shows which descriptor is currently used by the GRSPW. For each new descriptor read, the selector will increase with 8 and eventually wrap to zero again. Reset value: 0.
2: 0	RESERVED

Table 104. GRSPW dma channel address register

31	16	15	8	7	0
RESERVED		MASK		ADDR	

31: 8	RESERVED
15: 8	Mask (MASK) - Mask used for node identification on the SpaceWire network. This field is used for masking the address before comparison. Both the received address and the ADDR field are anded with the inverse of MASK before the address check.
7: 0	Address (ADDR) - Address used for node identification on the SpaceWire network for the corresponding dma channel when the EN bit in the DMA control register is set. Reset value: 254.

16.10 Signal definitions

The signals are described in table 105.

*Table 105.*Signal definitions

Signal name	Type	Function	Active
SPWCLK	Input	SpaceWire receiver and transmitter clock	-
SPW_RXD[]	Input	Data input	High
SPW_RXS[]	Input	Strobe input	High
SPW_TXD[]	Output	Data output	High
SPW_TXS[]	Output	Strobe output	High

17   Ethernet Media Access Controller (MAC)

17.1   Overview

The Ethernet Media Access Controller (GRETH) provides an interface between an AMBA-AHB bus and an Ethernet network. It supports 10/100 Mbit speed in both full- and half-duplex. The AMBA interface consists of an APB interface for configuration and control and an AHB master interface for transmit and receive data. There is one DMA engine for the transmitter and one for the receiver. Both share the same AHB master interface. The ethernet interface supports the RMII interfaces which should be connected to an external PHY. The GRETH also provides access to the RMII Management interface which is used to configure the PHY.
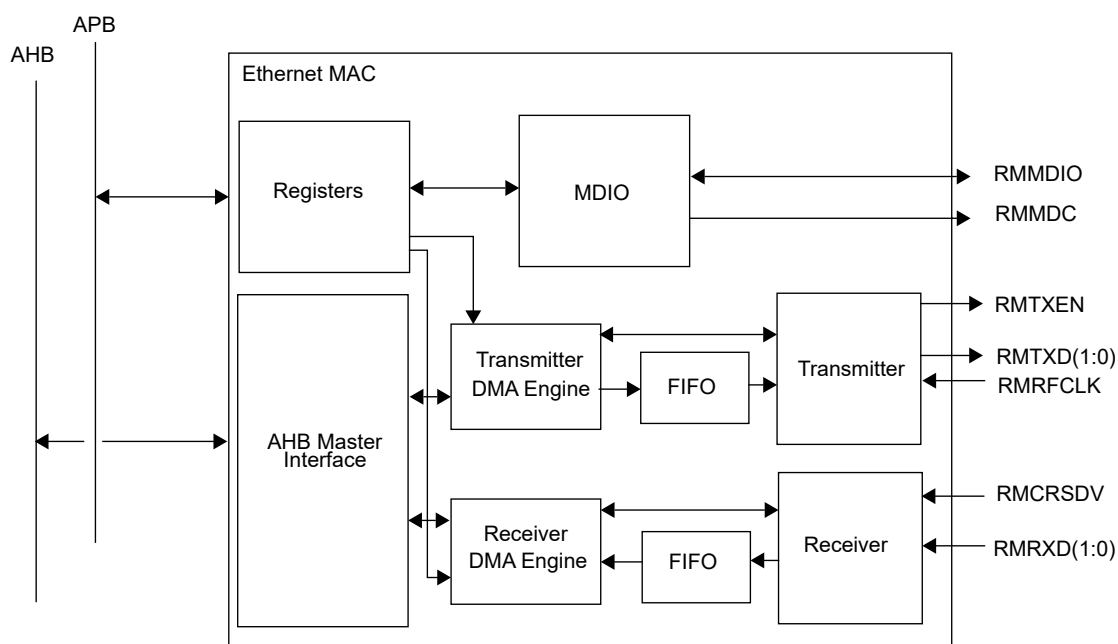
Figure 64. Block diagram of the internal structure of the GRETH with RMII

17.2   Operation

17.2.1   Protocol support

The GRETH is implemented according to IEEE standard 802.3-2002. There is no support for the optional control sublayer and no multicast addresses can be assigned to the MAC. This means that packets with type 0x8808 (the only currently defined ctrl packets) are discarded.

17.2.2   Clocking

GRETH has two clock domains: The AHB clock and the Ethernet PHY clock. The ethernet transmitter and receiver clock is generated by the external ethernet PHY, and is an input to the core through the REFCLK signal. The two clock domains are unrelated to each other and all signals crossing the clock regions are fully synchronized inside the core.

Both full-duplex and half-duplex operating modes are supported and both can be run in either 10 or 100 Mbit. The minimum AHB clock for 10 Mbit operation is 2.5 MHz, while 18 MHz is needed for 100 Mbit. Using a lower AHB clock than specified will lead to excessive packet loss.

17.3 Tx DMA interface

The transmitter DMA interface is used for transmitting data on an Ethernet network. The transmission is done using descriptors located in memory.

17.3.1 Setting up a descriptor.

A single descriptor is shown in table 106 and 107. The number of bytes to be sent should be set in the length field and the address field should point to the data. The address must be word-aligned. If the interrupt enable (IE) bit is set, an interrupt will be generated when the packet has been sent (this requires that the transmitter interrupt bit in the control register is also set). The interrupt will be generated regardless of whether the packet was transmitted successfully or not. The Wrap (WR) bit is also a control bit that should be set before transmission and it will be explained later in this section.

Table 106. GRETH transmit descriptor word 0 (address offset 0x0)

31 16	15	14	13	12	11	10 0
RESERVED	AL	UE	IE	WR	EN	LENGTH

Bits	Description
31: 16	RESERVED
15	Attempt Limit Error (AL) - The packet was not transmitted because the maximum number of attempts was reached.
14	Underrun Error (UE) - The packet was incorrectly transmitted due to a FIFO underrun error.
13	Interrupt Enable (IE) - Enable Interrupts. An interrupt will be generated when the packet from this descriptor has been sent provided that the transmitter interrupt enable bit in the control register is set. The interrupt is generated regardless if the packet was transmitted successfully or if it terminated with an error.
12	Wrap (WR) - Set to one to make the descriptor pointer wrap to zero after this descriptor has been used. If this bit is not set the pointer will increment by 8. The pointer automatically wraps to zero when the 1 kB boundary of the descriptor table is reached.
11	Enable (EN) - Set to one to enable the descriptor. Should always be set last of all the descriptor fields.
10: 0	LENGTH - The number of bytes to be transmitted.

Table 107. GRETH transmit descriptor word 1 (address offset 0x4)

31 2	1	0
ADDRESS		RES

Bits	Description
31: 2	Address (ADDRESS) - Pointer to the buffer area from where the packet data will be loaded.
1: 0	RESERVED

To enable a descriptor the enable (EN) bit should be set and after this is done, the descriptor should not be touched until the enable bit has been cleared by the GRETH.

17.3.2 Starting transmissions

Enabling a descriptor is not enough to start a transmission. A pointer to the memory area holding the descriptors must first be set in the GRETH. This is done in the transmitter descriptor pointer register. The address must be aligned to a 1 kB boundary. Bits 31 to 10 hold the base address of descriptor area while bits 9 to 3 form a pointer to an individual descriptor.The first descriptor should be located at the base address and when it has been used by the GRETH the pointer field is incremented by 8 to point at the next descriptor. The pointer will automatically wrap back to zero when the next 1 kB boundary has been reached (the descriptor at address offset 0x3F8 has been used). The WR bit in the descriptors can be set to make the pointer wrap back to zero before the 1 kB boundary.

The pointer field has also been made writable for maximum flexibility but care should be taken when writing to the descriptor pointer register. It should never be touched when a transmission is active.

The final step to activate the transmission is to set the transmit enable bit in the control register. This tells the GRETH that there are more active descriptors in the descriptor table. This bit should always be set when new descriptors are enabled, even if transmissions are already active. The descriptors must always be enabled before the transmit enable bit is set.

17.3.3 Descriptor handling after transmission

When a transmission of a packet has finished, status is written to the first word in the corresponding descriptor. The Underrun Error bit is set if the FIFO became empty before the packet was completely transmitted while the Alignment Error bit is set if more collisions occurred than allowed. The packet was successfully transmitted only if both of these bits are zero. The other bits in the first descriptor word are set to zero after transmission while the second word is left untouched.

The enable bit should be used as the indicator when a descriptor can be used again, which is when it has been cleared by the GRETH. There are three bits in the GRETH status register that hold transmission status. The Transmitter Error (TE) bit is set each time an transmission ended with an error (when at least one of the two status bits in the transmit descriptor has been set). The Transmitter Interrupt (TI) is set each time a transmission ended successfully.

The transmitter AHB error (TA) bit is set when an AHB error was encountered either when reading a descriptor or when reading packet data. Any active transmissions were aborted and the transmitter was disabled. The transmitter can be activated again by setting the transmit enable register.

17.3.4 Setting up the data for transmission

The data to be transmitted should be placed beginning at the address pointed by the descriptor address field. The GRETH does not add the Ethernet address and type fields so they must also be stored in the data buffer. The 4 B Ethernet CRC is automatically appended at the end of each packet. Each descriptor will be sent as a single Ethernet packet. If the size field in a descriptor is greater than 1514 B, the packet will not be sent.

17.4 Rx DMA interface

The receiver DMA interface is used for receiving data from an Ethernet network. The reception is done using descriptors located in memory.

17.4.1 Setting up descriptors

A single descriptor is shown in table 108 and 109. The address field should point to a word-aligned buffer where the received data should be stored. The GRETH will never store more than 1514 B to the buffer. If the interrupt enable (IE) bit is set, an interrupt will be generated when a packet has been received to this buffer (this requires that the receiver interrupt bit in the control register is also set). The interrupt will be generated regardless of whether the packet was received successfully or not. The Wrap (WR) bit is also a control bit that should be set before the descriptor is enabled and it will be explained later in this section.

Table 108. GRETH receive descriptor word 0 (address offset 0x0)

31	19	18	17	16	15	14	13	12	11	10	0
RESERVED		LE	OE	CE	FT	AE	IE	WR	EN	LENGTH	

31: 19	RESERVED
18	Length error (LE) - The length/type field of the packet did not match the actual number of received bytes.
17	Overrun error (OE) - The frame was incorrectly received due to a FIFO overrun.

Table 108. GRETH receive descriptor word 0 (address offset 0x0)

16	CRC error (CE) - A CRC error was detected in this frame.
15	Frame too long (FT) - A frame larger than the maximum size was received. The excessive part was truncated.
14	Alignment error (AE) - An odd number of nibbles were received.
13	Interrupt Enable (IE) - Enable Interrupts. An interrupt will be generated when a packet has been received to this descriptor provided that the receiver interrupt enable bit in the control register is set. The interrupt is generated regardless if the packet was received successfully or if it terminated with an error.
12	Wrap (WR) - Set to one to make the descriptor pointer wrap to zero after this descriptor has been used. If this bit is not set the pointer will increment by 8. The pointer automatically wraps to zero when the 1 kB boundary of the descriptor table is reached.
11	Enable (EN) - Set to one to enable the descriptor. Should always be set last of all the descriptor fields.
10: 0	LENGTH - The number of bytes received to this descriptor.

Table 109. GRETH receive descriptor word 1 (address offset 0x4)

31	2	1	0
ADDRESS		RES	

31: 2	Address (ADDRESS) - Pointer to the buffer area from where the packet data will be loaded.
1: 0	RESERVED

17.4.2 Starting reception

Enabling a descriptor is not enough to start reception. A pointer to the memory area holding the descriptors must first be set in the GRETH. This is done in the receiver descriptor pointer register. The address must be aligned to a 1 kB boundary. Bits 31 to 10 hold the base address of descriptor area while bits 9 to 3 form a pointer to an individual descriptor. The first descriptor should be located at the base address and when it has been used by the GRETH the pointer field is incremented by 8 to point at the next descriptor. The pointer will automatically wrap back to zero when the next 1 kB boundary has been reached (the descriptor at address offset 0x3F8 has been used). The WR bit in the descriptors can be set to make the pointer wrap back to zero before the 1 kB boundary.

The pointer field has also been made writable for maximum flexibility but care should be taken when writing to the descriptor pointer register. It should never be touched when reception is active.

The final step to activate reception is to set the receiver enable bit in the control register. This will make the GRETH read the first descriptor and wait for an incoming packet.

17.4.3 Descriptor handling after reception

The GRETH indicates a completed reception by clearing the descriptor enable bit. The other control bits (WR, IE) are also cleared. The number of received bytes is shown in the length field. The parts of the Ethernet frame stored are the destination address, source address, type and data fields. Bits 17-14 in the first descriptor word are status bits indicating different receive errors. All four bits are zero after a reception without errors. The status bits are described in table 108.

Packets arriving that are smaller than the minimum Ethernet size of 64 B are not considered as a reception and are discarded. The current receive descriptor will be left untouched an used for the first packet arriving with an accepted size. The TS bit in the status register is set each time this event occurs.

If a packet is received with an address not accepted by the MAC, the IA status register bit will be set.

Packets larger than maximum size cause the FT bit in the receive descriptor to be set. The length field is not guaranteed to hold the correct value of received bytes. The counting stops after the word containing the last byte up to the maximum size limit has been written to memory.

The address word of the descriptor is never touched by the GRETH.

17.4.4 Reception with AHB errors

If an AHB error occurs during a descriptor read or data store, the Receiver AHB Error (RA) bit in the status register will be set and the receiver is disabled. The current reception is aborted. The receiver can be enabled again by setting the Receive Enable bit in the control register.

17.5 MDIO Interface

The MDIO interface provides access to PHY configuration and status registers through a two-wire interface which is included in the RMII interface. The GRETH provided full support for the MDIO interface.

The MDIO interface can be used to access from 1 to 32 PHY containing 1 to 32 16-bit registers. A read transfer i set up by writing the PHY and register addresses to the MDIO Control register and setting the read bit. This caused the Busy bit to be set and the operation is finished when the Busy bit is cleared. If the operation was successful the Linkfail bit is zero and the data field contains the read data. An unsuccessful operation is indicated by the Linkfail bit being set. The data field is undefined in this case.

A write operation is started by writing the 16-bit data, PHY address and register address to the MDIO Control register and setting the write bit. The operation is finished when the busy bit is cleared and it was successful if the Linkfail bit is zero.

17.6 Reduced Media Independent Interfaces (RMII)

The GRETH is configured to use the Reduced Media Independent Interface (RMII) for interfacing the external PHY. It uses 7 signals which are a subset of the MII signals. The table below shows the signals and their function.

*Table 110.*RMII signals

Signal	Function
RMTXD[1:0]	Transmit data
RMTXEN	Transmit enable
RMCRSDV	Carrier sense and data valid
RMRXD[1:0]	Receiver data
RMRFCLK	Reference clock (50 MHz)

17.7 Registers

The core is programmed through registers mapped into APB address space.

*Table 111.*GRETH registers

APB address	Register
0x80000E00	Control register
0x80000E04	Status/Interrupt-source register
0x80000E08	MAC Address MSB
0x80000E0C	MAC Address LSB
0x80000E10	MDIO Control/Status
0x80000E14	Transmit descriptor pointer
0x80000E18	Receiver descriptor pointer

Table 112. GRETH control register

31	28	27		8	7	6	5	4	3	2	1	0
0000		RESERVED			SP	RS	PM	FD	RI	TI	RE	TE

27: 8	RESERVED
7	Speed (SP) - Sets the current speed mode. 0 = 10 Mbit, 1 = 100 Mbit. A default value is automatically read from the PHY after reset. Reset value: '1'.
6	Reset (RS) - A one written to this bit resets the GRETH core. Self clearing.
5	Promiscuous mode (PM) - If set, the GRETH operates in promiscuous mode which means it will receive all packets regardless of the destination address. Reset value: '0'.
4	Full duplex (FD) - If set, the GRETH operates in full-duplex mode otherwise it operates in half-duplex. Reset value: '0'.
3	Receiver interrupt (RI) - Enable Receiver Interrupts. An interrupt will be generated each time a packet is received when this bit is set. The interrupt is generated regardless if the packet was received successfully or if it terminated with an error. Reset value: '0'.
2	Transmitter interrupt (TI) - Enable Transmitter Interrupts. An interrupt will be generated each time a packet is transmitted when this bit is set. The interrupt is generated regardless if the packet was transmitted successfully or if it terminated with an error. Reset value: '0'.
1	Receive enable (RE) - Should be written with a one each time new descriptors are enabled. As long as this bit is one the GRETH will read new descriptors and as soon as it encounters a disabled descriptor it will stop until TE is set again. This bit should be written with a one after the new descriptors have been enabled. Reset value: '0'.
0	Transmit enable (TE) - Should be written with a one each time new descriptors are enabled. As long as this bit is one the GRETH will read new descriptors and as soon as it encounters a disabled descriptor it will stop until TE is set again. This bit should be written with a one after the new descriptors have been enabled. Reset value: '0'.

Table 113. GRETH status register

31		8	7	6	5	4	3	2	1	0
RESERVED			IA	TS	TA	RA	TI	RI	TE	RE

7	Invalid address (IA) - A packet with an address not accepted by the MAC was received. Cleared when written with a one. Reset value: '0'.
6	Too small (TS) - A packet smaller than the minimum size was received. Cleared when written with a one. Reset value: '0'.
5	Transmitter AHB error (TA) - An AHB error was encountered in transmitter DMA engine. Cleared when written with a one. Not Reset.
4	Receiver AHB error (RA) - An AHB error was encountered in receiver DMA engine. Cleared when written with a one. Not Reset.
3	Transmitter interrupt (TI) - A packet was transmitted without errors. Cleared when written with a one. Not Reset.
2	Receiver interrupt (RI) - A packet was received without errors. Cleared when written with a one. Not Reset.
1	Transmitter error (TE) - A packet was transmitted which terminated with an error. Cleared when written with a one. Not Reset.
0	Receiver error (RE) - A packet has been received which terminated with an error. Cleared when written with a one. Not Reset.

Table 114. GRETH MAC address MSB.

31	16	15	0
RESERVED		Bit 47 downto 32 of the MAC address	

31: 16   RESERVED

15: 0   The two most significant bytes of the MAC Address. Not Reset.

Table 115. GRETH MAC address LSB.

31	0
Bit 31 downto 0 of the MAC address	

31: 0   The four least significant bytes of the MAC Address. Not Reset.

Table 116. GRETH MDIO ctrl/status register.

31	16	15	11	10	6	5	4	3	2	1	0
DATA		PHYADDR		REGADDR			NV	BU	LF	RD	WR

31: 16   Data (DATA) - Contains data read during a read operation and data that is transmitted is taken from this field. Reset value: 0x0000.

15: 11   PHY address (PHYADDR) - This field contains the address of the PHY that should be accessed during a write or read operation. Reset value: "00001".

10: 6   Register address (REGADDR) - This field contains the address of the register that should be accessed during a write or read operation. Reset value: "00000".

5   RESERVED

4   Not valid (NV) - When an operation is finished (BUSY = 0) this bit indicates whether valid data has been received that is, the data field contains correct data. Reset value: '0'.

3   Busy (BU) - When an operation is performed this bit is set to one. As soon as the operation is finished and the management link is idle this bit is cleared. Reset value: '0'.

2   Linkfail (LF) - When an operation completes (BUSY = 0) this bit is set if a functional management link was not detected. Reset value: '1'.

1   Read (RD) - Start a read operation on the management interface. Data is stored in the data field. Reset value: '0'.

0   Write (WR) - Start a write operation on the management interface. Data is taken from the Data field. Reset value: '0'.

Table 117. GRETH transmitter descriptor table base address register.

31	10	9	3	2	0
BASEADDR		DESCPNT		RES	

31: 10   Transmitter descriptor table base address (BASEADDR) - Base address to the transmitter descriptor table.Not Reset.

9: 3   Descriptor pointer (DESCPNT) - Pointer to individual descriptors. Automatically incremented by the Ethernet MAC.

2: 0   RESERVED

Table 118. GRETH receiver descriptor table base address register.

31	10	9	3	2	0
BASEADDR		DESCPNT		RES	

31: 10   Receiver descriptor table base address (BASEADDR) - Base address to the receiver descriptor table.Not Reset.

9: 3   Descriptor pointer (DESCPNT) - Pointer to individual descriptors. Automatically incremented by the Ethernet MAC.

2: 0   RESERVED

17.8   Signal definitions

The signals are described in table 119.

*Table 119.*Signal definitions

Signal name	Type	Function	Active
RMTXEN	Output	Ethernet Transmit Enable	High
RMTXD[0]	Output	Ethernet Transmit Data 0	-
RMTXD[1]	Output	Ethernet Transmit Data 1	-
RMRXD[0]	Input	Ethernet Receive Data 0	-
RMRXD[1]	Input	Ethernet Receive Data 1	-
RMCRSDV	Input	Ethernet Carrier Sense / Data Valid	High
RMINTN	Input	Ethernet Management Interrupt	Low
RMMDIO	Input/Output	Ethernet Media Interface Data	-
RMMDC	Output	Ethernet Media Interface Clock	High
RMRFCLK	Input	Ethernet Reference Clock	-

18   CAN Interface

18.1   Overview

The GR712RC includes two identical CAN-2.0 interfaces. The CAN interfaces implement the CAN 20.A and 2.0B protocols. It is based on the Philips SJA1000 and has a compatible register map with a few exceptions.
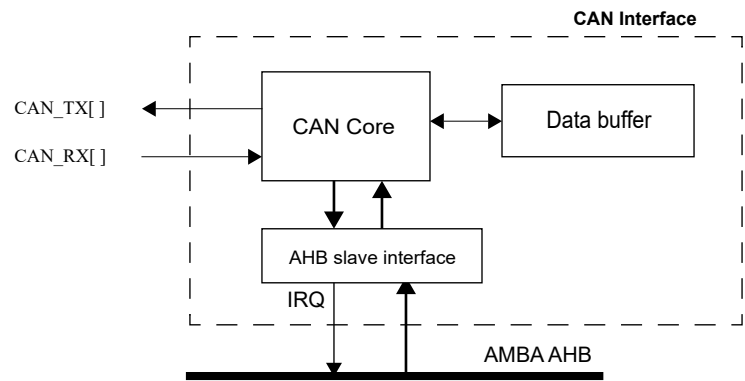
Figure 65. Block diagram

Note: The CAN bus multiplexer has to be programmed to activate OC-CAN1 and OC-CAN2 on the GR712RC can bus A and can bus B. The OC-CAN1 and OC-CAN2 are disabled at reset by CAN-MUX (see section 20 for further information).

18.2   CAN controller overview

This CAN controller is based on the Philips SJA1000 and has a compatible register map with a few exceptions. It also supports both BasicCAN (PCA82C200 like) and PeliCAN mode. In PeliCAN mode the extended features of CAN 2.0B is supported. The mode of operation is chosen through the Clock Divider register.

This document will list the registers and their functionality. The Philips SJA1000 data sheet can be used as a reference for further clarification. See also the Design considerations chapter for differences between this core and the SJA1000.

The register map and functionality is different between the two modes of operation. First the Basic-CAN mode will be described followed by PeliCAN. Common registers (clock divisor and bus timing) are described in a separate chapter. The register map also differs depending on whether the core is in operating mode or in reset mode. When reset the core starts in reset mode awaiting configuration. Operating mode is entered by clearing the reset request bit in the command register. To re-enter reset mode set this bit high again.

18.3   AHB interface

All registers are one byte wide and the addresses specified in this document are byte addresses. Byte reads and writes should be used when interfacing with this core. The read byte is duplicated on all byte lanes of the AHB bus. The wrapper is big endian so the core expects the MSB at the lowest address. The bit numbering in this document uses bit 7 as MSB and bit 0 as LSB.

The table below shows the AHB base address and interrupt for the two CAN cores.

TABLE 120. CAN AHB base address

Core	Address range	Interrupt
CAN core 1	0xFFF30000 - 0xFFF3001F	5
CAN core 2	0xFFF30100 - 0xFFF3011F	6

18.4 BasicCAN mode

18.4.1 BasicCAN register map

*Table 121.*BasicCAN address allocation

Address	Operating mode		Reset mode	
	Read	**Write**	**Read**	**Write**
0	Control	Control	Control	Control
1	(0xFF)	Command	(0xFF)	Command
2	Status	-	Status	-
3	Interrupt	-	Interrupt	-
4	(0xFF)	-	Acceptance code	Acceptance code
5	(0xFF)	-	Acceptance mask	Acceptance mask
6	(0xFF)	-	Bus timing 0	Bus timing 0
7	(0xFF)	-	Bus timing 1	Bus timing 1
8	(0x00)	-	(0x00)	-
9	(0x00)	-	(0x00)	-
10	TX id1	TX id1	(0xFF)	-
11	TX id2, rtr, dlc	TX id2, rtr, dlc	(0xFF)	-
12	TX data byte 1	TX data byte 1	(0xFF)	-
13	TX data byte 2	TX data byte 2	(0xFF)	-
14	TX data byte 3	TX data byte 3	(0xFF)	-
15	TX data byte 4	TX data byte 4	(0xFF)	-
16	TX data byte 5	TX data byte 5	(0xFF)	-
17	TX data byte 6	TX data byte 6	(0xFF)	-
18	TX data byte 7	TX data byte 7	(0xFF)	-
19	TX data byte 8	TX data byte 8	(0xFF)	-
20	RX id1	-	RX id1	-
21	RX id2, rtr, dlc	-	RX id2, rtr, dlc	-
22	RX data byte 1	-	RX data byte 1	-
23	RX data byte 2	-	RX data byte 2	-
24	RX data byte 3	-	RX data byte 3	-
25	RX data byte 4	-	RX data byte 4	-
26	RX data byte 5	-	RX data byte 5	-
27	RX data byte 6	-	RX data byte 6	-
28	RX data byte 7	-	RX data byte 7	-
29	RX data byte 8	-	RX data byte 8	-
30	(0x00)	-	(0x00)	-
31	Clock divider	Clock divider	Clock divider	Clock divider

18.4.2 Control register

The control register contains interrupt enable bits as well as the reset request bit.

Table 122. Bit interpretation of control register (CR) (address 0)

Bit	Name	Description
CR.7	-	reserved
CR.6	-	reserved
CR.5	-	reserved
CR.4	Overrun Interrupt Enable	1 - enabled, 0 - disabled
CR.3	Error Interrupt Enable	1 - enabled, 0 - disabled
CR.2	Transmit Interrupt Enable	1 - enabled, 0 - disabled
CR.1	Receive Interrupt Enable	1 - enabled, 0 - disabled
CR.0	Reset request	Writing 1 to this bit aborts any ongoing transfer and enters reset mode. Writing 0 returns to operating mode.

18.4.3 Command register

Writing a one to the corresponding bit in this register initiates an action supported by the core.

Table 123. Bit interpretation of command register (CMR) (address 1)

Bit	Name	Description
CMR.7	-	reserved
CMR.6	-	reserved
CMR.5	-	reserved
CMR.4	-	not used (go to sleep in SJA1000 core)
CMR.3	Clear data overrun	Clear the data overrun status bit
CMR.2	Release receive buffer	Free the current receive buffer for new reception
CMR.1	Abort transmission	Aborts a not yet started transmission.
CMR.0	Transmission request	Starts the transfer of the message in the TX buffer

A transmission is started by writing 1 to CMR.0. It can only be aborted by writing 1 to CMR.1 and only if the transfer has not yet started. If the transmission has started it will not be aborted when setting CMR.1 but it will not be retransmitted if an error occurs.

Giving the Release receive buffer command should be done after reading the contents of the receive buffer in order to release this memory. If there is another message waiting in the FIFO a new receive interrupt will be generated (if enabled) and the receive buffer status bit will be set again.

To clear the Data overrun status bit CMR.3 must be written with 1.

18.4.4 Status register

The status register is read only and reflects the current status of the core.

Table 124. Bit interpretation of status register (SR) (address 2)

Bit	Name	Description
SR.7	Bus status	1 when the core is in bus-off and not involved in bus activities
SR.6	Error status	At least one of the error counters have reached or exceeded the CPU warning limit (96).
SR.5	Transmit status	1 when transmitting a message
SR.4	Receive status	1 when receiving a message
SR.3	Transmission complete	1 indicates the last message was successfully transferred.
SR.2	Transmit buffer status	1 means CPU can write into the transmit buffer
SR.1	Data overrun status	1 if a message was lost because no space in fifo.
SR.0	Receive buffer status	1 if messages available in the receive fifo.

Receive buffer status is cleared when the Release receive buffer command is given and set high if there are more messages available in the fifo.

The data overrun status signals that a message which was accepted could not be placed in the fifo because not enough space left. NOTE: This bit differs from the SJA1000 behavior and is set first when the fifo has been read out.

When the transmit buffer status is high the transmit buffer is available to be written into by the CPU. During an on-going transmission the buffer is locked and this bit is 0.

The transmission complete bit is set to 0 when a transmission request has been issued and will not be set to 1 again until a message has successfully been transmitted.

18.4.5 Interrupt register

The interrupt register signals to CPU what caused the interrupt. The interrupt bits are only set if the corresponding interrupt enable bit is set in the control register.

Table 125. Bit interpretation of interrupt register (IR) (address 3)

Bit	Name	Description
IR.7	-	reserved
IR.6	-	reserved
IR.5	-	reserved
IR.4	-	not used (wake-up interrupt of SJA1000)
IR.3	Data overrun interrupt	Set when SR.1 goes from 0 to 1.
IR.2	Error interrupt	Set when the error status or bus status are changed.
IR.1	Transmit interrupt	Set when the transmit buffer is released (status bit 0->1)
IR.0	Receive interrupt	This bit is set while there are more messages in the fifo.

This register is reset on read with the exception of IR.0. Note that this differs from the SJA1000 behavior where all bits are reset on read in BasicCAN mode. This core resets the receive interrupt bit when the release receive buffer command is given (like in PeliCAN mode).

Also note that bit IR.5 through IR.7 reads as 1 but IR.4 is 0.

18.4.6 Transmit buffer

The table below shows the layout of the transmit buffer. In BasicCAN only standard frame messages can be transmitted and received (EFF messages on the bus are ignored).

Table 126. Transmit buffer layout

Addr	Name	Bits							
		7	6	5	4	3	2	1	0
10	ID byte 1	ID.10	ID.9	ID.8	ID.7	ID.6	ID.5	ID.4	ID.3
11	ID byte 2	ID.2	ID.1	ID.0	RTR	DLC.3	DLC.2	DLC.1	DLC.0
12	TX data 1	TX byte 1							
13	TX data 2	TX byte 2							
14	TX data 3	TX byte 3							
15	TX data 4	TX byte 4							
16	TX data 5	TX byte 5							
17	TX data 6	TX byte 6							
18	TX data 7	TX byte 7							
19	TX data 8	TX byte 8							

If the RTR bit is set no data bytes will be sent but DLC is still part of the frame and must be specified according to the requested frame. Note that it is possible to specify a DLC larger than 8 bytes but should not be done for compatibility reasons. If DLC > 8 still only 8 bytes can be sent.

18.4.7 Receive buffer

The receive buffer on address 20 through 29 is the visible part of the 64 byte RX FIFO. Its layout is identical to that of the transmit buffer.

18.4.8 Acceptance filter

Messages can be filtered based on their identifiers using the acceptance code and acceptance mask registers. The top 8 bits of the 11 bit identifier are compared with the acceptance code register only comparing the bits set to zero in the acceptance mask register. If a match is detected the message is stored to the fifo.

18.5 PeliCAN mode

18.5.1 PeliCAN register map

*Table 127.*PeliCAN address allocation

#	Operating mode			Reset mode	
	Read	Write		Read	Write
0	Mode	Mode		Mode	Mode
1	(0x00)	Command		(0x00)	Command
2	Status	-		Status	-
3	Interrupt	-		Interrupt	-
4	Interrupt enable	Interrupt enable		Interrupt enable	Interrupt enable
5	reserved (0x00)	-		reserved (0x00)	-
6	Bus timing 0	-		Bus timing 0	Bus timing 0
7	Bus timing 1	-		Bus timing 1	Bus timing 1
8	(0x00)	-		(0x00)	-
9	(0x00)	-		(0x00)	-
10	reserved (0x00)	-		reserved (0x00)	-
11	Arbitration lost capture	-		Arbitration lost capture	-
12	Error code capture	-		Error code capture	-
13	Error warning limit	-		Error warning limit	Error warning limit
14	RX error counter	-		RX error counter	RX error counter
15	TX error counter	-		TX error counter	TX error counter
16	RX FI SFF	RX FI EFF	TX FI SFF	TX FI EFF	Acceptance code 0	Acceptance code 0
17	RX ID 1	RX ID 1	TX ID 1	TX ID 1	Acceptance code 1	Acceptance code 1
18	RX ID 2	RX ID 2	TX ID 2	TX ID 2	Acceptance code 2	Acceptance code 2
19	RX data 1	RX ID 3	TX data 1	TX ID 3	Acceptance code 3	Acceptance code 3
20	RX data 2	RX ID 4	TX data 2	TX ID 4	Acceptance mask 0	Acceptance mask 0
21	RX data 3	RX data 1	TX data 3	TX data 1	Acceptance mask 1	Acceptance mask 1
22	RX data 4	RX data 2	TX data 4	TX data 2	Acceptance mask 2	Acceptance mask 2
23	RX data 5	RX data 3	TX data 5	TX data 3	Acceptance mask 3	Acceptance mask 3
24	RX data 6	RX data 4	TX data 6	TX data 4	reserved (0x00)	-
25	RX data 7	RX data 5	TX data 7	TX data 5	reserved (0x00)	-
26	RX data 8	RX data 6	TX data 8	TX data 6	reserved (0x00)	-
27	FIFO	RX data 7	-	TX data 7	reserved (0x00)	-
28	FIFO	RX data 8	-	TX data 8	reserved (0x00)	-
29	RX message counter	-		RX msg counter	-
30	(0x00)	-		(0x00)	-
31	Clock divider	Clock divider		Clock divider	Clock divider

The transmit and receive buffers have different layout depending on if standard frame format (SFF) or extended frame format (EFF) is to be transmitted/received. See the specific section below.

18.5.2 Mode register

*Table 128.*Bit interpretation of mode register (MOD) (address 0)

Bit	Name	Description
MOD.7	-	reserved
MOD.6	-	reserved
MOD.5	-	reserved
MOD.4	-	not used (sleep mode in SJA1000)
MOD.3	Acceptance filter mode	1 - single filter mode, 0 - dual filter mode
MOD.2	Self test mode	If set the controller is in self test mode
MOD.1	Listen only mode	If set the controller is in listen only mode
MOD.0	Reset mode	Writing 1 to this bit aborts any ongoing transfer and enters reset mode. Writing 0 returns to operating mode

Writing to MOD.1-3 can only be done when reset mode has been entered previously.

In Listen only mode the core will not send any acknowledgements. Note that unlike the SJA1000 the Opencores core does not become error passive and active error frames are still sent!

When in Self test mode the core can complete a successful transmission without getting an acknowledgement if given the Self reception request command. Note that the core must still be connected to a real bus, it does not do an internal loopback.

18.5.3 Command register

Writing a one to the corresponding bit in this register initiates an action supported by the core.

*Table 129.*Bit interpretation of command register (CMR) (address 1)

Bit	Name	Description
CMR.7	-	reserved
CMR.6	-	reserved
CMR.5	-	reserved
CMR.4	Self reception request	Transmits and simultaneously receives a message
CMR.3	Clear data overrun	Clears the data overrun status bit
CMR.2	Release receive buffer	Free the current receive buffer for new reception
CMR.1	Abort transmission	Aborts a not yet started transmission.
CMR.0	Transmission request	Starts the transfer of the message in the TX buffer

A transmission is started by writing 1 to CMR.0. It can only be aborted by writing 1 to CMR.1 and only if the transfer has not yet started. Setting CMR.0 and CMR.1 simultaneously will result in a so called single shot transfer, i.e. the core will not try to retransmit the message if not successful the first time.

Giving the Release receive buffer command should be done after reading the contents of the receive buffer in order to release this memory. If there is another message waiting in the FIFO a new receive interrupt will be generated (if enabled) and the receive buffer status bit will be set again.

The Self reception request bit together with the self test mode makes it possible to do a self test of the core without any other cores on the bus. A message will simultaneously be transmitted and received and both receive and transmit interrupt will be generated.

18.5.4 Status register

The status register is read only and reflects the current status of the core.

Table 130. Bit interpretation of command register (SR) (address 2)

Bit	Name	Description
SR.7	Bus status	1 when the core is in bus-off and not involved in bus activities
SR.6	Error status	At least one of the error counters have reached or exceeded the error warning limit.
SR.5	Transmit status	1 when transmitting a message
SR.4	Receive status	1 when receiving a message
SR.3	Transmission complete	1 indicates the last message was successfully transferred.
SR.2	Transmit buffer status	1 means CPU can write into the transmit buffer
SR.1	Data overrun status	1 if a message was lost because no space in fifo.
SR.0	Receive buffer status	1 if messages available in the receive fifo.

Receive buffer status is cleared when there are no more messages in the fifo. The data overrun status signals that a message which was accepted could not be placed in the fifo because not enough space left. NOTE: This bit differs from the SJA1000 behavior and is set first when the fifo has been read out.

When the transmit buffer status is high the transmit buffer is available to be written into by the CPU. During an on-going transmission the buffer is locked and this bit is 0.

The transmission complete bit is set to 0 when a transmission request or self reception request has been issued and will not be set to 1 again until a message has successfully been transmitted.

18.5.5 Interrupt register

The interrupt register signals to CPU what caused the interrupt. The interrupt bits are only set if the corresponding interrupt enable bit is set in the interrupt enable register.

Table 131. Bit interpretation of interrupt register (IR) (address 3)

Bit	Name	Description
IR.7	Bus error interrupt	Set if an error on the bus has been detected
IR.6	Arbitration lost interrupt	Set when the core has lost arbitration
IR.5	Error passive interrupt	Set when the core goes between error active and error passive
IR.4	-	not used (wake-up interrupt of SJA1000)
IR.3	Data overrun interrupt	Set when data overrun status bit is set
IR.2	Error warning interrupt	Set on every change of the error status or bus status
IR.1	Transmit interrupt	Set when the transmit buffer is released
IR.0	Receive interrupt	Set while the fifo is not empty.

This register is reset on read with the exception of IR.0 which is reset when the fifo has been emptied.

18.5.6 Interrupt enable register

In the interrupt enable register the separate interrupt sources can be enabled/disabled. If enabled the corresponding bit in the interrupt register can be set and an interrupt generated.

Table 132. Bit interpretation of interrupt enable register (IER) (address 4)

Bit	Name	Description
IR.7	Bus error interrupt	1 - enabled, 0 - disabled
IR.6	Arbitration lost interrupt	1 - enabled, 0 - disabled
IR.5	Error passive interrupt	1 - enabled, 0 - disabled
IR.4	-	not used (wake-up interrupt of SJA1000)
IR.3	Data overrun interrupt	1 - enabled, 0 - disabled
IR.2	Error warning interrupt	1 - enabled, 0 - disabled.
IR.1	Transmit interrupt	1 - enabled, 0 - disabled
IR.0	Receive interrupt	1 - enabled, 0 - disabled

18.5.7 Arbitration lost capture register

Table 133. Bit interpretation of arbitration lost capture register (ALC) (address 11)

Bit	Name	Description
ALC.7-5	-	reserved
ALC.4-0	Bit number	Bit where arbitration is lost

When the core loses arbitration the bit position of the bit stream processor is captured into arbitration lost capture register. The register will not change content again until read out.

18.5.8 Error code capture register

Table 134. Bit interpretation of error code capture register (ECC) (address 12)

Bit	Name	Description
ECC.7-6	Error code	Error code number
ECC.5	Direction	1 - Reception, 0 - transmission error
ECC.4-0	Segment	Where in the frame the error occurred

When a bus error occurs the error code capture register is set according to what kind of error occurred, if it was while transmitting or receiving and where in the frame it happened. As with the ALC register the ECC register will not change value until it has been read out. The table below shows how to interpret bit 7-6 of ECC.

Table 135. Error code interpretation

ECC.7-6	Description
0	Bit error
1	Form error
2	Stuff error
3	Other

Bit 4 downto 0 of the ECC register is interpreted as below

Table 136. Bit interpretation of ECC.4-0

ECC.4-0	Description
0x03	Start of frame
0x02	ID.28 - ID.21
0x06	ID.20 - ID.18
0x04	Bit SRTR
0x05	Bit IDE
0x07	ID.17 - ID.13
0x0F	ID.12 - ID.5
0x0E	ID.4 - ID.0
0x0C	Bit RTR
0x0D	Reserved bit 1
0x09	Reserved bit 0
0x0B	Data length code
0x0A	Data field
0x08	CRC sequence
0x18	CRC delimiter
0x19	Acknowledge slot
0x1B	Acknowledge delimiter
0x1A	End of frame
0x12	Intermission
0x11	Active error flag
0x16	Passive error flag
0x13	Tolerate dominant bits
0x17	Error delimiter
0x1C	Overload flag

18.5.9 Error warning limit register

This registers allows for setting the CPU error warning limit. It defaults to 96. Note that this register is only writable in reset mode.

18.5.10 RX error counter register (address 14)

This register shows the value of the rx error counter. It is writable in reset mode. A bus-off event resets this counter to 0.

18.5.11 TX error counter register (address 15)

This register shows the value of the tx error counter. It is writable in reset mode. If a bus-off event occurs this register is initialized as to count down the protocol defined 128 occurrences of the bus-free signal and the status of the bus-off recovery can be read out from this register. The CPU can force a bus-off by writing 255 to this register. Note that unlike the SJA1000 this core will signal bus-off immediately and not first when entering operating mode. The bus-off recovery sequence starts when entering operating mode after writing 255 to this register in reset mode.

18.5.12 Transmit buffer

The transmit buffer is write-only and mapped on address 16 to 28. Reading of this area is mapped to the receive buffer described in the next section. The layout of the transmit buffer depends on whether a standard frame (SFF) or an extended frame (EFF) is to be sent as seen below.

Table 137.

#	Write (SFF)	Write(EFF)
16	TX frame information	TX frame information
17	TX ID 1	TX ID 1
18	TX ID 2	TX ID 2
19	TX data 1	TX ID 3
20	TX data 2	TX ID 4
21	TX data 3	TX data 1
22	TX data 4	TX data 2
23	TX data 5	TX data 3
24	TX data 6	TX data 4
25	TX data 7	TX data 5
26	TX data 8	TX data 6
27	-	TX data 7
28	-	TX data 8

TX frame information (this field has the same layout for both SFF and EFF frames)

*Table 138.*TX frame information address 16

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
FF	RTR	-	-	DLC.3	DLC.2	DLC.1	DLC.0

Bit 7 - FF selects the frame format, i.e. whether this is to be interpreted as an extended or standard frame. 1 = EFF, 0 = SFF.
Bit 6 - RTR should be set to 1 for an remote transmission request frame.
Bit 5:4 - are don't care.
Bit 3:0 - DLC specifies the Data Length Code and should be a value between 0 and 8. If a value greater than 8 is used 8 bytes will be transmitted.

TX identifier 1 (this field is the same for both SFF and EFF frames)

*Table 139.*TX identifier 1 address 17

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.28	ID.27	ID.26	ID.25	ID.24	ID.23	ID.22	ID.21

Bit 7:0 - The top eight bits of the identifier.

TX identifier 2, SFF frame

*Table 140.*TX identifier 2 address 18

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.20	ID.19	ID.18	-	-	-	-	-

Bit 7:5 - Bottom three bits of an SFF identifier.
Bit 4:0 - Don't care.

TX identifier 2, EFF frame

Table 141. TX identifier 2 address 18

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.20	ID.19	ID.18	ID.17	ID.16	ID.15	ID.14	ID.13

Bit 7:0 - Bit 20 downto 13 of 29 bit EFF identifier.

TX identifier 3, EFF frame

Table 142. TX identifier 3 address 19

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.12	ID.11	ID.10	ID.9	ID.8	ID.7	ID.6	ID.5

Bit 7:0 - Bit 12 downto 5 of 29 bit EFF identifier.

TX identifier 4, EFF frame

Table 143. TX identifier 4 address 20

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.4	ID.3	ID.2	ID.1	ID.0	-	-	-

Bit 7:3 - Bit 4 downto 0 of 29 bit EFF identifier
Bit 2:0 -   Don't care

Data field

For SFF frames the data field is located at address 19 to 26 and for EFF frames at 21 to 28. The data is transmitted starting from the MSB at the lowest address.

18.5.13 Receive buffer

Table 144.

#	Read (SFF)	Read (EFF)
16	RX frame information	RX frame information
17	RX ID 1	RX ID 1
18	RX ID 2	RX ID 2
19	RX data 1	RX ID 3
20	RX data 2	RX ID 4
21	RX data 3	RX data 1
22	RX data 4	RX data 2
23	RX data 5	RX data 3
24	RX data 6	RX data 4
25	RX data 7	RX data 5
26	RX data 8	RX data 6
27	RX FI of next message in fifo	RX data 7
28	RX ID1 of next message in fifo	RX data 8

RX frame information (this field has the same layout for both SFF and EFF frames)

*Table 145.*RX frame information address 16

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
FF	RTR	0	0	DLC.3	DLC.2	DLC.1	DLC.0

Bit 7 -   Frame format of received message. 1 = EFF, 0 = SFF.
Bit 6 -   1 if RTR frame.
Bit 5:4 - Always 0.
Bit 3:0 - DLC specifies the Data Length Code.

RX identifier 1(this field is the same for both SFF and EFF frames)

*Table 146.*RX identifier 1 address 17

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.28	ID.27	ID.26	ID.25	ID.24	ID.23	ID.22	ID.21

Bit 7:0 - The top eight bits of the identifier.

RX identifier 2, SFF frame

*Table 147.*RX identifier 2 address 18

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.20	ID.19	ID.18	RTR	0	0	0	0

Bit 7:5 - Bottom three bits of an SFF identifier.
Bit 4 -   1 if RTR frame.
Bit 3:0 - Always 0.

RX identifier 2, EFF frame

*Table 148.*RX identifier 2 address 18

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.20	ID.19	ID.18	ID.17	ID.16	ID.15	ID.14	ID.13

Bit 7:0 - Bit 20 downto 13 of 29 bit EFF identifier.

RX identifier 3, EFF frame

*Table 149.*RX identifier 3 address 19

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.12	ID.11	ID.10	ID.9	ID.8	ID.7	ID.6	ID.5

Bit 7:0 - Bit 12 downto 5 of 29 bit EFF identifier.

RX identifier 4, EFF frame

*Table 150.*RX identifier 4 address 20

Bit 7	Bit 6	Bit 5	Bit 4	Bit 3	Bit 2	Bit 1	Bit 0
ID.4	ID.3	ID.2	ID.1	ID.0	RTR	0	0

Bit 7:3 - Bit 4 downto 0 of 29 bit EFF identifier
Bit 2-   1 if RTR frame
Bit 1:0 - Don't care

Data field

For received SFF frames the data field is located at address 19 to 26 and for EFF frames at 21 to 28.

18.5.14 Acceptance filter

The acceptance filter can be used to filter out messages not meeting certain demands. If a message is filtered out it will not be put into the receive fifo and the CPU will not have to deal with it.

There are two different filtering modes, single and dual filter. Which one is used is controlled by bit 3 in the mode register. In single filter mode only one 4 byte filter is used. In dual filter two smaller filters are used and if either of these signals a match the message is accepted. Each filter consists of two parts the acceptance code and the acceptance mask. The code registers are used for specifying the pattern to match and the mask registers specify don't care bits. In total eight registers are used for the acceptance filter as shown in the table below. Note that they are only read/writable in reset mode.

*Table 151.*Acceptance filter registers

Address	Description
16	Acceptance code 0 (ACR0)
17	Acceptance code 1 (ACR1)
18	Acceptance code 2 (ACR2)
19	Acceptance code 3 (ACR3)
20	Acceptance mask 0 (AMR0)
21	Acceptance mask 1 (AMR1)
22	Acceptance mask 2 (AMR2)
23	Acceptance mask 3 (AMR3)

Single filter mode, standard frame

When receiving a standard frame in single filter mode the registers ACR0-3 are compared against the incoming message in the following way:

ACR0.7-0 & ACR1.7-5 are compared to ID.28-18
ACR1.4 is compared to the RTR bit.
ACR1.3-0 are unused.
ACR2 & ACR3 are compared to data byte 1 & 2.

The corresponding bits in the AMR registers selects if the results of the comparison doesn't matter. A set bit in the mask register means don't care.

Single filter mode, extended frame

When receiving an extended frame in single filter mode the registers ACR0-3 are compared against the incoming message in the following way:

ACR0.7-0 & ACR1.7-0 are compared to ID.28-13
ACR2.7-0 & ACR3.7-3 are compared to ID.12-0
ACR3.2 are compared to the RTR bit
ACR3.1-0 are unused.

The corresponding bits in the AMR registers selects if the results of the comparison doesn't matter. A set bit in the mask register means don't care.

Dual filter mode, standard frame

When receiving a standard frame in dual filter mode the registers ACR0-3 are compared against the incoming message in the following way:

Filter 1
ACR0.7-0 & ACR1.7-5 are compared to ID.28-18
ACR1.4 is compared to the RTR bit.
ACR1.3-0 are compared against upper nibble of data byte 1
ACR3.3-0 are compared against lower nibble of data byte 1
Filter 2
ACR2.7-0 & ACR3.7-5 are compared to ID.28-18
ACR3.4 is compared to the RTR bit.

The corresponding bits in the AMR registers selects if the results of the comparison doesn't matter. A set bit in the mask register means don't care.

Dual filter mode, extended frame

When receiving a standard frame in dual filter mode the registers ACR0-3 are compared against the incoming message in the following way:

Filter 1:   ACR0.7-0 & ACR1.7-0 are compared to ID.28-13

Filter 2:   ACR2.7-0 & ACR3.7-0 are compared to ID.28-13

The corresponding bits in the AMR registers selects if the results of the comparison doesn't matter. A set bit in the mask register means don't care.

18.5.15 RX message counter

The RX message counter register at address 29 holds the number of messages currently stored in the receive fifo. The top three bits are always 0.

18.6   Common registers

There are three common registers with the same addresses and the same functionality in both Basic-CAN and PeliCAN mode. These are the clock divider register and bus timing register 0 and 1.

18.6.1   Clock divider register

The only real function of this register in the GRLIB version of the Opencores CAN is to choose between PeliCAN and BasicCAN. The clkout output of the Opencore CAN core is not connected and it is its frequency that can be controlled with this register.

*Table 152.*Bit interpretation of clock divider register (CDR) (address 31)

Bit	Name	Description
CDR.7	CAN mode	1 - PeliCAN, 0 - BasicCAN
CDR.6	-	unused (cbp bit of SJA1000)
CDR.5	-	unused (rxinten bit of SJA1000)
CDR.4	-	reserved
CDR.3	Clock off	Disable the clkout output
CDR.2-0	Clock divisor	Frequency selector

18.6.2   Bus timing 0

*Table 153.*Bit interpretation of bus timing 0 register (BTR0) (address 6)

Bit	Name	Description
BTR0.7-6	SJW	Synchronization jump width
BTR0.5-0	BRP	Baud rate prescaler

The CAN core system clock is calculated as:

$$t_{scl} = 2 * t_{clk} * (BRP+1)$$

where t_{clk} is the system clock.

The sync jump width defines how many clock cycles (t_{scl}) a bit period may be adjusted with by one re-synchronization.

18.6.3 Bus timing 1

*Table 154.*Bit interpretation of bus timing 1 register (BTR1) (address 7)

Bit	Name	Description
BTR1.7	SAM	1 - The bus is sampled three times, 0 - single sample point
BTR1.6-4	TSEG2	Time segment 2
BTR1.3-0	TSEG1	Time segment 1

The CAN bus bit period is determined by the CAN system clock and time segment 1 and 2 as shown in the equations below:

$$t_{tseg1} = t_{scl} * (TSEG1+1)$$
$$t_{tseg2} = t_{scl} * (TSEG2+1)$$
$$t_{bit} = t_{tseg1} + t_{tseg2} + t_{scl}$$

The additional t_{scl} term comes from the initial sync segment. Sampling is done between TSEG1 and TSEG2 in the bit period.

18.7 Design considerations

This section lists known differences between this CAN controller and SJA1000 on which is it based:

- All bits related to sleep mode are unavailable

- Output control and test registers do not exist (reads 0x00)

- Clock divisor register bit 6 (CBP) and 5 (RXINTEN) are not implemented

- Overrun irq and status not set until fifo is read out

BasicCAN specific differences:

- The receive irq bit is not reset on read, works like in PeliCAN mode

- Bit CR.6 always reads 0 and is not a flip flop with no effect as in SJA1000

PeliCAN specific differences:

- Writing 256 to tx error counter gives immediate bus-off when still in reset mode

- Read Buffer Start Address register does not exist

- Addresses above 31 are not implemented (i.e. the internal RAM/FIFO access)

- The core transmits active error frames in Listen only mode

18.8 Signal definitions

The signals are described in table 155.

*Table 155.*Signal definitions

Signal name	Type	Function	Active
CANTX[]	Output	CAN transmit data	Low
CANRX[]	Input	CAN receive data	Low

19   Obsolete

Proprietary function not supported.

19.1   Signal definitions

The signals are described in table 156.

*Table 156.*Signal definitions

Signal name	Type	Function	Active
-	Output	Proprietary	
-	Output	Proprietary	
-	Output	Proprietary	
-	Output	Proprietary	
-	Output	Proprietary	

20 CAN Bus multiplexer

20.1 Overview

The CAN bus multiplexer provides a way to select which of the OC-CAN1 and OC-CAN2 cores should drive the two CAN bus interfaces. OC-CAN1 can only be connected to CAN bus interface A and OC-CAN2 only to CAN bus interface B. The CAN multiplexer control register determines which core that is active on each bus.
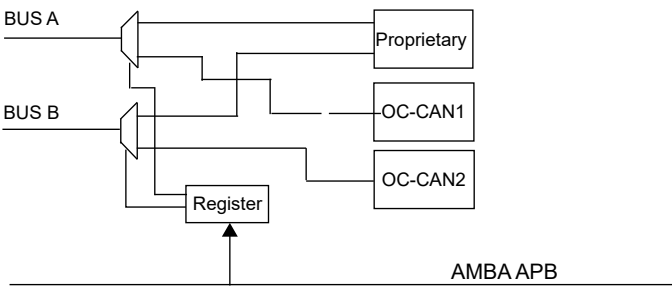
Figure 66. CAN multiplexer

20.2 Operation

The select signal for each multiplexer can be controlled through an APB-register. The following options are possible:

*Table 157.*CAN Multiplexer control register configuration

Option	Register value
Unused	0
OC-CAN1 on bus A	1
OC-CAN2 on bus B	2
OC-CAN1 on bus A, OC-CAN2 on bus B	3

20.3 Registers

*Table 158.*CAN multiplexer registers

APB address offset	Register
0x80000500	CAN Multiplexer control register

Table 159. CAN Multiplexer control register

31	2	1	0
RESERVED		BUSB	BUSA

31: 2   RESERVED

1   Bus B select (BUSB) - 0, unused. 1, OC-CAN2 on bus B. 0 on reset.

0   Bus A select (BUSA) - 0, unused. 1, OC-CAN1 on bus A. 0 on reset.

21 MIL-STD-1553B BC/RT/BM

21.1 Overview

The 1553 interface provides a complete Mil-Std-1553B Bus Controller (BC), Remote Terminal (RT) or Monitor Terminal (MT). The interface connects to the MIL-STD-1553B bus through external transceivers and transformers. The interface is based on the Actel Core1553BRM core version 2.16.

The interface consists of six main blocks: 1553 encoder, 1553B decoders, a protocol controller block, AMBA bus interface, command word legality interface, and a backend interface.

The interface can be configured to provide all three functions BC, RT and MT or any combination of the three. All variations use all six blocks except for the command legalization interface, which is only required on RT functions that implement RT legalization function externally.

A single 1553 encoder takes each word to be transmitted and serializes it using Manchester encoding. The encoder also includes independent logic to prevent the interface from transmitting for greater than the allowed period as well as loopback fail logic. The loopback logic monitors the received data and verifies that the interface has correctly received every word that it transmits. The output of the encoder is gated with the bus enable signals to select which buses the interface should be transmitting on. Two decoders take the serial Manchester received data from each bus and extract the received data words.

The decoder contains a digital phased lock loop (PLL) that generates a recovery clock used to sample the incoming serial data. The data is then de-serialized and the 16-bit word decoded. The decoder detects whether a command, status, or data word has been received, and checks that no Manchester encoding or parity errors occurred in the word.

The protocol controller block handles all the message sequencing and error recovery for all three operating modes, Bus Controller, Remote Terminal, and Bus Monitor. This is complex state machine that processes messages based on the message tables setup in memory, or reacts to incoming command words. The protocol controller implementation varies depending on which functions are implemented. The AMBA interface allows a system processor to access the control registers. It also allows the processor to directly access the memory connected to the backend interface, this simplifies the system design.

The interface comprises 33 16-bit registers. Of the 33 registers, 17 are used for control function and 16 for RT command legalization. The core generates interrupt 14.
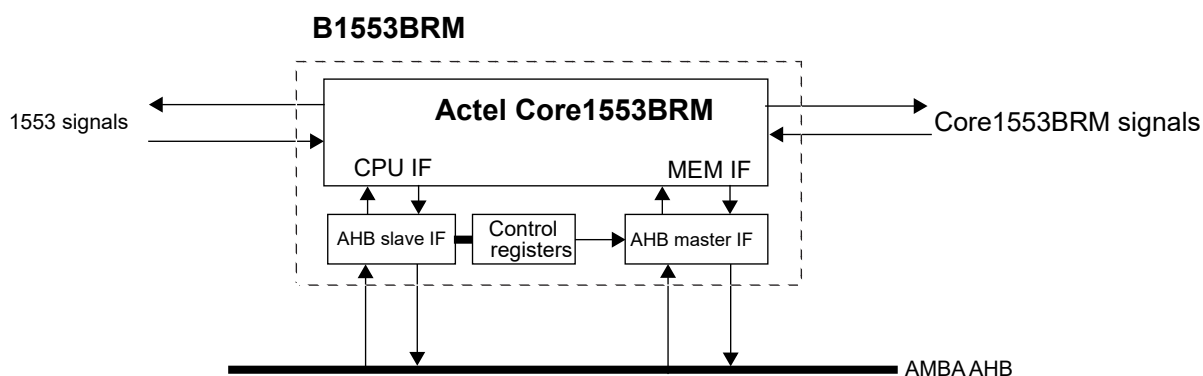

Figure 67. Block diagram

21.2   AHB interface

The amount of memory that the Mil-Std-1553B interface can address is 128 KiB. The base address of this memory area must be aligned to a boundary of its own size and written into the AHB page address register.

The 16 bit address provided by the Core1553BRM core is shifted left one bit, and forms the AHB address together with the AHB page address register. Note that all pointers given to the Core1553BRM core needs to be right shifted one bit because of this.

When the Core1553BRM core has been granted access to the bus it expects to be able to do a series of uninterrupted accesses. To handle this requirement the AHB master locks the bus during these transfers. In the worst case, the Core1553BRM can do up to 7 writes in one such access and each write takes 2 plus the number of waitstate cycles with 4 idle cycles between each write strobe. This means care has to be taken if using two simultaneous active Core1553BRM cores on the same AHB bus.All AHB accesses are done as half word single transfers.

The AMBA AHB protection control signal is driven permanently with "0011" i.e a not cacheable, not bufferable, privileged data access. During all AHB accesses the AMBA AHB lock signal is driven with `1' and `0' otherwise.

21.3   1553 Clock generation

The B1553BRM core needs a 24, 20 or 16 MHz clock to generate the 1553 waveforms. The frequency used can be configured through a register in the core. This clock can either be supplied directly to the B1553BRM core through the 1553CK pin, the system clock or it can be generated through a clock divider that divides the system clock and is programmable through the GPREG.

If the system clock is used to generate the BRM clock it must be one of the frequencies in the table below.

TABLE 160. BRM frequencies

System clock (MHz) [1]	Division	BRM frequency (MHz)
32	2	16 [2]
40	2	20
48	2	24
64	4	16
80	4	20
96	4	24
120	6	20
144	6	24

Note 1: The system clock frequency must always be higher than BRM clock frequency, regardless of how the BRM clock is generated.

Note 2: This specific case has not been validated.

The access time requirements of the B1553BRM core from the information in the Core1553BRM handbook [1553BRM] has been translated into requirements on arbitration latency and access time for the B1553BRM core on the AMBA on-chip bus. This takes into account the transfer between clock domains and translation into AMBA accesses that are done inside the B1553BRM core.

The BRM frequency of 12 MHz allows much less time for the access to complete than the other configurations, and for this reason this frequency is not been listed as supported.

The B1553BRM and GRETH cores have elevated priority in the GR712RC arbiter over all other AMBA masters, meaning that for latency calculation the B1553BRM will only have to wait for the currently active master to complete its access before getting access to the AMBA bus. For the other AMBA masters a round-robin arbitration scheme is used. The B1553BRM and GRETH cores are never active at the same time since they conflict in the I/O matrix.

TABLE 161. System versus BRM clock frequencies and latency times

System clock	BRM clock	Max arbitration latency		Max access time	
(MHz)	(MHz)	(system periods)	(ns)	(system periods)	(ns)
24 [1]	12	91	3792	6	250
32 [2]	16	123	3844	14	251
40	20	165	4125	20	500
48	24	231	4812	22	458
64	16	249	3891	30	469
80	20	333	4162	42	525
96	24	333	4844	46	525
120	20	501	4175	64	533
144	24	699	4175	70	533

Note 1: This specific case is not supported.

Note 2: This specific case has not been validated.

GR712RC

21.4 Registers

The core is programmed through registers mapped into AHB I/O address space. The internal registers of Core1553BRM are mapped on the 33 lowest AHB addresses. These addresses are 32-bit word aligned although only the lowest 16 bits are used. Refer to the *Actel Core1553BRM MIL-STD-1553 BC, RT, and MT* data sheet for detailed information.

*Table 162.*B1553BRM registers

AHB address	Register
0x0xFFF00000 - 0xFFF00084	Core1553BRM registers
0xFFF00100	B1553BRM status/control
0xFFF00104	B1553BRM interrupt settings
0xFFF00108	AHB page address register

B1553BRM status/control register

31	14	13	12	11 9	8 5	4	3	2	1	0
RESERVED		busrst	*reset*	RES	RES	rtaderr	memfail	busy	active	ssysfn

Figure 68. B1553BRM status/control register

13   Bus reset. If set a bus reset mode code has been received. Generates an irq when set.

12   Reset. For asynchronous toplevel only. Software reset. Self clearing.

11:9   Reserved.

8:5   Reserved.

4:        Address error. Shows the value of the rtaderr output from Core1553BRM.

3:        Memory failure. Shows the value of the memfail output from Core1553BRM.

2:        Busy. Shows the value of the busy output from Core1553BRM.

1:        Active. Show the value of the active output from Core1553BRM.

0:        Ssyfn. Connects directly to the ssyfn input of the Core1553BRM core. Resets to 0.

Note: Ssyfn is reset to '0', which causes the BRM to set the subsystem flag bit in the status words sent out.

B1553BRM interrupt register

31	2	1	0
RESERVED	intackm	intackh	intlevel

Figure 69. B1553RM interrupt register

2:        Message interrupt acknowledge. Controls the intackm input signal of the Core1553BRM core.

1:        Hardware interrupt acknowledge. Controls the intackh input signal of the Core1553BRM core.

0:        Interrupt level. Controls the intlevel input signal of the Core1553BRM core.

AHB page address register

31	abits	0
ahbaddr		RESERVED

Figure 70. AHB page address register

[31:17]: Holds the top most bits of the AHB address of the allocated memory area.

GR712RC

21.5 Signal definitions

The signals are described in table 163.

Table 163. Signal definitions

Signal name	Type	Function	Active
1553CK	Input	Clock	
1553RXENA	Output	Enable for the A receiver	High
1553RXA	Input	Positive data input from the A receiver	High
1553RXNA	Input	Negative data input from the A receiver	Low
1553RXENB	Output	Enable for the B receiver	High
1553RXB	Input	Positive data to the B receiver	High
1553RXNB	Input	Negative data to the B receiver	Low
1553TXINHA	Output	Inhibit for the A transmitter	High
1553TXA	Output	Positive data to the A transmitter	High
1553TXNA	Output	Negative data to the A transmitter	Low
1553TXINHB	Output	Inhibit for the B transmitter	High
1553TXB	Output	Positive output to the B transmitter	High
1553TXNB	Output	Negative output to the B transmitter	Low

22   I^{2}C-master

22.1   Overview

The I^{2}C-master core is compatible with Philips I^{2}C standard and supports 7- and 10-bit addressing. Standard-mode (100 kb/s) and Fast-mode (400 kb/s) operation are supported directly. External pull-up resistors must be supplied for both bus lines.
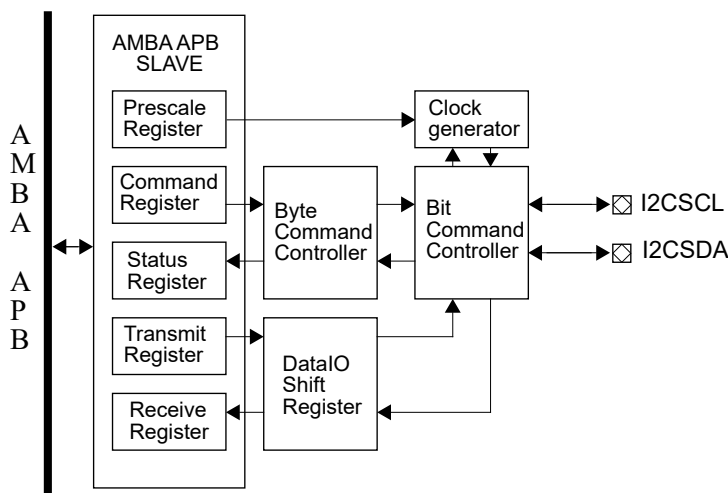

Figure 71. Block diagram

22.2   Operation

22.2.1   Transmission protocol

The I^{2}C-bus is a simple 2-wire serial multi-master bus with collision detection and arbitration. The bus consists of a serial data line (I2CSDA) and a serial clock line (I2CSCL). The I^{2}C standard defines three transmission speeds; Standard (100 kb/s), Fast (400 kb/s) and High speed (3.4 Mb/s).

A transfer on the I^{2}C-bus begins with a START condition. A START condition is defined as a high to low transition of the I2CSDA line while I2CSCL is high. Transfers end with a STOP condition, defined as a low to high transition of the I2CSDA line while I2CSCL is high. These conditions are always generated by a master. The bus is considered to be busy after the START condition and is free after a certain amount of time following a STOP condition. The bus free time required between a STOP and a START condition is defined in the I^{2}C-bus specification and is dependent on the bus bit rate.

Figure 72 shows a data transfer taking place over the I^{2}C-bus. The master first generates a START condition and then transmits the 7-bit slave address. The bit following the slave address is the R/$\overline{\text{W}}$ bit which determines the direction of the data transfer. In this case the R/$\overline{\text{W}}$ bit is zero indicating a write operation. After the master has transmitted the address and the R/$\overline{\text{W}}$ bit it releases the I2CSDA line. The receiver pulls the I2CSDA line low to acknowledge the transfer. If the receiver does not acknowledge the transfer, the master may generate a STOP condition to abort the transfer or start a new transfer by generating a repeated START condition.

After the first byte has been acknowledged the master transmits the data byte. If the R/$\overline{\text{W}}$ bit had been set to '1' the master would have acted as a receiver during this phase of the transfer. After the data byte has been transferred the receiver acknowledges the byte and the master generates a STOP condition to complete the transfer. Section 22.2.4 contains three more example transfers from the perspective of a software driver.

Figure 72. Complete I²C data transfer

If the data bitrate is too high for a slave device, it may stretch the clock period by keeping I2CSCL low after the master has driven I2CSCL low.

22.2.2 Clock generation

The core uses the prescale register to determine the frequency of the I2CSCL clock line and of the 5*I2CSCL clock that the core uses internally. To calculate the prescale value use the formula:

$$Prescale = \frac{AMBAclockfrequency}{5 \cdot SCLfrequency} - 1$$

The *SCLfrequency* is 100 kHz for Standard-mode operation (100 kb/s) and 400 kHz for Fast mode operation. To use the core in Standard-mode in a system with a 60 MHz clock driving the AMBA bus the required prescale value is:

$$Prescale = \frac{60Mhz}{5 \cdot 100kHz} - 1 = 119 = 0x77$$

Note that the prescale register should only be changed when the core is disabled. The minimum recommended prescale value is 3 due to synchronization issues. Lower values may cause the master to violate I²C timing requirements. This limits the minimum system frequency to 2 MHz for operation in Standard-mode.

22.2.3 Interrupts

The core generates interrupt 28. The interrupt can be enabled by setting the EN bit in the control register.

22.2.4 Software operational model

The core is initialized by writing an appropriate value to the clock prescale register and then setting the enable (EN) bit in the control register. Interrupts are enabled via the interrupt enable (IEN) bit in the control register.

To write a byte to a slave the I²C-master must generate a START condition and send the slave address with the R/W̄ bit set to '0'. After the slave has acknowledged the address, the master transmits the

data, waits for an acknowledge and generates a STOP condition. The sequence below instructs the core to perform a write:

1. Left-shift the I^{2}C-device address one position and write the result to the transmit register. The least significant bit of the transmit register (R/$\overline{W}$) is set to '0'.

2. Generate START condition and send contents of transmit register by setting the STA and WR bits in the command register.

3. Wait for interrupt, or for TIP bit in the status register to go low.

4. Read RxACK bit in status register. If RxACK is low the slave has acknowledged the transfer, proceed to step 5. If RxACK is set the device did not acknowledge the transfer, go to step 1.

5. Write the slave-data to the transmit register.

6. Send the data to the slave and generate a stop condition by setting STO and WR in the command register.

7. Wait for interrupt, or for TIP bit in the status register to go low.

8. Verify that the slave has acknowledged the data by reading the RxACK bit in the status register. RxACK should not be set.

To read a byte from an I^{2}C-connected memory much of the sequence above is repeated. The data written in this case is the memory location on the I^{2}C slave. After the address has been written the master generates a repeated START condition and reads the data from the slave. The sequence that software should perform to read from a memory device:

1. Left-shift the I^{2}C-device address one position and write the result to the transmit register. The least significant bit of the transmit register (R/W) is set to '0'.

2. Generate START condition and send contents of transmit register by setting the STA and WR bits in the command register.

3. Wait for interrupt or for TIP bit in the status register to go low.

4. Read RxACK bit in status register. If RxACK is low the slave has acknowledged the transfer, proceed to step 5. If RxACK is set the device did not acknowledge the transfer, go to step 1.

5. Write the memory location to be read from the slave to the transmit register.

6. Set the WR bit in the command register. Note that a STOP condition is not generated here.

7. Wait for interrupt, or for TIP bit in the status register to go low.

8. Read RxACK bit in the status register. RxACK should be low.

9. Address the I^{2}C-slave again by writing its left-shifted address into the transmit register. Set the least significant bit of the transmit register (R/W) to '1' to read from the slave.

10. Set the STA and WR bits in the command register to generate a repeated START condition.

11. Wait for interrupt, or for TIP bit in the status register to go low.

12. Read RxACK bit in the status register. The slave should acknowledge the transfer.

13. Prepare to receive the data read from the I^{2}C-connected memory. Set bits RD, ACK and STO on the command register. Setting the ACK bit NAKs the received data and signifies the end of the transfer.

14. Wait for interrupt, or for TIP in the status register to go low.

15. The received data can now be read from the receive register.

To perform sequential reads the master can iterate over steps 13 - 15 by not setting the ACK and STO bits in step 13. To end the sequential reads the ACK and STO bits are set. Consult the documentation of the I^{2}C-slave to see if sequential reads are supported.

The final sequence illustrates how to write one byte to an I^{2}C-slave which requires addressing. First the slave is addressed and the memory location on the slave is transmitted. After the slave has acknowledged the memory location the data to be written is transmitted without a generating a new START condition:

1. Left-shift the I^{2}C-device address one position and write the result to the transmit register. The least significant bit of the transmit register (R/W) is set to '0'.

2. Generate START condition and send contents of transmit register by setting the STA and WR bits in the command register.

3. Wait for interrupt or for TIP bit in the status register to go low.

4. Read RxACK bit in status register. If RxACK is low the slave has acknowledged the transfer, proceed to step 5. If RxACK is set the device did not acknowledge the transfer, go to step 1.

5. Write the memory location to be written from the slave to the transmit register.

6. Set the WR bit in the command register.

7. Wait for interrupt, or for TIP bit in the status register to go low.

8. Read RxACK bit in the status register. RxACK should be low.

9. Write the data byte to the transmit register.

10. Set WR and STO in the command register to send the data byte and then generate a STOP condition.

11. Wait for interrupt, or for TIP bit in the status register to go low.

12. Check RxACK bit in the status register. If the write succeeded the slave should acknowledge the data byte transfer.

The example sequences presented here can be generally applied to I^{2}C-slaves. However, some devices may deviate from the protocol above, please consult the documentation of the I^{2}C-slave in question. Note that a software driver should also monitor the arbitration lost (AL) bit in the status register.

22.3   Registers

The core is programmed through registers mapped into APB address space.

*Table 164.*I^{2}C-master registers

APB address	Register
0x80000C00	Clock prescale register
0x80000C04	Control register
0x80000C08	Transmit register*
0x80000C08	Receive register**
0x80000C0C	Command register*
0x80000C0C	Status register**

* Write only

** Read only

Table 165. I^{2}C-master Clock prescale register

31				16	15	7	6	5	4	3	2	1	0
RESERVED					Clock prescale								

Table 165. I²C-master Clock prescale register

31: 16	RESERVED
15:0	Clock prescale - Value is used to prescale the I2CSCL clock line. Do not change the value of this register unless the EN field of the control register is set to '0'. The minimum recommended value of this register is 0x0003. Lower values may cause the master to violate I²C timing requirements due to synchronization issues.

Table 166. I²C-master control register

31			8	7	6	5		0
RESERVED				EN	IEN	RESERVED		

31: 8	RESERVED
7	Enable (EN) - Enable I²C core. The core is enabled when this bit is set to '1'.
6	Interrupt enable (IEN) - When this bit is set to '1' the core will generate interrupts upon transfer completion.
5:0	RESERVED

Table 167. I²C-master transmit register

31		8	7		1	0
RESERVED			TDATA			RW

31: 8	RESERVED
7:1	Transmit data (TDATA) - Most significant bits of next byte to transmit via I²C
0	Read/Write (RW) - In a data transfer this is the data's least significant bit. In a slave address transfer this is the RW bit. '1' reads from the slave and '0' writes to the slave.

Table 168. I²C-master receive register

31		8	7		0
RESERVED			RDATA		

31: 8	RESERVED
7:0	Receive data (RDATA) - Last byte received over I²C-bus.

Table 169. I²C-master command register

31		8	7	6	5	4	3	2	1	0
RESERVED			STA	STO	RD	WR	ACK	RESERVED		IACK

31: 8	RESERVED
7	Start (STA) - Generate START condition on I²C-bus. This bit is also used to generate repeated START conditions.
6	Stop (STO) - Generate STOP condition
5	Read (RD) - Read from slave
4	Write (WR) - Write to slave
3	Acknowledge (ACK) - Used when acting as a receiver. '0' sends an ACK, '1' sends a NACK.
2:1	RESERVED
0	Interrupt acknowledge (IACK) - Clears interrupt flag (IF) in status register.

Table 170. I²C-master status register

31		8	7	6	5	4	3	2	1	0
RESERVED			RxACK	BUSY	AL	RESERVED			TIP	IF

31: 8	RESERVED

Table 170. I^{2}C-master status register

7	Receive acknowledge (RxACK) - Received acknowledge from slave. '1' when no acknowledge is received, '0' when slave has acknowledged the transfer.
6	I^{2}C-bus busy (BUSY) - This bit is set to '1' when a start signal is detected and reset to '0' when a stop signal is detected.
5	Arbitration lost (AL) - Set to '1' when the core has lost arbitration. This happens when a stop signal is detected but not requested or when the master drives I2CSDA high but I2CSDA is low.
4:2	RESERVED
1	Transfer in progress (TIP) - '1' when transferring data and '0' when the transfer is complete. This bit is also set when the core will generate a STOP condition.
0	Interrupt flag (IF) - This bit is set when a byte transfer has been completed and when arbitration is lost. If IEN in the control register is set an interrupt will be generated. New interrupts will be generated even if this bit has not been cleared.

22.4 Signal definitions

The signals are described in table 171.

Table 171. Signal definitions

Signal name	Type	Function	Active
I2CSCL	Input/Output	I^{2}C clock line	-
I2CSDA	Input/Output	I^{2}C data line	-

23 SPI Controller

23.1 Overview

The SPI controller provides a link between the AMBA APB bus and the Serial Peripheral Interface (SPI) bus. The SPI core is controlled through the APB address space, and can only work as master. The SPI bus parameters are highly configurable via registers, and the controller has configurable word length, bit ordering and clock gap insertion. To facilitate reuse of existing software drivers, the controller's interface is compatible with the SPI controller interface in the Freescale MPC83xx series.
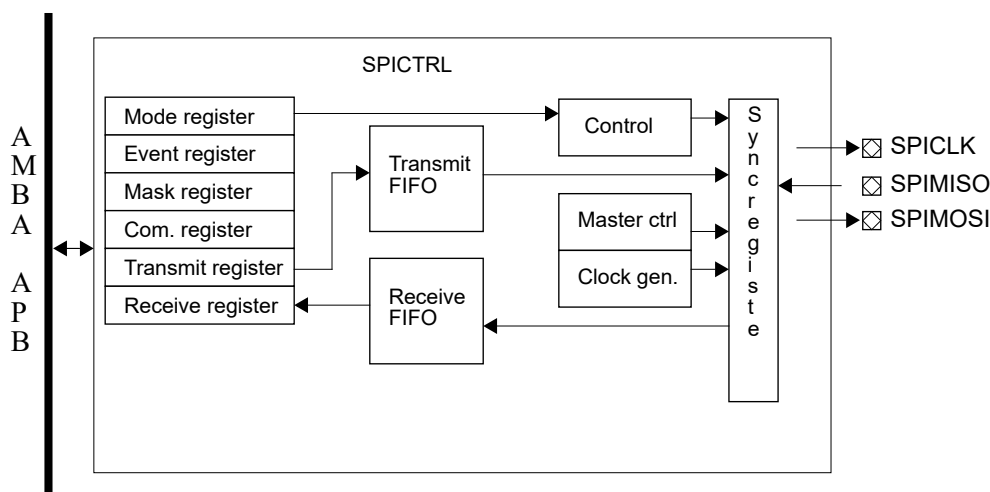
Figure 73. Block diagram

23.2 Operation

23.2.1 SPI transmission protocol

The SPI bus is a full-duplex synchronous serial bus. Transmission starts when the clock line SPICLK transitions from its idle state. Data is transferred from the master through the Master-Output-Slave-Input (SPIMOSI) signal and from the slave through the Master-Input-Slave-Output (SPIMISO) signal. In a system with only one master and one slave, the Slave Select input of the slave may be always active and the master does not need to have a slave select output.

During a transmission on the SPI bus data is either changed or read at a transition of SPICLK. If data has been read at edge n, data is changed at edge n+1. If data is read at the first transition of SPICLK the bus is said to have clock phase 0, and if data is changed at the first transition of SPICLK the bus has clock phase 1. The idle state of SPICLK may be either high or low. If the idle state of SPICLK is low, the bus has clock polarity 0 and if the idle state is high the clock polarity is 1. The combined values of clock polarity (CPOL) and clock phase (CPHA) determine the mode of the SPI bus. Figure 74 shows one byte (0x55) being transferred MSb first over the SPI bus under the four different modes. Note that the idle state of the SPIMOSI line is '1' and that CPHA = 0 means that the devices must have data ready before the first transition of SPICLK. The figure does not include the SPIMISO signal, the behavior of this line is the same as for the SPIMOSI signal. However, due to synchronization issues the SPIMISO signal will be delayed when the core is operating in slave mode.
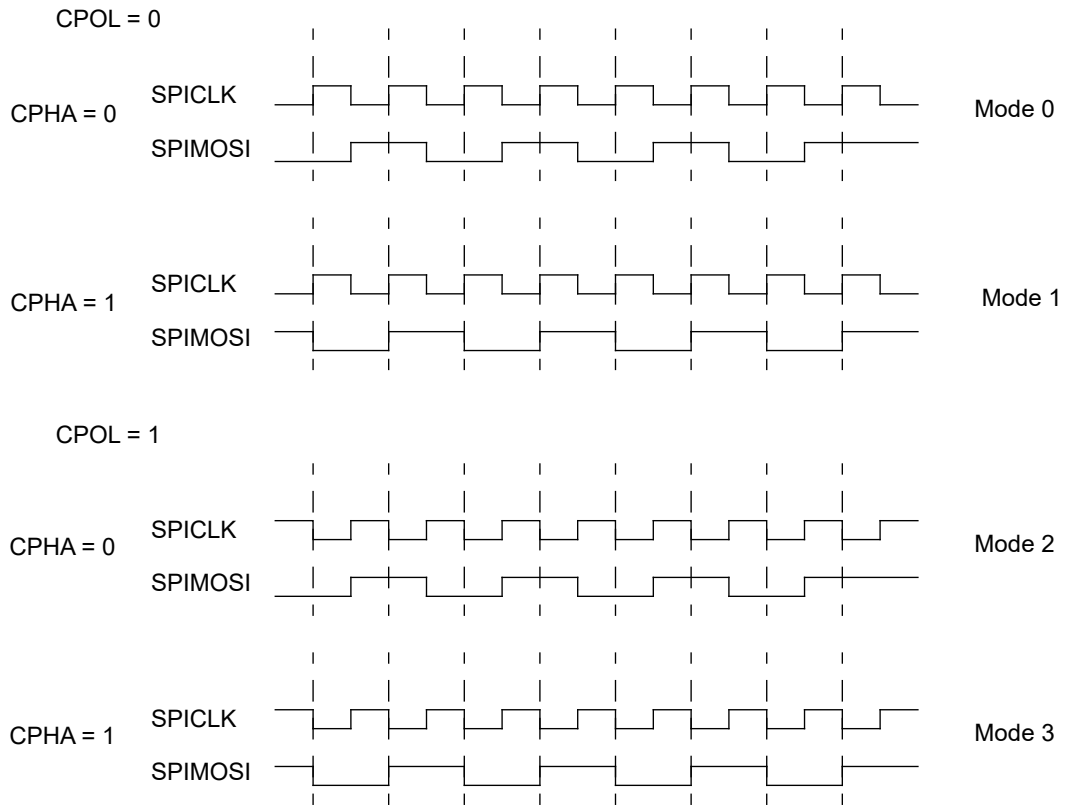
Figure 74. SPI transfer of byte 0x55 in all modes

23.2.2 Receive and transmit queues

The core's transmit queue consists of the transmit register and the transmit FIFO. The receive queue consists of the receive register and the receive FIFO. The total number of words that can exist in each queue is thus the FIFO depth plus one. When the core has one or more free slots in the transmit queue it will assert the Not full (NF) bit in the event register. Software may only write to the transmit register when this bit is asserted. When the core has received a word, as defined by word length (LEN) in the Mode register, it will place the data in the receive queue. When the receive queue has one or more elements stored the Event register bit Not empty (NE) will be asserted. The receive register will only contain valid data if the Not empty bit is asserted and software should not access the receive register unless this bit is set. If the receive queue is full and the core receives a new word, an overrun condition will occur. The received data will be discarded and the Overrun (OV) bit in the Event register will be set.

23.2.3 Clock generation

The core only generates the clock in master mode, the generated frequency depends on the system clock frequency and the Mode register fields DIV16, FACT, and PM. Without DIV16 the SPICLK frequency is:

$$SCKFrequency = \frac{AMBAclockfrequency}{(4 - (2 \cdot FACT)) \cdot (PM + 1)}$$

With DIV16 enabled the frequency of SPICLK is derived through:

$$SCKFrequency = \frac{AMBAclockfrequency}{16 \cdot (4 - (2 \cdot FACT)) \cdot (PM + 1)}$$

Note that the fields of the Mode register, which includes DIV16, FACT and PM, should not be changed when the core is enabled. If the FACT field is set to 0 the core's register interface is compatible with the register interface found in MPC83xx SoCs. If the FACT field is set to 1, the core can generate an SPICLK clock with higher frequency.

23.2.4 Operation (master-only)

Master operation will transmit a word when there is data available in the transmit queue. When the transmit queue is empty the core will drive SPICLK to its idle state.

The core does not implement any slave select outputs. General purpose I/O pins can be used as slave select signals, which will be controlled via the General Purpose I/O Port(s).

23.3 Registers

The core is programmed through registers mapped into APB address space.

*Table 172.*SPI controller registers

APB address	Register
0x80000400	Capability register
0x80000420	Mode register
0x80000424	Event register
0x80000428	Mask register
0x8000042C	Command register
0x80000430	Transmit register
0x80000434	Receive register

Table 173. SPI controller Capability register

31			24	23		20	19	18	17	16
	SSSZ				RESERVED		0	0	0	SSEN

15		8	7		0
	FDEPTH			REV	

31: 24	Slave Select register size (SSSZ) - If the core has been configured with slave select signals this field contains the number of available signals. This value in this field (1) is not applicable as this implementation does not have dedicated slave select signals.
23: 20	RESERVED - The value of this field should not be used
16	Slave Select Enable (SSEN) - If the core has a slave select register, and corresponding slave select lines, the value of this field is one. Otherwise the value of this field is zero. The value in this implementation is zero.
15: 8	FIFO depth (FDEPTH) - This field contains the depth of the core's internal FIFOs, which is 16 for this implementation. The number of words the core can store in each queue is FDEPTH+1, since the transmit and receive registers can contain one word each.
6: 0	Core revision (REV) - This manual applies to core revision 2.

Table 174. SPI controller Mode register

31	30	29	28	27	26	25	24	23		20	19		16
0	LOOP	CPOL	CPHA	DIV16	REV	MS	EN		LEN			PM	

15	14	13	12	11		7	6				0
RESERVED	FACT	R		CG				RESERVED			

Table 174. SPI controller Mode register

30	Loop mode (LOOP) - When this bit is set, and the core is enabled, the core's transmitter and receiver are interconnected and the core will operate in loopback mode. The core will still detect, and will be disabled, on Multiple-master errors.
29	Clock polarity (CPOL) - Determines the polarity (idle state) of the SPICLK clock.
28	Clock phase (CPHA) - When CPHA is '0' data will be read on the first transition of SPICLK. When CPHA is '1' data will be read on the second transition of SPICLK.
27	Divide by 16 (DIV16) - Divide system clock by 16, see description of PM field below and see section 23.2.3 on clock generation. This bit has no significance in slave mode.
26	Reverse data (REV) - When this bit is '0' data is transmitted LSB first, when this bit is '1' data is transmitted MSB first. This bit affects the layout of the transmit and receive registers.
25	Master/Slave (MS) - When this bit is set to '1' the core will act as a master, when this bit is set to '0' the core will operate in slave mode. This implementation only supports operation in master mode. Software must set this bit to '1'.
24	Enable core (EN) - When this bit is set to '1' the core is enabled. No fields in the mode register should be changed while the core is enabled. This can bit can be set to '0' by software, or by the core if a multiple-master error occurs.
23: 20	Word length (LEN) - The value of this field determines the length in bits of a transfer on the SPI bus. Values are interpreted as: 0b0000 - 32-bit word length 0b0001-0b0010 - Illegal values 0b0011-0b1111 - Word length is LEN+1, allows words of length 4-16 bits.
19: 16	Prescale modulus (PM) - This value is used in master mode to divide the system clock and generate the SPI SPICLK clock. The value in this field depends on the value of the FACT bit. If bit 13 (FACT) is '0':The system clock is divided by 4*(PM+1) if the DIV16 field is '0' and 16*4*(PM+1) if the DIV16 field is set to '1'. The highest SPICLK frequency is attained when PM is set to 0b0000 and DIV16 to '0', this configuration will give a SPICLK frequency that is (system clock)/4. With this setting the core is compatible with the SPI register interface found in MPC83xx SoCs. If bit 13 (FACT) is '1': The system clock is divided by 2*(PM+1) if the DIV16 field is '0' and 16*2*(PM+1) if the DIV16 field is set to '1'. The highest SPICLK frequency is attained when PM is set to 0b0000 and DIV16 to '0', this configuration will give a SPICLK frequency that is (system clock)/2.
15:14	RESERVED
13	PM factor (FACT) - If this bit is 1 the core's register interface is no longer compatible with the MPC83xx register interface. The value of this bit affects how the PM field is utilized to scale the SPI clock. See the description of the PM field.
12	RESERVED
11: 7	Clock gap (CG) - The value of this field is only significant in master mode. The core will insert CG SPICLK clock cycles between each consecutive word. This only applies when the transmit queue is kept non-empty. After the last word of the transmit queue has been sent the core will go into an idle state and will continue to transmit data as soon as a new word is written to the transmit register, regardless of the value in CG. A value of 0b00000 in this field enables back-to-back transfers.
6: 0	RESERVED (R) - Read as zero and should be written as zero to ensure forward compatibility.

Table 175. SPI controller Event register

31	30			15	14	13	12	11	10	9	8	7	0
TIP		R			LT	R	OV	UN	R	NE	NF	R	

31	Transfer in progress (TIP) - This bit is '1' when the core has a transfer in progress. Writes have no effect.
30: 15	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.
14	Last character (LT) - This bit is set when a transfer completes if the transmit queue is empty and the LST bit in the Command register has been written. This bit is cleared by writing '1', writes of '0' have no effect.
13	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.

Table 175. SPI controller Event register

12	Overrun (OV) - This bit gets set when the receive queue is full and the core receives new data. The core continues communicating over the SPI bus but discards the new data. This bit is cleared by writing '1', writes of '0' have no effect.
11	Underrun (UN) - This bit is only set when the core is operating in slave mode. The bit is set if the core's transmit queue is empty when a master initiates a transfer. When this happens the core will respond with a word where all bits are set to '1'. This bit is cleared by writing '1', writes of '0' have no effect.
10	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.
9	Not empty (NE) - This bit is set when the receive queue contains one or more elements. It is cleared automatically by the core, writes have no effect.
8	Not full (NF) - This bit is set when the transmit queue has room for one or more words. It is cleared automatically by the core when the queue is full, writes have no effect.
7: 0	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.

Table 176. SPI controller Mask register

31	30		15	14	13	12	11	10	9	8	7	0
TIPE	R			LTE	R	OVE	UNE	MMEE	NEE	NFE	R	

31	Transfer in progress enable (TIPE) - When this bit is set the core will generate an interrupt when the TIP bit in the Event register transitions from '0' to '1'.
30: 15	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.
14	Last character enable (LTE) - When this bit is set the core will generate an interrupt when the LT bit in the Event register transitions from '0' to '1'.
13	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.
12	Overrun enable (OVE) - When this bit is set the core will generate an interrupt when the OV bit in the Event register transitions from '0' to '1'.
11	Underrun enable (UNE) - When this bit is set the core will generate an interrupt when the UN bit in the Event register transitions from '0' to '1'.
10	Multiple-master error enable (MMEE) - When this bit is set the core will generate an interrupt when the MME bit in the Event register transitions from '0' to '1'. This event is not applicable for this implementation,
9	Not empty enable (NEE) - When this bit is set the core will generate an interrupt when the NE bit in the Event register transitions from '0' to '1'.
8	Not full enable (NFE) - When this bit is set the core will generate an interrupt when the NF bit in the Event register transitions from '0' to '1'.
7: 0	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.

Table 177. SPI controller Command register

31	23	22	21	0
R		LST	R	

31: 23	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.
22	Last (LST) - After this bit has been written to '1' the core will set the Event register bit LT when a character has been transmitted and the transmit queue is empty. This bit is automatically cleared when the Event register bit has been set and is always read as zero.
21: 0	RESERVED (R) - Read as zero and should be written to zero to ensure forward compatibility.

Table 178. SPI controller Transmit register

31	0
TDATA	

Table 178. SPI controller Transmit register

31: 0	Transmit data (TDATA) - Writing a word into this register places the word in the transmit queue. This register will only react to writes if the Not full (NF) bit in the Event register is set. The layout of this register depends on the value of the REV field in the Mode register:
	Rev = '0': The word to transmit should be written with its least significant bit at bit 0.
	Rev = '1': The word to transmit should be written with its most significant bit at bit 31.

Table 179. SPI controller Receive register

31	0
RDATA	

31: 0	Receive data (RDATA) - This register contains valid receive data when the Not empty (NE) bit of the Event register is set. The placement of the received word depends on the Mode register fields LEN and REV:
	For LEN = 0b0000 - The data is placed with its MSb in bit 31 and its LSb in bit 0.
	For other lengths and REV = '0' - The data is placed with its MSB in bit 15.
	For other lengths and REV = '1' - The data is placed with its LSB in bit 16.
	To illustrate this, a transfer of a word with eight bits (LEN = 7) that are all set to one will have the following placement:
	REV = '0' - 0x0000FF00
	REV = '1' - 0x00FF0000

23.4 Signal definitions

The signals are described in table 180.

Table 180. Signal definitions

Signal name	Type	Function	Active
SPICLK	Output	SPI clock	-
SPIMISO	Input	Master-Input-Slave-Output	-
SPIMOSI	Output	Master-Output-Slave-Input	-

24   SLINK Serial Bus Based Real-Time Network Master

24.1   Overview

The core provides a link between the AMBA bus and the SLINK Serial Bus Based Real-Time Network. The core is configured through registers mapped into AMBA APB address space and can also perform DMA operations via AMBA AHB. A DMA engine enables it to perform sequences of SLINK operations without CPU intervention. Single operations are performed via the register interface. The core generates the SLINK SCLK clock and SYNC signals by dividing the system clock.
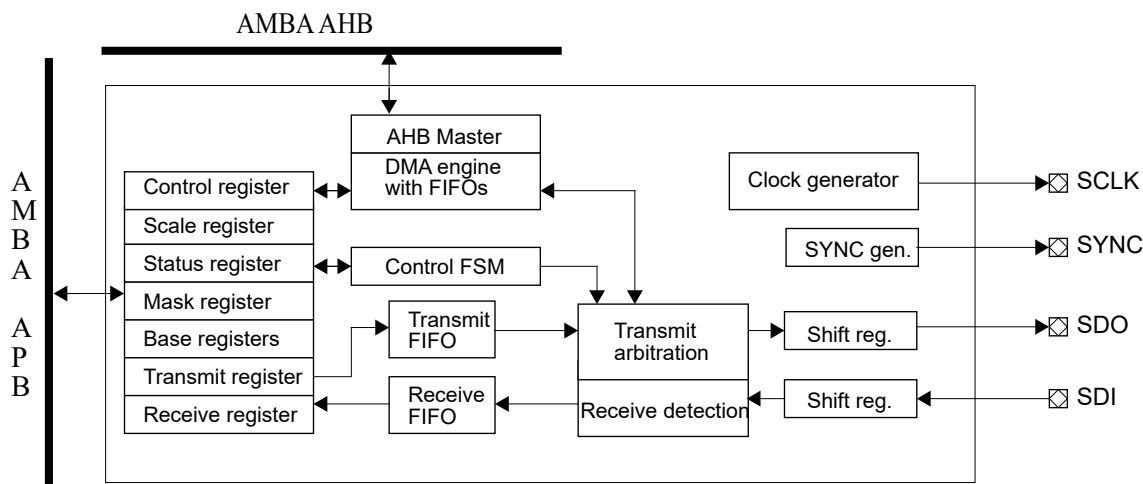

Figure 75. Block diagram

24.2   Operation

24.2.1   Transmission protocol

The SLINK bus is a full duplex synchronous bus that connects one master to seven slaves. The bus contains four differential lines, data to the master (SDI), data from the master (SDO), a periodic synchronization signal (SYNC) and a clock line (SCLK). The SCLK clock operates at 6 MHz continuously, one cycle of this clock constitutes a SLINK bit time. Words sent on the SLINK bus are 25 bits long and are always followed by a SYNC pulse that lasts one SLINK bit time, which results in a word time of 26 SLINK bit times. The data word formats are shown in figures 182 and 183.

When the bus is inactive the master sends NULL messages, see figure 184, and listens to the SDI lines for traffic from the slaves. The NULL word shown in figure 184 is the default word with odd parity, the contents of the NULL word is configurable via the NULL word register and the core's use of parity is configurable via the Control register.

The SLINK specification defines four types of data transfers:

*Table 181.*SLINK data transfer types

Transfer type	Description
MASTER-WORD-SEND	The master sends one READ or WRITE command. In case of a READ command a slave will respond in a later SYNC cycle.
INTERRUPT	A slave sends a word where the channel # field is set to 0.
SLAVE-WORD-SEND	An unsolicited word sent from a slave that is not an INTERRUPT nor a response to a READ command.
SEQUENCE	The master performs a sequence of operations described by an in-memory array which has a maximum size of 1024 elements.

Core operation for the first three types of transfers is described in section 24.2.4 and SEQUENCE operation is described in section 24.2.3. Conflicts on the bus are resolved by the slaves, the master performs no arbitration nor collision detection.

The core considers the data on the SDI line to be a SLINK word if bit 21 of the data word is '0'.

Table 182. Data word format for master to slaves

24	23		17	16		1	0	SYNC
RD/$\overline{\text{WR}}$		CHAN. NO			DATA / ADDRESS		PARITY	X

Table 183. Data word format for slave to master

24	23	22	21	20	17	16		1	0	SYNC
Hi-Z	SLAVE #		0	CHAN. NO			DATA		PARITY	1

Table 184. NULL word sent by master

24	23		17	16		1	0	SYNC
0		0000000			0000000000000000		1	X

24.2.2 Receive and transmit queues

The core uses receive and transmit queues to guarantee that it is able to transmit word k+1 directly after word k has been sent and that word m+1 can be received directly after word m. The core also has receive and transmit queues for SEQUENCE transfers, this section describes the queues used for MASTER-WORD-SEND, INTERRUPT and SLAVE-WORD-SEND transfers.

The receive and transmit queues are implemented as FIFOs with room for three SLINK data words each. The back of the transmit queue is the Transmit register and the front of the receive queue is the Receive register. Software can monitor the status of these queues via the Transmit Not Full (TNF) and Receive Not Empty (RNE) bits in the status register. The core can be configured to generate interrupts when these bits transition to one.

If the receive queue is full when a word is received the core will discard the received word and assert the Status register bit Receive Overflow (ROV). This is a critical error since the master system must guarantee that words can be received without delay.

24.2.3 SEQUENCE operations

During a SEQUENCE of operations the core performs READ/WRITE operations described by an in-memory array A. Responses to these operations are written back to an array named B. The core executes one operation and waits for the response before executing the operation described by the next element in A. To determine if a received word is a response to a SEQUENCE operation the core compares the received word's channel field to the value of the Sequence Channel Number (SCN) field in the Control register. If the values match, the received word is considered to be a response to the current operation and the core will execute the next pending SEQUENCE operation in the next SYNC cycle. If the core has pending MASTER-WORD-SEND transfers it will interleave these among the SEQUENCE operations. However, SEQUENCE operations have higher priority.

Software initiates a SEQUENCE transfer by allocating room for both arrays and filling array A with words that match the format of the Transmit register. The base address of this array is then written to the Array A Base Address register. The base address of array B is written to the Array B Base Address Register. Software enables processing of the SEQUENCE operations by setting the SEQUENCE Enable (SE) bit and initializing the SEQUENCE Length (SLEN) and SEQUENCE Channel Number (SCN) fields in the Control register.

Software should only change the base address registers when the SEQUENCE Active (SA) bit in the Status register is zero. The core starts processing the elements of array A at the address set in the Array Base Address register and writes responses to the B array. The response to an operation described by an element at A[i] is placed at position B[i]. The core maintains a queue of SEQUENCE operations which enables it to start operation A[i+1] in the SYNC cycle after the response to operation A[i] has been received.

When SLEN+1 elements in the A array have been processed and corresponding responses have been written to the B array, the core will clear the Control register's SEQUENCE Enable (SE) bit and set SEQUENCE Completed (SC) in the Status register. The SEQUENCE Completed bit can be used as an interrupt source by setting the corresponding bit in the Mask register.

A SEQUENCE may be prematurely terminated by an error event or by software intervention. If the core receives an ERROR response to a transfer on the AMBA AHB bus the AMBA ERROR (AERR) bit in the status register will be set and the SEQUENCE Completed (SC) bit will not be set.

If software aborts a SEQUENCE operation by setting the Abort SEQUENCE (AS) bit in the Control register the core will write back the currently completed operations and will then clear the Control register's SE and SA bits. If the core completed all SLEN operations the SEQUENCE Completed (SC) bit will be set, otherwise the Sequence Aborted (SA) bit will be set. If the Sequence Aborted bit is set, the SEQUENCE Index (SI) field in the Status register can be used to determine the number successful transfers that were written back to the B array. If a SEQUENCE is aborted while a slave is preparing a response, and the response arrives after the abort procedure has finished, the last response will be interpreted as an unsolicited request and placed in the receive queue. The core will clear the Abort SEQUENCE (AS) bit when the SEQUENCE has been aborted and all the completed operations have been written back to memory.

The core is also capable of receive-only SEQUENCE operation. When the SEQUENCE Receive Only (SRO) bit is set in the Control register the core will not access the elements in the A array and will, when the SEQUENCE Enable bit is set, store incoming data words with a channel matching the SEQUENCE channel number into the B array. All other behavior is identical to a normal SEQUENCE operation.

24.2.4 MASTER-WORD-SEND, SLAVE-WORD-SEND and INTERRUPT requests

Single READ/WRITE operations of data transfer type MASTER-WORD-SEND are performed by writing to the core's Transmit register. This register may only be written when the Status register bit Transmit Not Full (TNF) is set. Writes to the transmit register when the TNF bit is not set may overwrite a previously written value.

The core does not distinguish between the three transfer types MASTER-WORD-SEND, INTERRUPT and SLAVE-WORD-SEND. All received words, that do not belong to a ongoing SEQUENCE operation, are placed in the receive queue. Software can detect and identify the type of incoming words by monitoring the Receive Not Empty (RNE) and SLINK Received (SRX) bits in the status field. Both bits can be used as interrupt sources. It is recommended that software makes use of Receive Not Empty as an interrupt source and identifies the transfer type of the word by reading the Receive register and examining the appropriate fields. Even if SLINK Received is set and used as the interrupt source, software must never read the Receive register unless the Receive Not Empty bit is set.

Note that if the core receives a word that has a channel field value that matches the value of SEQUENCE Channel Number (SCN) and there is no ongoing SEQUENCE operation the core will regard this to be an unsolicited request and will place the word in the receive queue.

24.2.5 Parity and Parity Error

Parity is calculated on bits 1 to 24 of both received and sent SLINK words. The Control register bit Parity (PAR) determines if the core uses odd or even parity. If the core detects a parity error, the received word is discarded and the Parity Error (PERR) bit in the status register is set. If the Parity

Error Enable (PERRE) bit in the Mask register is set this event will cause an interrupt. The incoming word may have been a response to an ongoing SEQUENCE operation or to a MASTER-WORD-SEND READ transfer. Software should abort any ongoing SEQUENCE operation and possibly stop waiting for the response to a READ operation when the core reports a parity error.

24.2.6 Clock and SYNC generation

To avoid multiple clock domains in the design the core generates the SCLK clock by dividing the system clock. The clock divisor is selected by the fields CLKSCALEH and CLKSCALEL in the Clock Scaler register. The SCLK frequency is calculated with the formula:

$$SCLK = \frac{SystemClockFrequency}{((CLKSCALEH + 1) + (CLKSCALEL + 1))}$$

The CLKSCALEH determines the number of system clock cycles, plus one, that the SCLK clock is kept HIGH and the CLKSCALEL value determines how many clock cycles, plus one, that the SCLK clock is kept low. This scaling constrains the system clock frequency. To generate a 6 MHz, SCLK the system clock frequency must be a multiple of 6 MHz. To attain a 6 MHz SCLK with 50/50 duty cycle, the multiple must be dividable by two. The lowest possible system frequency that can be used to generate a 6 MHz clock is 12 MHz, with both scaler register fields set to zero:

$$SCLK = \frac{SystemClockFrequency}{((CLKSCALEH + 1) + (CLKSCALEL + 1))} = \frac{12MHz}{((0 + 1) + (0 + 1))} = \frac{12MHz}{2} = 6MHz$$

The next frequency that will generate a SCLK clock with a 6 MHz period is a system clock of 18 MHz. This system frequency does not allow a SCLK clock with a 50/50 duty cycle.

24.3 Registers

The core is programmed through registers at address. 0x80000800

*Table 185.*GRSLINK registers

APB address offset	Register
0x08000080	Clock Scaler register
0x80000804	Control register
0x80000808	NULL word register
0x8000080C	Status register
0x80000810	IRQ mask register
0x80000814	Array A base address register
0x80000818	Array B base address register
0x8000081C	Transmit register
0x80000820	Receive register

Table 186. GRSLINK Clock Scaler register

31	21	20	16	15	5	4	0
RESERVED		CLKSCALEH		RESERVED		CLKSCALEL	

31:21   RESERVED

Table 186. GRSLINK Clock Scaler register

Bits	Description
20:16	Clock Scale HIGH (CLKSCALEH) - This value determines how many system clock cycles the SCLK clock is high. The SCLK clock's high time is CLKSCALEH+1 system clock cycles. Please see section 24.2.6 for a description of clock scaling.
15:5	RESERVED
4:0	Clock scale LOW (CLKSCALEL) - This value determines how many system clock cycles the SCLK clock is low. The SCLK clock's low time is CLKSCALEL+1 system clock cycles. Please see section 24.2.6 for a description of clock scaling.

Reset value: 0x00000000

Table 187. GRSLINK Control register

31	30	29		26	25								16
ODEL		RESERVED			SLEN								

15			9	8	7		4	3	2	1	0
RESERVED				SRO	SCN			PAR	AS	SE	SLE

Bits	Description
31:30	Output delay (ODEL) - This field determines the delay applied to transitions on the SDO and SYNC lines relative to the SLINK clock SCLK. The SDO and SYNC signals will transition ODEL+1 system clock cycles after the rising edge on SCLK. This means that the delay can be adjusted between 1 and 4 system clock cycles.
29:26	RESERVED
25:16	SEQUENCE Length (SLEN) - Number of elements in a SEQUENCE. The value is encoded as number of elements - 1; when this field is set to 0 the core will perform one operation from the A array. Or, if the SRO bit is set, only transfer one element to the B array.
15:9	RESERVED
8	SEQUENCE Receive Only (SRO) - When this bit is set the core will not transmit any words as part of a SEQUENCE operation. The core will only access the B array.
7:4	SEQUENCE Channel Number (SCN) - When the SLINK core is performing a SEQUENCE of operations it compares the channel number in the incoming packets with the value of this field to determine if the received word is a response to a SEQUENCE operation.
3	Parity (PAR) - This bit determines the method used to calculate the parity bit in the SLINK data words. When this bit is set to '1' the core uses odd parity, when the bit is set to '0' the core uses even parity. The core reads the value of this bit during each SYNC cycle, however the bit's value should only be changed when the core is disabled. Default value: '1' - odd parity.
2	Abort SEQUENCE (AS) - Setting this bit aborts an ongoing SEQUENCE. This bit is automatically cleared when the SEQUENCE has stopped.
1	SEQUENCE Enable (SE) - When this bit is set to '1' the core will use the Array base addresses to perform a SEQUENCE of SLEN operations. This bit is automatically cleared by the core after the SEQUENCE has completed and can not be reset by writes to this register.
0	SLINK Enable (SLE) - When this bit is set the core will accept transmit requests and react to incoming data on the bus. When this bit is '0' the controller will not generate the SCLK clock and SYNC pulses and will not perform, or react to, any communication on the bus. When the core is enabled it always outputs one NULL word before issuing the first SYNC pulse. When the core is disabled by writing '0' to this bit, the SCLK clock stops within one SCLK period.

Reset value: 0x00000008

Table 188. GRSLINK NULL word register

31		24	23		0
RESERVED			NULLWORD		

Bits	Description
31:24	RESERVED

Table 188. GRSLINK NULL word register

23:0	SLINK NULL Word (NULLWORD) - This register contains the value of the NULL word that the core transmits during inactivity. Bits 23:0 in this register constitutes bits 24:1 in the SLINK data word. The parity bit is calculated by the core and appended at the end of the word.

Reset value: 0x00000000

Table 189. GRSLINK Status register

31		26	25								16
RESERVED			SI								

15				8	7	6	5	4	3	2	1	0
RESERVED					PERR	AERR	ROV	RNE	TNF	SC	SA	SRX

31:26	RESERVED
25:16	SEQUENCE Index (SI) - This fields contains the index of the next element to write to the B array during SEQUENCE operation. With the help of this field software can monitor the progress SEQUENCE operations. Software can also determine that SI elements were written back if the SEQUENCE was aborted. Note that the field holds a maximum value of 1023. If the SEQUENCE is 1024 operations long this counter will wrap on the last element. In this case software must check if the SC bit is set to determine if the core transferred 0 or 1024 elements. This field is read only.
15: 8	RESERVED
7	Parity Error (PERR) - This bit is set when a parity error is detected. The bit is cleared by writing '1' to this position.
6	AMBA ERROR (AERR) - This bit is set to '1' when the core receives an ERROR response to a transfer on the AMBA AHB bus. This event aborts any ongoing SEQUENCE and clears bit 1 in the Control register and bit 1 in the Status register. This bit is considered to get set even if it already has the value '1' and an AMBA ERROR event will always generate a interrupt if ARRE (bit 6) in the mask register is set.
5	Receive Overflow (ROV) - This bit is set to '1' when the core receives a word and the receive queue is full. The incoming word is discarded. This is a critical error since a master system must guarantee that words can be received.
4	Receive Not Empty (RNE) - This bit is '1' when the receive queue contains a word. This read only bit is automatically cleared when the receive queue is empty.
3	Transmit Not Full (TNF) - This bit is '1' when the core has one or more free slots in the transmit queue. When this bit is set software may write a new word into the Transmit register. This bit is read only.
2	SEQUENCE Completed (SC) - This bit is set to '1' when the core has performed the READ/WRITE operation described at the end of the SEQUENCE array. This bit is cleared by writing '1' to this position.
1	SEQUENCE Aborted (SA) - This bit is set to '1' if SEQUENCE operation was aborted by software setting the Abort SEQUENCE (AS) bit in the control register. If this bit is set the core did not write back results for all SLEN operations. This bit is cleared by writing '1' to this position.
0	SLINK Word Received (SRX) - This bit is set when the master has accepted a single word from an IO card and is not set if the core receives a word with parity error or if the reception of a word leads to a receive overflow. This bit is cleared by writing '1' to this position.

Reset value: 0x00000008

Table 190. GRSLINK Interrupt Mask register

31		8	7	6	5	4	3	2	1	0
RESERVED			SERRE	AERRE	ROVE	RNEE	TNFE	SCE	SAE	SRXE

31:7	RESERVED
7	Parity Error Enable (PERRE) - When this bit is set to '1' the core will generate an interrupt when a parity error is detected.

Table 190. GRSLINK Interrupt Mask register

6	AMBA ERROR Enable (AERRE) - When this bit is set to '1' the core will generate an interrupt when an AMBA ERROR event occurs.
5	Receive Overflow Enable (ROVE) - When this bit is set to '1' the core will generate an interrupt when bit 5 of the Status register is set.
4	Receive Not Empty Enable (RNEE) - When this bit is set to '1' the core will generate an interrupt when bit 4 of the Status register is set.
3	Transmit Not Full Enable (TNFE) - When this bit is set to '1' the core will generate an interrupt when bit 3 of the Status register is set.
2	SEQUENCE Completed Enable (SCE) - When this bit is set to '1' the core will generate an interrupt when bit 2 of the Status register is set.
1	SEQUENCE Aborted Enable (SCE) - When this bit is set to '1' the core will generate an interrupt when bit 1 of the Status register is set.
0	SLINK Word Received Enable (SRXE) - When this bit is set to '1' the core will generate an interrupt when bit 0 of the Status register is set.

Reset value: 0x000000UU, where U is undefined

Table 191. GRSLINK Array A Base Address register

31	2	1	0
ABASE		RESERVED	

31:2	Array A Base Address (ABASE) - This field contains bits 31:2 of the address of the fast sequence array's first element. This field may only be written when the SEQUENCE Enable bit in the Control register is '0'.
1:0	RESERVED - This field must be written with only zeroes and is always read as zero.

Reset value: Undefined

Table 192. GRSLINK Array B Base Address register

31	2	1	0
BBASE		RESERVED	

31:2	Array B Base Address (BBASE) - This field contains bits 31:2 of the address of the fast sequence array B's first element. This field may only be written when the SEQUENCE Enable in the Control register is '0'.
1:0	RESERVED - This field must be written with only zeroes and is always read as zero.

Reset value: Undefined

Table 193. GRSLINK Transmit register

31	24	23	22	16	15	0
RESERVED		R$\overline{W}$	CHAN. NO		Data / Address	

31:24	RESERVED - This field is always read as zero, writes have no effect.
23:0	Fields in SLINK data word from master to slave. Reads of this field always return zero. This register must not be written unless the Transmit Not Full (TNF) bit in the Status register is set.

Reset value: 0x00000000

Table 194. GRSLINK Receive register

31	23	22	21	20	19	16	15	0
RESERVED		IO CARD #	SPARE		CHAN. NO		Data	

31:23	RESERVED - This field is always read as zero, writes have no effect.
22:0	Fields in SLINK data word from slave to master. This register must not be read unless the Receive Not Empty (RNE) bit in the Status register is set.

Reset value: Undefined

24.4 Signal definitions

The signals are described in table 195.

Table 195. Signal definitions

Signal name	Type	Function	Active
SLI	Input	Data into the master	-
SLO	Output	Data out of the master	-
SLSYNC	Output	SLINK SYNC signal	Logical 1
SLCLK	Output	SLINK clock	Logical 1

25 ASCS Controller

25.1 Overview

The ASCS core provides a serial link that performs data reads (telemetry) and data writes (telecommand). The core supports 2 slaves, data word lengths of 16 bits, and different frequencies for the synchronization link (in the 1 - 100 Hz range, for all supported system frequencies).

25.2 Operation

25.2.1 Serial link

The task of the serial link is to perform data writes (telecommand, TC) and data reads (telemetry, TM). The interface of the serial link consists of five signals: hs, dcs, mcs, mas, das. The hs signal is the serial clock (500 kHz) signal used to control the TC or TM; dcs is the data output from the host to the slaves; mcs indicates that a TC is about to start or finish; mas indicates that a TM is about to start or finish; das is the data input from the slaves to the host. The core supports both a configuration where each slave has its own data input and a configuration where only one data input is used, and where the slaves are assumed to tri-state their output. For a more detailed description of how a single TC/TM is performed and the timing requirements that apply please refer to the ASCS specification.

The activity on the serial link is controlled through the core's command and control register, status register, TC data register, and TM data register (see 25.3). The following also needs to be taken into consideration:

- For transactions to be timed correctly the *us1* field of the command register must be setup.
- A TC and TM will never be performed at the same time.
- The serial link must be running for a TC or TM to be carried out. If the serial link is stopped and software tells the core to start a transfer, the transfer will be delayed until the serial link is started.
- The GRASCS has no FIFO where it buffers transfer requests or data. However since separate registers are used for TC and TM the core can handle that software starts one TC and one TM at the same time. In such a case the TC will always be performed before the TM.
- Once a transfer has been started by software the only way to abort it is to reset the core.
- There might be a delay between the time software starts a transfer and the time the core performs it. This delay could be because the core is still waiting for the minimum delay time between transfers to pass. A TM can also be delay up to 150 microseconds if a synchronization pulse is being generated.

25.2.2 Synchronization link

The task of the synchronization link is to issue synchronization pulses to the slaves. The interface of the synchronization link consists of one signal called etr. When the link is running the core generates pulses on the etr signal. The frequency of the etr pulses can be configured through the core's ETR scale register. The synchronization link is controlled through the core's command and control register and status register. The following also needs to be taken into consideration:

- For synchronization pulses to be generated correctly the ETR scale register must be set properly.
- When software starts the synchronization interface the core might delay the actual start of the interface. The reason for this could be that a TM is in progress. It could also be because the internal ETR timer, which is always running, is in the middle of a pulse. The core then delays the start to be sure not to generate an external pulse that is too short.
- The first pulse on the *etr* signal might be delayed with up to one period from the time that the core starts the synchronization interface, depending on the source used to generate the pulse.

25.3 Registers

The core is programmed through registers at address 0x80000700.

*Table 196.*GRASCS registers

APB address offset	Register
0x80000700	Command and control register
0x80000704	ETR scale register
0x80000708	Status and capability register
0x8000070C	TC data and control register
0x80000710	TM data register

Table 197. Command and control register

31		25	24	23		16	15	14	13	12	11		8	7	6		4	3	2	1	0
OEN		R	us1c	us1			R		tmd	tcd	slave_sel			R	etr_ctrl			stm	estst	stst	res

31	Output enable. Reset value 0. (Read/Write)
30:25	Reserved, always zero
24	This bit can't be written and it always reads zero. But if the us1 field should be updated then this bit needs to be set to a '1' in the same write cycle.
23:16	Number of system clock cycles in a micro second. In order to preserve correct timing of transactions the core prevents these bits from being written if the *run* bit in the status register is set to '1'. (Read/Write)
15:14	Reserved, always zero
13	Interrupt control mask bit for TMs. When set to '1' an interrupt is generated when a TM is done, when set to '0' no interrupt is generated. (Read/Write)
12	Interrupt control mask bit for TCs. When set to '1' an interrupt is generated when a TC is done, when set to '0' no interrupt is generated. (Read/Write)
11:8	Slave select bits. Decides which slave input the core listens to when performing a TM. In order to prevent faulty data to be stored in the TM register the core prevents these bits from being written if the *tma* bit in the status register is '1'. Only the number of bits required to represent the number of slaves are used, the rest are always zero. (Read/Write)
7	Reserved, always zero
6:4	ETR source control. Decides which source that should be used for the ETR synchronization pulse. A value of 0 means that an internal counter is used while a value of 1 - 6 means that external time marker 1 - 6 will be used. In order to prevent an invalid ETR period or pulse the core prevents these bits from being written when the *erun* bit in the status register is '1'. (Read/Write)
3	Send TM bit. When this bit is set to '1' the core will perform a TM the next chance it gets and at the same time clear this bit. (Read/Write)
2	Synchronization interface start/stop. This bit is set/cleared immediately but the actual start/stop of the interface might be delayed. The real status of the synchronization interface is reported in the status register. (Read/Write)
1	Serial interface start/stop. This bit is set/cleared immediately but the actual start/stop of the interface might be delayed. The real status of the serial interface is reported in the status register. (Read/Write)
0	Reset. Writing a '1' to this bit is equivalent with a hardware reset with the following exceptions: In order to be sure that the core never generates a invalid ETR pulse or transactions with too short delays in between them, the ETR scale register and internal timers are not reset, and the *etr_running* bit in the status register might not be cleared immediately. This bit is cleared by the core the cycle after is set. (Read/Write)

Table 198. ETR scale register

31		28	27		0
	R				

31:28	Reserved, always zero
27:0	This field should be set to the number of system clock cycles in one ETR period. (Read/Write)

GR712RC

Table 199. Status and capability register

31			20	19	18	17	16		13	12		8	7	6	5	4	3	2	1	0
	R			tmc	usc	sdc	nslaves			dbits			R		tmd	tcd	tma	tca	erun	run

31:20	Reserved, always zero.
19	Capability bit '1': five external time markers enabled (Read only)
18	Capability bit '0': using internal us counter (Read only)
17	Capability bit '1': separate input signals (Read only)
16:13	Capability bits "01": number of slaves minus one, 2 slaves. (Read only)
12:8	Capability bits "1111": number of bits in data word minus one, 16 bits. (Read only)
7:6	Reserved, always zero.
5	TM done bit. Core sets this bit to '1' when a TM is done. Software can clear this bit by writing '1' to it. (Read/Write clear)
4	TC done bit. Core sets this bit to '1' when a TC is done. Software can clear this bit by writing '1' to it. (Read/Write clear)
3	TM active. Set to '1' if a TM is in progress. (Read only)
2	TC active. Set to '1' if a TC is in progress. (Read only)
1	ETR running bit. This bit is set to '1'when the synchronization interface is running. (Read only)
0	Serial running bit. This bit is set to '1' when the serial interface is running. (Read only)

Table 200. TC data and control register

31	16	15	0
R		tcd	

31:16	Reserved, always zero.
15: 0	TC data. When these bits are written the core will perform a TC at the next chance it gets and send the data written. In order to not send wrong TC data the core prevents these bits from being written if the *tca* bit in the status register is set to '1'. (Read/Write)

Table 201. TM data register

31		16	15		0
	R			tmd	

31:16   Reserved, always zero.

15:0   TM data. Received TM data is stored here. (Read only)

25.4 Signal definitions

The signals are described in table 202.

Table 202. Signal definitions

Signal name	Type	Function	Active
A16DASA, A16DASB	Input	ASCS slave data in	-
A16MCS	Output	TC start/stop signal	Logical 1
A16HS	Output	Serial clock used for TC/TM	Logical 1
A16DCS	Output	ASCS slave data out	-
A16MAS	Output	TM start/stop signal	Logical 1
A61ETR	Output	Synchronization signal	Logical 1

26 GRTC - Telecommand Decoder

26.1 Overview

The Telecommand Decoder (GRTC) is compliant with the Packet Telecommand protocol and specification defined by [ECSS-E-ST-50-04C]. The decoder is compatible with the [PSS-04-107] and [PSS-04-151] standards. The decoder is compatible with the CCSDS recommendations [CCSDS-231.0-B-2], [CCSDS-232.0-B-2] and [CCSDS-232.1-B-2]. The Telecommand Decoder (GRTC) only implements the Coding Layer (CL).

In the Coding Layer (CL), the telecommand decoder receives bit streams on multiple channel inputs. The streams are assumed to have been generated in accordance with the Physical Layer specifications. In the Coding Layer, the decoder searches all input streams simultaneously until a start sequence is detected. Only one of the channel inputs is selected for further reception. The selected stream is bit-error corrected and the resulting corrected information is passed to the user. The corrected information received in the CL is transfer by means of Direct Memory Access (DMA) to the on-board processor.

The Command Link Control Word (CLCW) and the Frame Analysis Report (FAR) can be read and written as registers via the AMBA AHB bus. Parts of the two registers are generated by the Coding Layer (CL). Note that most parts of the CLCW and FAR are not produced by the Telecommand Decoder (GRTC) hardware portion. This is instead to be done in software. The CLCW register contents need to be transferred from the Telecommand Decoder (GRTC) to the Telemetry Encoder (GRTM) by means of software. The FAR register contents need also to be transferred from the Telecommand Decoder to the Telemetry Encoder by means of software via the creation of a (Telemetry) Space Packet. There is thus no automatic hardware connection between the Telecommand Decoder and the Telemetry Encoder.
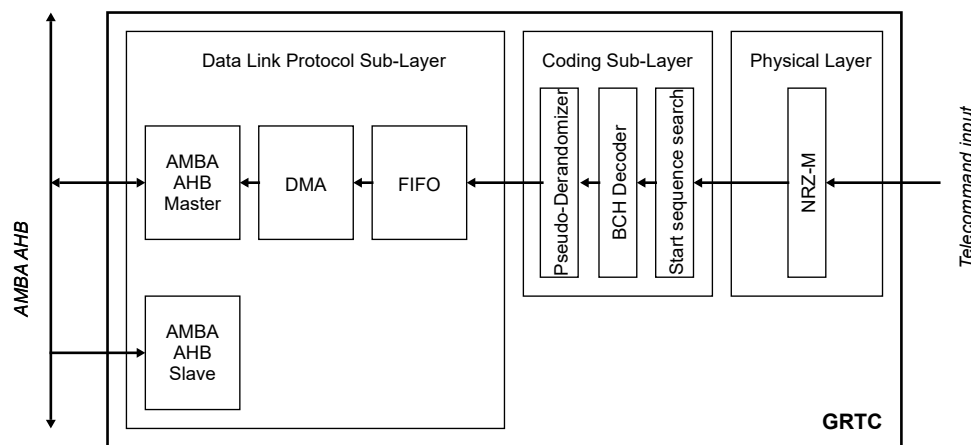

Figure 76. Block diagram

26.1.1 Concept

A telecommand decoder in this concept is mainly implemented by software in the on-board processor. The supporting hardware in the Telecommand Decoder (GRTC) implements the Coding Layer, which includes synchronisation pattern detection, channel selection, codeblock decoding, Direct Memory Access (DMA) capability and buffering of corrected codeblocks.

The GRTC has been split into several clock domains to facilitate higher bit rates and partitioning. The two resulting sub-cores have been named Telecommand Channel Layer (TCC) and the Telecommand Interface (TCI). Note that TCI is called AHB2TCI. A complete ECSS/CCSDS packet telecommand decoder can be realized at software level according to the latest available standards, staring from the Transfer Layer.
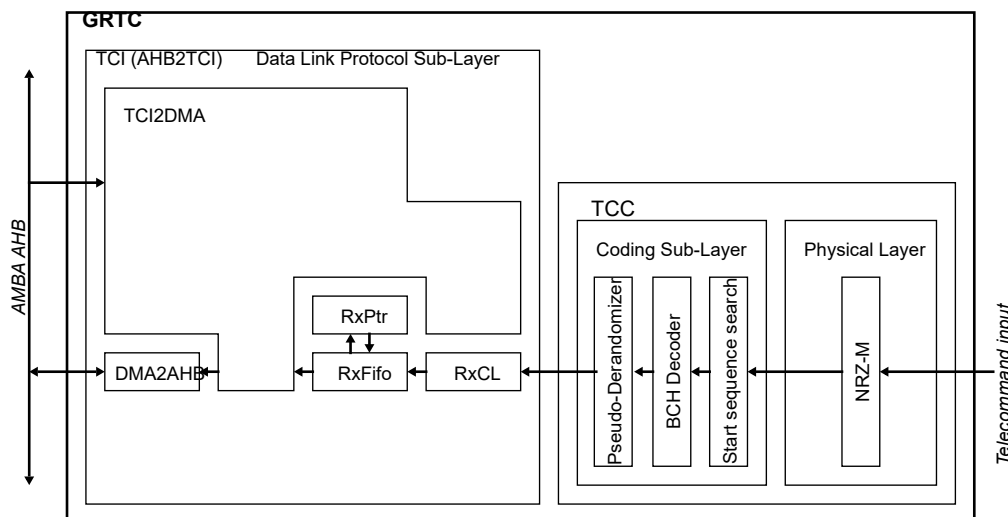
Figure 77. Detailed block diagram showing the internal structure

26.1.2   Functions and options

The Telecommand Decoder (GRTC) only implements the Coding Layer of the telecommand protocol standard [ECSS-E-ST-50-04C]. Other layers are to be implemented in software, e.g. Authentication Unit (AU). The Command Pulse Distribution Unit (CPDU) is not implemented.

The following functions of the GRTC are programmable by means of registers:

- Pseudo-De-Randomisation
- Non-Return-to-Zero – Mark decoding

The following functions of the GRTC are fixed:

- Polarity of RF Available and Bit Lock inputs is active high
- Edge selection for input channel clock is rising edge

26.2   Data formats

26.2.1   Reference documents

| [PSS-04-107]	Packet Telecommand Standard, Issue 2
[PSS-04-151]	Telecommand Decoder Standard, Issue 1
[CCSDS 231.0-B-2]	TC Synchronization and Channel Coding
[CCSDS 232.0-B-2]	TC Space Data Link Protocol
[CCSDS 232.1-B-2]	Communications Operation Procedure-1
[ECSS-E-ST-50-04C]	Space engineering - Space data links - Telecommand protocols, synchronization and channel coding

26.2.2 Waveforms

Figure 78. Telecommand input protocol

26.3 Coding Layer (CL)

The Coding Layer synchronises the incoming bit stream and provides an error correction capability for the Command Link Transmission Unit (CLTU). The Coding Layer receives a dirty bit stream together with control information on whether the physical channel is active or inactive for the multiple input channels.

The bit stream is assumed to be NRZ-L encoded, as the standards specify for the Physical Layer. As an option, it can also be NRZ-M encoded. There are no assumptions made regarding the periodicity or continuity of the input clock signal while an input channel is inactive. The most significant bit (Bit 0 according to [ECSS-E-ST-50-04C]) is received first.

Searching for the Start Sequence, the Coding Layer finds the beginning of a CLTU and decodes the subsequent codeblocks. As long as no errors are detected, or errors are detected and corrected, the Coding Layer passes clean blocks of data to the Transfer Layer which is to be implemented in software. When a codeblock with an uncorrectable error is encountered, it is considered as the Tail Sequence, its contents are discarded and the Coding Layer returns to the Start Sequence search mode.

The Coding Layer also provides status information for the FAR, and it is possible to enable an optional de-randomiser according to [ECSS-E-ST-50-04C].

26.3.1 Synchronisation and selection of input channel

Synchronisation is performed by means of bit-by-bit search for a Start Sequence on the channel inputs. The detection of the Start Sequence is tolerant to a single bit error anywhere in the Start Sequence pattern. The Coding Layer searches both for the specified pattern as well as the inverted pattern. When an inverted Start Sequence pattern is detected, the subsequent bit-stream is inverted till the detection of the Tail Sequence.

The detection is accomplished by a simultaneous search on all active channels. The first input channel where the Start Sequence is found is selected for the CLTU decoding. The selection mechanism is restarted on any of the following events:

- The input channel active signal is de-asserted, or
- a Tail Sequence is detected, or
- a Codeblock rejection is detected, or
- an abandoned CLTU is detected, or the clock time-out expires.

As a protection mechanism in case of input failure, a clock time-out is provided for all selection modes. The clock time-out expires when no edge on the bit clock input of the selected input channel in decode mode has been detected for a specified period.

26.3.2 Codeblock decoding

The received Codeblocks are decoded using the standard (63,56) modified BCH code. Any single bit error in a received Codeblock is corrected. A Codeblock is rejected as a Tail Sequence if more than one bit error is detected. Information regarding Count of Single Error Corrections and Count of Accept Codeblocks is provided to the FAR. Information regarding Selected Channel Input is provided via a register.

26.3.3 De-Randomiser

In order to maintain bit synchronisation with the received telecommand signal, the incoming signal must have a minimum bit transition density. If a sufficient bit transition density is not ensured for the channel by other methods, the randomiser is required. Its use is optional otherwise. The presence or absence of randomisation is fixed for a physical channel and is managed (i.e., its presence or absence is not signalled but must be known a priori by the spacecraft and ground system). A random sequence is exclusively OR-ed with the input data to increase the frequency of bit transitions. On the receiving end, the same random sequence is exclusively OR-ed with the decoded data, restoring the original data form. At the receiving end, the de-randomisation is applied to the successfully decoded data. The de-randomiser remains in the "all-ones" state until the Start Sequence has been detected. The pattern is exclusively OR-ed, bit by bit, to the successfully decoded data (after the Error Control Bits have been removed). The de-randomiser is reset to the "all-ones" state following a failure of the decoder to successfully decode a codeblock or other loss of input channel.

26.3.4 Non-Return-to-Zero – Mark

An optional Non-Return-to-Zero – Mark decoder can be enabled by means of a register.

26.3.5 Design specifics

The coding layer is supporting channel inputs ([PSS-04-151] requires at least 4).

A codeblock is fixed to 56 information bits (as per CCSDS/ECSS).

The CCSDS/ECSS (1024 octets) or [PSS-04-151] (256 octets) standard maximum frame lengths are supported, being programmable via bit PSS in the GCR register. The former allows more than 37 codeblocks to be received.

The Frame Analysis Report (FAR) interface supports 8 bit CAC field, as well as the 6 bit CAC field specified in [PSS-04-151]. When the PSS bit is cleared to '0', the two most significant bits of the CAC will spill over into the "LEGAL/ILLEGAL" FRAME QUALIFIER field in the FAR. These bits will however be all-zero when [PSS-04-151] compatible frame lengths are received or the PSS bit is set to '1'. The saturation is done at 6 bits when PSS bit is set to '1' and at 8 bits when PSS bit is cleared to '0'.

The Pseudo-Randomiser decoder is included (as per CCSDS/ECSS), its usage being programmable.

The Physical Layer input can be NRZ-L or NRZ-M modulated, allowing for polarity ambiguity. NRZ-L/M selection is programmable. This is an extension to ECSS: Non-Return to Zero - Mark decoder added, with its internal state reset to zero when channel is deactivated.

Note: If input clock disappears, it will also affect the codeblock acquired immediately before the codeblock just being decoded (accepted by [PSS-04-151]).

In state S1, all active inputs are searched for start sequence, there is no priority search, only round robin search. The search for the start sequence is sequential over all inputs: maximum input frequency = system frequency / 7.

The [PSS-04-151] specified CASE-1 and CASE-2 actions are implemented according to aforementioned specification, not leading to aborted frames.

Extended E2 handling is implemented:

- E2b   Channel Deactivation - selected input becomes inactive in S3
- E2c   Channel Deactivation - too many codeblocks received in S3
- E2d   Channel Deactivation - selected input is timed-out in S3
  (design choice being: S3 => S1, abandoned frame)

26.3.6 Direct Memory Access (DMA)

This interface provides Direct Memory Access (DMA) capability between the AMBA bus and the Coding Layer. The DMA operation is programmed via an AHB slave interface.

The DMA interface is an element in a communication concept that contains several levels of buffering. The first level is performed in the Coding Layer where a complete codeblock is received and kept until it can be corrected and sent to the next level of the decoding chain. This is done by inserting each correct information octet of the codeblock in an on-chip local First-In-First-Out (FIFO) memory which is used for providing improved burst capabilities. The data is then transferred from the FIFO to a system level ring buffer in the user memory which is accessed by means of DMA.

The following storage elements can thus be found in this design:

The shift and hold registers in the Coding Layer

The local FIFO (parallel; 32-bit; 4 words deep)

The system ring buffer (external memory; 32-bit; 1 to 256 KiB deep).

26.4 Transmission

The transmission of data from the Coding Layer to the system buffer is described hereafter.

The serial data is received and shifted in a shift register in the Coding Layer when the reception is enabled. After correction, the information content of the shift register is put into a hold register.

When space is available in the peripheral FIFO, the content of the hold register is transferred to the FIFO. The FIFO is of 32-bit width and the byte must thus be placed on the next free byte location in the word.

When the FIFO is filled for 50%, a request is done to transfer the available data towards the system level buffer.

If the system level ring buffer isn't full, the data is transported from the FIFO, via the AHB master interface towards the main processor and stored in external memory. If no place is available in the system level ring buffer, the data is held in the FIFO.

When the GRTC keeps receiving data, the FIFO will fill up and when it reaches 100% of data, and the hold and shift registers are full, a receiver overrun interrupt will be generated (IRQ_RX_OVERRUN). All new incoming data is rejected until space is available in the peripheral FIFO.

When the receiving data stream is stopped (e.g. when a complete data block is received), and some bytes are still in the peripheral FIFO, then these bytes will be transmitted to the system level ring buffer automatically. Received bytes in the shift and hold register are always directly transferred to the peripheral FIFO.

The FIFO is automatically emptied when a CLTU is either ready or has been abandoned. The reason for the latter can be codeblock error, time out etc. as described in CLTU decoding state diagram.

The operational state machine is shown in figure 79.
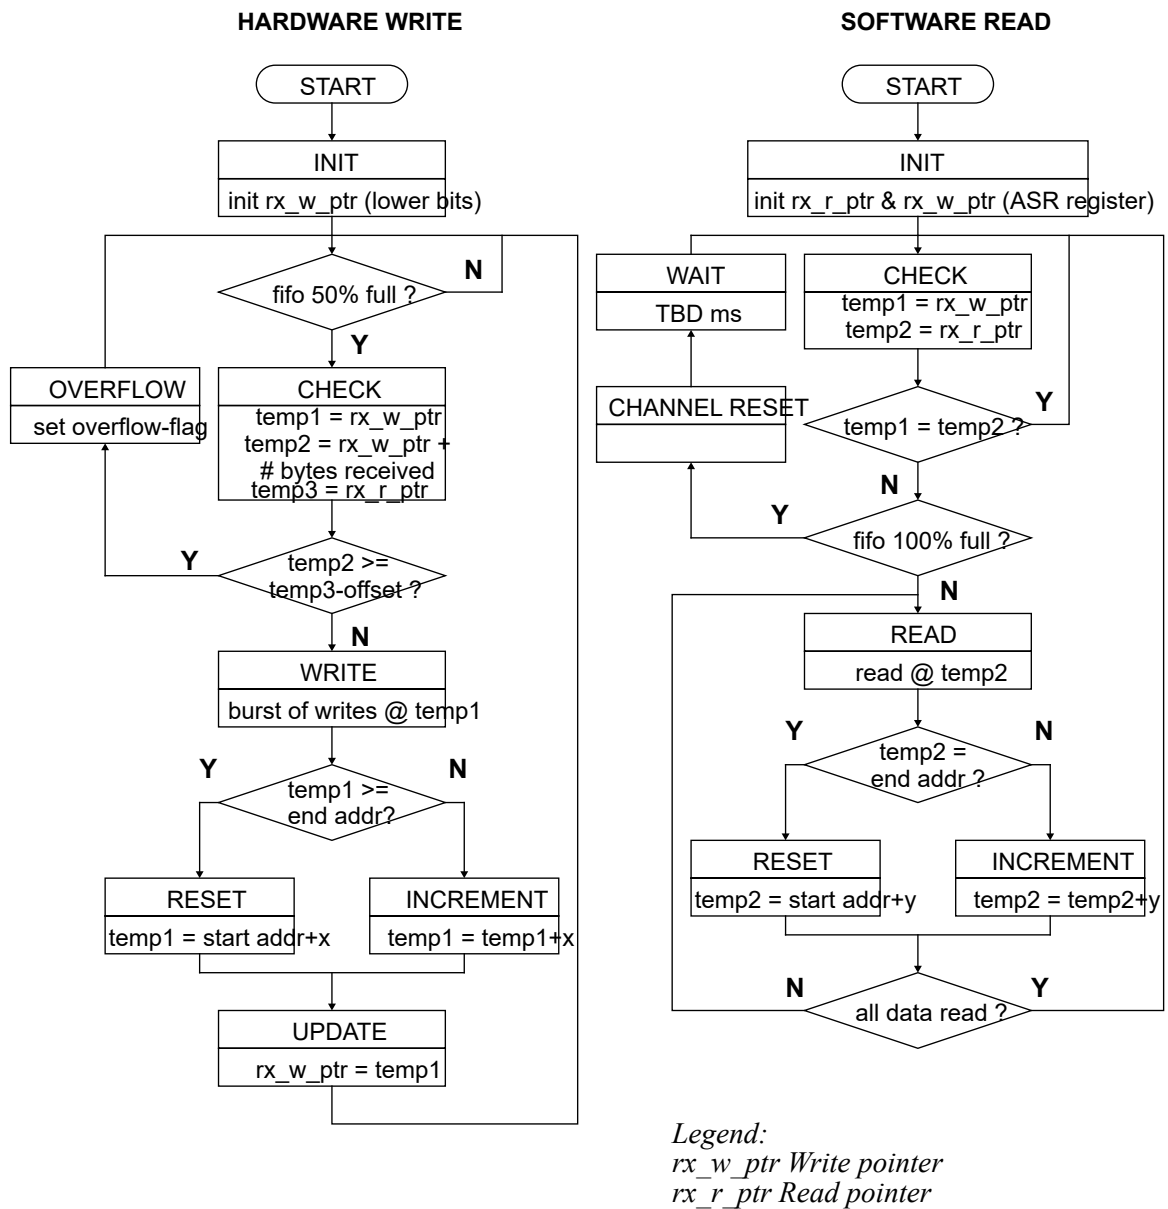
HARDWARE WRITE

SOFTWARE READD

Figure 79. Direct Memory Access

26.4.1 Data formatting

When in the decode state, each candidate codeblock is decoded in single error correction mode as described hereafter.

26.4.2 CLTU Decoder State Diagram

Note that the diagram has been improved with explicit handling of different E2 events listed below.

Figure 80. Decoder state diagram

State   Definition:

    S1   Inactive

    S2   Search

    S3   Decode

Event   Definition:

    E1   Channel Activation

    E2a   Channel Deactivation - all inputs are inactive

    E2b   Channel Deactivation - selected becomes inactive (CB=0 -> frame abandoned)

    E2c   Channel Deactivation - too many codeblocks received (all -> frame abandoned)

    E2d   Channel Deactivation - selected is timed-out (all -> frame abandoned)

    E3   Start Sequence Found

    E4   Codeblock Rejection (CB=0 -> frame abandoned)

26.4.3 Nominal

A: When the first "Candidate Codeblock" (i.e. "Candidate Codeblock" 0, which follows Event 3 (E3):START SEQUENCE FOUND) is found to be error free, or if it contained an error which has been corrected, its information octets are transferred to the remote ring buffer in table 203. At the same time, a "Start of Candidate Frame" flag is written to bit 0 or 16, indicating the beginning of a transfer of a block of octets that make up a "Candidate Frame". There are two cases that are handled differently as described in the next sections.

Table 203. Data format

	Bit[31…………24]	Bit[23…………16]	Bit[15….………8]	Bit[7……...……0]
0x40000000	information octet0	0x01	information octet1	0x00
0x40000004	information octet2	0x00	information octet3	0x00
0x40000008	information octet4	0x00	end of frame	0x02
…	…		…	
0x400000xx	information octet6	0x01	information octet7	0x00
0x400000xx	information octet8	0x00	abandoned frame	0x03

Legend: Bit [17:16] or [1:0]:

    "00" = continuing octet

    "01" = Start of Candidate Frame

    "10" = End of Candidate Frame

    "11" = Candidate Frame Abandoned

26.4.4 CASE 1

When an Event 4 – (E4): CODEBLOCK REJECTION – occurs for any of the 37 possible "Candidate Codeblocks" that can follow Codeblock 0 (possibly the tail sequence), the decoder returns to the SEARCH state (S2), with the following actions:

- The codeblock is abandoned (erased)
- No information octets are transferred to the remote ring buffer
- An "End of Candidate Frame" flag is written, indicating the end of the transfer of a block of octets that make up a "Candidate Frame".

26.4.5 CASE 2

When an Event 2 – (E2): CHANNEL DEACTIVATION – occurs which affects any of the 37 possible "Candidate Codeblocks" that can follow Codeblock 0, the decoder returns to the INACTIVE state (S1), with the following actions:

- The codeblock is abandoned (erased)
- No information octets are transferred to the remote ring buffer
- An "End of Candidate Frame" flag is written, indicating the end of the transfer of a block of octets that make up a "Candidate Frame"

26.4.6 Abandoned

- B: When an Event 4 (E4), or an Event 2 (E2), occurs which affects the first candidate codeblock 0, the CLTU shall be abandoned. No candidate frame octets have been transferred.
- C: If and when more than 37 Codeblocks have been accepted in one CLTU, the decoder returns to the SEARCH state (S2). The CLTU is effectively aborted and this is will be reported to the software by writing the "Candidate Frame Abandoned flag" to bit 1 or 17, indicating to the software to erase the "Candidate frame".

26.5 Relationship between buffers and FIFOs

The conversion from the peripheral data width (8 bit for the coding layer receiver), to 32 bit system word width, is done in the peripheral FIFO.

All access towards the system ring buffer are 32-bit aligned. When the amount of received bytes is odd or not 32-bit aligned, the FIFO will keep track of this and automatically solve this problem. For the reception data path, the 32 bit aligned accesses could result in incomplete words being written to the ring buffer. This means that some bytes aren't correct (because not yet received), but this is no problem due to the fact that the hardware write pointer (rx_w_ptr) always points to the last, correct, data byte.

The local FIFO ensures that DMA transfer on the AMBA AHB bus can be made by means of 2-word bursts. If the FIFO is not yet filled and no new data is being received this shall generate a combination of single accesses to the AMBA AHB bus if the last access was indicating an end of frame or an abandoned frame.

If the last single access is not 32-bit aligned, this shall generate a 32-bit access anyhow, but the receive-write-pointer shall only be incremented with the correct number of bytes. Also in case the previous access was not 32-bit aligned, then the start address to write to will also not be 32-bit aligned. Here the previous 32-bit access will be repeated including the bytes that were previously missing, in order to fill-up the 32-bit remote memory-controller without gaps between the bytes.

The receive-write-pointer shall be incremented according to the number of bytes being written to the remote memory controller.

26.5.1 Buffer full

The receiving buffer is full when the hardware has filled the complete buffer space while the software didn't read it out. Due to hardware implementation and safety, the buffer can't be filled completely without interaction of the software side. A space (offset) between the software read pointer (rx_r_ptr) and the hardware write pointer (rx_w_ptr) is used as safety buffer. When the write pointer (rx_w_ptr) would enter this region (due to a write request from the receiver), a buffer full signal is generated and all hardware writes to the buffer are suppressed. The offset is hard coded to 8 bytes.

Warning: If the software wants to receive a complete 1 KiB block (when RXLEN = 0), then it must read out at least 8 bytes of data from the buffer. In this case, the hardware can write the 1024 bytes without being stopped by the rx buffer full signal.

Figure 81. Buffer full situation

26.5.2 Buffer full interrupt

The buffer full interrupt is given when the difference between the hardware write pointer (rx_w_ptr) and the software read pointer (rx_r_ptr) is less than 1/8 of the buffer size. The way it works is the same as with the buffer full situation, only is the interrupt active when the security zone is entered. The buffer full interrupt is active for 1 system clock cycle. When the software reads out data from the buffer, the security zone shifts together with the read pointer (rx_r_ptr) pointer. Each time the hardware write pointer (rx_w_ptr) enters the security zone, a single interrupt is given.

Figure 82. Buffer full interrupt (buffers size is 2 KiB in this example)

26.6 Command Link Control Word interface (CLCW)

The Command Link Control Word (CLCW) is inserted in the Telemetry Transfer Frame by the Telemetry Encoder (TM) when the Operational Control Field (OCF) use is enabled. The CLCW needs to be created by the software part of the telecommand decoder and to be written by software to the GRTM OCF register in the Telemetry Encoder.
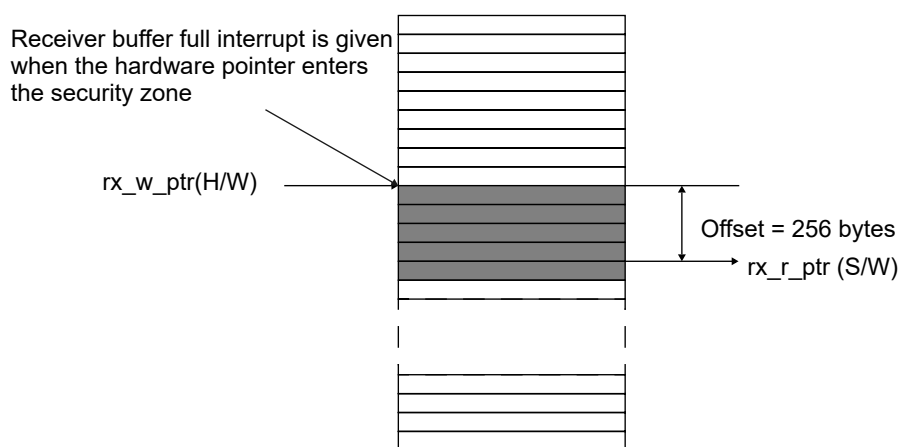
The telecommand decoder hardware provides two CLCW registers through which bit 16 (No RF Available) and 17 (No Bit Lock) of the CLCW can be read by the telecommand decoder software. The information carried in these bits is based on the discrete inputs TCRFAVL[4:0] and TCACT[4:0] respectively, where the multiple inputs are OR-ed before being assigned to the two bits. Note that if these telecommand input signals are shared with other functions via the I/O switch matrix, then the values of bit 16 and 17 might reflect the status of the transponders incorrectly. The other bits of the CLCW registers have no impact on any hardware more than providing a temporary storage place that software may utilize.

Note that the two corresponding bits can also be automatically overwritten in the OCF part of the Telemetry Transfer Frame by the Telemetry Encoder (GRTM) hardware if enabled through the Over Write OCF (OW) bit in the GRTM master frame generation register. This provides the same functionality as reading the bits through the CLCW registers in the telecommand decoder and then writing them to the GRTM OCF register in the Telemetry Encoder.

26.7 Configuration Interface (AMBA AHB slave)

The AMBA AHB slave interface supports 32 bit wide data input and output. Since each access is a word access, the two least significant address bits are assumed always to be zero, address bits 23:0 are decoded. Note that address bits 31:24 are not decoded and should thus be handled by the AHB arbiter/ decoder. The address input of the AHB slave interfaces is thus incompletely decoded. Misaligned addressing is not supported. For read accesses, unmapped bits are always driven to zero.

The AMBA AHB slave interface has been reduced in function to support only what is required for the TC. The following AMBA AHB features are constrained:

- Only supports HSIZE=WORD, HRESP_ERROR generated otherwise
- Only supports HMASTLOCK='0', HRESP_ERROR generated otherwise
- Only support HBURST=SINGLE or INCR, HRESP_ERROR generated otherwise
- No HPROT decoding
- No HSPLIT generated
- HRETRY is generated if a register is inaccessible due to an ongoing reset.
- HRESP_ERROR is generated for unmapped addresses, and for write accesses to register without any writeable bits
- Only big-endianness is supported.

During a channel reset the RRP and RWP registers are temporary unavailable. The duration of this reset-inactivity is 8 HCLK clock periods and the AHB-slave generates a HRETRY response during this period if an access is made to these registers.

If the channel reset is initiated by or during a burst-access the reset will execute correctly but a part of the burst could be answered with a HRETRY response. It is therefore not recommended to initiate write bursts to the register.

GRTC has one interrupt output, that is asserted on the occurrence of one of following events:

- 'CLTU stored' (generated when CLTU has been stored towards the AMBA bus, also issued for abandoned CLTUs)
- 'Receive buffer full' (generated when the buffer has less than 1/8 free) (note that this interrupt is issued on a static state of the buffer, and can thus be re-issued immediately after the correspond-

ing register has been read out by software, it should be masked in the interrupt controller to avoid an immediate second interrupt).

- 'Receiver overrun' (generated when received data is dropped due to a reception overrun)
- 'CLTU ready' (note that this interrupt is also issued for abandoned CLTUs)
- FAR interrupt 'Status Survey Data'
- CLCW interrupt 'Bit Lock'
- CLCW interrupt 'RF Available'

26.8   Miscellaneous

26.8.1   Numbering and naming conventions

Convention according to the CCSDS recommendations, applying to time structures:

- The most significant bit of an array is located to the left, carrying index number zero.
- An octet comprises eight bits.

*Table 204.*CCSDS n-bit field definition

CCSDS n-bit field		
most significant		least significant
0	1 to n-2	n-1

Convention according to AMBA specification, applying to the APB/AHB interfaces:

- Signal names are in upper case, except for the following:
- A lower case 'n' in the name indicates that the signal is active low.
- Constant names are in upper case.
- The least significant bit of an array is located to the right, carrying index number zero.
- Big-endian support.

*Table 205.*AMBA n-bit field definition

AMBA n-bit field		
most significant		least significant
n-1	n-2 down to 1	0

General convention, applying to all other signals and interfaces:

- Signal names are in mixed case.
- An upper case '_N' suffix in the name indicates that the signal is active low.

26.9 Registers

The core is programmed through registers mapped into AHB I/O address space.

Table 206. GRTC registers

AHB address	Register
0xFFF10000	Global Reset Register (GRR)
0xFFF10004	Global Control Register (GCR)
0xFFF10008	Physical Interface Mask Register (PMR)
0xFFF1000C	Spacecraft Identifier Register (SIR)
0xFFF10010	Frame Acceptance Report Register (FAR)
0xFFF10014	CLCW Register 1 (CLCWR1)
0xFFF10018	CLCW Register 2 (CLCWR2)
0xFFF1001C	Physical Interface Register (PHIR)
0xFFF10020	Control Register (COR)
0xFFF10024	Status Register (STR)
0xFFF10028	Address Space Register (ASR)
0xFFF1002C	Receive Read Pointer Register (RRP)
0xFFF10030	Receive Write Pointer Register (RWP)
0xFFF10060	Pending Interrupt Masked Status Register (PIMSR)
0xFFF10064	Pending Interrupt Masked Register (PIMR)
0xFFF10068	Pending Interrupt Status Register (PISR)
0xFFF1006C	Pending Interrupt Register (PIR)
0xFFF10070	Interrupt Mask Register (IMR)
0xFFF10074	Pending Interrupt Clear Register (PICR)

Table 207. Global Reset Register (GRR)

31	24	23	1	0
SEB		RESERVED		SRST

31: 24   SEB (Security Byte):

        Write: '0x55'= the write will have effect (the register will be updated).
                        Any other value= the write will have no effect on the register.

        Read: All zero.

23: 1   RESERVED

        Write: Don't care.

        Read: All zero.

0   System reset (SRST): **[1]**

        Write: '1'= initiate reset, '0'= do nothing

        Read: '1'= unsuccessful reset, '0'= successful reset

Power-up default: 0x00000000

Table 208. Global Control Register (GCR)

31	24	23	13	12	11	10	9	0
SEB		RESERVED		PSS	NRZM	PSR	RESERVED	

31: 24   SEB (Security Byte):

        Write: '0x55'= the write will have effect (the register will be updated).
                        Any other value= the write will have no effect on the register.

        Read: All zero.

Table 208. Global Control Register (GCR)

23: 13	RESERVED	
	Write:	Don't care.
	Read:	All zero.
12	PSS (ESA/PSS enable) [11]	
	Write/Read:	'0'= disable, '1'= enable
11	NRZM (Non-Return-to-Zero Mark Decoder enable)	
	Write/Read:	'0'= disable, '1'= enable
10	PSR (Pseudo-De-Randomiser enable)	
	Write/Read:	'0'= disable, '1'= enable
9: 0	RESERVED	
	Write:	Don't care.
	Read:	All zero.

Power-up default: 0x00000000

Table 209. Physical Interface Mask Register (PMR)

31	8	7	0
RESERVED		MASK	

31: 8	RESERVED	
	Write:	Don't care.
	Read:	All zero.
7: 0	RESERVED	
	Write:	Mask TC input when set, bit 0 corresponds to TC input 0
	Read:	Current mask

Power-up default: 0x00000000

Table 210. Spacecraft Identifier Register (SIR) [7]

31	10	9	0
RESERVED		SCID	

31: 10	RESERVED	
	Write:	Don't care.
	Read:	All zero.
9: 0	SCID (Spacecraft Identifier)	
	Write:	Don't care.
	Read:	Bit[9]=MSB, Bit[0]=LSB

Power-up default: 0x00000000

Table 211. Frame Acceptance Report Register (FAR) [7]

31	30	25	24	19	18	16	15	14	13	11	10	0
SSD	RESERVED		CAC		CSEC		RESERVED		SCI		RESERVED	

31	SSD (Status of Survey Data) (see [PSS-04-151])	
	Write:	Don't care.
	Read:	Automatically cleared to 0 when any other field is updated by the coding layer. Automatically set to 1 upon a read.
30: 25	RESERVED	
	Write:	Don't care.
	Read:	All zero.
24: 19	CAC (Count of Accept Codeblocks) (see [PSS-04-151])	
	Write:	Don't care.

Table 211. Frame Acceptance Report Register (FAR) [7]

		Read:	Information obtained from coding layer. [2]
18: 16	CSEC (Count of Single Error Corrections) (see [PSS-04-151])		
		Write:	Don't care.
		Read:	Information obtained from coding layer.
15: 14	RESERVED		
		Write:	Don't care.
		Read:	All zero.
13: 11	SCI (Selected Channel Input) (see [PSS-04-151])		
		Write:	Don't care.
		Read:	Information obtained from coding layer.
10: 0	RESERVED		
		Write:	Don't care.
		Read:	All zero.

Power-up default: 0x00003800

Table 212. CLCW Register (CLCWRx) (see [PSS-04-107])

31	30 29	28 26	25 24	23 18	17 16	15	14	13	12	11	10 9	8	7 0
CWTY	VNUM	STAF	CIE	VCI	RESERVED	NRFA	NBLO	LOUT	WAIT	RTMI	FBCO	RTYPE	RVAL

31	CWTY (Control Word Type)		
30: 29	VNUM (CLCW Version Number)		
28: 26	STAF (Status Fields)		
25: 24	CIE (COP In Effect)		
23: 18	VCI (Virtual Channel Identifier)		
17: 16	Reserved (PSS/ECSS requires "00")		
15	NRFA (No RF Available)		
		Write:	Don't care.
		Read:	Based on discrete inputs TCRFAVL[4:0], where the inputs are OR-ed
14	NBLO (No Bit Lock)		
		Write:	Don't care.
		Read:	Based on discrete inputs TCACT[4:0], where the inputs are OR-ed
13	LOUT (Lock Out)		
12	WAIT (Wait)		
11	RTMI (Retransmit)		
10: 9	FBCO (FARM-B Counter)		
8	RTYPE (Report Type)		
7: 0	RVAL (Report Value)		

Power-up default: 0x00000000

Table 213. Physical Interface Register (PHIR) [7]

31 16	15 8	7 0
RESERVED	RFA	BLO

31: 16	RESERVED		
		Write:	Don't care.
		Read:	All zero.
15: 8	RFA (RF Available) [3]		

Only the implemented inputs 0 through 4 are taken into account. All other bits are zero.

		Write:	Don't care.
		Read:	Bit[8] = input 0, Bit[15] = input 7

Table 213. Physical Interface Register (PHIR) [7]

7: 0   BLO (Bit Lock) [3]

Only the implemented inputs 0 through 4 are taken into account. All other bits are zero.

Write: Don't care.

Read: Bit[0] = input 0, Bit[7] = input 7

Power-up default: Depends on inputs.

Table 214. Control Register (COR)

31		24	23		10	9	8		1	0
	SEB			RESERVED		CRST		RESERVED		RE

31: 24   SEB (Security Byte):

Write: '0x55'= the write will have effect (the register will be updated).
            Any other value= the write will have no effect on the register.

Read: All zero.

23: 10   RESERVED

Write: Don't care.

Read: All zero.

9   CRST (Channel reset) [4]

Write: '1'= initiate channel reset,'0'= do nothing

Read: '1'= unsuccessful reset, '0'= successful reset

8: 1   RESERVED

Write: Don't care.

Read: All zero.

0   RE (Receiver Enable) TCACT[4:0] inputs of the receiver are masked when the RE bit is disabled.

Read/Write: '0'= disabled, '1'= enabled

Power-up default: 0x00000000

Table 215. Status Register (STR) [7]

31		11	10	9	8	7	6	5	4	3	1	0
	RESERVED		RBF	RESERVED		RFF	RESERVED		OV	RESERVED		CR

31: 11   RESERVED

Write: Don't care.

Read: All zero.

10   RBF (RX BUFFER Full)

Write: Don't care.

Read: '0' = Buffer not full,

            '1'= Buffer full (this bit is set if the buffer has less then 1/8 of free space)

9: 8   RESERVED

Write: Don't care.

Read: All zero.

7   RFF (RX FIFO Full)

Write: Don't care.

Read: '0' = FIFO not full, '1' = FIFO full

6: 5   RESERVED

Write: Don't care.

Read: All zero.

4   OV (Overrun) [5]

Write: Don't care.

Read: '0'= nominal, '1'= data lost

Table 215. Status Register (STR) [7]

3: 1   RESERVED

         Write: Don't care.

         Read: All zero.

0   CR (CLTU Ready) [5]

There is a worst case delay from the CR bit being asserted, until the data has actually been transferred from the receiver FIFO to the ring buffer. This depends on the bus load etc.

         Write: Don't care.

         Read: '1'= new CLTU in ring buffer. '0'= no new CLTU in ring buffer.

Power-up default: 0x00000000

Table 216. Address Space Register (ASR) [8]

31	10 9	8 7	0
BUFST	RESERVED	RXLEN	

31: 10   BUFST (Buffer Start Address)

22-bit address pointer

This pointer contains the start address of the allocated buffer space for this channel.

Register has to be initialized by software before DMA capability can be enabled.

9: 8   RESERVED

         Write: Don't care.

         Read: All zero.

7: 0   RXLEN *(RX buffer length)*

Number of 1kB-blocks reserved for the RX buffer.
(Min.   1 KiB = 0x00, Max. 256 KiB = 0xFF)

Power-up default: 0x00000000

Table 217. Receive Read Pointer Register (RRP) [6] [9][10]

31	24 23	0
RxRd Ptr Upper	RxRd Ptr Lower	

31: 24   10-bit upper address pointer

         Write: Don't care.

         Read: This pointer = ASR[31..24].

23: 0   24-bit lower address pointer.

This pointer contains the current RX read address. This register is to be incremented with the actual amount of bytes read.

Power-up default: 0x00000000

Table 218. Receive Write Pointer Register (RWP) [6] [9]

31	24 23	0
RxWr Ptr Upper	RxWr Ptr Lower	

31: 24   10-bit upper address pointer

         Write: Don't care.

         Read: This pointer = ASR[31..24].

23: 0   24-bit lower address pointer.

This pointer contains the current RX write address. This register is incremented with the actual amount of bytes written.

Power-up default: 0x00000000

Legend:

[1]   The global system reset caused by the SRST-bit in the GRR-register results in the following actions:

- Initiated by writing a '1", gives '0' on read-back when the reset was successful.

- No need to write a '0' to remove the reset.

- Unconditionally, means no need to check/disable something in order for this reset-function to correctly execute.

- Could of course lead to data-corruption coming/going from/to the reset core.

- Resets the complete core (all logic, buffers & register values)

- Behaviour is similar to a power-up.

[2]   The FAR register supports the CCSDS/ECSS standard frame lengths (1024 octets), requiring an 8 bit CAC field instead of the 6 bits specified in [PSS-04-151]. The two most significant bits of the CAC will thus spill over into the "LEGAL/ILLEGAL" FRAME QUALIFIER field, Bit [26:25]. This is only the case when the PSS bit is set to '0'.

[3]   Only inputs 0 trough 4 are implemented.

[4]   The channel reset caused by the CRST-bit in the COR-register results in the following actions:

- Initiated by writing a '1", gives '0' on read-back when the reset was successful.

- No need to write a '0' to remove the reset.

- All other bit's in the COR are neglected (not looked at) when the CRST-bit is set during a write, meaning that the value of these bits has no impact on the register-value after the reset.

- Unconditionally, means no need to check/disable something in order for this reset-function to correctly execute.

- Could of course lead to data-corruption coming/going from/to the reset channel.

- Resets the complete channel (all logic, buffers & register values)

- Except the ASR-register of that channel which remains it's value.

- All read- and write-pointers are automatically re-initialized and point to the start of the ASR-address.

- All registers of the channel (except the ones described above) get their power-up value.

- This reset shall not cause any spurious interrupts.

[5]   These bits are sticky bits which means that they remain present until the register is read and that they are cleared automatically by reading the register.

[6]   The value of the pointers depends on the content of the corresponding Address Space Register (ASR).

During a system reset, a channel reset or a change of the ASR register, the pointers are recalculated based on the values in the ASR register.

The software has to take care (when programming the ASR register) that the pointers never have to cross a 16 MiB boundary (because this would cause an overflow of the 24-bit pointers).

It is not possible to write an out of range value to the RRP register. Such access will be ignored with an HERROR.

[7]   An AMBA AHB ERROR response is generated if a write access is attempted to a register without any writeable bits.

[8]   The channel reset caused by a write to the ASR-register results in the following actions:

- Initiated by writing an updated value into the ASR-register.

- Unconditionally, means no need to check/disable something in order for this reset-function to correctly execute.

- Could of course lead to data-corruption coming/going from/to the reset channel.

- Resets the complete channel (all logic & buffers) but not all register values, only the following:

- COR-register, TE & RE bits get their power-up value, other bits remain their value.

- STR-register, all bits get their power-up value

- Other registers remain their value

- Updates the ASR-register of that channel with the written value

- All read- and write-pointers are automatically re-initialized and point to the start of the ASR-address.

- This reset shall not cause any spurious interrupts

[9]   During a channel reset the register is temporarily unavailable and HRETRY response is generated if accessed.

[10] It is not possible to write an out of range value to the RRP register. Such access will be ignored without an error.

[11] The PSS bit usage is supported.

26.9.1 Interrupt registers

The interrupt registers give complete freedom to the software, by providing means to mask interrupts, clear interrupts, force interrupts and read interrupt status.

When an interrupt occurs the corresponding bit in the Pending Interrupt Register is set. The normal sequence to initialize and handle a module interrupt is:

- Set up the software interrupt-handler to accept an interrupt from the module.

- Read the Pending Interrupt Register to clear any spurious interrupts.

- Initialize the Interrupt Mask Register, unmasking each bit that should generate the module interrupt.

- When an interrupt occurs, read the Pending Interrupt Status Register in the software interrupt-handler to determine the causes of the interrupt.

- Handle the interrupt, taking into account all causes of the interrupt.

- Clear the handled interrupt using Pending Interrupt Clear Register.

Masking interrupts: After reset, all interrupt bits are masked, since the Interrupt Mask Register is zero. To enable generation of a module interrupt for an interrupt bit, set the corresponding bit in the Interrupt Mask Register.

Clearing interrupts: All bits of the Pending Interrupt Register are cleared when it is read or when the Pending Interrupt Masked Register is read. Reading the Pending Interrupt Masked Register yields the contents of the Pending Interrupt Register masked with the contents of the Interrupt Mask Register. Selected bits can be cleared by writing ones to the bits that shall be cleared to the Pending Interrupt Clear Register.

Forcing interrupts: When the Pending Interrupt Register is written, the resulting value is the original contents of the register logically OR-ed with the write data. This means that writing the register can force (set) an interrupt bit, but never clear it.

Reading interrupt status: Reading the Pending Interrupt Status Register yields the same data as a read of the Pending Interrupt Register, but without clearing the contents.

Reading interrupt status of unmasked bits: Reading the Pending Interrupt Masked Status Register yields the contents of the Pending Interrupt Register masked with the contents of the Interrupt Mask Register, but without clearing the contents.

The interrupt registers comprise the following:

- Pending Interrupt Masked Status Register   [PIMSR]   R
- Pending Interrupt Masked Register   [PIMR]   R
- Pending Interrupt Status Register   [PISR]   R
- Pending Interrupt Register   [PIR]   R/W
- Interrupt Mask Register   [IMR]   R/W
- Pending Interrupt Clear Register   [PICR]   W

*Table 219.*Interrupt registers

31 7	6	5	4	3	2	1	0
-	CS	OV	RBF	CR	FAR	BLO	RFA

6:	CS	CLTU stored
5:	OV	Input data overrun
4:	RBF	Output buffer full
3:	CR	CLTU ready/aborted
2:	FAR	FAR available
1:	BLO	Bit Lock changed

0:         RFA   RF Available Changed

All bits in all interrupt registers are reset to 0b after reset.

26.10 Signal definitions

The signals are described in table 220.

Table 220. Signal definitions

Signal name	Type	Function	Active
TCACT[4:0]	Input, async	Active	High
TCLK[4:0]	Input, async	Bit clock	Rising edge
TCD[4:0]	Input, async	Data	-
TCRFAVL[4:0]	Input, async	RF Available for CLCW	High

27 GRTM - CCSDS Telemetry Encoder

27.1 Overview

The CCSDS/ECSS/PSS Telemetry Encoder implements part of the Data Link Layer, covering the Protocol Sub-layer and the Frame Synchronization and Coding Sub-layer and part of the Physical Layer of the packet telemetry encoder protocol.

The operation of the Telemetry Encoder is highly programmable by means of control registers.

The Telemetry Encoder comprises several encoders and modulators implementing the Consultative Committee for Space Data Systems (CCSDS) recommendations, European Cooperation on Space Standardization (ECSS) and the European Space Agency (ESA) Procedures, Standards and Specifications (PSS) for telemetry and channel coding. The Telemetry Encoder comprises the following:

- Packet Telemetry and/or Advanced Orbiting Systems (AOS) Encoder
- Reed-Solomon Encoder
- Pseudo-Randomiser (PSR)
- Non-Return-to-Zero Mark encoder (NRZ)
- Convolutional Encoder (CE)
- Split-Phase Level modulator (SP)
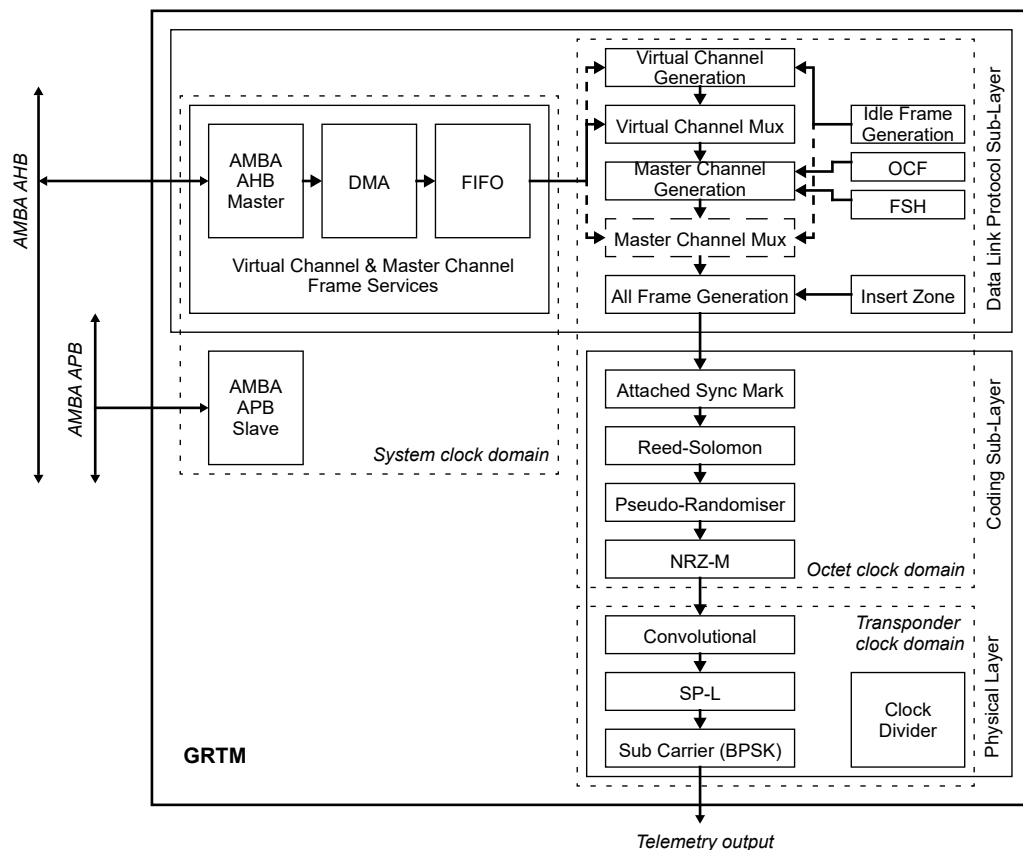- Sub-Carrier modulator (SC)
- Clock Divider (CD)

Figure 83. Block diagram

27.2 References

27.2.1 Documents

[CCSDS-131.0-B-1]   TM Synchronization and Channel Coding

[CCSDS-132.0-B-1]   TM Space Data Link Protocol

[CCSDS-133.0-B-1]   Space Packet Protocol

[CCCSDS-732.0-B-2]   AOS Space Data Link Protocol

[ECSS-E-ST-50-01C]   Space engineering - Space data links - Telemetry synchronization and channel coding

[ECSS-E-ST-50-03C]   Space engineering - Space data links - Telemetry transfer frame protocol

[ECSS-E-ST-50-05C]   Space engineering - Radio frequency and modulation

[PSS-04-103]   Telemetry channel coding standard

[PSS-04-105]   Radio frequency and modulation standard

[PSS-04-106]   Packet telemetry standard

27.2.2 Acronyms and abbreviations

AOS   Advanced Orbiting Systems

ASM   Attached Synchronization Marker

CCSDS   Consultative Committee for Space Data Systems

CLCW   Command Link Control Word

CRC   Cyclic Redundancy Code

DMA   Direct Memory Access

ECSS   European Cooperation for Space Standardization

ESA   European Space Agency

FECF   Frame Error Control Field

FHEC   Frame Header Error Control

FHP   First Header Pointer

GF   Galois Field

LFSR   Linear Feedback Shift Register

MC   Master Channel

NRZ   Non Return to Zero

OCF   Operational Control Field

PSR   Pseudo Randomiser

PSS   Procedures, Standards and Specifications

RS   Reed-Solomon

SP   Split-Phase

TE   Turbo Encoder

TM   Telemetry

VC   Virtual Channel

27.3 Layers

27.3.1 Introduction

The Packet Telemetry (or simply Telemetry or TM) and Advanced Orbiting System (AOS) standards are similar in their format, with only some minor variations. The AOS part covered here is the down-link or transmitter, not the uplink or receiver.

The relationship between these standards and the Open Systems Interconnection (OSI) reference model is such that the OSI Data Link Layer corresponds to two separate layer, namely the Data Link Protocol Sub-layer and Synchronization and Channel Coding Sub-Layer. The OSI Data Link Layer is covered here.

The OSI Physical Layer is also covered here to some extended, as specified in [ECSS-E-ST-50-05C] and [PSS-04-105].

The OSI Network Layer or higher layers are not covered here.

27.3.2 Data Link Protocol Sub-layer

The Data Link Protocol Sub-layer differs somewhat between TM and AOS. Differences are pointed out where needed in the subsequent descriptions.

The following functionality is not implemented in the core:

• Packet Processing

• Bitstream Processing (applies to AOS only)

The following functionality is implemented in the core:

• Virtual Channel Generation (for Idle Frame generation only)

• Virtual Channel Multiplexing (for Idle Frame generation only)

• Master Channel Generation (applies to Packet Telemetry only)

• Master Channel Multiplexing (including Idle Frame generation)

• All Frame Generation

27.3.3 Synchronization and Channel Coding Sub-Layer

The Synchronization and Channel Coding Sub-Layer does not differ between TM and AOS.

The following functionality is implemented in the core:

• Attached Synchronization Marker

• Reed-Solomon coding

• Turbo coding (future option)

• Pseudo-Randomiser

• Convolutional coding

27.3.4 Physical Layer

The Physical Layer does not differ between TM and AOS.

The following functionality is implemented in the core:

• Non-Return-to-Zero modulation

• Split-Phase modulation

• Sub-Carrier modulation

27.4 Data Link Protocol Sub-Layer

27.4.1 Physical Channel

The configuration of a Physical Channel covers the following parameters:

• Transfer Frame Length (in number of octets)

• Transfer Frame Version Number

Note that there are other parameters that need to be configured for a Physical Channel, as listed in section 27.4.8, covering the All Frame Generation functionality.

The Transfer Frame Length can be programmed by means of the DMA length register.

The Transfer Frame Version Number can be programmed by means of a register, and can take one of two legal values: 00b for Telemetry and 01b for AOS.

27.4.2 Virtual Channel Frame Service

The Virtual Channel Frame Service is implemented by means of a DMA interface, providing the user with a means for inserting Transfer Frames into the Telemetry Encoder. Transfer Frames are automatically fetched from memory, for which the user configures a descriptor table with descriptors that point to each individual Transfer Frame. For each individual Transfer Frame the descriptor also provides means for bypassing functions in the Telemetry Encoder. This includes the following:

• Virtual Channel Counter generation can be enabled in the Virtual Channel Generation function (this function is normally only used for Idle Frame generation but can be used for the Virtual Channel Frame Service when sharing a Virtual Channel)

• Master Channel Counter generation can be bypassed in the Master Channel Generation function (TM only)

• Frame Secondary Header (FSH) is not supported

• Operational Control Field (OCF) generation can be bypassed in the Master Channel Generation function (TM only)

• Frame Error Header Control (FECH) generation can be bypassed in the All Frame Generation function (AOS only)

• Insert Zone (IZ) generation is not supported

• Frame Error Control Field (FECF) generation can be bypassed in the All Frame Generation function

• A Time Strobe can be generated for the Transfer Frame.

Note that the above features can only be bypassed for each Transfer Frame, the overall enabling of the features is done for the corresponding functions in the Telemetry Encoder, as described in the subsequent sections.

The detailed operation of the DMA interface is described in section 27.8.

27.4.3 Virtual Channel Generation

The Virtual Channel Generation function is used to generate the Virtual Channel Counter for Idle Frames as described hereafter. The function can however also be enabled for any Transfer Frame inserted via the Virtual Channel Frame Service described above, allowing a Virtual Channel to be shared between the two services. In this case the Virtual Channel Counter, the Extended Virtual Channel Counter (only for TM, as defined for ECSS and PSS, including the complete Transfer Frame Secondary Header) and the Virtual Channel Counter Cycle (only for AOS) fields will be inserted and incremented automatically when enabled as described hereafter.

27.4.4 Virtual Channel Multiplexing

The Virtual Channel Multiplexing Function is used to multiplex Transfer Frames of different Virtual Channels of a Master Channel. Virtual Channel Multiplexing in the core is performed between two sources: Transfer Frames provided through the Virtual Channel Frame Service and Idle Frames. Note that multiplexing between different Virtual Channels is assumed to be done as part of the Virtual Channel Frame Service outside the core.

The Virtual Channel Frame Service user interface is described above. The Idle Frame generation is described hereafter.

Idle Frame generation can be enabled and disabled by means of a register. The Spacecraft ID to be used for Idle Frames is programmable by means of a register. The Virtual Channel ID to be used for Idle Frames is programmable by means of a register.

Master Channel Counter generation for Idle Frames can be enabled and disabled by means of a register (only for TM). Note that it is also possible to generate the Master Channel Counter field as part of the Master Channel Generation function described in the next section. When Master Channel Counter generation is enabled for Idle Frames, then the generation in the Master Channel Generation function is bypassed.

The Virtual Channel Counter generation for Idle Frames is always enabled (both for TM and AOS) and generated in the Virtual Channel Generation function described above.

Extended Virtual Channel Counter generation for Idle Frames can be enabled and disabled by means of a register (only for TM, as defined for ECSS and PSS). This includes the complete Transfer Frame Secondary Header.

Virtual Channel Counter Cycle generation for Idle Frames can be enabled and disabled by means of a register (only for AOS).

If Frame Secondary Header generation is not supported.

If Operation Control Field generation is enabled in the Master Channel Generation function described in the next section, it can be bypassed for Idle Frames, programmable by means of a register. This allows VC_OCF or MC_OCF realization.

27.4.5 Master Channel Generation

The Master Channel Counter can be generated for all frames on a master channel (only for TM). It can be can be enabled and disabled by means of a register. The generation can also be bypassed for Idle Frames or Transfer Frames provided via the DMA interface.

The Frame Secondary Header (FSH) is not supported.

The Operational Control Field (OCF) can be generated from a 32-bit register. This can be done for all frames on an master channel (MC_OCF) or be bypassed for Idle Frames or Transfer Frames provided via the DMA interface, effectively implementing OCF on a per virtual channel basis (VC_OCF).

27.4.6 Master Channel Frame Service

The Master Channel Frame Service user interface is equivalent to the previously described Virtual Channel Frame Service user interface, using the same DMA interface. The interface can thus be used for inserting both Master Channel Transfer Frame and Virtual Channel Transfer Frames.

27.4.7 Master Channel Multiplexing

The Master Channel Multiplexing Function is used to multiplex Transfer Frames of different Master Channels of a Physical Channel. Master Channel Multiplexing is performed between three sources: Master Channel Generation Service, Master Channel Frame Service and Idle Frames.

Bypassing all the functionality of the Master Channel Generation functionality described above effectively establishes the master channel frame service. The same holds for the Idle Frame generation described above, allowing the core to generate Idle Frame on the level of the Physical Channel.

27.4.8 All Frame Generation

The All Frame Generation functionality operates on all transfer frames of a Physical Channel. Each of the individual functions can be bypassed for each frame coming from the DMA interface or idle frame generation functionality.

The Frame Header Error Control (FHEC) generation can be enabled and disabled by means of a register (AOS only).

The Insert Zone is not supported.

Frame Error Control Field (FECF) generation can be enabled and disabled by means of a register.

27.5 Synchronization and Channel Coding Sub-Layer

27.5.1 Attached Synchronization Marker

The 32-bit Attached Synchronization Marker is placed in front of each Transfer Frame as per [CCSDS-131.0-B-1] and [ECSS-E-ST-50-03C].

An alternative Attached Synchronization Marker for embedded data streams can also be used, its enabling and bit pattern being programmable via a configuration register.

27.5.2 Reed-Solomon Encoder

The CCSDS recommendation [CCSDS-131.0-B-1] and ECSS standard [ECSS-E-ST-50-03C] specify Reed-Solomon codes, one (255, 223) code and one (255, 239) code. The ESA PSS standard [PSS013] only specifies the former code. Although the definition style differs between the documents, the (255, 223) code is the same in all three documents. The definition used in this document is based on the PSS standard [PSS013].

The Reed-Solomon Encoder implements both codes.

The Reed-Solomon encoder is compliant with the coding algorithms in [CCSDS-131.0-B-1] and [ECSS-E-ST-50-03C]:

- there are 8 bits per symbol;

- there are 255 symbols per codeword;

- the encoding is systematic:

- for E=8 or (255, 239), the first 239 symbols transmitted are information symbols, and the last 16 symbols transmitted are check symbols;

- for E=16 or (255, 223), the first 223 symbols transmitted are information symbols, and the last 32 symbols transmitted are check symbols;

- the E=8 code can correct up to 8 symbol errors per codeword;

- the E=16 code can correct up to 16 symbol errors per codeword;

- the field polynomial is

$$f_{esa}(x) = x^8 + x^6 + x^4 + x^3 + x^2 + x + 1$$

- the code generator polynomial for E=8 is

    for which the highest power of x is transmitted first;

- the code generator polynomial for E=16 is

$$g_{esa}(x) = \prod_{i=120}^{135} (x + \alpha^i) = \sum_{j=0}^{16} g_j \cdot x^j$$

$$g_{esa}(x) = \prod_{i=112}^{143} (x + \alpha^i) = \sum_{j=0}^{32} g_j \cdot x^j$$

for which the highest power of x is transmitted first;

- interleaving is supported for depth $I = \{1 \text{ to } 8\}$, where information symbols are encoded as I codewords with symbol numbers $i + j*I$ belonging to codeword i {where $0 \leq i < I$ and $0 \leq j < 255$};

- shortened codeword lengths are supported;

- the input and output data from the encoder are in the representation specified by the following transformation matrix T_{esa}, where i_0 is transferred first

$$\begin{bmatrix} \iota_0 & \iota_1 & \iota_2 & \iota_3 & \iota_4 & \iota_5 & \iota_6 & \iota_7 \end{bmatrix} = \begin{bmatrix} \alpha_7 & \alpha_6 & \alpha_5 & \alpha_4 & \alpha_3 & \alpha_2 & \alpha_1 & \alpha_0 \end{bmatrix} \times \begin{bmatrix} 0 & 0 & 1 & 1 & 0 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 & 1 & 1 & 1 & 1 \\ 1 & 0 & 0 & 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 & 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 0 & 1 & 0 & 1 & 1 \\ 0 & 1 & 1 & 1 & 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 1 & 0 & 1 & 1 \end{bmatrix}$$

- the following matrix T^{-1}_{esa} specifying the reverse transformation

$$\begin{bmatrix} \alpha_7 & \alpha_6 & \alpha_5 & \alpha_4 & \alpha_3 & \alpha_2 & \alpha_1 & \alpha_0 \end{bmatrix} = \begin{bmatrix} \iota_0 & \iota_1 & \iota_2 & \iota_3 & \iota_4 & \iota_5 & \iota_6 & \iota_7 \end{bmatrix} \times \begin{bmatrix} 1 & 1 & 1 & 0 & 1 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 1 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 & 1 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 & 0 & 0 \end{bmatrix}$$

- the Reed-Solomon output is non-return-to-zero level encoded.

The Reed-Solomon Encoder encodes a bit stream from preceding encoders and the resulting symbol stream is output to subsequent encoder and modulators. The encoder generates codeblocks by receiving information symbols from the preceding encoders which are transmitted unmodified while calculating the corresponding check symbols which in turn are transmitted after the information symbols. The check symbol calculation is disabled during reception and transmission of unmodified data not related to the encoding. The calculation is independent of any previous codeblock and is perform correctly on the reception of the first information symbol after a reset.

Each information symbol corresponds to an 8 bit symbol. The symbol is fed to a binary network in which parallel multiplication with the coefficients of a generator polynomial is performed. The products are added to the values contained in the check symbol memory and the sum is then fed back to the check symbol memory while shifted one step. This addition is performed octet wise per symbol. This cycle is repeated until all information symbols have been received. The contents of the check symbol memory are then output from the encoder. The encoder is based on a parallel architecture, including parallel multiplier and adder.

The encoder supports E=16 (255, 223) and E=8 (255, 239) codes. The choice between the E=16 and E=8 coding can be performed during operation by means of a configuration register.

The maximum number of supported interleave depths I_{max} is 8. The choice of any interleave depth ranging from 1 to the chosen I_{max} is supported during operation. The interleave depth is chosen during operation by means of a configuration register.

27.5.3 Pseudo-Randomiser

The Pseudo-Randomiser (PSR) generates a bit sequence according to [CCSDS-131.0-B-1] and [ECSS-E-ST-50-03C] which is xor-ed with the data output of preceding encoders. This function allows the required bit transition density to be obtained on a channel in order to permit the receiver on ground to maintain bit synchronization.

The polynomial for the Pseudo-Randomiser is h (x) = $x^8 + x^7 + x^5 + x^3 + 1$ and is implemented as a Fibonacci version (many-to-one implementation) of a Linear Feedback Shift Register (LFSR). The registers of the LFSR are initialized to all ones between Transfer Frames. The Attached Synchronization Marker (ASM) is not effected by the encoding.


Figure 84. Pseudo-randomiser

27.5.4 Convolutional Encoder

The Convolutional Encoder (CE) implements two convolutional encoding schemes. The ESA PSS standard [PSS-04-103] specifies a basic convolutional code without puncturing. This basic convolutional code is also specified in the CCSDS recommendation [CCSDS-131.0-B-1] and ECSS standard [ECSS-E-ST-50-03C], which in addition specifies a punctured convolutional code.

The basic convolutional code has a code rate of 1/2, a constraint length of 7, and the connection vectors G1 = 1111001_b (171 octal) and G2 = 1011011_b (133 octal) with symbol inversion on output path, where G1 is associated with the first symbol output.

The punctured convolutional code has a code rate of 1/2 which is punctured to 2/3, 3/4, 5/6 or 7/8, a constraint length of 7, and the connection vectors G1 = 1111001_b (171 octal) and G2 = 1011011_b (133 octal) without any symbol inversion. The puncturing and output sequences are defined in [CCSDS-131.0-B-1]. The encoder also supports rate 1/2 unpunctured coding with aforementioned connection vectors and no symbol inversion.


*Figure 85.* Unpuctured convolutional encoder

Figure 86. Punctured convolutional encoder

27.6 Physical Layer

27.6.1 Non-Return-to-Zero Mark encoder

The Non-Return-to-Zero Mark encoder (NRZ) encodes differentially a bit stream from preceding encoders according to [ECSS-E-ST-50-05C]. The waveform is shown in figure 87

Both data and the Attached Synchronization Marker (ASM) are affected by the coding. When the encoder is not enabled, the bit stream is by default non-return-to-zero level encoded.


Figure 87. NRZ-L and NRZ-M waveform

27.6.2 Split-Phase Level modulator

The Split-Phase Level modulator (SP) modulates a bit stream from preceding encoders according to [ECSS-E-ST-50-05C]. The waveform is shown in figure 88.

Both data and the Attached Synchronization Marker (ASM) are effected by the modulator. The modulator will increase the output bit rate with a factor of two.


Figure 88. SP-L waveform

27.6.3 Sub-Carrier modulator

The Sub-Carrier modulator (SC) modulates a bit stream from preceding encoders according to [ECSS-E-ST-50-05C], which is Binary Phase Shift Key modulation (BPSK) or Phase Shift Key Square.

The sub-carrier modulation frequency is programmable. The symbol rate clock be divided to a degree 2^{15}. The divider can be configured during operation to divide the symbol rate clock frequency from $1/2$ to $1/2^{15}$. The phase of the sub-carrier is programmable, selecting which phase 0^o or 180^o should correspond to a logical one on the input.

27.6.4 Clock Divider

The Clock Divider (CD) provides clock enable signals for the telemetry and channel encoding chain. The clock enable signals are used for controlling the bit rates of the different encoder and modulators.

The source for the bit rate frequency is the dedicated bit rate clock input TCLKI or the system clock. The bit rate clock input can be divided to a degree 2^{15}. The divider can be configured during operation to divide the bit rate clock frequency from 1/1 to $1/2^{15}$. In addition, the Sub-Carrier modulator can divide the above resulting clock frequency from 1/2 to $1/2^{15}$. The divider in the sub-carrier modulator can be used without enabling actual sub-carrier modulation, allowing division up to $1/2^{30}$.

The bit rate frequency is based on the output frequency of the last encoder in a coding chain, except for the sub-carrier modulator. No actual clock division is performed, since clock enable signals are used. No clock multiplexing is performed in the core.

The Clock Divider (CD) supports clock rate increases for the following encoders and rates:

* Convolutional Encoder (CE), 1/2, 2/3, 3/4, 5/6 and 7/8;

* Split-Phase Level modulator (SP-L), rate 1/2;

* Sub-Carrier modulator (SC), rate 1/2 to $1/2^{15}$.

The resulting symbol rate and telemetry rate are depended on what encoders and modulators are enabled. The following variables are used in the tables hereafter: f = input bit frequency, n = SYMBOLRATE+1 (GRTM physical layer register field +1), and m = SUBRATE+1 (physical layer register field +1), c = convolutional coding rate {1/2, 2/3, 3/4, 5/6, 7/8} (see CERATE field in GRTM coding sub-layer register).

*Table 221.*Data rates without sub-carrier modulation (SUB=0)

Coding & Modulation	Telemetry rate	Convolutional rate	Split-Phase rate	Sub-carrier frequency	Output symbol rate	Output clock frequency
-	f / n / m	-	-	-	f / n / m	f / n / m
CE	(f / n / m) * c	f / n / m	-	-	f / n / m	f / n / m
SP-L	f / n / m / 2	-	f / n / m	-	f / n / m	f / n / m
CE + SP-L	(f / n / m / 2) * c	f / n / m / 2	f / n / m	-	f / n / m	f / n / m

For n = 1, no output symbol clock is generated, i.e. SYMBOLRATE register field equals 0.
m should be an even number, i.e. SUBRATE register field should be uneven and > 0 to generate an output symbol clock with 50% duty cycle.
If m > 1 then also n must be > 1, i.e. if SUBRATE register field is > 0 then SYMBOLRATE register field must be > 0.

*Table 222.*Data rates with sub-carrier modulation (SUB=1)

Coding & Modulation	Telemetry rate	Convolutional rate	Split-Phase rate	Sub-carrier frequency	Output symbol rate [1]	Output clock frequency
SC	f / n / m	-	-	f / n / 2	f / n	f / n
CE + SC	(f / n / m) * c	f / n / m	-	f / n / 2	f / n	f / n
SP-L+ SC	f / n / m / 2	-	f / n / m	f / n / 2	f / n	f / n
CE + SP-L + SC	(f / n / m / 2) * c	f / n / m / 2	f / n / m	f / n / 2	f / n	f / n

n = 1 or m = 1 are invalid settings for sub-carrier modulation, i.e SYMBOLRATE and SUBRATE register fields must be > 0.
m must be an even number, i.e. SUBRATE register field must be uneven and > 0.
m defines number of sub-carrier phases per input bit from preceding encoder or modulator.
Note 1: The output symbol rate for sub-carrier modulation corresponds to the rate of phases, not the frequency. Sub-carrier frequency is half the symbol rate.

27.7 Connectivity

The output from the Packet Telemetry and AOS encoder can be connected to:

- Reed-Solomon encoder
- Pseudo-Randomiser
- Non-Return-to-Zero Mark encoder
- Convolutional encoder
- Split-Phase Level modulator
- Sub-Carrier modulator

The input to the Reed-Solomon encoder can be connected to:

- Packet Telemetry and AOS encoder

The output from the Reed-Solomon encoder can be connected to:

- Pseudo-Randomiser
- Non-Return-to-Zero Mark modulator
- Convolutional encoder
- Split-Phase Level modulator
- Sub-Carrier modulator

The input to the Pseudo-Randomiser (PSR) can be connected to:

- Packet Telemetry and AOS encoder
- Reed-Solomon encoder

The output from the Pseudo-Randomiser (PSR) can be connected to:

- Non-Return-to-Zero Mark modulator
- Convolutional encoder
- Split-Phase Level modulator
- Sub-Carrier modulator

The input to the Non-Return-to-Zero Mark encoder (NRZ) can be connected to:

- Packet Telemetry and AOS encoder
- Reed-Solomon encoder
- Pseudo-Randomiser

The output from the Non-Return-to-Zero Mark encoder (NRZ) can be connected to:

- Convolutional encoder
- Split-Phase Level modulator
- Sub-Carrier modulator

The input to the Convolutional Encoder (CE) can be connected to:

- Packet Telemetry and AOS encoder
- Reed-Solomon encoder

- Pseudo-Randomiser
- Non-Return-to-Zero Mark encoder

The output from the Convolutional Encoder (CE) can be connected to:

- Split-Phase Level modulator
- Sub-Carrier modulator

The input to the Split-Phase Level modulator (SP) can be connected to:

- Packet Telemetry and AOS encoder
- Reed-Solomon encoder
- Pseudo-Randomiser
- Non-Return-to-Zero Mark encoder
- Convolutional encoder

The output from the Split-Phase Level modulator (SP) can be connected to:

- Sub-Carrier modulator

The input to the Sub-Carrier modulator (SC) can be connected to:

- Packet Telemetry and AOS encoder
- Reed-Solomon encoder
- Pseudo-Randomiser
- Non-Return-to-Zero Mark encode
- Convolutional encoder
- Split-Phase Level modulator

27.8 Operation

27.8.1 Introduction

The DMA interface provides a means for the user to insert Transfer Frames in the Packet Telemetry and AOS Encoder. Depending on which functions are enabled in the encoder, the various fields of the Transfer Frame are overwritten by the encoder. It is also possible to bypass some of these functions for each Transfer Frame by means of the control bits in the descriptor associated to each Transfer Frame. The DMA interface allows the implementation of Virtual Channel Frame Service and Master Channel Frame Service, or a mixture of both, depending on what functions are enabled or bypassed.

27.8.2 Descriptor setup

The transmitter DMA interface is used for transmitting transfer frames on the downlink. The transmission is done using descriptors located in memory.

A single descriptor is shown in table 223 and 224. The number of bytes to be sent is set globally for all transfer frames in the length field in register DMA length register. The the address field of the descriptor should point to the start of the transfer frame. The address must be word-aligned. If the interrupt enable (IE) bit is set, an interrupt will be generated when the transfer frame has been sent (this requires that the transmitter interrupt enable bit in the control register is also set). The interrupt will be generated regardless of whether the transfer frame was transmitted successfully or not. The wrap (WR) bit is also a control bit that should be set before transmission and it will be explained later

in this section.

Table 223. GRTM transmit descriptor word 0 (address offset 0x0)

31	16	15	14	13 10	9	8	7	6	5	4	3	2	1	0
RESERVED		UE	TS	0000	VCE	MCB	-	OCFB	FHECB	-	FECFB	IE	WR	EN

31: 16	RESERVED
15	Underrun Error (UE) - underrun occurred while transmitting frame (status bit only)
14	Time Strobe (TS) - generate a time strobe for this frame
13: 10	RESERVED
9	Virtual Channel Counter Enable (VCE) - enable virtual channel counter generation (using the Idle Frame virtual channel counter)
8	Master Channel Counter Bypass (MCB) - bypass master channel counter generation (TM only)
7	RESERVED
6	Operational Control Field Bypass (OCFB) - bypass operational control field generation
5	Frame Error Header Control Bypass (FECHB) - bypass frame error header control generation (AOS)
4	RESERVED
3	Frame Error Control Field Bypass (FECFB) - bypass frame error control field generation
2	Interrupt Enable (IE) - an interrupt will be generated when the frame from this descriptor has been sent provided that the transmitter interrupt enable bit in the control register is set. The interrupt is generated regardless if the frame was transmitted successfully or if it terminated with an error.
1	Wrap (WR) - Set to one to make the descriptor pointer wrap to zero after this descriptor has been used. If this bit is not set the pointer will increment by 8. The pointer automatically wraps to zero when the 1 KiB boundary of the descriptor table is reached.
0	Enable (EN) - Set to one to enable the descriptor. Should always be set last of all the descriptor fields.

Table 224. GRTM transmit descriptor word 1 (address offset 0x4)

31	2	1	0
ADDRESS			RES

31: 2	Address (ADDRESS) - Pointer to the buffer area from where the packet data will be loaded.
1: 0	RESERVED

To enable a descriptor the enable (EN) bit should be set and after this is done, the descriptor should not be touched until the enable bit has been cleared by the core.

27.8.3 Starting transmissions

Enabling a descriptor is not enough to start a transmission. A pointer to the memory area holding the descriptors must first be set in the core. This is done in the transmitter descriptor pointer register. The address must be aligned to a 1 KiB boundary. Bits 31 to 10 hold the base address of descriptor area while bits 9 to 3 form a pointer to an individual descriptor.The first descriptor should be located at the base address and when it has been used by the core, the pointer field is incremented by 8 to point at the next descriptor. The pointer will automatically wrap back to zero when the next 1 KiB boundary has been reached (the descriptor at address offset 0x3F8 has been used). The WR bit in the descriptors can be set to make the pointer wrap back to zero before the 1 KiB boundary.

The pointer field has also been made writable for maximum flexibility but care should be taken when writing to the descriptor pointer register. It should never be touched when a transmission is active.

The final step to activate the transmission is to set the transmit enable bit in the DMA control register. This tells the core that there are more active descriptors in the descriptor table. This bit should always be set when new descriptors are enabled, even if transmissions are already active. The descriptors must always be enabled before the transmit enable bit is set.

27.8.4 Descriptor handling after transmission

When a transmission of a frame has finished, status is written to the first word in the corresponding descriptor. The Underrun Error bit is set if the FIFO became empty before the frame was completely transmitted. The other bits in the first descriptor word are set to zero after transmission while the second word is left untouched. The enable bit should be used as the indicator when a descriptor can be used again, which is when it has been cleared by the core.

There are multiple bits in the DMA status register that hold transmission status.

The Transmitter Interrupt (TI) bit is set each time a DMA transmission of a transfer frame ended successfully. An interrupt is generated for transfer frames for which the Interrupt Enable (IE) was set in the descriptor. The interrupt is maskable with the Interrupt Enable (IE) bit in the control register.

If a DMA transmission of a transfer frame ends with an underrun error, the Underrun Error (UE) bit in the descriptor will be written to 1 and the Enable (EN) bit will be written to 0; the Transmitter Error (TE) bit in the DMA status register will be set and the Enable (EN) bit in the DMA control register will be cleared; an interrupt will be generated if not masked with the Interrupt Enable (IE) bit in the control register. The Telemetry Encoder will continue sending idle transfer frames after the failed transfer frame. The Telemetry Encoder requires a complete reset after an underrun error before new user transfer frames can be sent.

The Transmitter AMBA error (TA) bit is set when an AMBA AHB error was encountered either when reading a descriptor or when reading transfer frame data. Any active transmissions were aborted and the DMA channel was disabled. It is recommended that the Telemetry Encoder is reset after an AMBA AHB error. The interrupt is maskable with the Interrupt Enable (IE) bit in the control register.

The Transfer Frame Sent (TFS) bit is set whenever a transfer frame has been sent, independently if it was sent via the DMA interface or generated by the core. The interrupt is maskable with the Transfer Frame Interrupt Enable (TFIE) bit in the control register.

The Transfer Frame Failure (TFF) bit is set whenever a transfer frame has failed for other reasons, such as when Idle Frame generation is not enabled and no user Transfer Frame is ready for transmission, independently if it was sent via the DMA interface or generated by the core. The interrupt is maskable with the Transfer Frame Interrupt Enable (TFIE) bit in the control register.

The Transfer Frame Ongoing (TFO) bit is set when DMA transfers are enabled, and is not cleared until all DMA induced transfer frames have been transmitted after DMA transfers are disabled.

27.8.5 Interrupts

The Transfer Frame Sent (TFS) and Transfer Frame Failure (TFF) interrupts are maskable with the Transfer Frame Interrupt Enable (TFIE) bit in the DMA control register, and can be observed via the DMA status register.

The Transmitter Interrupt (TI), Transmitter Error (TE) and Transmitter AMBA Error (TA) interrupts are maskable with the Interrupt Enable (IE) bit in the DMA control register, and can be observed via the DMA status register.

The Time Strobe Interrupt (TSI) is maskable with the Transfer Frame Interrupt Enable (TFIE) bit in the DMA control register.

All interrupts except Time Strobe Interrupt (TSI) are output on interrupt number 29.

The Time Strobe Interrupt (TSI) is output on interrupt number 30.

27.9   Registers

The core is programmed through registers mapped into APB address space.

*Table 225.*GRTM registers

APB address offset	Register
0x80000B00	GRTM DMA Control register
0x80000B04	GRTM DMA Status register
0x80000B08	GRTM DMA Length register
0x80000B0C	GRTM DMA Descriptor Pointer register
0x80000B10	GRTM DMA Configuration register
0x80000B14	GRTM DMA Revision register
0x80000B80	GRTM Control register
0x80000B88	GRTM Configuration register
0x80000B90	GRTM Physical Layer register
0x80000B94	GRTM Coding Sub-Layer register
0x80000B98	GRTM Attached Synchronization Marker
0x80000BA0	GRTM All Frames Generation register
0x80000BA4	GRTM Master Frame Generation register
0x80000BA8	GRTM Idle Frame Generation register
0x80000BD0	GRTM Operational Control Field register

Table 226. GRTM DMA control register

31						5	4	3	2	1	0
RESERVED							TFIE	RST	TXRST	IE	EN

31: 5	RESERVED
4	Transfer Frame Interrupt Enable (TFIE) - enable telemetry frame sent (TFS) and failure (TFF) interrupt, and time strobe interrupt
3	Reset (RST) - reset DMA and telemetry transmitter
2	Reset Transmitter (TXRST) - reset telemetry transmitter
1	Interrupt Enable (IE) - enable DMA interrupt (TI), (TE) and (TA)
0	Enable (EN) - enable DMA transfers

Table 227. GRTM DMA status register

31			8	7	6	5	4	3	2	1	0
RESERVED				TXSTAT	TXRDY	TFO	TFS	TFF	TA	TI	TE

31: 8	RESERVED
7	Transmitter Reset Status (TXSTAT) - telemetry transmitter is in reset mode when set (read-only)
6	Transmitter Ready (TXRDY) - telemetry transmitter ready for operation after setting the TE bit in GRTM control register (read-only)
5	Transfer Frame Ongoing (TFO) - telemetry frames via DMA transfer are on-going (read-only)
4	Transfer Frame Sent (TFS) - telemetry frame interrupt, cleared by writing a logical 1
3	Transfer Frame Failure (TFF) - telemetry transmitter failure, cleared by writing a logical 1
2	Transmitter AMBA Error (TA) - DMA AMBA AHB error, cleared by writing a logical 1
1	Transmitter Interrupt (TI) - DMA interrupt, cleared by writing a logical 1
0	Transmitter Error (TE) - DMA transmitter underrun, cleared by writing a logical 1

Table 228. GRTM DMA length register

31	27 26	16	15	11 10	0
RESERVED	LIMIT-1		RESERVED	LENGTH-1	

31: 27	RESERVED
26: 16	Transfer Limit (LIMIT)- length-1 of data to be fetched by DMA before transfer starts.
	Note: LIMIT must be equal to or less than LENGTH.
	LIMIT must be equal to or less than 128.
	LIMIT must be equal to or larger than 64 for LENGTH > 64
15: 11	RESERVED
10: 0	Transfer Length (LENGTH) - length-1 of data to be transferred by DMA

Table 229. GRTM DMA descriptor pointer register

31	10 9	3 2	0
BASE	INDEX	"000"	

31: 10	Descriptor base (BASE) - base address of descriptor table
9: 3	Descriptor index (INDEX) - index of active descriptor in descriptor table
2: 0	Reserved - fixed to "00"

Table 230. GRTM DMA configuration register (read-only)

31	16 15	0
FIFOSZ	BLOCKSZ	

31: 16	FIFO size (FIFOSZ) - size of FIFO memory in number of bytes (read-only), 128 bytes
15: 0	Block size (BLOCKSZ) - size of block in number of bytes (read-only), 32 bytes

Table 231. GRTM DMA revision register (read-only)

31	17 16	15	8 7	0
RESERVED	TIRQ	REVISION	SUB REVISION	

31: 17	RESERVED
16	Time Strobe Interrupt (TIRQ) - Separate time strobe interrupt supported, set to 1
15: 8	REVISION - Main revision number
	0x00: Initial release
7: 0	SUB REVISION - Sub revision number
	0x00: Initial release
	0x01: Added time interrupt, moved TXRDY bit, added TXSTAT bit, added this revision register

Table 232. GRTM control register

31	1	0
RESERVED		TE

31: 1	RESERVED
0:	Transmitter Enable (TE) - enables telemetry transmitter (should be done after the complete configuration of the telemetry transmitter, including the LENGTH field in the GRTM DMA length register)

Table 233. GRTM configuration register (read-only)

31	21	20	19	18	17	16	15	14	13	12	11	10 9	8	6	5	4	3	2	1	0
RESERVED		AOS	FHEC	-	MCG	-	IDLE	EVC	OCF	FECF	AASM	RS	RS DEPTH		-	PSR	NRZ	CE	SP	SC

31: 21	RESERVED

Table 233. GRTM configuration register (read-only)

20	Advanced Orbiting Systems (AOS) - AOS transfer frame generation implemented
19	Frame Header Error Control (FHEC) - frame header error control implemented, only if AOS also set
18	RESERVED
17	Master Channel Generation (MCG) - master channel counter generation implemented
16	RESERVED
15	Idle Frame Generation (IDLE) - idle frame generation implemented
14	Extended VC Cntr (EVC) - extended virtual channel counter implemented (ECSS)
13	Operational Control Field (OCF) - CLCW implemented
12	Frame Error Control Field (FECF) - transfer frame CRC implemented
11	Alternative ASM (AASM) - alternative attached synchronization marker implemented
10: 9	Reed-Solomon (RS) - reed-solomon encoder implemented, set to "11" E=16 & 8
8: 6	Reed-Solomon Depth (RSDEPTH) - reed-solomon interleave depth -1 implemented, set to 7
5	RESERVED
4	Pseudo-Randomiser (PSR) - pseudo-Randomiser implemented
3	Non-Return-to-Zero (NRZ) - non-return-to-zero - mark encoding implemented
2	Convolutional Encoding (CE) - convolutional encoding implemented
1	Split-Phase Level (SP) - split-phase level modulation implemented
0	Sub Carrier (SC) - sub carrier modulation implemented

Table 234. GRTM physical layer register

31	30	16	15	14	0
SF	SYMBOLRATE		SCF	SUBRATE	

31	Symbol Fall (SF) - symbol clock has a falling edge at start of symbol bit
30: 16	Symbol Rate (SYMBOLRATE) - symbol rate division factor - 1
15	Sub Carrier Fall (SCF) -sub carrier output start with a falling edge for logical 1
14: 0	Sub Carrier Rate (SUBRATE) - sub carrier division factor - 1

Table 235. GRTM coding sub-layer register

31	17	16	15	14	12	11	10	8	7	6	5	4	2	1	0
RESERVED		AASM	RS	RSDEPTH		RS8	RESERVED		PSR	NRZ	CE	CE RATE		SP	SC

31: 17	RESERVED
16	Alternative ASM (AASM) - alternative attached synchronization marker enable. When enabled the value from the GRTM Attached Synchronization Marker register is used, else the standardized ASM value 0x1ACFFC1D is used
15	Reed-Solomon (RS) - reed-solomon encoder enable
14: 12	Reed-Solomon Depth (RSDEPTH) - reed-solomon interleave depth -1
11	Reed-Solomon Rate (RS8) - '0' E=16, '1' E=8
10: 8	RESERVED
7	Pseudo-Randomiser (PSR) - pseudo-Randomiser enable
6	Non-Return-to-Zero (NRZ) - non-return-to-zero - mark encoding enable
5	Convolutional Encoding (CE) - convolutional encoding enable
4: 2	Convolutional Encoding Rate (CERATE): "00-" rate 1/2, no puncturing "01-" rate 1/2, punctured "100" rate 2/3, punctured "101" rate 3/4, punctured "110" rate 5/6, punctured "111" rate 7/8, punctured
1	Split-Phase Level (SP) - split-phase level modulation enable
0	Sub Carrier (SC) - sub carrier modulation enable

Table 236. GRTM attached synchronization marker register

31	0
ASM	

31: 0	Attached Synchronization Marker (ASM) - pattern for alternative ASM, (bit 31 MSB sent first, bit 0 LSB sent last) (The reset value is the standardized alternative ASM value 0x352EF853.)

Table 237. GRTM all frames generation register

31	16	15	14	13	12	11	0
RESERVED		FECF	FHEC	VER		RESERVED	

31: 16	RESERVED
15	Frame Error Control Field (FECF) - transfer frame CRC enabled
14	Frame Header Error Control (FHEC) - frame header error control enabled, only with AOS
13: 12	Version (VER) - Transfer Frame Version - "00" Packet Telemetry, "01" AOS
11: 0	RESERVED

Table 238. GRTM master frame generation register

31				4	3	2	1	0
RESERVED					MC	-	OCF	OW

31: 4	RESERVED
3	Master Channel (MC) - enable master channel counter generation (TM only)
2	RESERVED
1	Operation Control Field (OCF) - enable MC_OCF for master channel
0	Over Write OCF (OW) - overwrite OCF bits 16 and 17 when set.
	When overwritten, bit 16 (No RF Available) and 17 (No Bit Lock) of the OCF are based on the discrete inputs TCRFAVL[4:0] and TCACT[4:0] respectively, where the multiple inputs are OR-ed before assigned to the two bits. Note that if these telecommand input signals are shared with other function via the I/O switch matrix, then the values of it 16 and 17 might reflect the status of the transponders incorrectly.

Table 239. GRTM idle frame generation register

31	22	21	20	19	18	17	16	15	10	9	0
RESERVED		ID LE	O C F	E V C	-	V C C	MC	VCID		SCID	

31: 22	RESERVED
21	Idle Frames (IDLE) - enable idle frame generation
20	Operation Control Field (OCF) - enable OCF for idle frames
19	Extended Virtual Channel Counter (EVC) - enable extended virtual channel counter generation for idle frames (TM only, ECSS)
18	RESERVED
17	Virtual Channel Counter Cycle (VCC) - enable virtual channel counter cycle generation for idle frames (AOS only)
16	Master Channel (MC) - enable separate master channel counter generation for idle frames (TM only)
15: 10	Virtual Channel Identifier (VCID) - virtual channel identifier for idle frames
9: 0	Spacecraft Identifier (SCID) - spacecraft identifier for idle frames

Table 240. GRTM OCF register

31	0
CLCW	

31: 0	Operational Control Field (OCF) - CLCW data (bit 31 MSB, bit 0 LSB)

27.10 Signal definitions

The signals are described in table 241.

Table 241. Signal definitions

Signal name	Type	Function	Active
TCACT[4:0]	Input	Bit Lock (the signals are OR-ed internally)	High
TCRFAVL[4:0]	Input	RF Available (the signals are OR-ed internally)	High
TMDO	Output	Serial bit data	-
TMCLKO	Output	Serial bit data clock	-
TMCLKI	Input	Transponder clock	Rising

28 Clock Gating Unit

28.1 Overview

The CLKGATE clock gating unit provides a mean to save power by disabling the clock to unused functional blocks.

28.2 Operation

The operation of the clock gating unit is controlled through three registers: the unlock, clock enable and core reset registers. The clock enable register defines if a clock is enabled or disabled. A '1' in a bit location will enable the corresponding clock, while a '0' will disable the clock. The core reset register allows to generate a reset signal for each generated clock. A reset will be generated as long as the corresponding bit is set to '1'. The bits in clock enable and core reset registers can only be written when the corresponding bit in the unlock register is 1. If a bit in the unlock register is 0, the corresponding bits in the clock enable and core reset registers cannot be written.

To gate the clock for a core, the following procedure should be applied:

1. Disable the core through software to make sure it does not initialize any AHB accesses

2. Write a 1 to the corresponding bit in the unlock register

3. Write a 1 to the corresponding bit in the core reset register

4. Write a 0 to the corresponding bit in the clock enable register

5. Write a 0 to the corresponding bit in the unlock register

To enable the clock for a core, the following procedure should be applied

1. Write a 1 to the corresponding bit in the unlock register

2. Write a 1 to the corresponding bit in the core reset register

3. Write a 1 to the corresponding bit in the clock enable register

4. Write a 0 to the corresponding bit in the core reset register

5. Write a 0 to the corresponding bit in the unlock register

The cores connected to the clock gating unit are defined in the table below:

*Table 242.*Clocks controlled by CLKGATE unit in the GR712RC

Bit	Functional module
0	GRETH 10/100 Mbit Ethernet MAC
1	GRSPW SpaceWire link 0
2	GRSPW SpaceWire link 1
3	GRSPW SpaceWire link 2
4	GRSPW SpaceWire link 3
5	GRSPW SpaceWire link 4
6	GRSPW SpaceWire link 5
7	CAN core 0 & 1
8	Proprietary, never enable this function
9	CCSDS Telemetry encoder
10	CCSDS Telecommand decoder
11	MIL-STD-1553 BC/RT/MT

28.3 Registers

Table 30 shows the clock gating unit registers. The base address for the registers is 0x80000D00.

*Table 243.*Clock unit control registers

Address	Functional module	Reset value (binary)
0x80000D00	Unlock register	000000000000
0x80000D04	Clock enable register	000000000110
0x80000D08	Core reset register	111111111001

After reset, the clocks for SpaceWire cores 0 and 1 are enabled to allow connection through RMAP without additional software. The remaining SpaceWire core clocks are disabled.

Cobham Gaisler AB
Kungsgatan 12
411 19 Göteborg
Sweden
www.cobham.com/gaisler
sales@gaisler.com
T: +46 31 7758650
F: +46 31 421407